

Lecture 11: Random Variables and Independence

Anup Rao

April 24, 2019

We discuss the use of hashing and pairwise independence.

RANDOM VARIABLES ARE A VERY USEFUL way to model many situations. Recall the law of total probability, that we discussed before:

Fact 1 (Law of Total Probability). *If A_1, A_2, \dots, A_n are disjoint events that form a partition of the whole sample space, and B is another event, then*

$$p(B) = p(A_1 \cap B) + p(A_2 \cap B) + \dots + p(A_n \cap B).$$

One consequence of the Law of Total Probability is that whenever you have a random variable X and an event E , we have

$$p(E) = \sum_x p(X = x) \cdot p(E|X = x).$$

Just like with events, we say that two random variables X, Y are *independent* if having information about one gives you no information about the other. Formally, X, Y are independent if for every x, y

$$p(X = x, Y = y) = p(X = x) \cdot p(Y = y).$$

This means that $p(X = x|Y = y) = p(X = x)$, and $p(Y = y|X = x) = p(Y = y)$.

Example: Even Number of Heads

Suppose you toss n coins. What is the probability that the number of heads is even? We already calculated this the *hard* way. Let us use an appropriate random variable to calculate this the easy way.

Let X denote the outcome of tossing the first $n - 1$ coins. Let E denote the event that the number of heads is even. Then:

$$p(E) = \sum_{x \in \{H,T\}^{n-1}} p(X = x) \cdot p(E|X = x).$$

Now observe that no matter what x is, $p(E|X = x) = 1/2$. This is because after you fix the first $n - 1$ coin-tosses, the last coin toss is still uniformly random. If x has an even number of heads, then the total number of tosses is even when the last coin toss is a tail and so $p(E|X = x) = 1/2$. If x has an odd number of heads, the opposite is true, and still $p(E|X = x) = 1/2$. So, we get:

See notes for Lecture 7 on April 15.

$$\begin{aligned}
 p(E) &= \sum_{x \in \{H,T\}^{n-1}} p(X = x) \cdot p(E|X = x) \\
 &= (1/2) \cdot \sum_{x \in \{H,T\}^{n-1}} p(X = x) = 1/2.
 \end{aligned}$$

Random variables X_1, \dots, X_n are said to be mutually independent if for every x_1, \dots, x_n , we have

$$\begin{aligned}
 &p(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n) \\
 &= p(X_1 = x_1) \cdot p(X_2 = x_2) \dots p(X_n = x_n).
 \end{aligned}$$

It is possible that the random variables are not independent, even though every pair of random variables is independent. For example, suppose $X, Y \in \{0, 1\}$ are uniformly random bits that are independent of each other. Set

$$Z = X + Y \pmod{2}.$$

Z is not independent of X, Y . However, we see that X, Z are mutually independent. Similarly, Y, Z are also mutually independent. Another example: let $X \in \{0, 1\}^n$ be a uniformly random string such that $X_1 + X_2 + \dots + X_n = 0 \pmod{2}$. Then X_1, \dots, X_n are not mutually independent. However, every $n - 1$ of those variables *are* mutually independent.

In general, we say that X_1, \dots, X_n are k -wise independent if every k variables are independent.

Some applications

Hashing

Here is an application of pairwise independence to a data structures problem. Suppose we have subsets $S \subseteq T \subseteq [n]$. We would like to store S and allow membership queries into S . We want to preprocess and store S so that at runtime, if we get an element $i \in T$, we can quickly deduce whether $i \in S$ or not.

An obvious way to solve this problem is to just store S as a sorted list. Then if we want to check whether or not $i \in S$, we can do binary search on the list. This takes time proportional to $\log |S|$, but creating the sorted list takes time about $|S|^2$. Can we speed this up?

Another way to solve this problem is to store the indicator vector of S as a vector in $\{0, 1\}^n$. Then membership can be checked in a single step. But this requires n bits of space! Do we really need n bits of space?

For example, imagine that T is the set of all students at UW, $[n]$ is the set of all people in the world, and S is the set of students enrolled in CSE312. We would like to be able to quickly check whether a student at UW is enrolled in CSE312 or not.

The idea of *hashing* is to map each element of $[n]$ to a much smaller set of hashes, in such a way that collisions are unlikely. Then membership can be checked quickly, and using small space at the same time. Let $h : [n] \rightarrow [m]$ be a uniformly random function, with $m \ll n$. Since h is uniformly random, $h(1), h(2), \dots, h(n)$ are mutually independent and uniformly random in $[m]$.

Now, here is the idea.

1. Pick a random function $h : [n] \rightarrow [m]$ as described above.
2. Initialize a bit vector $Z \in \{0, 1\}^m$ by setting all coordinates to 0.
3. For every element $j \in S$, set $Z_{h(j)} = 1$.

To check if $i \in T$ belongs to S or not, inspect $Z_{h(i)}$. If $Z_{h(i)} = 0$, we conclude that $i \notin S$. If $Z_{h(i)} = 1$, we conclude that $i \in S$.

The algorithm makes an error only if there are distinct i, j such that $i \in S, j \in T$ and $h(i) = h(j)$. In other words, an error can happen only if h is not an injective function on $S \cup T$.

Claim 2. $p(h \text{ is not injective on } S \cup T) \leq (1/m) \cdot \binom{|S \cup T|}{2}$

Proof. Let $E_{i,j}$ denote the event that $h(i) = h(j)$. Then $p(E_{i,j}) = 1/m$. If E denotes the event that h is not injective on T , then we see that $E = \bigcup_{i \neq j, i \in S, j \in T} E_{i,j}$. So

$$p(E) \leq \sum_{i \neq j, i \in S, j \in T} p(E_{i,j}) \leq (1/m) \cdot \binom{|S \cup T|}{2}.$$

□

So, as long as we set $m \gg \binom{|S \cup T|}{2}$, the probability that h is not injective on $S \cup T$ is extremely small. This means, we only need to store a bit vector whose length is about $O(|S \cup T|^2)$.

All of this is great, but it is not actually practical to sample a uniformly random function $h : [n] \rightarrow [m]$ and store it. Imagine trying to sample a uniformly random hash for every single possible human name. Luckily, it turns out that it is enough to sample a pairwise independent hash. We say that the hash function is pairwise independent if for every $i \neq j$, $h(i)$ and $h(j)$ are independent. If you look at the proof of the claim above, you see that we only need the hash function to be pairwise independent for the analysis to go through.

Moreover, there are several nice constructions that give pairwise independent functions. For example, define $h : \{0, 1\}^{k_1} \rightarrow \{0, 1\}^{k_2}$ by picking H_1, H_2 to be uniformly random $k_1 \times k_2$ matrices with 0/1 entries and setting $h(x) = H_1 x + H_2 \pmod{2}$. Then the outputs of h are pairwise independent. h can also be computed very quickly. So, setting $k_2 \gg \log |S \cup T|^2$ be a large enough, we get a hash function that allows us to store S and quickly answer membership queries.