CSE 331
Software Design & Implementation

Hal Perkins

Summer 2019

Specifications

Administrivia

- HW1: due Mon. night, 11 pm
  - Can use a late day, but you *really* don't want to ☺
- HW3 (logistics) due Wed. night, 11 pm
  - How's this one going?
- Next set of lecture slides for this week posted today and tomorrow
  - We're going to have to cover a fair amount because of 9-week quarter + missing July 4 section, so we may skip a few things in lecture that everyone can read themselves. Suggestion: look at slides before lectures
- Lots of new readings related to next few lectures – dig in if you haven't already
  - Readings on calendar are *sections* (items) in books

2 Goals of Software System Building

- Building the *right system*

    – Does the program meet the user's needs?

    – Determining this is usually called *validation*

- Building the *system right*

    – Does the program meet the specification?

    – Determining this is usually called *verification*

- CSE 331: the second goal is the focus – creating a correctly functioning artifact

    – Surprisingly hard to specify, design, implement, test, and debug even simple programs

Where we are

- We've started to see how to reason about code
- We'll build on those skills in many places:

    - *Specification*: What are we supposed to build?

    - *Design*: How do we decompose the job into manageable pieces? Which designs are "better"?

    - *Implementation*: Building code that meets the specification

    - *Testing*: Systematically finding problems

    - *Debugging*: Systematically fixing problems

    - *Maintenance*: How does the artifact adapt over time?

    - *Documentation*: What do we need to know to do these things? How/where do we write that down?

The challenge of scaling software

- Small programs are simple and malleable
    - Easy to write
    - Easy to change

- Big programs are (often) complex and inflexible
    - Hard to write
    - Hard to change

- Why does this happen?
    - Because *interactions* become unmanageable

- How do we keep things simple and malleable?
    - Divide and conquer!

A discipline of modularity

- Two ways to view a program:
  - The client's view (how to use it)
  - The implementer's view (how to build it)

- Apply implementer and client views to system parts:
  - While implementing one part, consider yourself a client of any other parts it depends on
  - Ignore the implementation of those other parts
  - Minimizes interactions between parts

- Formalized through the idea of a *specification*

A specification is a contract

- A set of requirements agreed to by the user and the manufacturer of the product

  – Describes their expectations of each other

- Facilitates simplicity via *two-way* isolation

  – Isolate client from implementation details

  – Isolate implementer from how the part is used

  – Discourages implicit, unwritten expectations

- Facilitates change

  – The code can change

  – The specification cannot

Isn't the interface sufficient?

The interface defines the boundary between implementers and users:

```java
public class List<E> {
    public E get(int x) { return null; }
    public void set(int x, E y){}
    public void add(E) {}
    public void add(int, E){}

    …
    public static <T> boolean isSub(List<T>, List<T>){
        return false;
    }
}
```

Interface provides the *syntax and types*

But nothing about the *behavior and effects*
- *Provides too little information to clients*

Note: Code above is right concept but might not be (completely) legal Java
- slides will often gloss over details to get main ideas to fit

Why not just read code?

```java
static <T> boolean sub(List<T> src, List<T> part) {
    int part_index = 0;
    for (T elt : src) {
        if (elt.equals(part.get(part_index))) {
            part_index++;
            if (part_index == part.size()) {
                return true;
            }
        } else {
            part_index = 0;
        }
    }
    return false;
}
```

Why are you better off with a specification?

Code is complicated

- Code gives more detail than needed by client

- Understanding or even reading every line of code is an excessive burden
  - Suppose you had to read source code of Java libraries to use them
  - Same applies to developers of different parts of the libraries

- Client cares only about *what* the code does, not *how* it does it

Code is ambiguous

- Code seems unambiguous and concrete
  - But which details of code's behavior are essential, and which are incidental?

- Code invariably gets rewritten
  - Clients need to know what they can rely on
    - What properties will be maintained over time?
    - What properties might be changed by future optimization, improved algorithms, or bug fixes?
  - Implementer needs to know what features the client depends on, and which can be changed

Comments are essential

Typical comments often convey only an informal, general idea of what that the code does:

```
// This method checks if "part" appears as a
// sub-sequence in "src"
static <T> boolean sub(List<T> src, List<T> part){
   ...
}
```

Problem: ambiguity remains
- – What if **src** and **part** are both empty lists?
- – When does the function return **true**?

From vague comments to specifications

- Roles of a specification:
  - Client agrees to rely *only* on information in the description in their use of the part
  - Implementer of the part promises to support everything in the description
    - Otherwise is perfectly at liberty

- Sadly, much code lacks a specification
  - Clients often work out what a method/class does in ambiguous cases by running it and depending on the results
  - Leads to bugs and programs with unclear dependencies, reducing simplicity and flexibility

Recall the sublist example

```
static <T> boolean sub(List<T> src, List<T> part) {
    int part_index = 0;
    for (T elt : src) {
        if (elt.equals(part.get(part_index))) {
            part_index++;
            if (part_index == part.size()) {
                return true;
            }
        } else {
            part_index = 0;
        }
    }
    return false;
}
```

A more careful description of `sub`

// Check whether "part" appears as a sub-sequence in "src"

needs to be given some caveats (why?):

// * src and part cannot be null

// * If src is empty list, always returns false

// * Results may be unexpected if partial matches

//   can happen right before a real match; e.g.,

//   list (1,2,1,3) will not be identified as a

//   sub sequence of (1,2,1,2,1,3).

or replaced with a more detailed description:

// This method scans the "src" list from beginning

// to end, building up a match for "part", and

// resetting that match every time that...

It's better to *simplify* than to *describe* complexity!

A complicated description suggests poor design
  – Rewrite **sub** to be more sensible, and easier to describe

```
// returns true iff possibly empty sequences A, B exist such that
//   src = A + part + B
// where "+" is sequence concatenation
static <T> boolean sub(List<T> src, List<T> part) {
```

- Mathematical flavor not always necessary, but often helps avoid ambiguity
- "Declarative" style is important
  – Avoid reciting or depending on implementation details

Sneaky fringe benefit of specs #1

- The discipline of writing specifications changes the incentive structure of coding
    - Rewards code that is easy to describe and understand
    - Punishes code that is hard to describe and understand
        - Even if it is shorter or easier to write

- If you find yourself writing complicated specifications, it is an incentive to redesign
    - In `sub`, code that does exactly the right thing may be slightly slower than a hack that assumes no partial matches before true matches, but cost of forcing client to understand the details is too high

Writing specifications with Javadoc

- Javadoc
  - Sometimes can be daunting; get used to using it

- Javadoc convention for writing specifications
  - Method signature (prototype – name, parameters, result type)
  - Text description of method
  - `@param`: description of what gets passed in
  - `@return`: description of what gets returned
  - `@throws`: exceptions that may occur

Example: Javadoc for `String.contains`

`public boolean contains(CharSequence s)`

Returns true if and only if this string contains the specified sequence of char values.

Parameters:

s - the sequence to search for

Returns:

true if this string contains s, false otherwise

Throws:

NullPointerException – if s is null

Since:

1.5

CSE 331 specifications

- The *precondition*: constraints that hold before the method is called (if not, all bets are off – no guarantees about what method will do)
  - **@requires**: spells out any obligations on client

- The *postcondition*: constraints that hold after the method is called (if the precondition held)
  - **@modifies**: lists objects that may be affected by method; any object not listed is guaranteed to be unchanged
  - **@throws**: lists possible exceptions and conditions under which they are thrown (Javadoc uses this too)
  - **@effects**: gives guarantees on final state of modified objects
  - **@returns**: describes return value (Javadoc uses this too)

Example 1

static **\<T\>** int change(List\<T\> lst, T oldelt, T newelt)

    requires   lst, oldelt, and newelt are non-null.
                       oldelt occurs in lst.

    modifies   lst

    effects   change the first occurrence of oldelt in lst to newelt
                       & makes no other changes to lst

    returns   the position of the element in lst that was oldelt and
                       is now newelt

```
static <T> int change(List<T> lst,
                      T oldelt, T newelt) {
    int i = 0;
    for (T curr : lst) {
        if (curr == oldelt) {
            lst.set(newelt, i);
            return i;
        }
        i = i + 1;
    }
    return -1;
}
```

Example 2

static List<Integer> zipSum(List<Integer> lst1, List<Integer> lst2)

requires	lst1 and lst2 are non-null. lst1 and lst2 are the same size.
modifies	none
effects	none
returns	a list of same size where the ith element is the sum of the ith elements of lst1 and lst2

```
static List<Integer> zipSum(List<Integer> lst1
                            List<Integer> lst2) {
  List<Integer> res = new ArrayList<Integer>();
  for(int i = 0; i < lst1.size(); i++) {
     res.add(lst1.get(i) + lst2.get(i));
  }
  return res;
}
```

Example 3

static void listAdd(List<Integer> lst1, List<Integer> lst2)

     requires   lst1 and lst2 are non-null.

                    lst1 and lst2 are the same size.

     modifies   lst1

     effects   ith element of lst2 is added to the ith element of lst1

     returns   none

```java
static void listAdd(List<Integer> lst1,
                    List<Integer> lst2) {
  for(int i = 0; i < lst1.size(); i++) {
     lst1.set(i, lst1.get(i) + lst2.get(i));
  }
}
```

Should requires clause be checked?

- If the client calls a method without meeting the precondition, the code is free to do *anything*
  - Including pass corrupted data back
  - It is polite, nevertheless, to *fail fast*: to provide an immediate error, rather than permitting mysterious/silent bad behavior

- Preconditions are common in "helper" methods/classes
  - In public libraries, it's friendlier to deal with all possible input
  - *But: binary search would normally impose a pre-condition rather than simply failing if list is not sorted. Why?*

- Rule of thumb: Check if cheap to do so
  - *Example: list has to be non-null → check*
  - *Example: list has to be sorted → skip*
  - A quality implementation will check preconditions whenever it is inexpensive and convenient to do so
    - Defensive programming

Sneaky fringe benefit of specs #2

- Specification means that client doesn't need to look at implementation
  - So the code might not even exist yet!

- Write specifications first, make sure system will fit together, and then assign separate implementers to different modules
  - Allows teamwork and parallel development
  - Also helps with testing (future topic)

Upgrading a library

- Your program uses a library
- Can you use a different library?
- Can you use a new version?

We want an objective test

You can upgrade if the specification of the new version is stronger

- It makes at least as many promises
- Example:
    - Weaker spec: returns the elements
    - Stronger spec: returns the elements in sorted order

Two specifications for find which is stronger?

```
int find(int[] a, int value) {
    for (int i=0; i<a.length; i++) {
        if (a[i]==value)
            return i;
    }
    return -1;
}
```

- Specification A
  - requires: value occurs in **a**
  - returns: **i** such that **a[i]** = **value**

- Specification B
  - requires: value occurs in **a**
  - returns: *smallest* **i** such that **a[i]** = **value**

Two specifications for find
Which is stronger?

```
int find(int[] a, int value) {
    for (int i=0; i<a.length; i++) {
        if (a[i]==value)
            return i;
    }
    return -1;
}
```

- Specification A
    - requires: value occurs in `a`
    - returns: `i` such that `a[i]` = `value`

- Specification C
    - returns: `i` such that `a[i]` = `value`, or `-1` if value is not in `a`

Stronger and weaker specifications

- A stronger specification:
  - Promises more
    - Effects clause is harder to satisfy and/or lists fewer objects in modifies clause
    - Consequence: may be harder to implement
  - Asks less of the client
    - Requires clause can be easier to satisfy (weaker!)
  - May be harder to implement but is likely easier to use
    - More guarantees, more predictable

Satisfaction of a specification

Let P be an implementation and S a specification

P satisfies S if and only if
  – Every behavior of P is permitted by S
  – "The behavior of P is a subset of S"

The statement "P is correct" is meaningless!
  – Though often made!

If P does not satisfy S, either (or both!) could be "wrong"
  – *"One person's feature is another person's bug."*
  – Usually better to change the program than the spec

Substitutability

- Suppose that
  - I_1 and I_2 satisfy specification S
  - P uses I_1 as a component (and relies only on S)
- Then P can use I_2

- Further, suppose that
  - I_3 satisfies S_3 which is stronger than S
- Then P can use I_3

- Fact: If specification S_3 is stronger than S_1, then for any implementation I, I satisfies S_3 => I satisfies S_1

Why compare specifications?

We wish to relate procedures to specifications
- Does the procedure satisfy the specification?
- Has the implementer succeeded?

We wish to compare specifications to one another
- Which specification (if either) is stronger?
- A procedure satisfying a stronger specification can be used anywhere that a weaker specification is required
  - Substitutability principle

Given two specifications, they may be *incomparable*
- Neither is weaker/stronger than the other
- *Some* implementations might still satisfy them both

"Strange" case: @throws

[Prior versions of course, including old exams, were clumsy/wrong about this]

Compare:

S1:

  @throws FooException if x<0

  @return x+3

S2:

  @return x+3

- These are *incomparable* because they promise different, incomparable things when x<0
- Both are *stronger* than @requires x>=0; @return x+3

Which is better?

- Stronger does not always mean better!

- Weaker does not always mean better!

- Strength of specification trades off:
  - Usefulness to client
  - Ease of simple, efficient, correct implementation
  - Promotion of reuse and modularity
  - Clarity of specification itself

- "It depends"

Review:
Ways to compare specifications

- A *stronger* specification is satisfied by fewer implementations

- A stronger specification has

  - *weaker* preconditions (note contravariance)

  - stronger postcondition

  - fewer modifications

  Can be checked by hand

- A stronger specification has a (logically) stronger formula – can be checked by tools

- A stronger specification has a smaller transition relation (intuition "stronger = smaller"; fewer choices)

Specification style

- The point of a specification is to be helpful
  - Formalism helps, excessive formalism doesn't
- A specification should be
  - coherent: not too many cases
  - informative: (bad example in Java library, `HashMap.get`; what does result of `null` mean?)
  - strong enough: to do something useful, to provide guarantees
  - weak enough: to permit (efficient) implementation

A method should do only one thing

- Either return a value, or have some *side effect* (modification of program state)
  - Poor style to have *both* @effects and @returns
  - Exception: methods like `HashMap.put` that return an old value as part of an update