

# Introduction to Data Management

## RA examples and Subqueries

Paul G. Allen School of Computer Science and Engineering  
University of Washington, Seattle

# Announcements

- HW 2 due tomorrow
- Office hours adjusted
  - Shana has office hours Thursday at 3:30pm, Gates 150
  - TAs and I often available by appointment
- Midterm 2 weeks from today – will be added to calendar

# Goals for Today

- Recap RA
- Use SQL queries to assist other SQL queries
- Subqueries give you 99% of the tools for queries you can think of

# Recap RA Operators

- These are all the operators you will see in this class
  - We'll profile these one at a time



Inner Join



Grouping &  
Aggregation



Sort



Cartesian Product



Union



Duplicate  
Elimination



Selection



Intersection



Projection



Difference

# Simple Example

## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

$\pi_{Job}$

|

*Payroll P*

# Simple Example

## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

$\pi_{Job}$

|

*Payroll P*

```
SELECT Job
FROM Payroll;
```

# Simple Example

## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

$\pi_{Job}$

|

*Payroll P*

```
SELECT Job
FROM Payroll;
```

Job
TA
TA
Prof
Prof

# Simple Example

## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job
TA
Prof

```
SELECT DISTINCT Job
FROM Payroll;
```



# Simple Example


## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

## Job

TA  
Prof

```
SELECT DISTINCT Job
FROM Payroll;
```



# Simple Example


## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

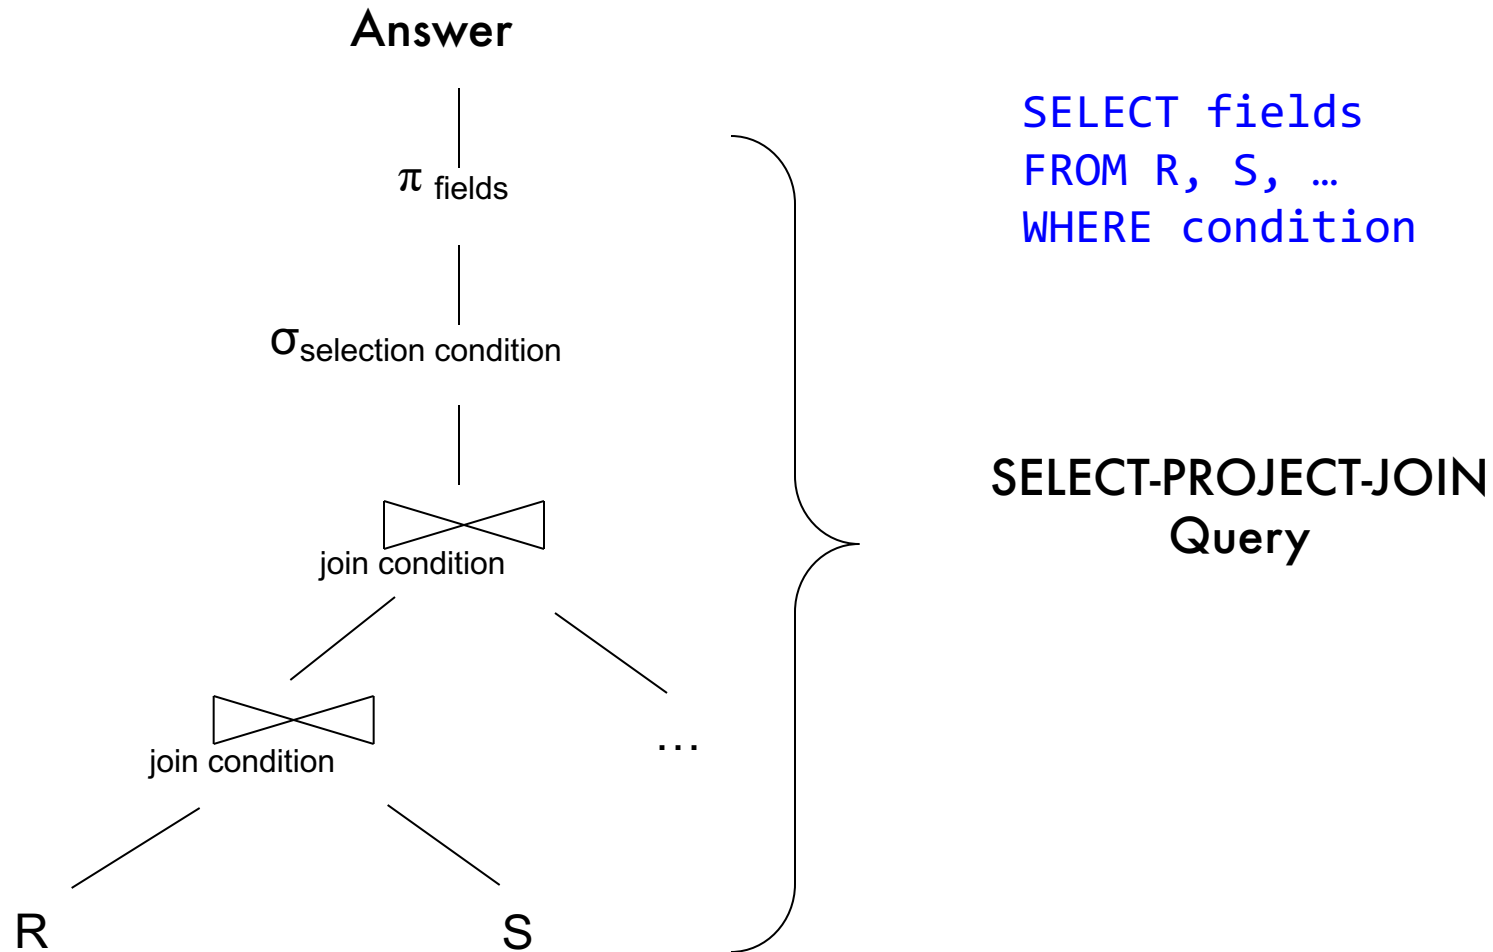
## Job

TA  
Prof

```
SELECT DISTINCT Job
FROM Payroll;
```



# Typical Plan for a Query



# Aggregation Order

How is aggregation processed internally?

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

# Aggregation RA

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
  FROM Payroll
 GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
  FROM Payroll
 GROUP BY Job
HAVING MIN(Salary) > 80000
```

Job	maxSal	minSal
TA	60000	50000
Prof	100000	90000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\sigma_{minSal > 80000}$

Job	maxSal	minSal
TA	60000	50000
Prof	100000	90000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...



# Aggregation RA

```
SELECT Job, MAX(Salary)
  FROM Payroll
 GROUP BY Job
HAVING MIN(Salary) > 80000
```

Job	maxSal	minSal
Prof	100000	90000

$\sigma_{minSal > 80000}$

Job	maxSal	minSal
TA	60000	50000
Prof	100000	90000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
  FROM Payroll
 GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\pi_{Job, maxSal}$

Job	maxSal	minSal
Prof	100000	90000

$\sigma_{minSal > 80000}$

Job	maxSal	minSal
TA	60000	50000
Prof	100000	90000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
  FROM Payroll
 GROUP BY Job
HAVING MIN(Salary) > 80000
```

Job	maxSal
Prof	100000

$\pi_{Job, maxSal}$

Job	maxSal	minSal
Prof	100000	90000

$\sigma_{minSal > 80000}$

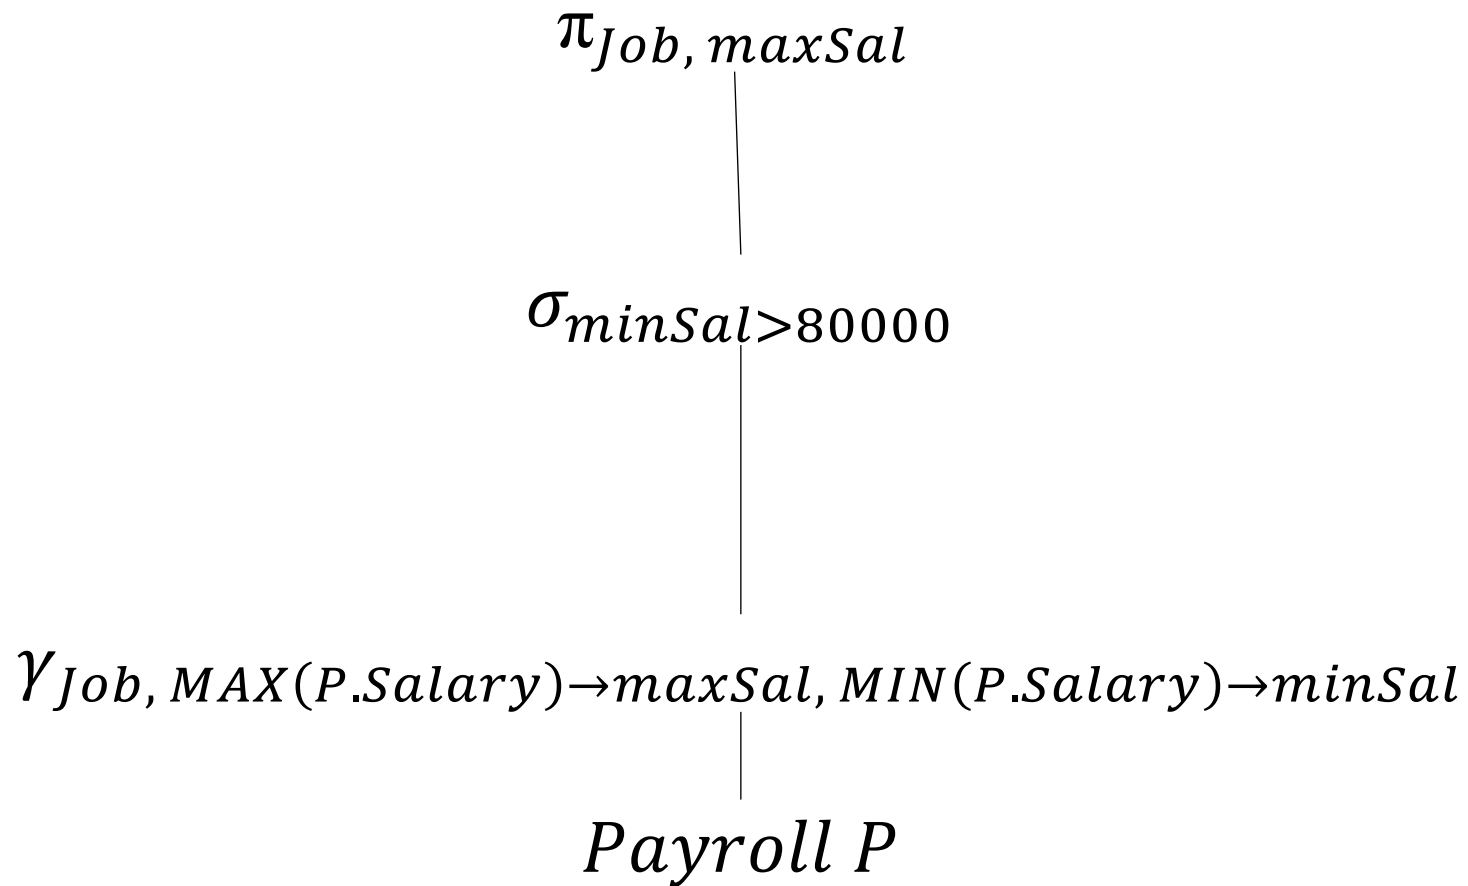
Job	maxSal	minSal
TA	60000	50000
Prof	100000	90000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

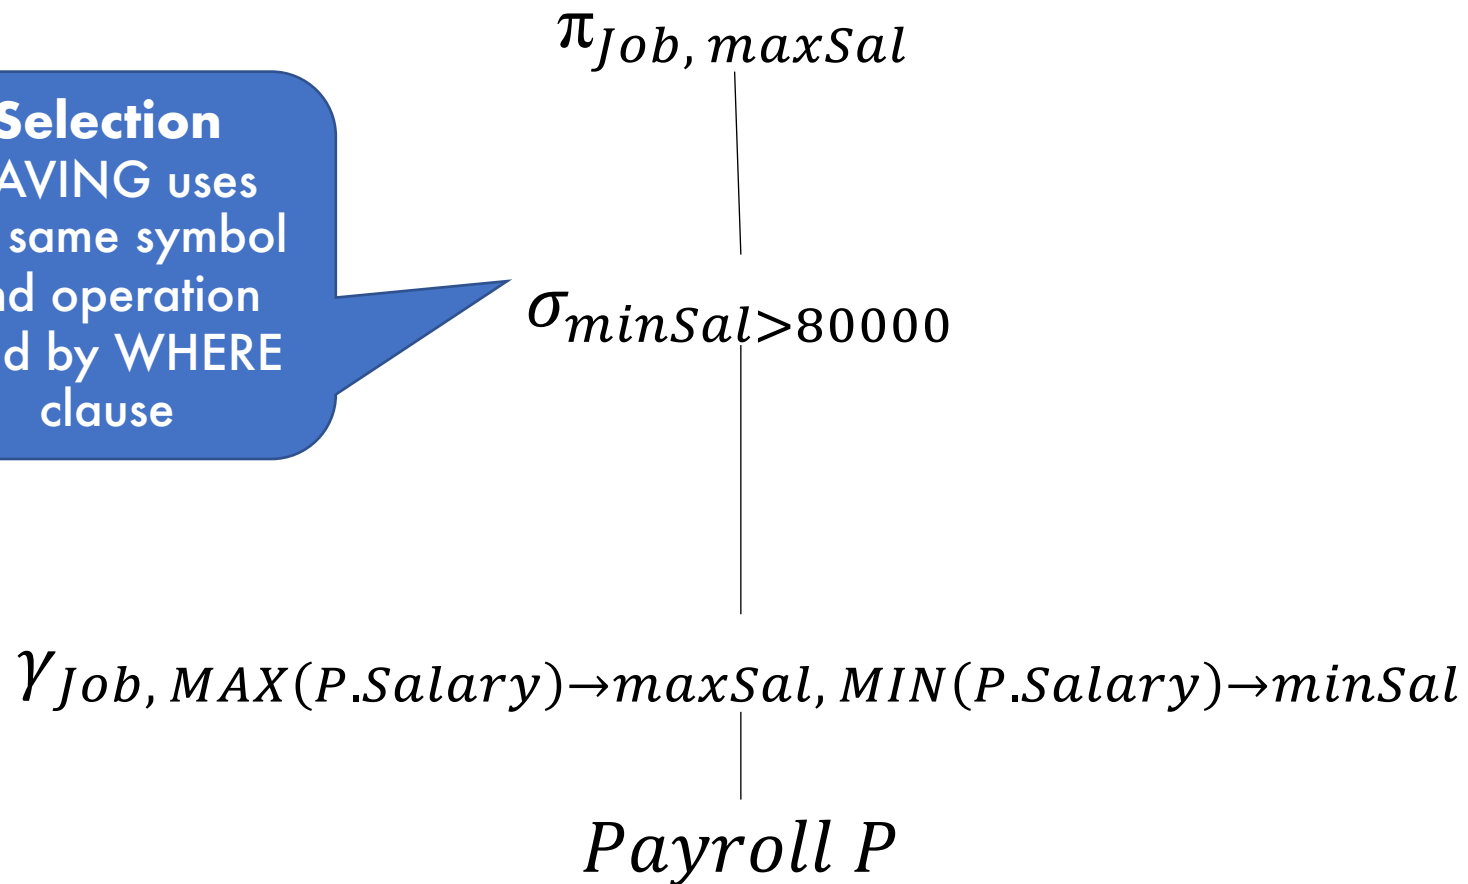
```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```




# Aggregation RA

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

**Selection**  
HAVING uses  
the same symbol  
and operation  
used by WHERE  
clause



# Typical Plan for an Aggregate Query



SELECT fields  
FROM R, S, ...  
WHERE condition  
GROUP BY fields  
HAVING condition

# Recap – The Witnessing Problem

- A question pattern that asks for data associated with a maxima of some value
  - Observed how to do it with grouping
  - “Self join” on values you find the maxima for
  - GROUP BY to deduplicate one side of the join
  - HAVING to compare values with respective maxima

# Outline

- Witnessing via subquery
- Subquery mechanics
  - Set/bag operations
  - SELECT
  - FROM
  - WHERE/HAVING
- Decorrelation and unnesting along the way



# The Witnessing Problem Simplified

- Wanted to join respective maxima
  - GROUP BY technique was interesting
  - Last time people suggested that we can **compute the maxima first then join**

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000

Return the person (or people) with the highest salary for each job type

# The Witnessing Problem - Previously

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000

Return the person with the highest salary for each job type

```
SELECT P1.Name, MAX(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P2.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(P2.Salary)
```

# The Witnessing Problem - Previously

```
SELECT P1.Name, MAX(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P2.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(P2.Salary)
```

Join on "original"  
grouping attributes

P1

P2

UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
567	Magda	Prof	90000	567	Magda	Prof	90000
567	Magda	Prof	90000	789	Dan	Prof	100000
789	Dan	Prof	100000	789	Dan	Prof	100000
789	Dan	Prof	100000	567	Magda	Prof	90000

# The Witnessing Problem - Previously

```
SELECT P1.Name, MAX(P2.Salary)
  FROM Payroll AS P1, Payroll AS P2
 WHERE P1.Job = P2.Job
  GROUP BY P2.Job, P1.Salary, P1.Name
  HAVING P1.Salary = MAX(P2.Salary)
```

Group on additional  
attributes that you are  
argmax-ing for

P1

P2

UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
567	Magda	Prof	90000	567	Magda	Prof	90000
567	Magda	Prof	90000	789	Dan	Prof	100000
789	Dan	Prof	100000	789	Dan	Prof	100000
789	Dan	Prof	100000	567	Magda	Prof	90000

# The Witnessing Problem - Previously

```
SELECT P1.Name, MAX(P2.Salary)
  FROM Payroll AS P1, Payroll AS P2
 WHERE P1.Job = P2.Job
  GROUP BY P2.Job, P1.Salary, P1.Name
  HAVING P1.Salary = MAX(P2.Salary)
```

Group on additional  
attributes that you are  
argmax-ing for

P1

P2

UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
567	Magda	Prof	90000	567	Magda	Prof	90000
567	Magda	Prof	90000	789	Dan	Prof	100000
789	Dan	Prof	100000	789	Dan	Prof	100000
789	Dan	Prof	100000	567	Magda	Prof	90000

# The Witnessing Problem - Previously

```

SELECT P1.Name, MAX(P2.Salary)
  FROM Payroll AS P1, Payroll AS P2
 WHERE P1.Job = P2.Job
 GROUP BY P2.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(P2.Salary)
    
```

P1

P2

UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
567	Magda	Prof	90000	567	Magda	Prof	90000
567	Magda	Prof	90000	789	Dan	Prof	100000
789	Dan	Prof	100000	789	Dan	Prof	100000
789	Dan	Prof	100000	567	Magda	Prof	90000

# The Witnessing Problem Simplified

```
SELECT P1.Name, MAX(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P2.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(P2.Salary)
```

We can compute  
the same thing!



```
WITH MaxPay AS
    (SELECT P1.Job AS Job,
            MAX(P1.Salary) AS Salary
     FROM Payroll AS P1
     GROUP BY P1.Job)
SELECT P.Name, P.Salary
FROM Payroll AS P, MaxPay AS MP
WHERE P.Job = MP.Job AND
        P.Salary = MP.Salary
```

# The Witnessing Problem Simplified

Useful intermediate result!

```
WITH MaxPay AS
    (SELECT P1.Job AS Job,
           MAX(P1.Salary) AS Salary
     FROM Payroll AS P1
    GROUP BY P1.Job)
```

```
SELECT P.Name, P.Salary
  FROM Payroll AS P, MaxPay AS MP
 WHERE P.Job = MP.Job AND
        P.Salary = MP.Salary
```

**MaxPay**

Job	Salary
TA	60000
Prof	100000



# The Witnessing Problem Simplified

```
WITH MaxPay AS
  (SELECT P1.Job AS Job,
         MAX(P1.Salary) AS Salary
   FROM Payroll AS P1
  GROUP BY P1.Job)
SELECT P.Name, P.Salary
  FROM Payroll AS P, MaxPay AS MP
 WHERE P.Job = MP.Job AND
       P.Salary = MP.Salary
```

**Payroll**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**MaxPay**

Job	Salary
TA	60000
Prof	100000

# The Witnessing Problem Simplified

```
WITH MaxPay AS
  (SELECT P1.Job AS Job,
         MAX(P1.Salary) AS Salary
   FROM Payroll AS P1
  GROUP BY P1.Job)
SELECT P.Name, P.Salary
  FROM Payroll AS P, MaxPay AS MP
 WHERE P.Job = MP.Job AND
       P.Salary = MP.Salary
```

**Join predicate**

**Payroll**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**MaxPay**

Job	Salary
TA	60000
Prof	100000

# The Witnessing Problem Simplified

```
WITH MaxPay AS
  (SELECT P1.Job AS Job,
         MAX(P1.Salary) AS Salary
   FROM Payroll AS P1
  GROUP BY P1.Job)
SELECT P.Name, P.Salary
  FROM Payroll AS P, MaxPay AS MP
 WHERE P.Job = MP.Job AND
       P.Salary = MP.Salary
```

## Selection Predicate

### Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

### MaxPay

Job	Salary
TA	60000
Prof	100000

# The Witnessing Problem Simplified

```
WITH MaxPay AS
  (SELECT P1.Job AS Job,
         MAX(P1.Salary) AS Salary
   FROM Payroll AS P1
  GROUP BY P1.Job)
SELECT P.Name, P.Salary
 FROM Payroll AS P, MaxPay AS MP
WHERE P.Job = MP.Job AND
      P.Salary = MP.Salary
```

Solving a subproblem  
can make your life easy

**Payroll**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**MaxPay**

Job	Salary
TA	60000
Prof	100000

# The Punchline about Subqueries

- Subqueries can be interpreted as **single values** or as **whole relations**
  - A single value (a 1x1 relation) can be returned as part of a tuple
  - A relation can be:
    - Used as input for another query
    - Checked for containment of a value

# Set Operations

- SQL mimics set theory in many ways
  - Bag = duplicates allowed
  - **UNION (ALL)** → set union (bag union)
  - **INTERSECT (ALL)** → set intersection (bag intersection)
  - **EXCEPT (ALL)** → set difference (bag difference)
- SQL Server Management Studio 2017
  - INTERSECT ALL not supported
  - EXCEPT ALL not supported



# Set Operations

- SQL set-like operators basically slap two queries together (not really a subquery...)

```
SELECT * FROM T1  
UNION  
SELECT * FROM T2;
```

# Subqueries in SELECT

- **Must return a single value**
- **Uses:**
  - Compute an associated value



# Subqueries in SELECT

- Must return a single value
- Uses:
  - Compute an associated value

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                  FROM Payroll AS P1
                  WHERE P.Job = P1.Job)
FROM Payroll AS P
```

# Subqueries in SELECT

- Must return a single value
- Uses:
  - Compute an associated value
- Example: For each employee, return their name and average job salary.

```
SELECT P.Name, ( SELECT AVG(P1.Salary)
                  FROM Payroll AS P1
                  WHERE P.Job = P1.Job)
FROM Payroll AS P
```

# Subqueries in SELECT

- Must return a single value
- Uses:
  - Compute an associated value
- Example: For each employee, return their name and average job salary.

```
SELECT P.Name, ( SELECT AVG(P1.Salary)
                  FROM Payroll AS P1
                  WHERE P.Job = P1.Job)
FROM Payroll AS P
```

# Subqueries in SELECT

- Must return a single value
- Uses:
  - Compute an associated value
- Example: For each employee, return their name and average job salary.

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                  FROM Payroll AS P1
                  WHERE P.Job = P1.Job)
FROM Payroll AS P
```

**Correlated subquery!**  
Semantics are that the  
entire subquery is  
recomputed for each tuple

# Subqueries in SELECT

- Must return a single value
- Uses:
  - Compute an associated value
- Example: For each employee, return their name and average job salary.

```
SELECT P.Name, ( SELECT AVG(P1.Salary)
                  FROM Payroll AS P1
                  WHERE P.Job = P1.Job)
FROM Payroll AS P
```

# Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```

**Payroll P**



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Payroll P1**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```

**Payroll P**



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Payroll P1**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



# Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```


**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Payroll P1**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

55000



# Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary	
123	Jack	TA	50000	55000
→ 345	Allison	TA	60000	
567	Magda	Prof	90000	
789	Dan	Prof	100000	

# Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary	
123	Jack	TA	50000	55000
→ 345	Allison	TA	60000	
567	Magda	Prof	90000	
789	Dan	Prof	100000	

**Payroll P1**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Payroll P1**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

55000



# Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Payroll P1**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

55000

55000



# Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary	
123	Jack	TA	50000	55000
345	Allison	TA	60000	55000
→ 567	Magda	Prof	90000	95000
789	Dan	Prof	100000	

# Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary	
123	Jack	TA	50000	55000
345	Allison	TA	60000	55000
567	Magda	Prof	90000	95000
→ 789	Dan	Prof	100000	95000

# Subqueries in SELECT

For each person find the average salary of their job

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                  FROM Payroll AS P1
                  WHERE P.Job = P1.Job)
FROM Payroll AS P
```



Same (decorrelated and unnested)

```
SELECT P1.Name, AVG(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P1.Name
```



# Subqueries in SELECT

For each person find the number of cars they drive

# Subqueries in SELECT

For each person find the number of cars they drive

```
SELECT P.Name, ( SELECT COUNT (R.Car)
                  FROM Regist AS R
                  WHERE P.UserID =
                      R.UserID)
FROM Payroll AS P
```

# Subqueries in SELECT

For each person find the number of cars they drive

```
SELECT P.Name, (SELECT COUNT(R.Car)
                  FROM Regist AS R
                  WHERE P.UserID =
                      R.UserID)
FROM Payroll AS P
```



Same? **Discuss!**

```
SELECT P.Name, COUNT(R.Car)
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
GROUP BY P.Name
```

# Subqueries in SELECT

For each person find the number of cars they drive

```
SELECT P.Name, (SELECT COUNT(R.Car)
                  FROM Regist AS R
                  WHERE P.UserID =
                      R.UserID)
FROM Payroll AS P
```

0-count case not covered!

```
SELECT P.Name, COUNT(R.Car)
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
GROUP BY P.Name
```

# Subqueries in SELECT

For each person find the number of cars they drive

```
SELECT P.Name, (SELECT COUNT(R.Car)
                  FROM Regist AS R
                  WHERE P.UserID =
                      R.UserID)
FROM Payroll AS P
```



Still possible to decorrelate and unnest

# Subqueries in SELECT

For each person find the number of cars they drive

```
SELECT P.Name, (SELECT COUNT(R.Car)
                  FROM Regist AS R
                  WHERE P.UserID =
                      R.UserID)
FROM Payroll AS P
```



Still possible to decorrelate and unnest

```
SELECT P.Name, COUNT(R.Car)
FROM Payroll AS P LEFT OUTER JOIN
    Regist AS R ON P.UserID = R.UserID
GROUP BY P.Name
```

# Announcements

- HW 3 out today
- Azure credits have been issued (\$75)
  - Sent to @uw.edu emails
  - Enter your @uw.edu email in sign-in
  - Post on Piazza if you have issues



## Accept your Azure lab assignment

You have a pending lab assignment. Please accept your assignment to get started with your course.

[Accept lab assignment >](#)

This email is generated from an unmonitored alias; please do not reply. If you have questions, please [submit a request](#).

# Subqueries in FROM

- Equivalent to a WITH subquery
- Uses:
  - Solve subproblems that can be later joined/evaluated

```
WITH MaxPay AS
    (SELECT P1.Job AS Job,
            MAX(P1.Salary) AS Salary
     FROM Payroll AS P1
     GROUP BY P1.Job)
SELECT P.Name, P.Salary
FROM Payroll AS P, MaxPay AS MP
WHERE P.Job = MP.Job AND
      P.Salary = MP.Salary
```

Syntactic sugar

```
SELECT P.Name, P.Salary
FROM Payroll AS P, (SELECT P1.Job AS Job,
                          MAX(P1.Salary) AS Salary
                     FROM Payroll AS P1
                     GROUP BY P1.Job) AS MP
WHERE P.Job = MP.Job AND
      P.Salary = MP.Salary
```



# Subqueries in WHERE/HAVING

## ■ Uses:

- ANY  $\rightarrow \exists$
- ALL  $\rightarrow \forall$
- (NOT) IN  $\rightarrow (\notin) \in$
- (NOT) EXISTS  $\rightarrow (\emptyset = \dots) \emptyset \neq \dots$

# Subqueries in WHERE/HAVING

## ■ Uses:

- ANY  $\rightarrow \exists$
- ALL  $\rightarrow \forall$
- (NOT) IN  $\rightarrow (\notin) \in$
- (NOT) EXISTS  $\rightarrow (\emptyset = \dots) \emptyset \neq \dots$

Find the name and salary of people who do not drive cars

```
SELECT P.Name, P.Salary
FROM Payroll AS P
WHERE P.UserID NOT IN (SELECT UserID
                        FROM Regist)
```

Decorrelated!

### 3. Subqueries in WHERE

- SELECT ..... WHERE EXISTS (subquery);
- SELECT ..... WHERE NOT EXISTS (subquery);
- SELECT ..... WHERE attribute IN (subquery);
- SELECT ..... WHERE attribute NOT IN (subquery);
- SELECT ..... WHERE constant > ANY (subquery );
- SELECT ..... WHERE constant > ALL (subquery);

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies that make some products with price < 200

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies that make some products with price < 200

Existential quantifiers

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies that make some products with price < 200

Existential quantifiers

Using **EXISTS**: EXISTS (subquery) returns true iff cardinality of subquery > 0

```
SELECT DISTINCT C.cname
FROM   Company C
WHERE  EXISTS (SELECT *
                FROM Product P
                WHERE C.cid = P.cid and P.price < 200)
```

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies that make some products with price < 200

Existential quantifiers

Using **IN**: attr IN (subquery) returns true iff value of attr is contained in subquery

```
SELECT DISTINCT C.cname
FROM Company C
WHERE C.cid IN (SELECT P.cid
                FROM Product P
                WHERE P.price < 200)
```

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies that make some products with price < 200

Existential quantifiers

Using **ANY**:

```
SELECT DISTINCT C.cname
FROM Company C
WHERE 200 > ANY (SELECT price
                  FROM Product P
                  WHERE P.cid = C.cid)
```



# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies that make some products with price < 200

Existential quantifiers

Using **ANY**:  $\text{const} > \text{ANY}(\text{sub})$  returns true if  $\text{const} > \text{value}$  for at least one value in sub

```
SELECT DISTINCT C.cname
FROM Company C
WHERE 200 > ANY (SELECT price
                  FROM Product P
                  WHERE P.cid = C.cid)
```

Not supported  
in sqlite

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies that make some products with price < 200

Existential quantifiers

Now let's unnest it:

```
SELECT DISTINCT C.cname
FROM Company C
WHERE 200 > ANY (SELECT price
                  FROM Product P
                  WHERE P.cid = C.cid)
```

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies that make some products with price < 200

Existential quantifiers

Now let's unnest it:

```
SELECT DISTINCT C.cname
FROM   Company C, Product P
WHERE  C.cid = P.cid and P.price < 200
```

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies that make some products with price < 200

Existential quantifiers

Now let's unnest it:

```
SELECT DISTINCT C.cname
FROM   Company C, Product P
WHERE  C.cid = P.cid and P.price < 200
```

Existential quantifiers are easy! 😊

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies s.t. all their products have price < 200

same as:

Find all companies that make only products with price < 200

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies s.t. all their products have price < 200

same as:

Find all companies that make only products with price < 200

Universal quantifiers

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies s.t. all their products have price < 200

same as:

Find all companies that make only products with price < 200

Universal quantifiers

Universal quantifiers are hard! ☹️

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies s.t. all their products have price < 200

1. Find *the other* companies that make some product  $\geq 200$

```
SELECT DISTINCT C.cname
FROM   Company C
WHERE  C.cid IN (SELECT P.cid
                  FROM   Product P
                  WHERE  P.price >= 200)
```



# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies s.t. all their products have price < 200

1. Find *the other* companies that make some product  $\geq 200$

```
SELECT DISTINCT C.cname
FROM   Company C
WHERE  C.cid IN (SELECT P.cid
                  FROM   Product P
                  WHERE  P.price >= 200)
```

2. Find all companies s.t. all their products have price < 200

```
SELECT DISTINCT C.cname
FROM   Company C
WHERE  C.cid NOT IN (SELECT P.cid
                     FROM   Product P
                     WHERE  P.price >= 200)
```

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies s.t. all their products have price < 200

Universal quantifiers

Using **EXISTS**:

```
SELECT DISTINCT C.cname
FROM   Company C
WHERE  NOT EXISTS (SELECT *
                   FROM Product P
                   WHERE P.cid = C.cid and P.price >= 200)
```

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies s.t. all their products have price < 200

Universal quantifiers

Using **ALL**:

```
SELECT DISTINCT C.cname
FROM   Company C
WHERE  200 >= ALL (SELECT price
                   FROM   Product P
                   WHERE  P.cid = C.cid)
```

# 3. Subqueries in WHERE

Product (pname, price, cid)

Company (cid, cname, city)

Find all companies s.t. all their products have price < 200

Universal quantifiers

Using **ALL**:

```
SELECT DISTINCT C.cname
FROM Company C
WHERE 200 >= ALL (SELECT price
                  FROM Product P
                  WHERE P.cid = C.cid)
```

Not supported  
in sqlite

# Subqueries in WHERE/HAVING

## ■ Uses:

- ANY  $\rightarrow \exists$
- ALL  $\rightarrow \forall$
- (NOT) IN  $\rightarrow (\notin) \in$
- (NOT) EXISTS  $\rightarrow (\emptyset = \dots) \emptyset \neq \dots$

Find the name and salary of people who do not drive cars

```
SELECT P.Name, P.Salary
FROM Payroll AS P
WHERE NOT EXISTS (SELECT *
                    FROM Regist AS R
                    WHERE P.UserID =
                        R.UserID)
```

# Subqueries in WHERE/HAVING

## ■ Uses:

- ANY  $\rightarrow \exists$
- ALL  $\rightarrow \forall$
- (NOT) IN  $\rightarrow (\notin) \in$
- (NOT) EXISTS  $\rightarrow (\emptyset = \dots) \emptyset \neq \dots$

Find the name and salary of people who do not drive cars

```
SELECT P.Name, P.Salary
FROM Payroll AS P
WHERE P.UserID NOT IN (SELECT UserID
                        FROM Regist)
```

Decorrelated!

# Encoding Universal Quantifiers

- Could we ever encode a universal quantifier with a `SELECT-FROM-WHERE` query with no subqueries or aggregates?


# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where  $I$  and  $J$  are data instances and  $q$  is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of  $I$ , the query over that superset must contain at least the query results of  $I$ .






# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where  $I$  and  $J$  are data instances and  $q$  is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of  $I$ , the query over that superset must contain at least the query results of  $I$ .



Monotone queries can be similar to monotonically increasing functions when considering cardinalities of results

# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name, P.Car  
      FROM Payroll AS P, Regist AS R  
      WHERE P.UserID = R.UserID
```

Is this query monotone?

# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name, P.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID
```

Is this query monotone? **Yes!**

# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name, P.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID
```

I can't add tuples to Payroll or Regist that would "remove" a previous result

Is this query monotone? **Yes!**

# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name
FROM Payroll AS P
WHERE P.Salary >= ALL (SELECT Salary
                        FROM Payroll)
```

Is this query monotone?

# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name
FROM Payroll AS P
WHERE P.Salary >= ALL (SELECT Salary
                        FROM Payroll)
```

Is this query monotone? **No!**

# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name
FROM Payroll AS P
WHERE P.Salary >= ALL (SELECT Salary
                        FROM Payroll)
```

I can add a tuple to Payroll that has a higher salary value than any other

Is this query monotone? **No!**

# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Job, COUNT (*)  
  FROM Payroll AS P  
GROUP BY P.Job
```

Is this query monotone?



# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Job, COUNT (*)  
  FROM Payroll AS P  
GROUP BY P.Job
```

Is this query monotone? **No!**

# Monotonicity

## Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

Aggregates generally are sensitive to any new tuples

```
SELECT P.Job, COUNT (*)  
FROM Payroll AS P  
GROUP BY P.Job
```

Is this query monotone? **No!**

# Monotonicity

- All SELECT-FROM-WHERE queries (without aggregates) are monotone
- **Queries with universal quantifiers are not generally monotone**
- You have to do something “complex” if you need to code a universal quantifier

# Takeaways

- SQL is able to mirror logic over sets more or less directly
- The internal interpretation of nested queries can be quite involved
  - But our DBMS is able to derive such interpretations automatically
- We can reason about expressive power of certain queries.

# Next Unit

- We are done with lectures on SQL queries!
- Up next: Database Design
  - Data modeling
  - Ethics and Security