

Introduction to Data Management

Database Design

Paul G. Allen School of Computer Science and Engineering
University of Washington, Seattle

Announcements

- HW 3 out
- Azure credits have been issued (\$75)
 - Go to this link (posted on message board too):
 - https://portal.azure.com/#blade/Microsoft_Azure_Education/EducationMenuBlade/overview
 - Enter your @uw.edu email in sign-in
- Make sure the setup is working by *yesterday*
- HW 1 grades will be back shortly
- HW 2 grades back by Friday

Recap – Relational Model

- SQL is parsed by the DBMS and translated into an RA plan that is more directly executable
- Both query types work on the assumption that you are **using relational data**
- The relational model specifies mechanics of how data can be organized
 - No prescription of how data should be organized

Goals for Today

- With some application in mind, we can use an entity relationship (ER) diagram to conceptualize and communicate
- And with an ER diagram, we can use SQL to realize the model

Outline

- Introduce Database Design
- ER Diagrams
- ER-to-SQL conversion along the way
- Integrity constraints along the way


Database Design

- Communication is Key
- Other people are involved in the design process
- Non-computer scientists have to interact with the data too

Database Design


Database Design

Database Design or **Logical Design** or **Relational Schema Design** is the process of organizing data into a database model. This is done by considering **what data needs to be stored** and the **interrelationship of the data**.



The Database Design Process


Conceptual Model



Relational Model


→ + Schema

→ + Constraints



Conceptual Schema


→ + Normalization



Physical Schema

→ + Partitioning


→ + Indexing



The Database Design Process

Today


Conceptual Model



Relational Model

→ + Schema

→ + Constraints



Conceptual Schema


→ + Normalization



Physical Schema


→ + Partitioning

→ + Indexing



The Database Design Process


Conceptual Model



Relational Model

→ + Schema

→ + Constraints



Conceptual Schema


→ + Normalization



Physical Schema


→ + Partitioning

→ + Indexing



The Database Design Process


Conceptual Model



Relational Model


→ + Schema

→ + Constraints



Conceptual Schema

→ + Normalization




Next Unit

Physical Schema

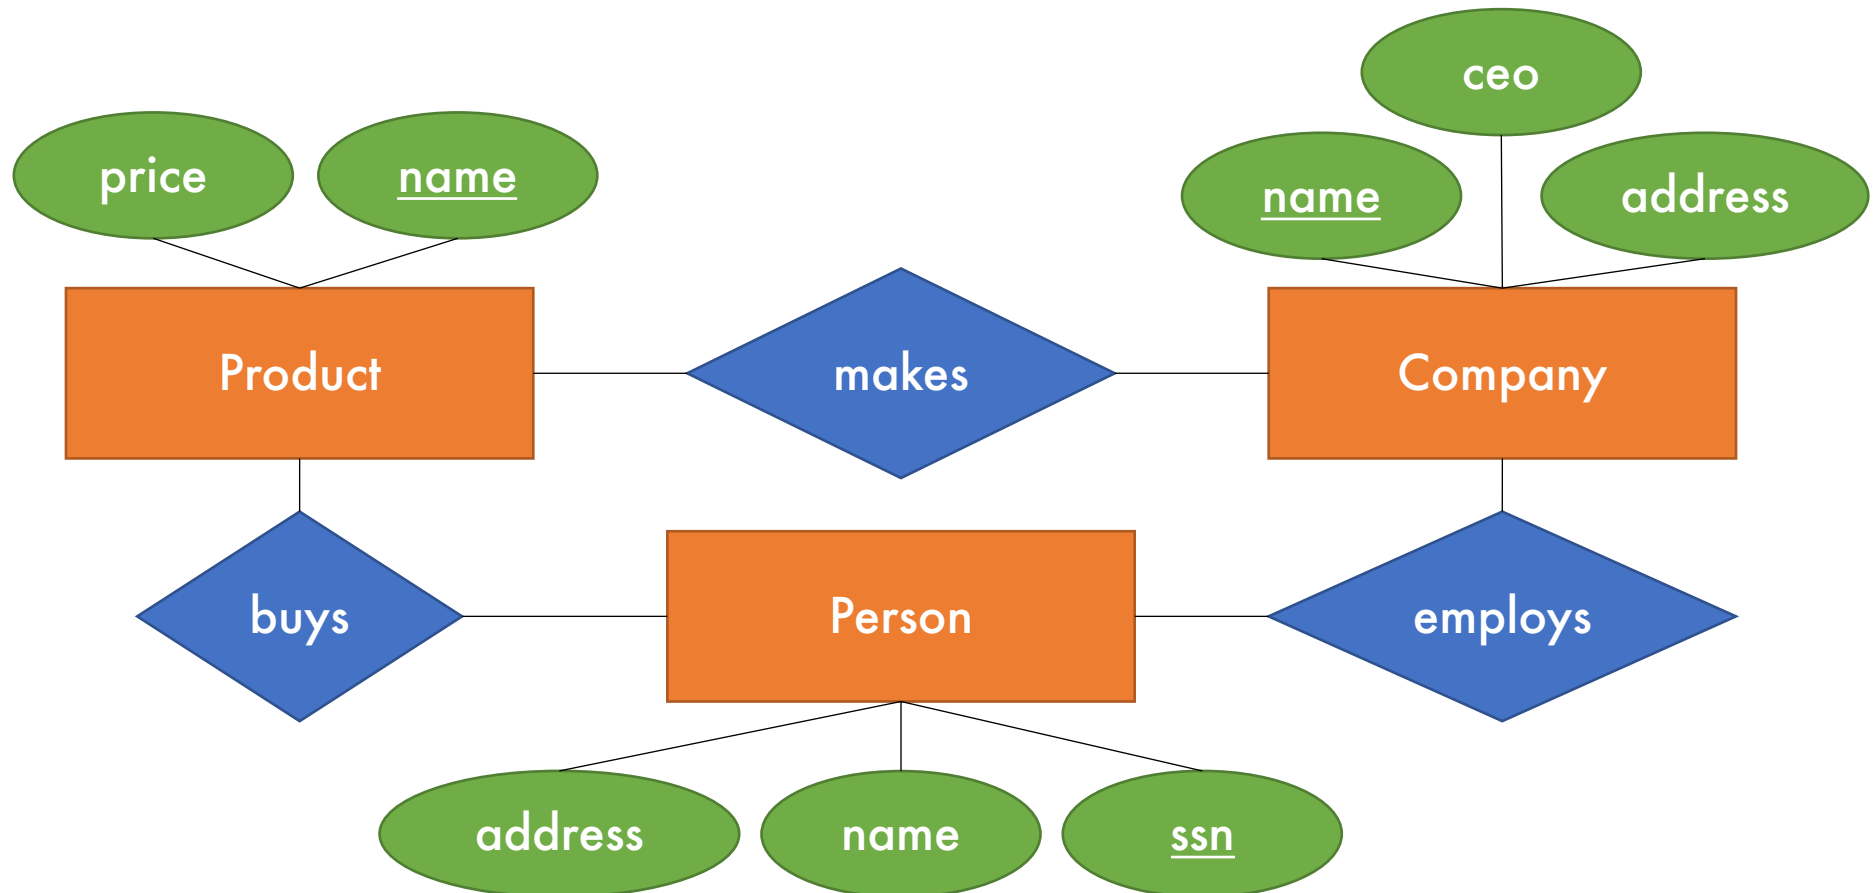
→ + Partitioning

→ + Indexing



ER Diagrams

- Humans are visual creatures so a visual model serves us best




ER Diagram Building Blocks

- These are all the blocks we will learn about


Entity set




Attribute




Relationship



Subclass




Weak Entity



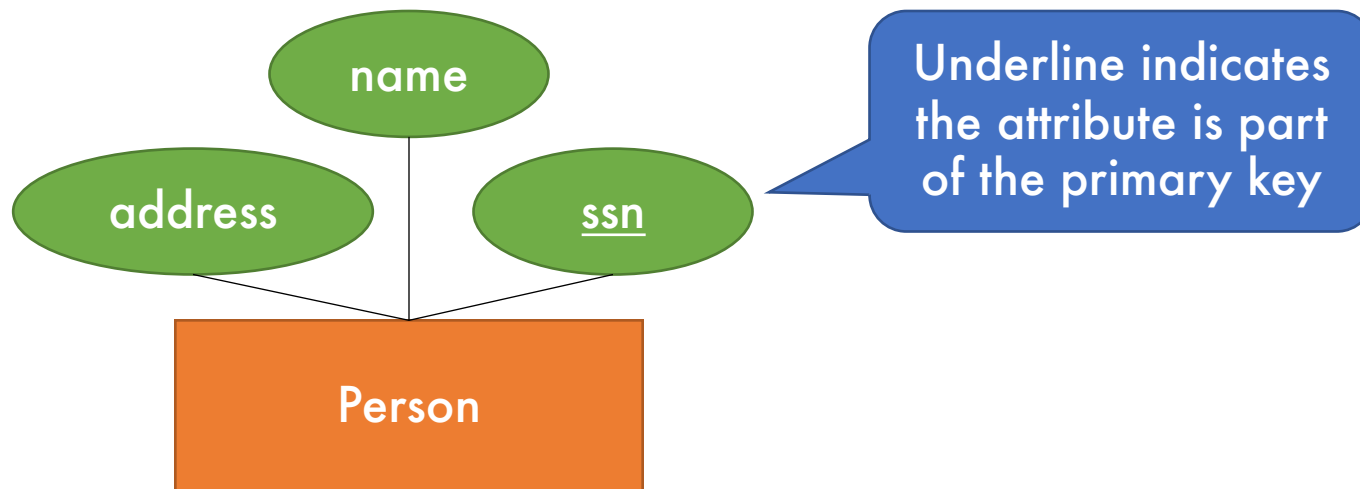
Entity Sets

- An **"entity set"** is like a **class**
- An **attribute** is like a **field**
- An **"entity"** is like a **object**
 - Corresponds to a row



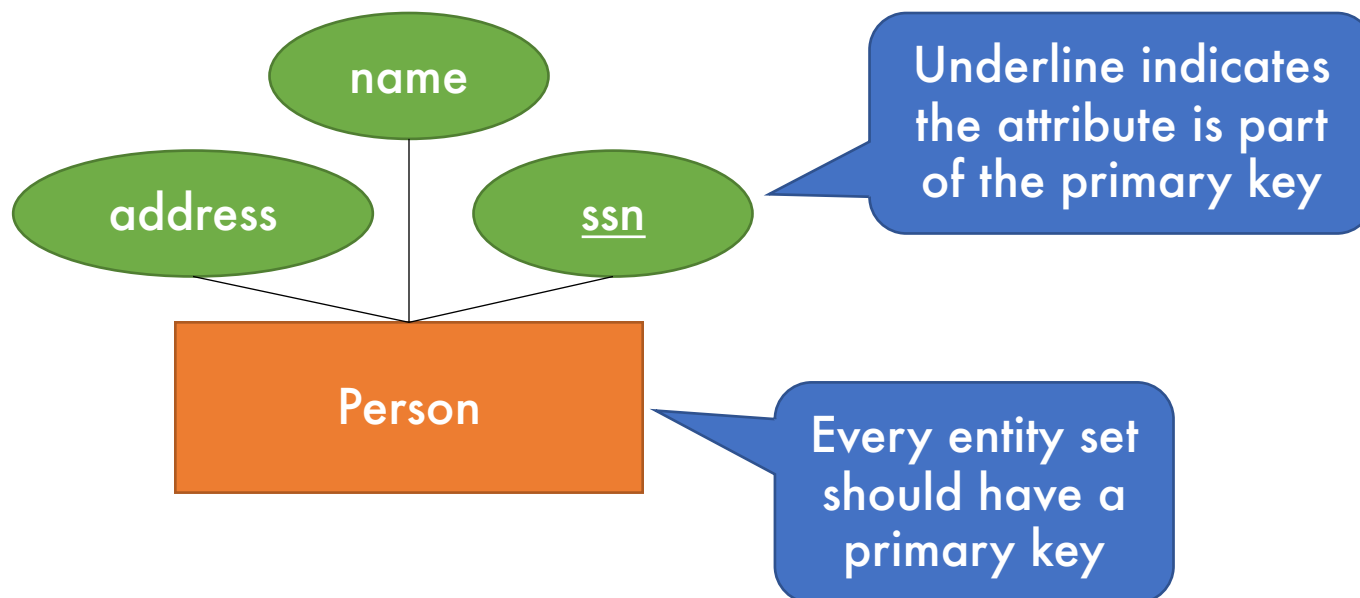
Entity Sets

- An “**entity set**” is like a **class**
- An **attribute** is like a **field**
- An “**entity**” is like a **object**
 - Corresponds to a row




Entity Sets

- An “**entity set**” is like a **class**
- An **attribute** is like a **field**
- An “**entity**” is like a **object**
 - Corresponds to a row



Entity Sets

- An “**entity set**” is like a **class**
- An **attribute** is like a **field**
- An “**entity**” is like a **object**
 - Corresponds to a row




```
CREATE TABLE Person (  
    ssn INT PRIMARY KEY,  
    name TEXT,  
    address TEXT);
```

Relations

Relationship



If A and B are sets, then a relation R is a subset of $A \times B$



Relations

Relationship



If A and B are sets, then a relation R is a subset of $A \times B$






Relations

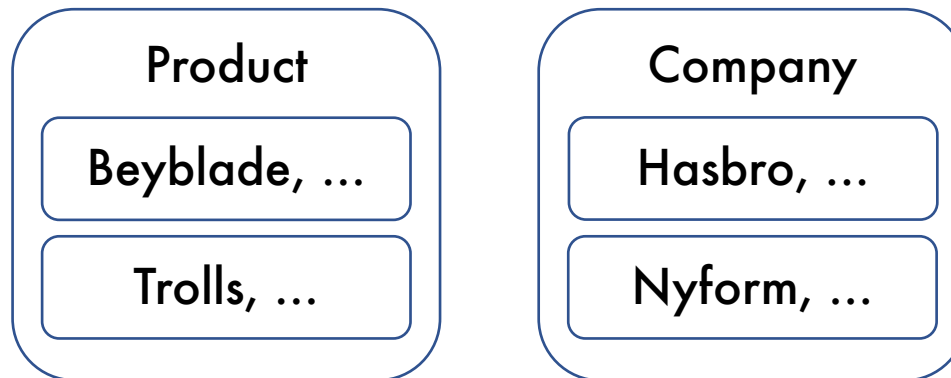
Relationship

If A and B are sets, then a relation R is a subset of $A \times B$



Relation Multiplicity

- One-to-one 
- Many-to-one 
- Many-to-many 



Relation Multiplicity


- **One-to-one**



- **Many-to-one**



- **Many-to-many**



Relation Multiplicity

- **One-to-one**



- **Many-to-one**



- **Many-to-many**




```
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    ...);  
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY,  
    ...);  
CREATE TABLE Makes (  
    cname VARCHAR(100) UNIQUE REFERENCES Company,  
    pname VARCHAR(100) UNIQUE REFERENCES Product,  
    ...);
```



Relation Multiplicity

- One-to-one 
- Many-to-one 
- **Many-to-many** 



Relation Multiplicity

- One-to-one



- Many-to-one



- **Many-to-many**





```
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    ...);  
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY,  
    ...);  
CREATE TABLE Makes (  
    cname VARCHAR(100) UNIQUE REFERENCES Company,  
    pname VARCHAR(100) UNIQUE REFERENCES Product,  
    PRIMARY KEY (cname, pname),  
    ...);
```






Relation Multiplicity

- One-to-one
- **Many-to-one**
- Many-to-many



Relation Multiplicity

- One-to-one
- **Many-to-one**
- Many-to-many




Relation Multiplicity

- One-to-one
- **Many-to-one**
- Many-to-many




Do I need a Makes table?





Relation Multiplicity

- One-to-one
- **Many-to-one**
- Many-to-many




Do I need a Makes table?
Key observation: In this many-to-one relationship, each company can make many products, but **each product can only be made by a one company**




Relation Multiplicity

- One-to-one
- **Many-to-one**
- Many-to-many



Do I need a Makes table?
Key observation: In this many-to-one relationship, each company can make many products, but **each product can only be made by a one company**



If we allow products to be made by multiple companies, we would have a many-to-many relationship





Relation Multiplicity

- One-to-one
- **Many-to-one**
- Many-to-many




```
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY,  
    ...);  
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    cname VARCHAR(100)  
        REFERENCES Company (name)  
    ...);
```




Relation Multiplicity

- One-to-one
- **Many-to-one**
- Many-to-many



```
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY,  
    ...);  
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    cname VARCHAR(100)  
    REFERENCES Company  
    ...);
```

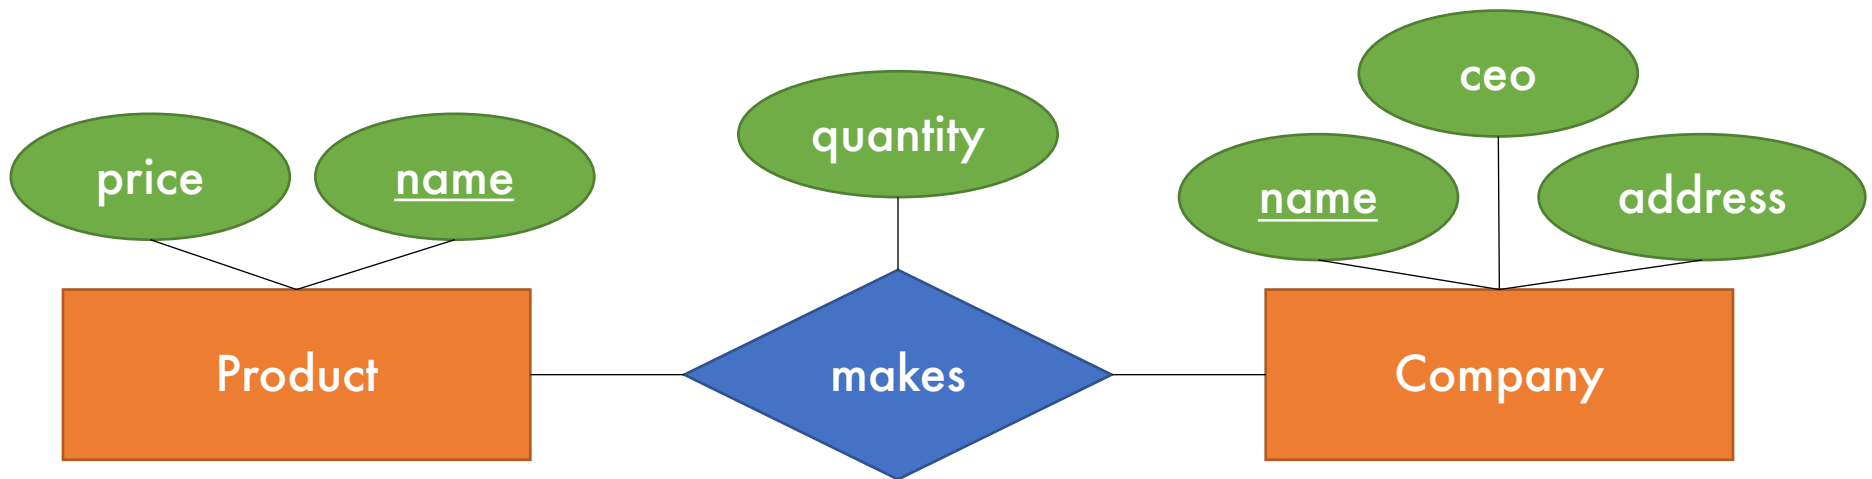


Foreign key alone is
able to encode the
Makes relationship




Relation Attributes

- Relations can have attributes too!



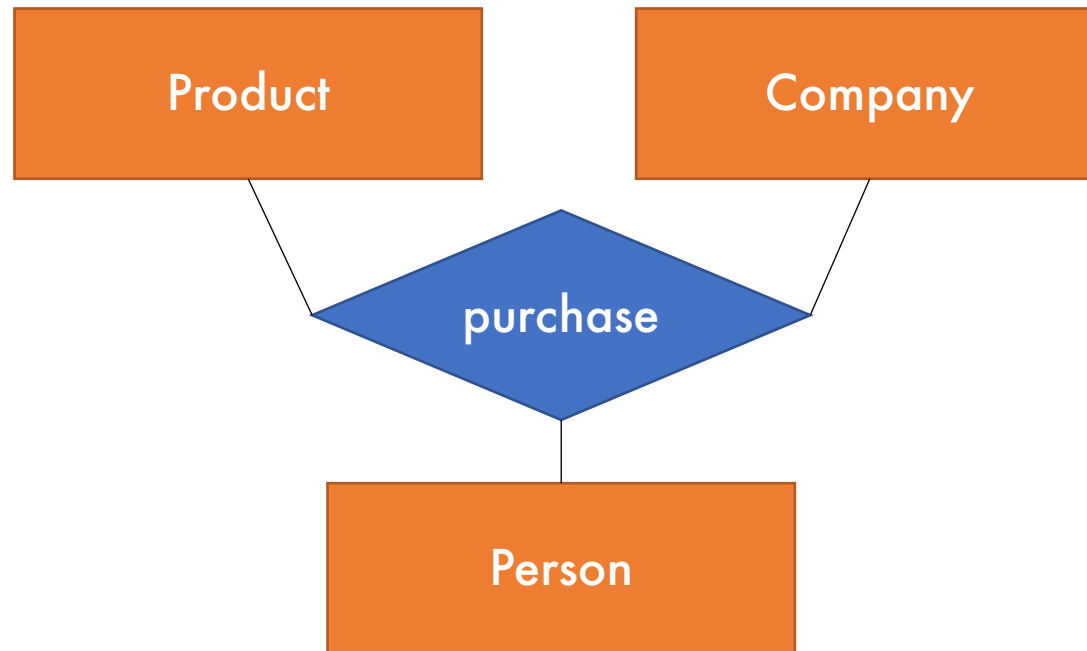
Exactly-One Reference

- Rounded arrow means the relationship is not optional (exactly one vs. at most one)

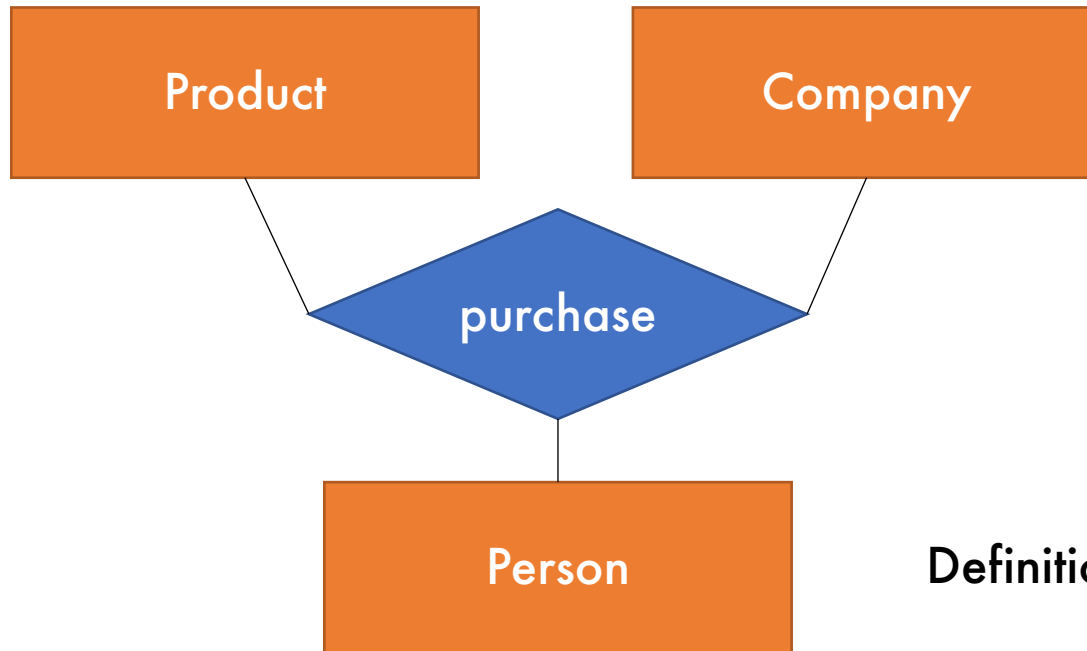


```
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY,  
    ...);  
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    cname VARCHAR(100) NOT NULL  
    REFERENCES Company  
    ...);
```

Multi-Way Relations




Multi-Way Relations



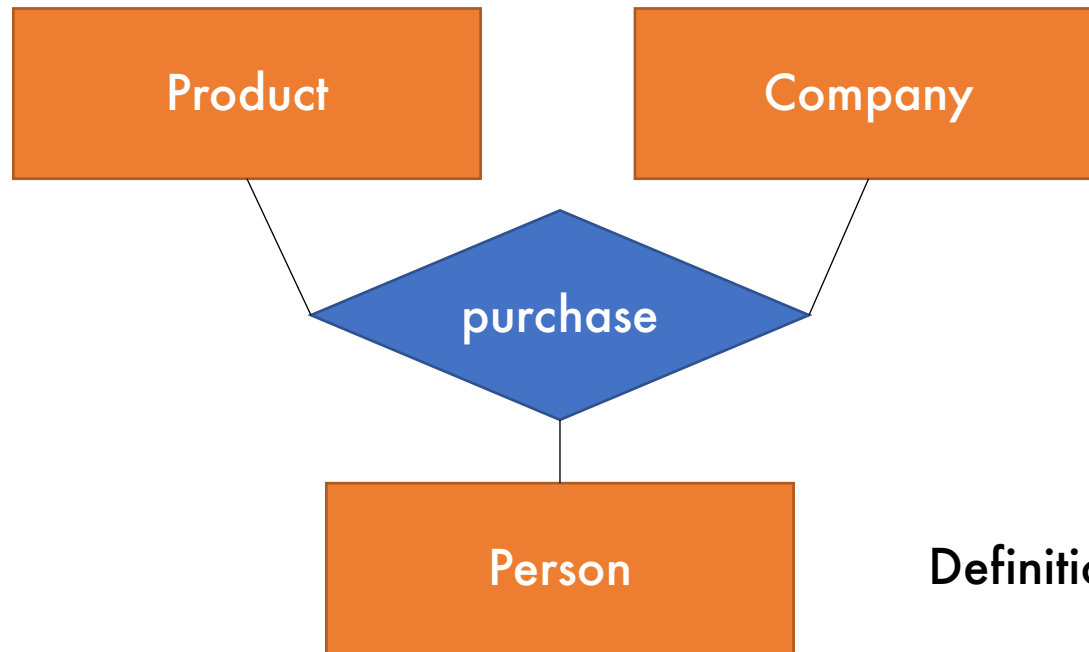
Definition of a relation generalizes!

Relationship

If A and B are sets, then a relation R is a subset of $A \times B$




Multi-Way Relations



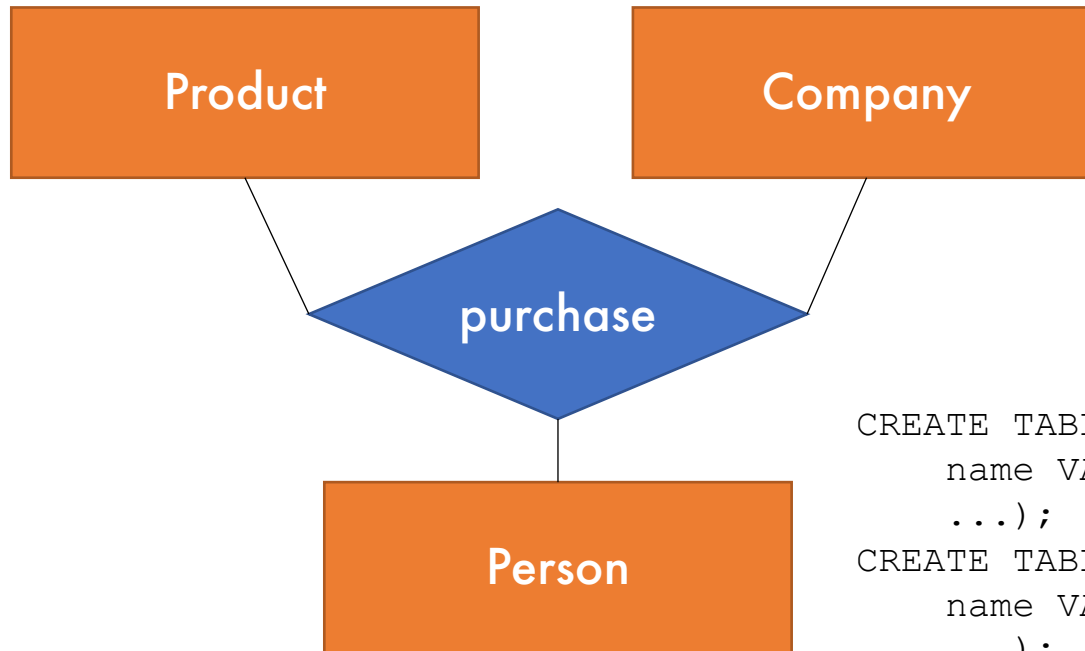
Definition of a relation generalizes!

Relationship

If A , B , and C are sets, then a relation R is a subset of $A \times B \times C$



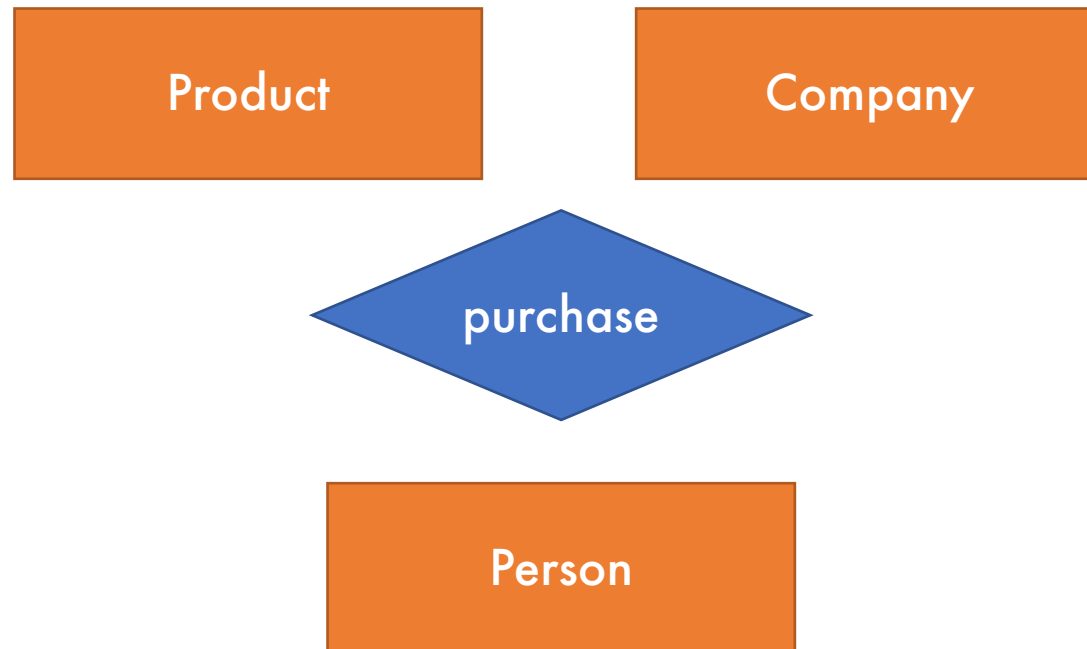
Multi-Way Relations



```
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    ...);  
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY,  
    ...);  
CREATE TABLE Person (  
    ssn INT PRIMARY KEY,  
    ...);  
CREATE TABLE Purchase (  
    cname VARCHAR(100) REFERENCES Company,  
    pname VARCHAR(100) REFERENCES Product,  
    ssn INT REFERENCES Person,  
    PRIMARY KEY (cname, pname, ssn),  
    ...);
```

It's Your Turn!

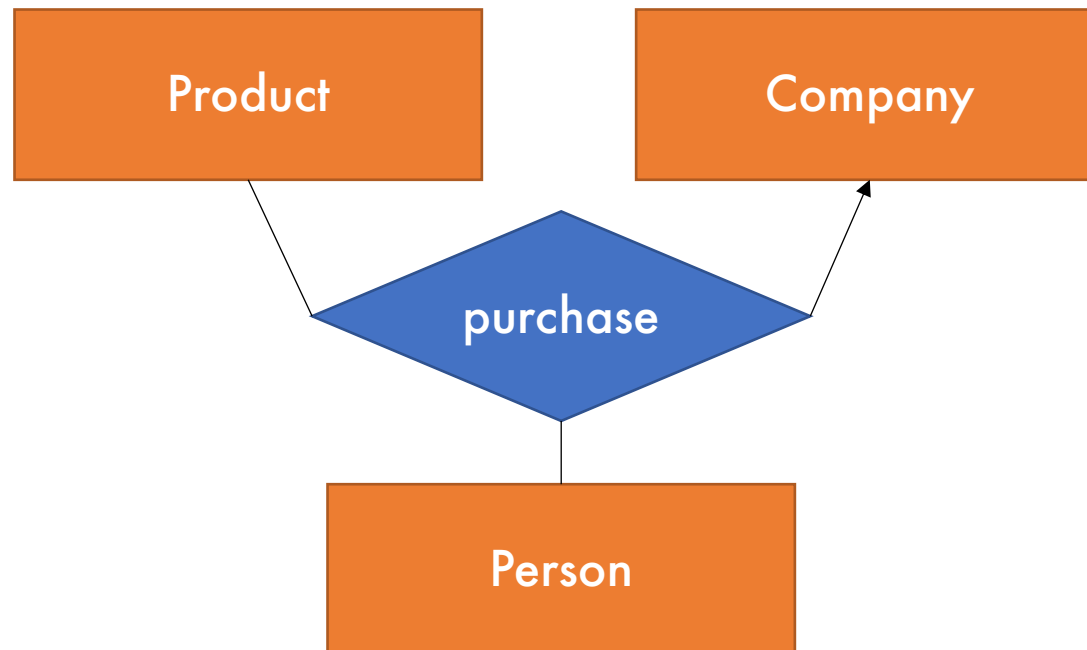
I want purchases to be such that a person will only buy each product from a single company.



How would you draw it?
Remember that the arrows read like an implication/function
Discuss!

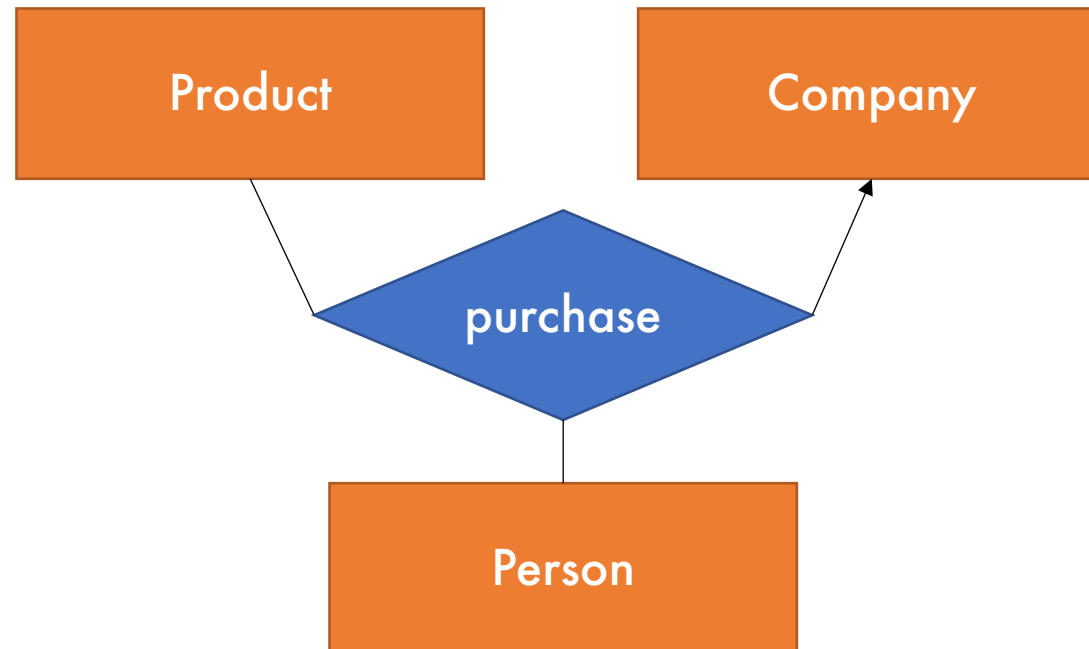
Multi-Way Relations

I want purchases to be such that a person will only buy each product from a single company.




Multi-Way Relations

Do I need a Purchase table?



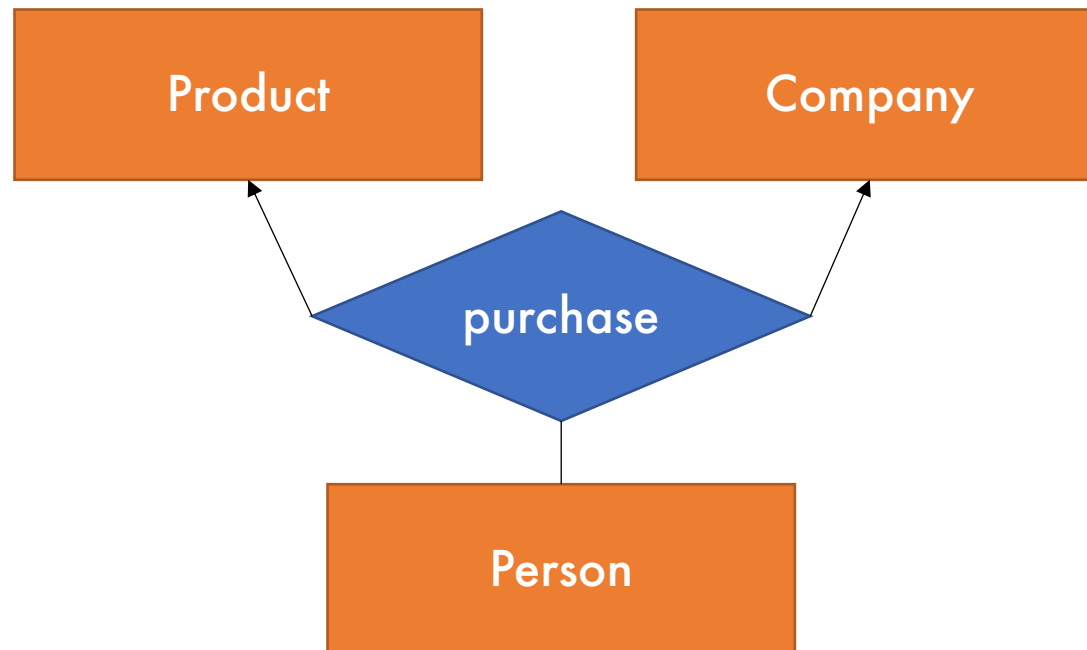
Multi-Way Relations

Do I need a Purchase table?
Probably a good idea




Multi-Way Relations

Now do I need a Purchase table?



Multi-Way Relations

Now do I need a Purchase table?
Nah



Rules of Thumb in Database Design

Design Principles (common sense):

- Pick the right entities
- Don't over complicate things
- **Follow the application spec**

Weak Entity Set

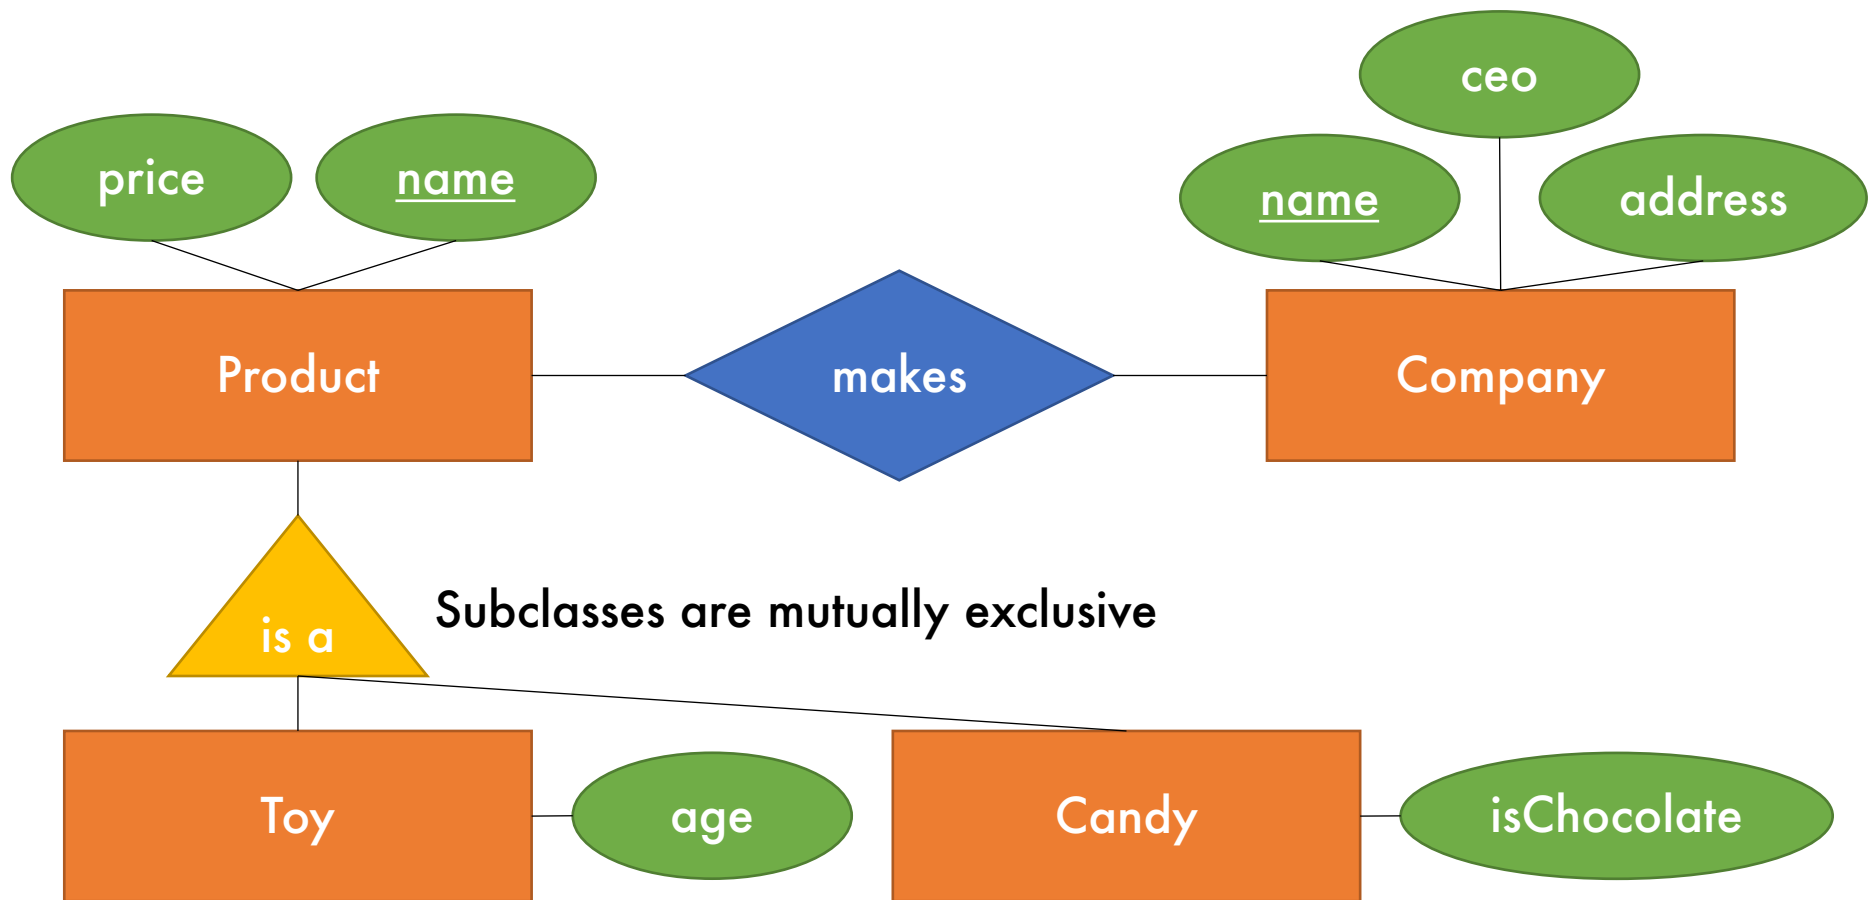
- A weak entity set has a key that is from another entity set



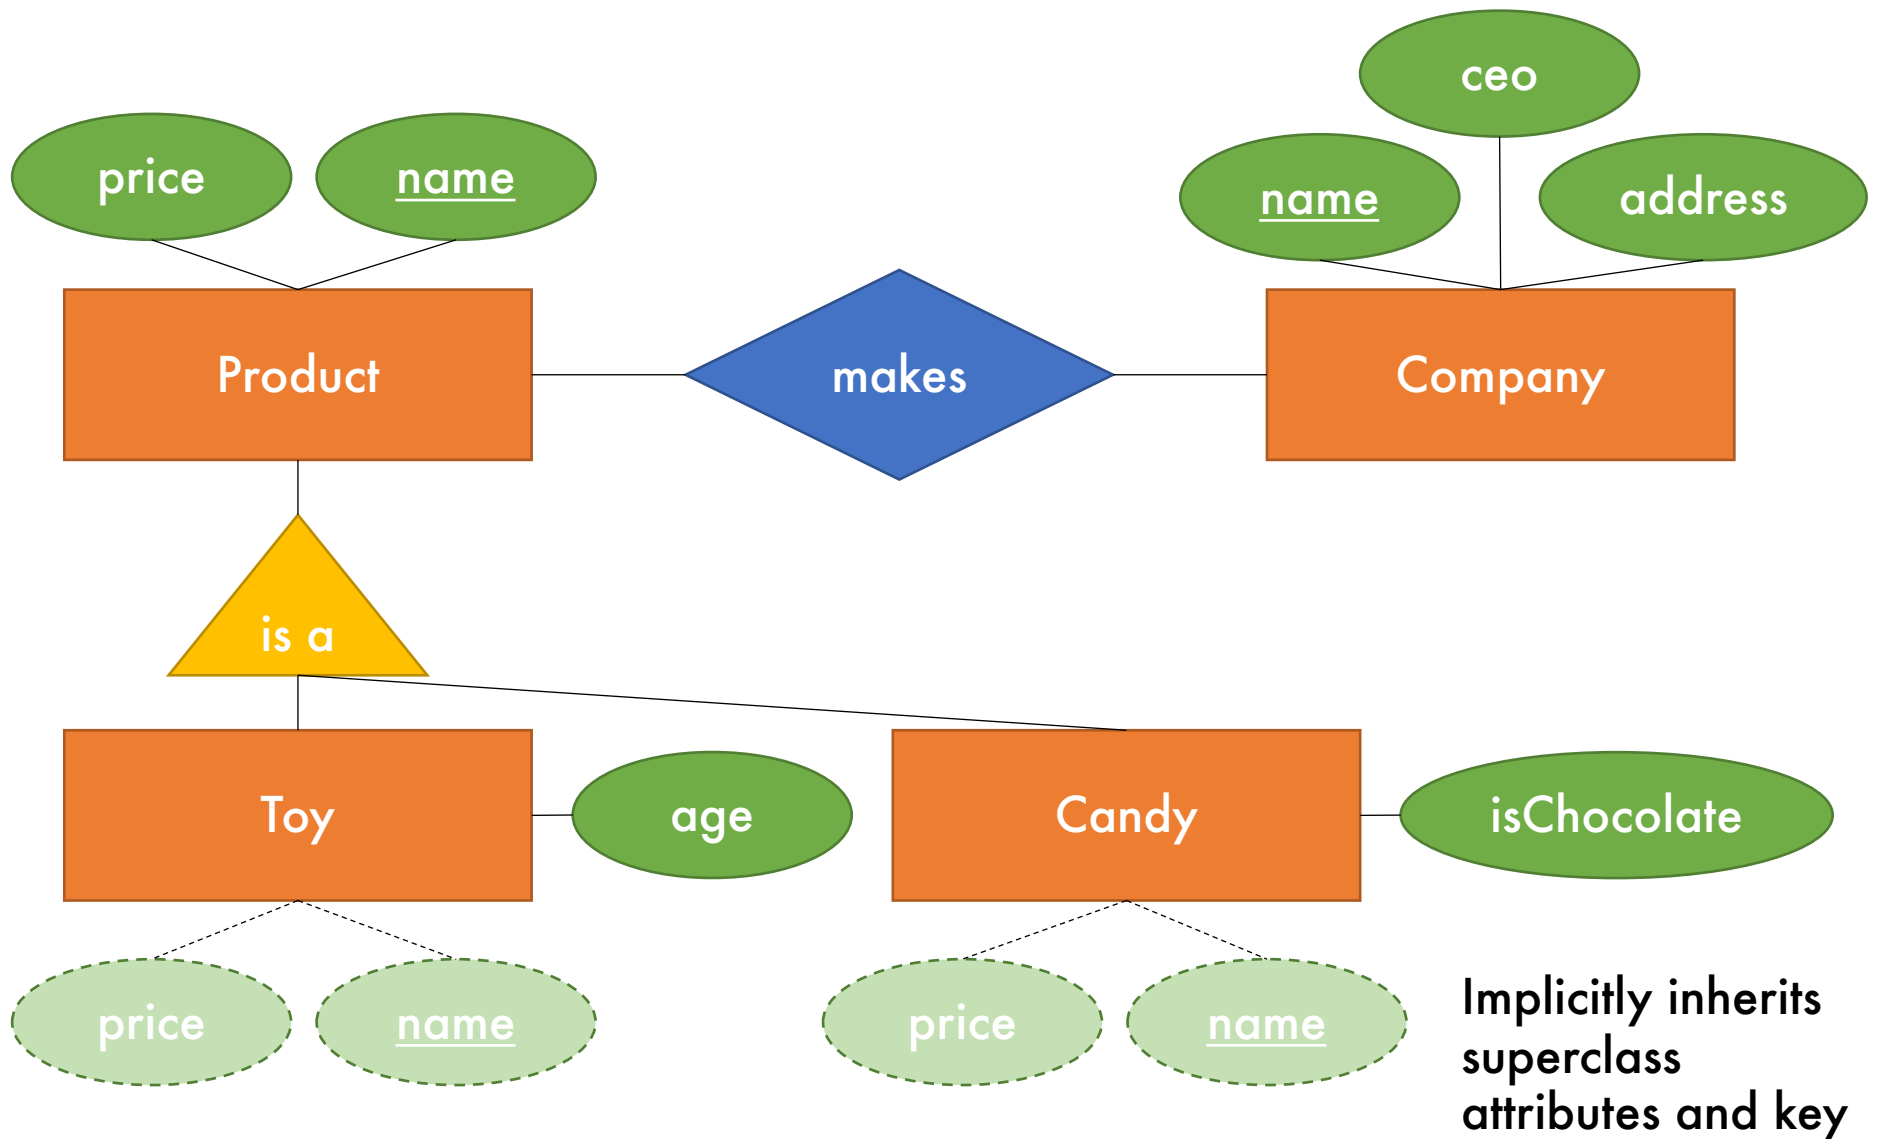
University(size, name)
Team(sport, name, uname)

Subclassing

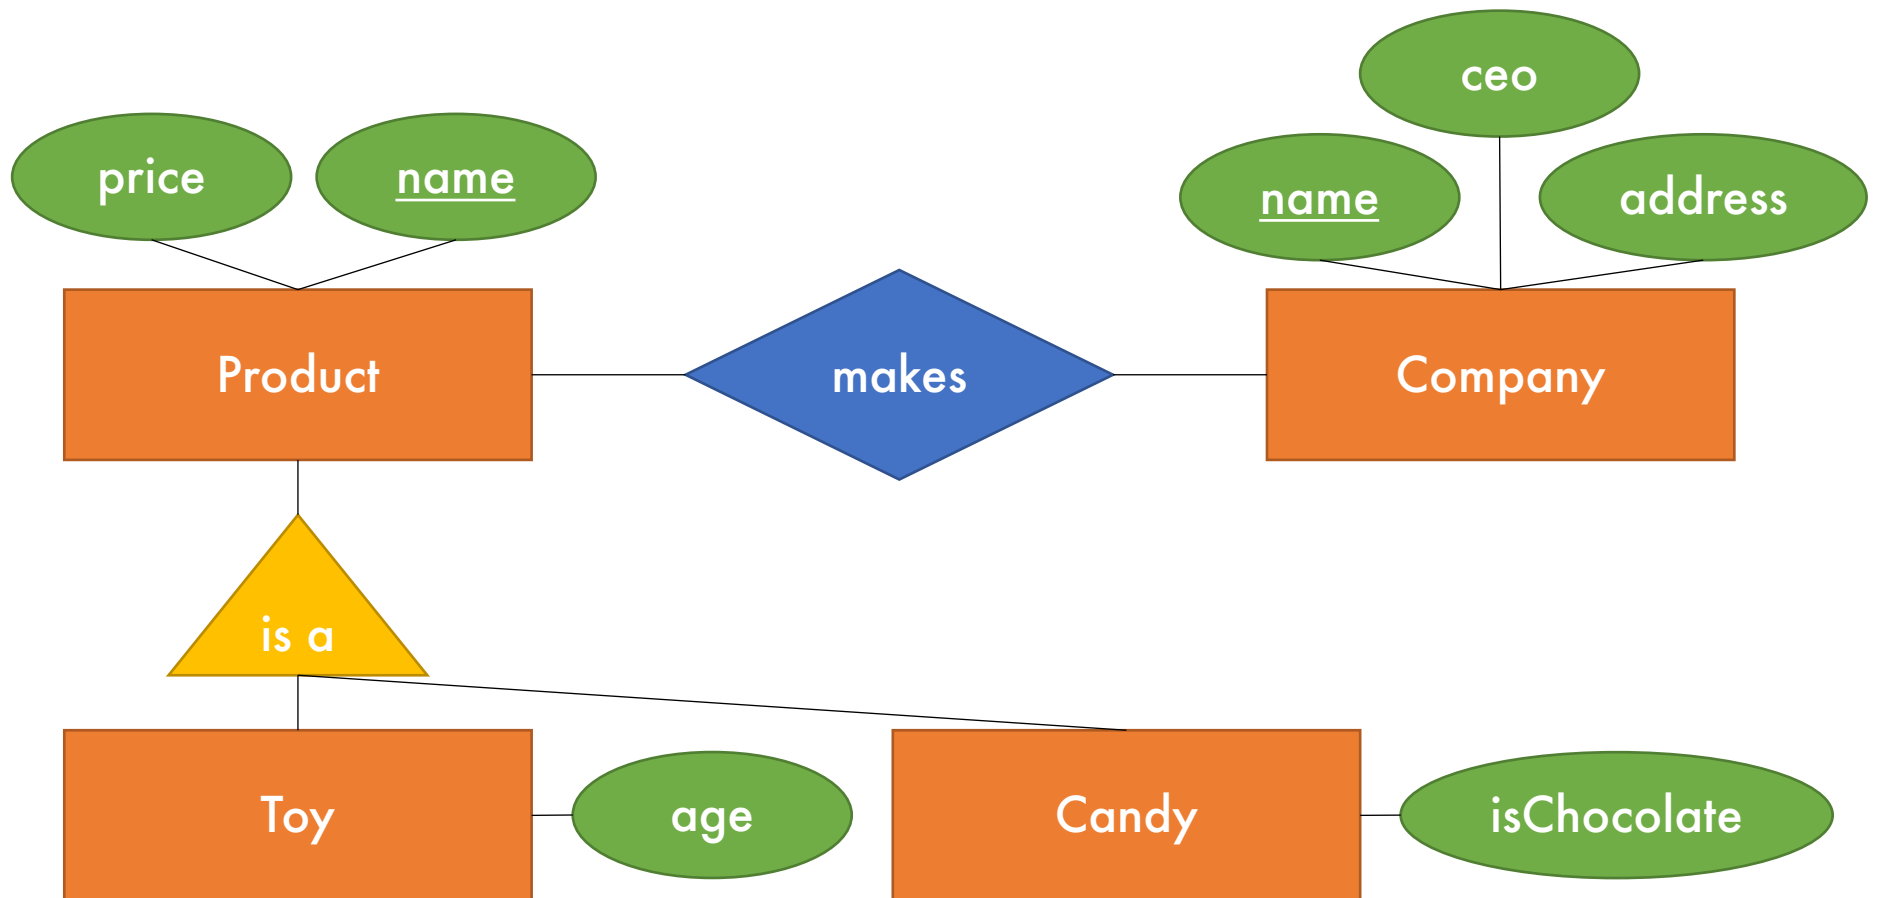
- Distinguish special entities in an entity set
- Mimics heuristics in object oriented programming



Subclassing



Subclassing




Company(...)
Makes(...)
Product(price, name)

Toy(name, age)
Candy(name, isChocolate)

Misc Constraints

- Normal arrows are shorthand versions of (≤ 1)
- Rounded arrows are shorthand versions of ($= 1$)



Each product can be made by, at most, 3 companies

Other Constraints

- CHECK (condition)
 - Single attribute
 - Single tuples

```
CREATE TABLE User (  
    uid INT PRIMARY KEY,  
    firstName TEXT,  
    lastName TEXT,  
    age INT CHECK (age > 12 AND age < 120),  
    email TEXT,  
    phone TEXT,  
    CHECK (email IS NOT NULL OR phone IS NOT NULL)  
);
```

Referential Constraint Maintenance

ON UPDATE/ON DELETE

- NO ACTION → (default) error out
- CASCADE → update/delete referencers
- SET NULL → set referencers' field to NULL
- SET DEFAULT → set referencers' field to default
 - **Assumes default was set, e.g.**

```
CREATE TABLE Table (  
    id INT DEFAULT 42 REFERENCES OtherTable,  
    ...  
);
```

Referential Constraint Maintenance

```
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY);  
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    cname VARCHAR(100)  
    REFERENCES Company  
    ON UPDATE CASCADE  
    ON DELETE SET NULL);
```

Company		Product	
name		name	cname
Hasbro		Beyblade	Hasbro
Nyform		Troll	Hasbro



Referential Constraint Maintenance

```
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY);  
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    cname VARCHAR(100)  
    REFERENCES Company  
    ON UPDATE CASCADE  
    ON DELETE SET NULL);
```

Company	Product	
name	name	cname
Hasbro	Beyblade	Hasbro
Nyform	Troll	Hasbro

```
UPDATE Company  
    SET name = 'lmao'  
    WHERE name = 'Hasbro';
```



Referential Constraint Maintenance

```
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY);  
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    cname VARCHAR(100)  
    REFERENCES Company  
    ON UPDATE CASCADE  
    ON DELETE SET NULL);
```


Company

name
lmao
Nyform

Product

name	cname
Beyblade	lmao
Troll	lmao

```
UPDATE Company  
    SET name = 'lmao'  
    WHERE name = 'Hasbro';
```



Referential Constraint Maintenance

```
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY);  
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    cname VARCHAR(100)  
    REFERENCES Company  
    ON UPDATE CASCADE  
    ON DELETE SET NULL);
```

Company	Product	
name	name	cname
lmao	Beyblade	lmao
Nyform	Troll	lmao

```
DELETE FROM Company  
WHERE name = 'lmao';
```



Referential Constraint Maintenance

```
CREATE TABLE Company (  
    name VARCHAR(100) PRIMARY KEY);  
CREATE TABLE Product (  
    name VARCHAR(100) PRIMARY KEY,  
    cname VARCHAR(100)  
    REFERENCES Company  
    ON UPDATE CASCADE  
    ON DELETE SET NULL);
```

Company		Product	
name		name	cname
Nyform		Beyblade	NULL
		Troll	NULL

```
DELETE FROM Company  
WHERE name = 'lmao';
```



Assertions

- Hard to support
- Usually impractical
- Usually not supported
 - Simulated with triggers

```
CREATE ASSERTION myAssert CHECK
  (NOT EXISTS (
    SELECT Product.name
      FROM Product, Purchase
     WHERE Product.name = Purchase.prodName
    GROUP BY Product.name
   HAVING count(*) > 200));
```

Triggers

▪ Triggers activate on a specified event

```
CREATE TRIGGER LowCredit ON Purchasing.PurchaseOrderHeader
AFTER INSERT AS
    IF (ROWCOUNT_BIG() = 0) RETURN;
    IF EXISTS (SELECT *
               FROM Purchasing.PurchaseOrderHeader AS p
               JOIN inserted AS i
               ON p.PurchaseOrderID = i.PurchaseOrderID
               JOIN Purchasing.Vendor AS v
               ON v.BusinessEntityID = p.VendorID
               WHERE v.CreditRating = 5
              )
    BEGIN
        RAISERROR ('A vendor''s credit rating is too
                    low to accept new purchase orders.', 16, 1);
        ROLLBACK TRANSACTION;
        RETURN
    END;
GO
```

= you don't need to study this for the class

Takeaways

- ER diagrams can sketch out **high-level designs**
- Certain rules of thumb for ER-to-SQL conversions help **preserve design semantics**
- SQL allows you to make **rules specific to your application**