

Introduction to Data Management

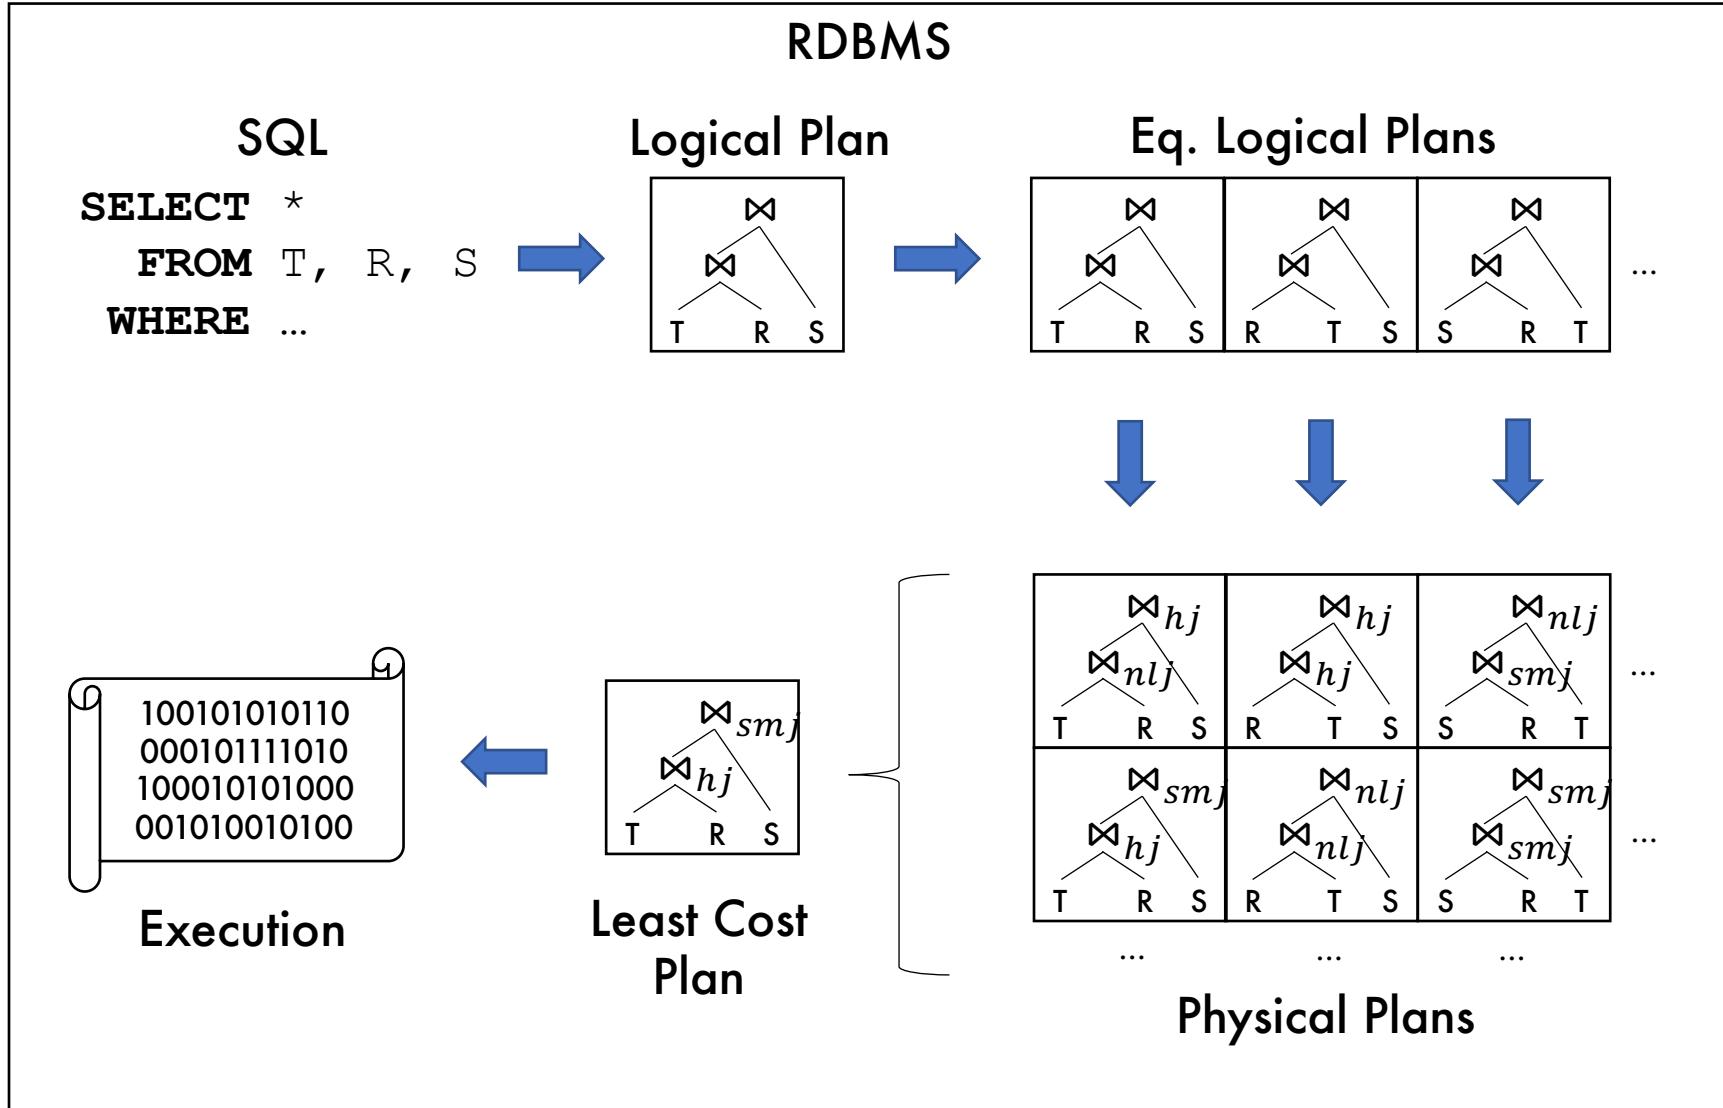
Database Tuning

Paul G. Allen School of Computer Science and Engineering
University of Washington, Seattle

Goals for Today

- We gave a baseline for what join algorithms (and respective costs) were possible
- Use DB structures to expand optimization options
- Key words to look out for:
 - “Clustered” index
 - “Unclustered” index

Recap – Plan Enumeration



Recap – Assumptions

For this class we make a lot of assumptions

- **Disk-based storage**

- HDD not SSD

- **Row-based storage**

- Tuples are stored contiguously

- **IO cost only**

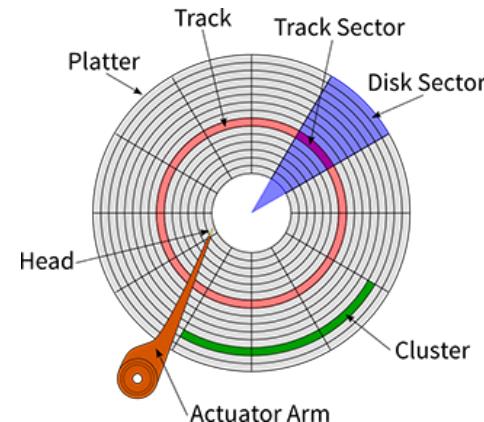
- One disk access is $\sim 100000x$ more expensive than one main memory access

- **Cold cache**

- No data preloaded into main memory

Recap – Disk Storage

- Can only read **1 block per read operation**
 - Usually 512B to 4kB
- **Sequential disk reads are faster than random ones**
 - Cost \sim 1-2% random scan = full sequential scan



Recap – Making Cost Estimations

- RDBMS keeps statistics about our tables
 - $B(R)$ = # of blocks in relation R
 - $T(R)$ = # of tuples in relation R
 - $V(attr, R)$ = # of distinct values of attr in R

Recap – Disk Storage

- **Tables are stored as files**
 - **Heap file** → Unsorted tuples
 - Nested-Loop Joins
 - Block-at-a-time Nested Loop Join ($\text{cost} = B(R) + B(R)*B(S)$)
 - Block-Nested-Loop Join ($\text{cost} = B(R) + B(R)/N * B(S)$)
 - Hash Join ($B(R) < M$, $\text{cost} = B(R) + B(S)$)
 - Sort-Merge Join ($B(R) + B(S) < M$, $\text{cost} = B(R) + B(S)$)
 - **Sequential file** → Sorted tuples

Outline

- Index structures
- Index join cost estimation
- Database tuning
- Multiple joins cost estimation
- Views

Indexing

- Indexes (for this class) can be assumed to be **already loaded into memory**
- An index does not have to contain all tuple data
 - Only key values are stored in the index
 - If an index contains all tuple data it is called a “covering index”
- **Indexes are access points for tables**

Index Structures

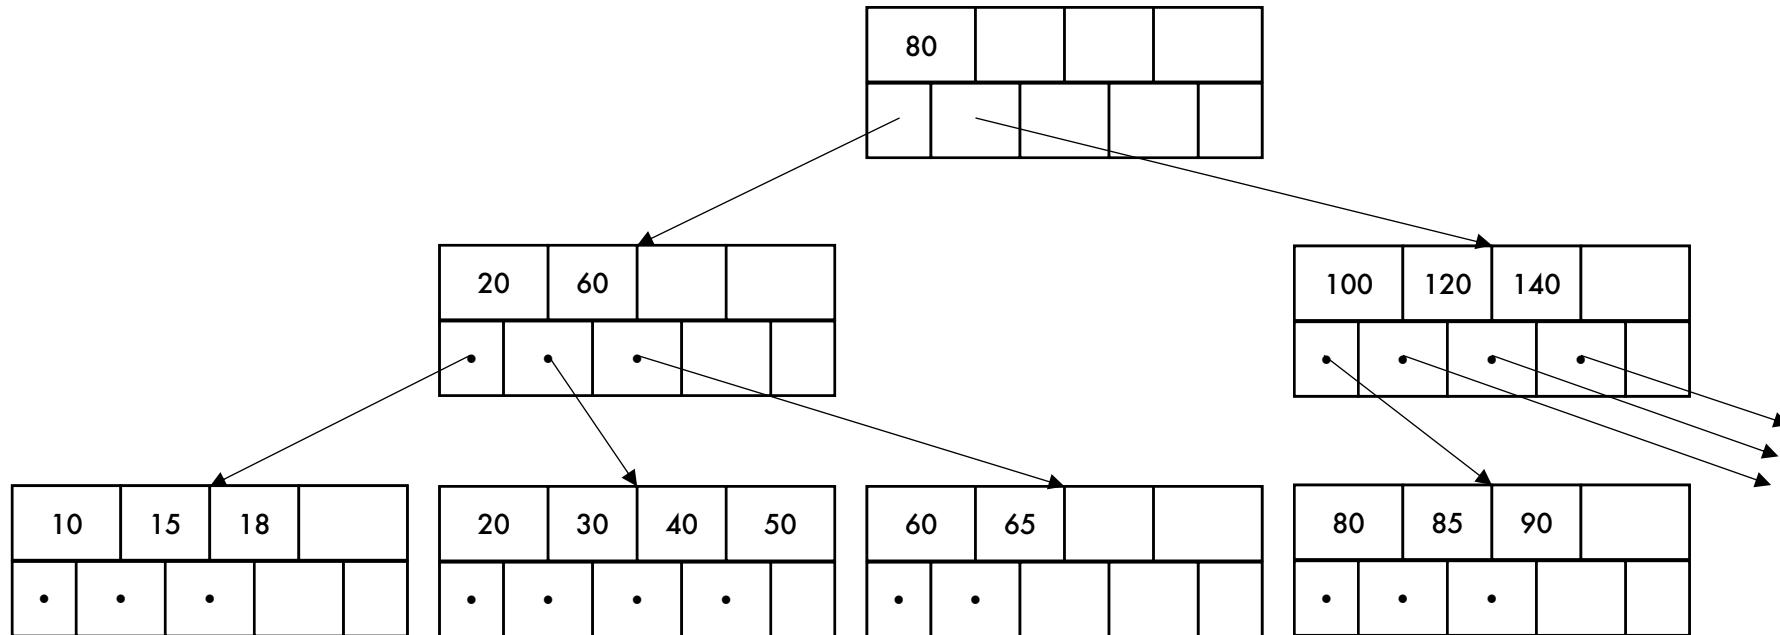
- **B+ Tree Index**
 - Clustered
 - Unclustered
- Hash Index
- R Tree
- Radix Tree
- Bloom Filter
- Hilbert Curves
- ...

What is a B Tree?

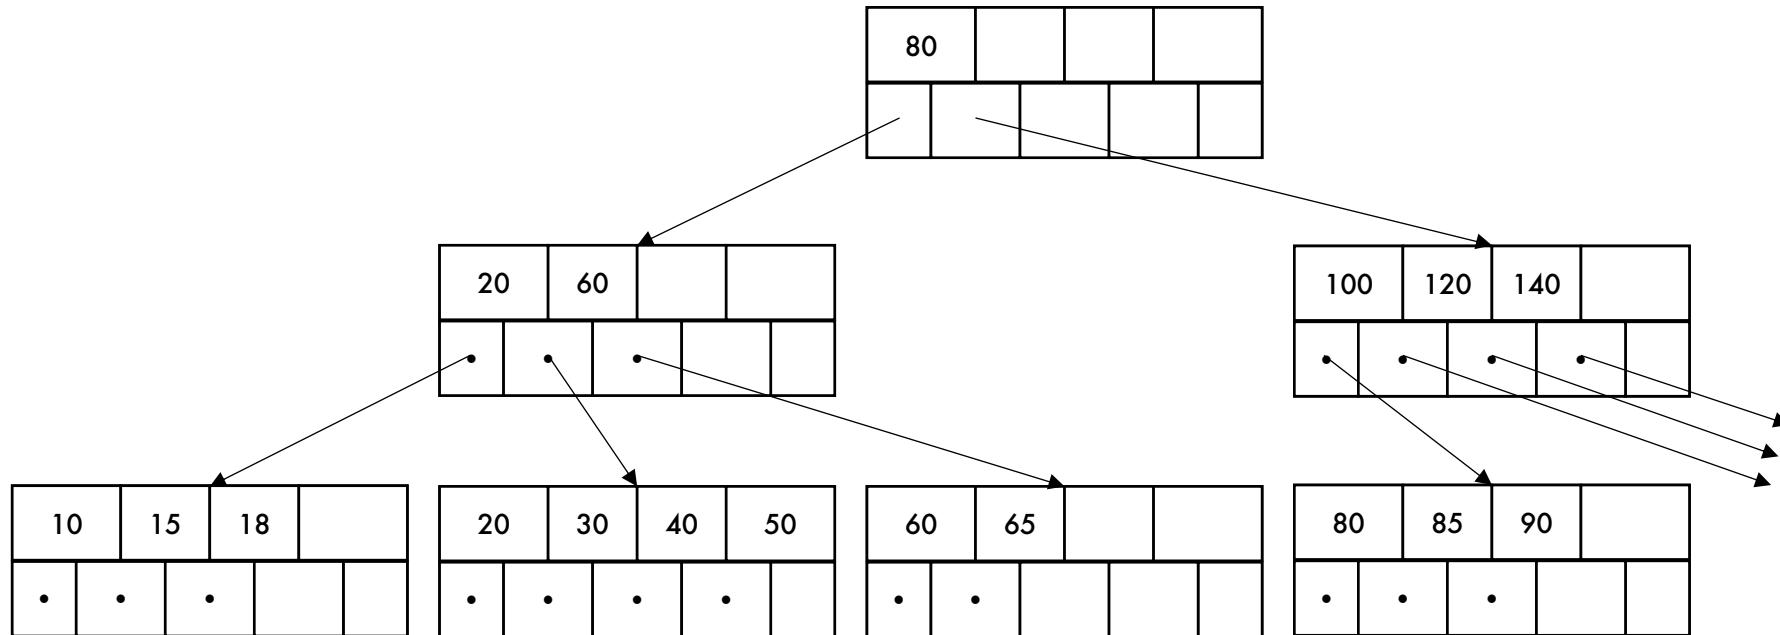
“What, if anything, the *B* stands for has never been established.” – [Wikipedia](#)

- **Search tree** (like a binary search tree)
 - Nodes annotate max values
 - Large number of children per node
- Tree/node structure that is memory efficient

What is a B Tree?

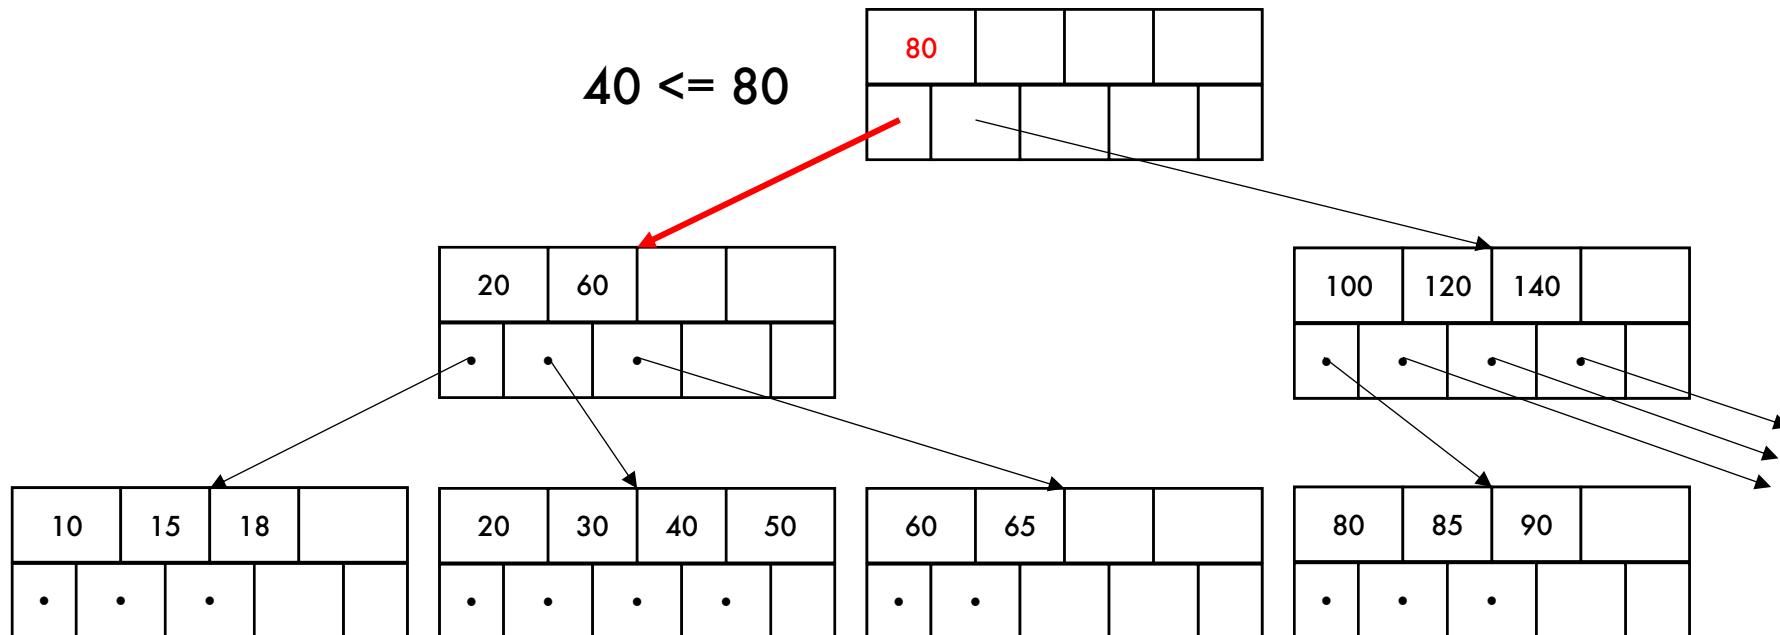


What is a B Tree?



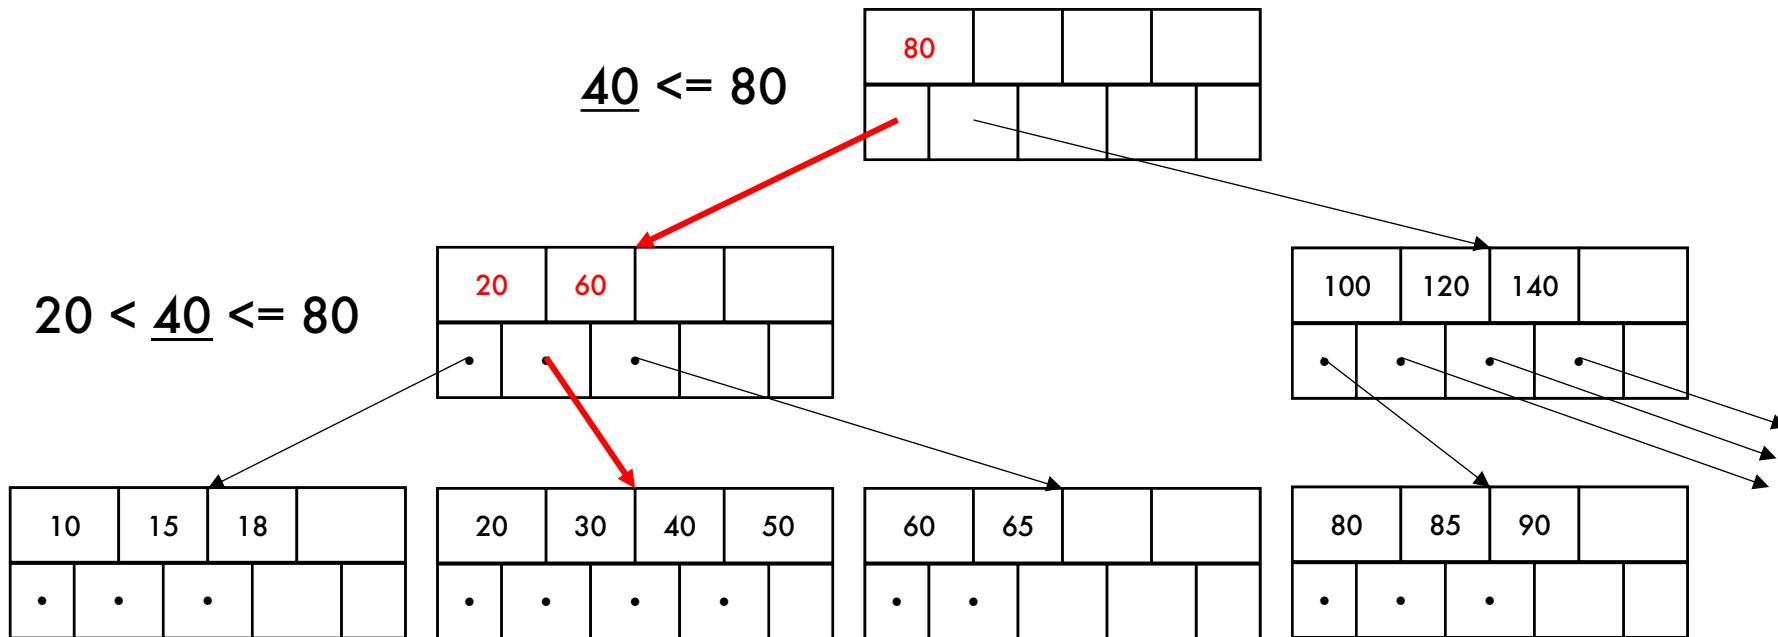
Find the value 40

What is a B Tree?



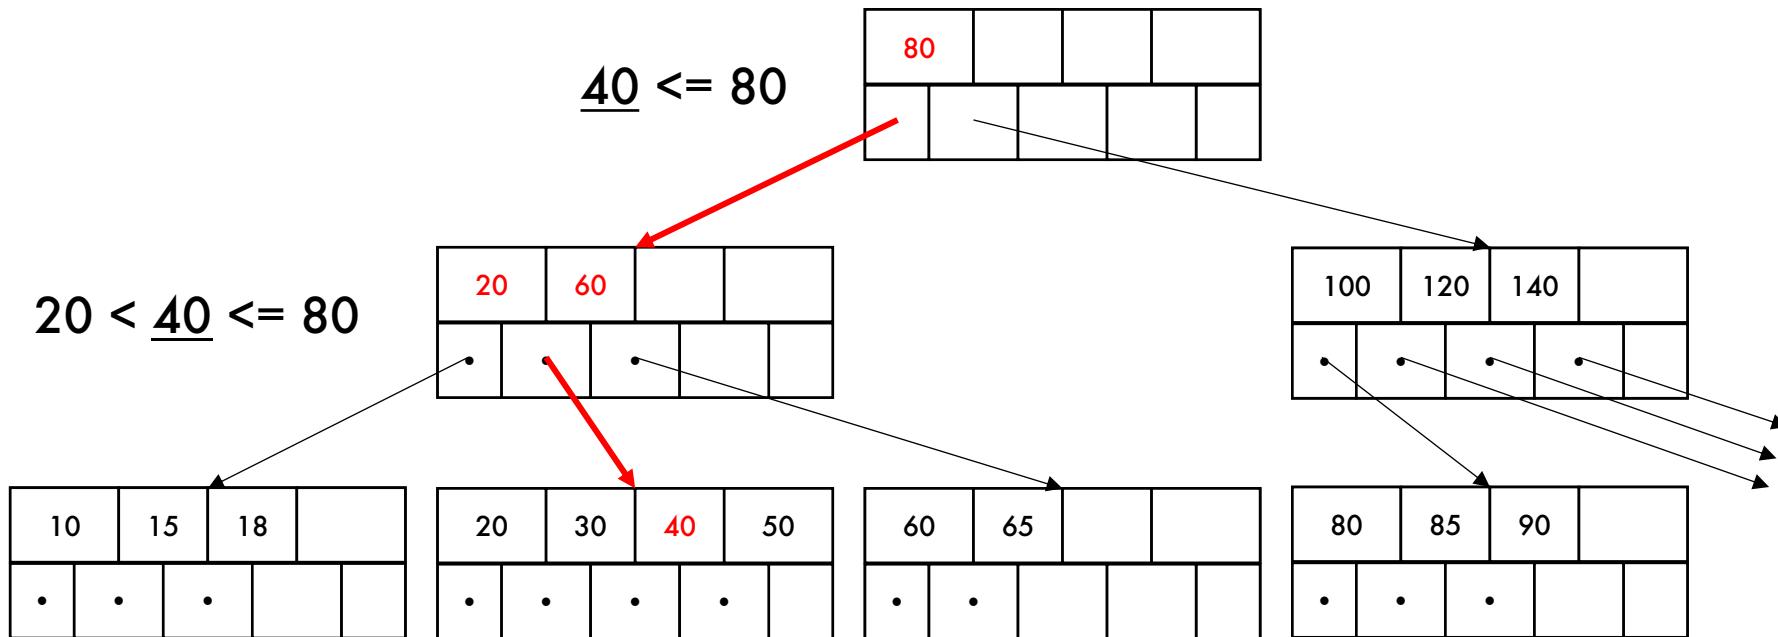
Find the value 40

What is a B Tree?



Find the value 40

What is a B Tree?

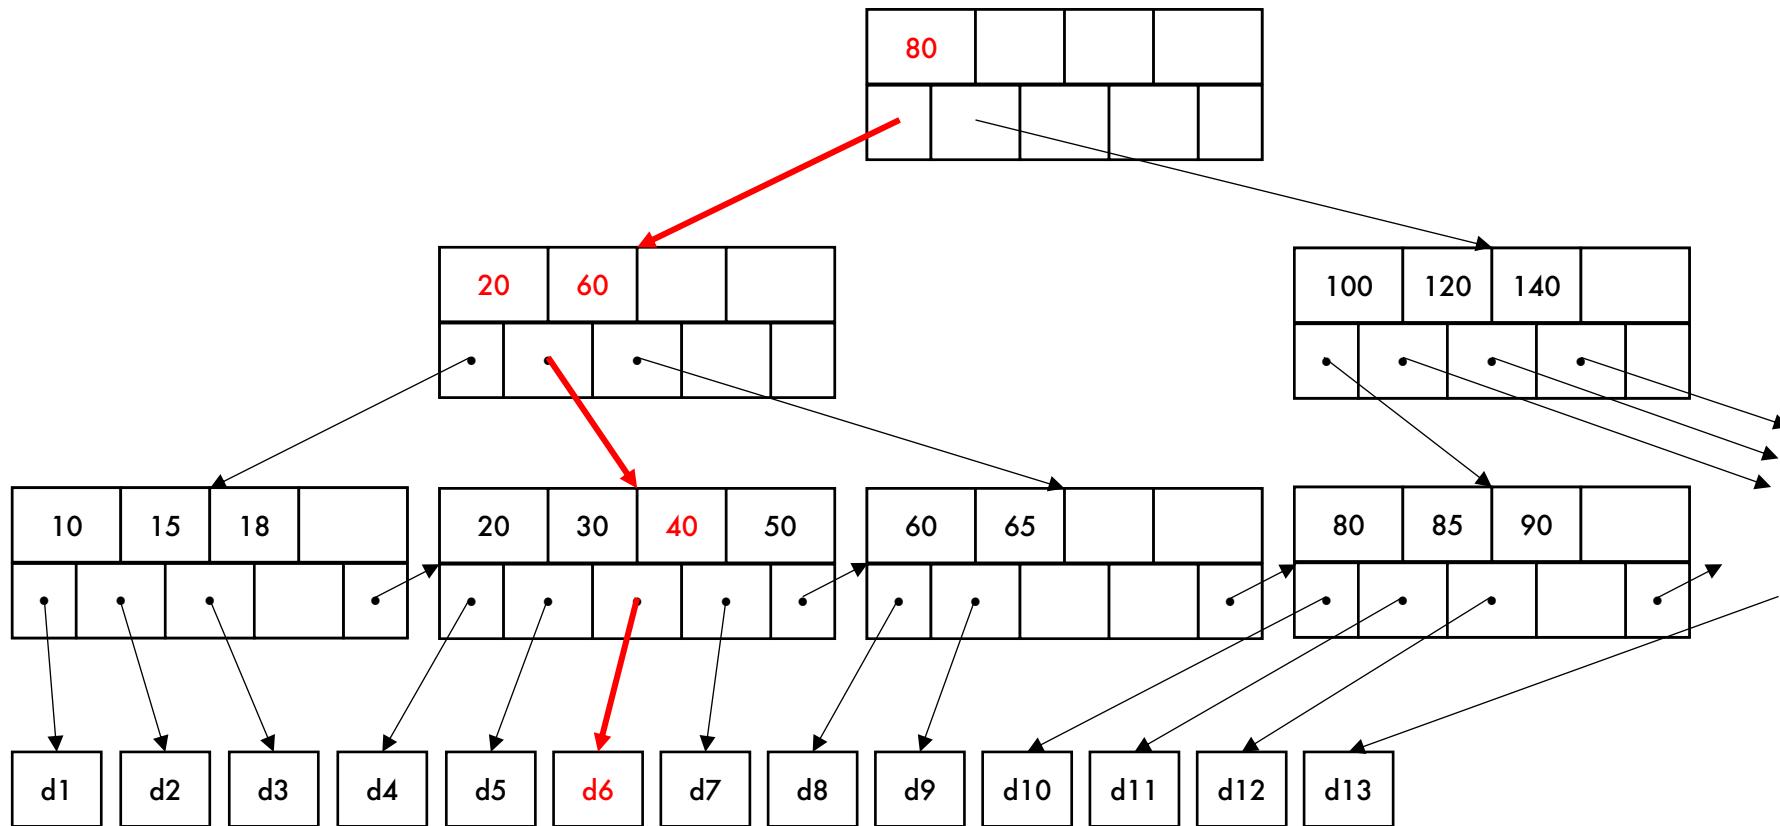


Find the value 40

How is a B+ Tree Different?

- Leaf nodes point to data
 - Data is searchable by **key** value annotated by the node labels
- Leaf nodes form a linked list

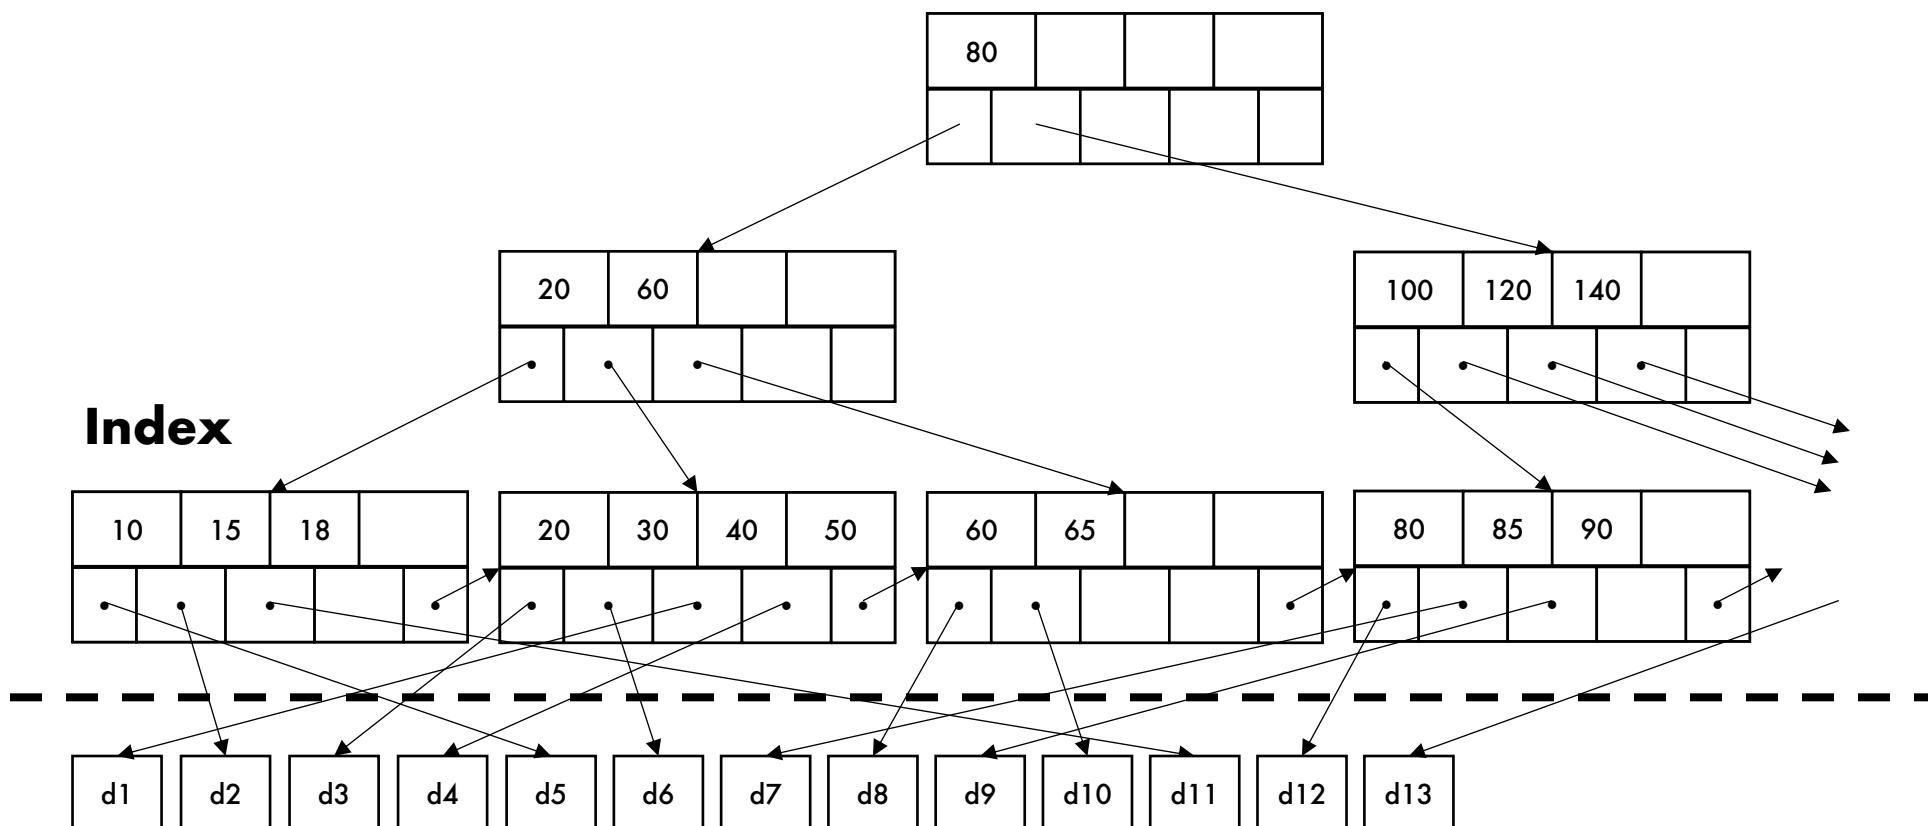
How is a B+ Tree Different?



Find the data associated with the key value 40
(same search process)

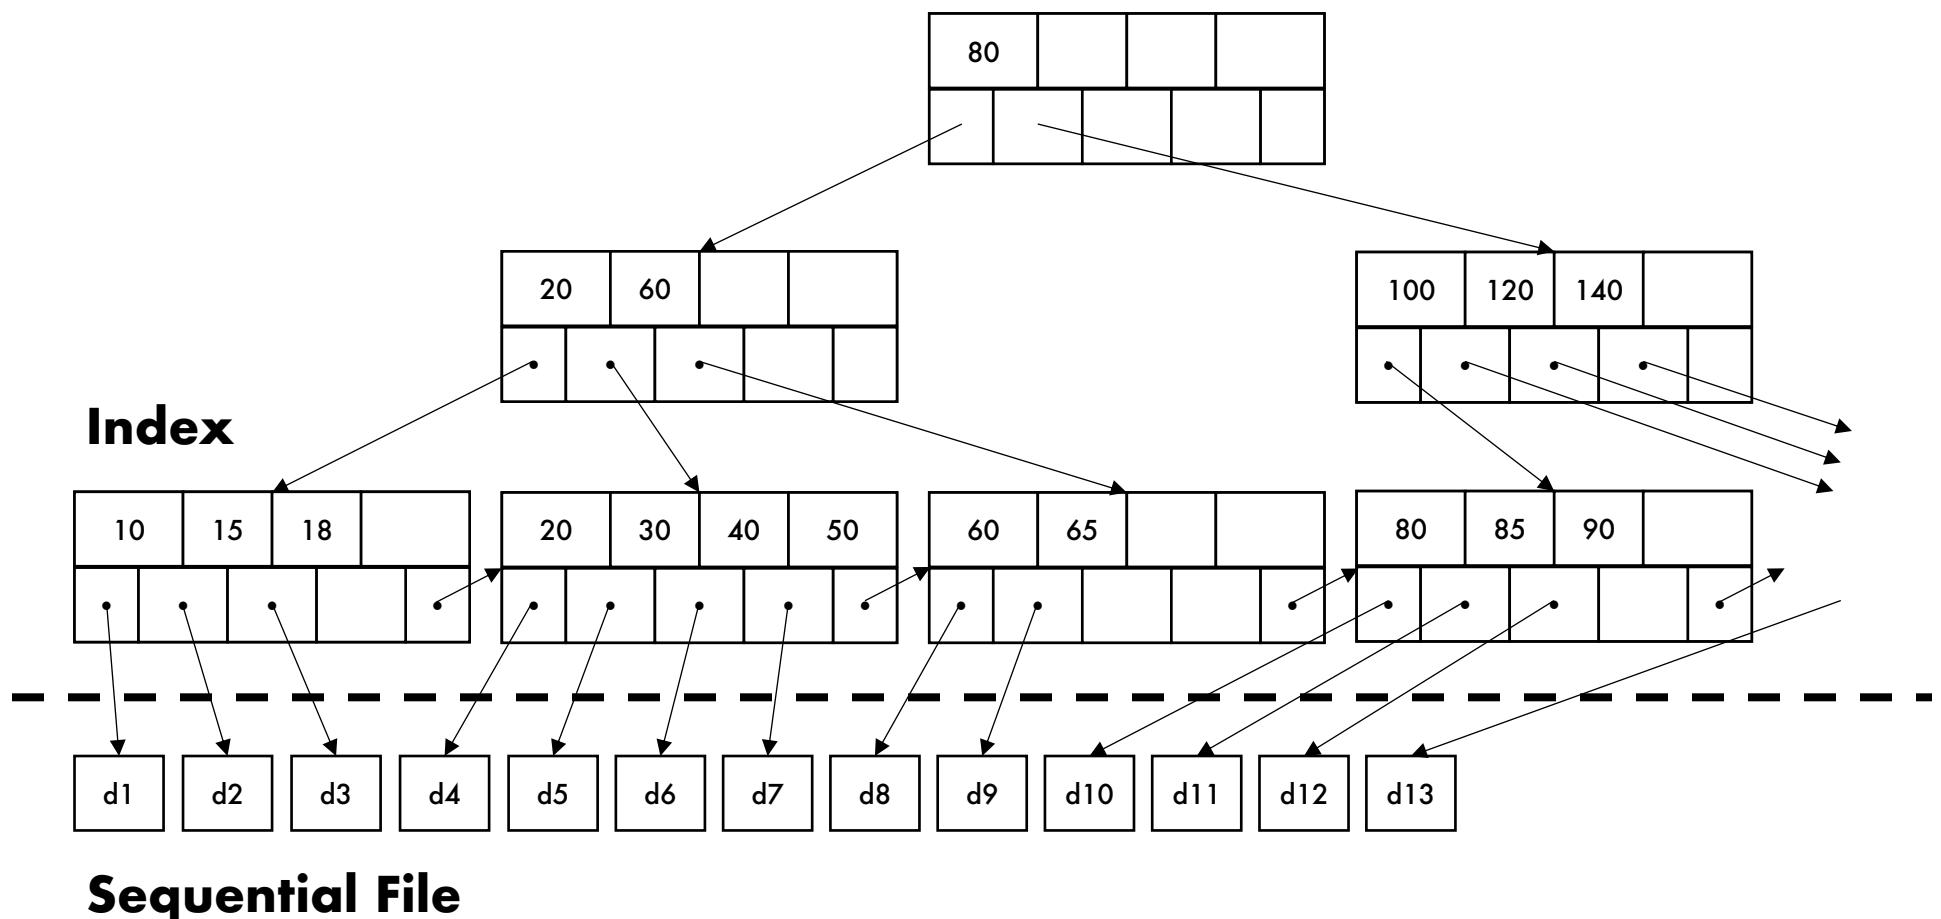
Clustered vs Unclustered Index

- An **unclustered index** may exist without any ordering on disk (any number per table)



Clustered vs Unclustered Index

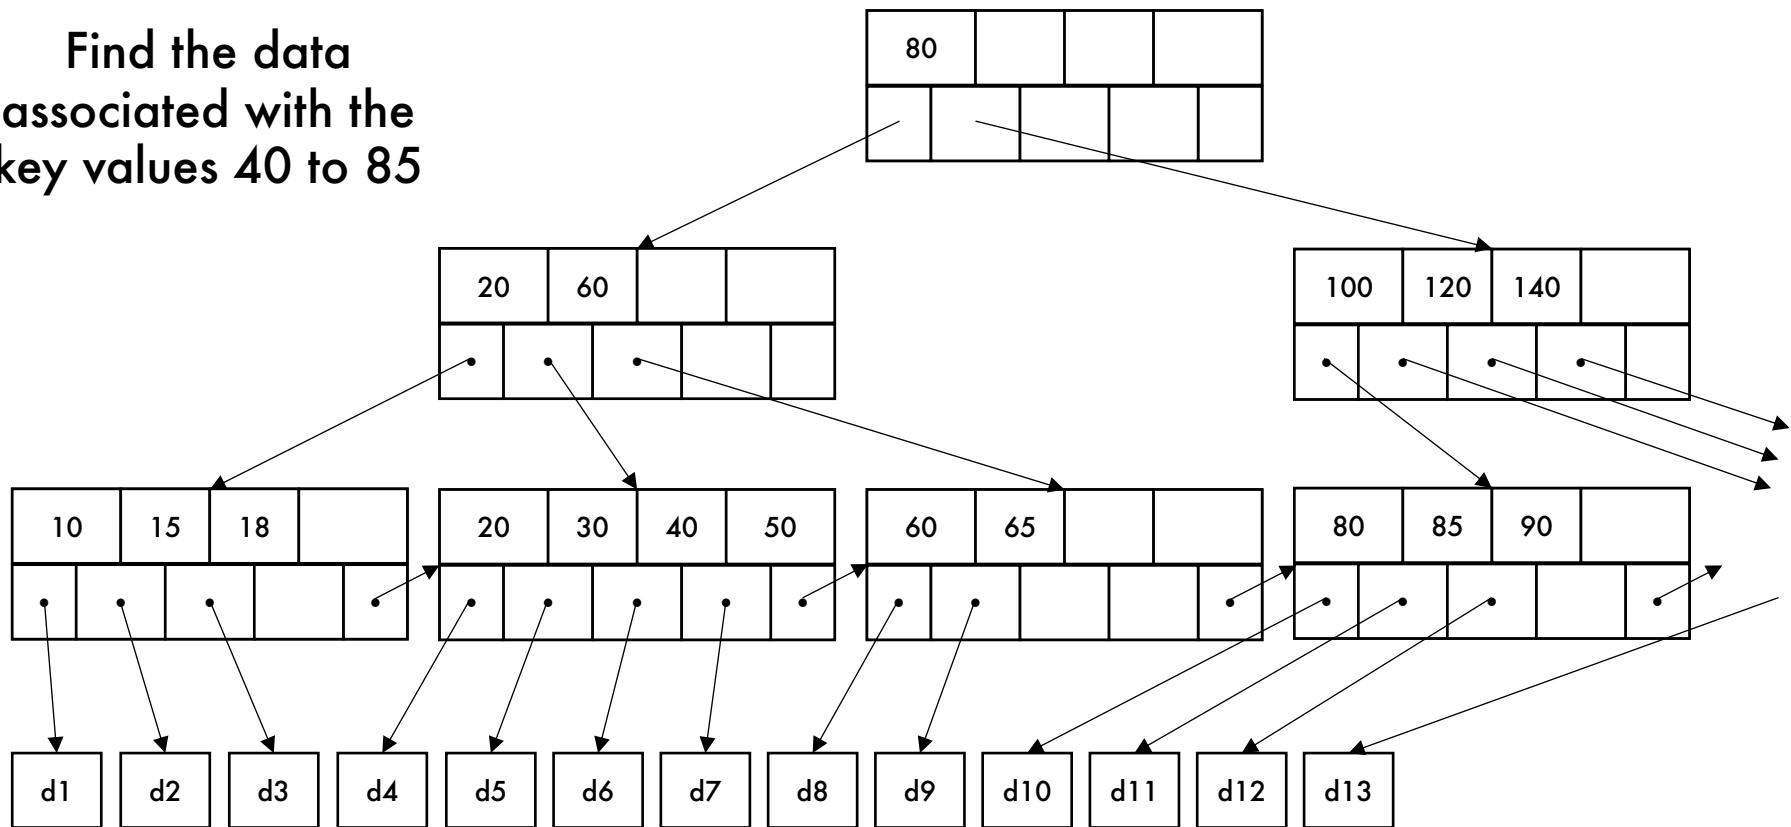
A **clustered index** is one that has the **same key ordering** as what is on disk (one per table)



Benefits of B+ Trees

- Range queries can be fast!
 - Filtering a value on a valid range is essentially looking up some portion of a B+ tree

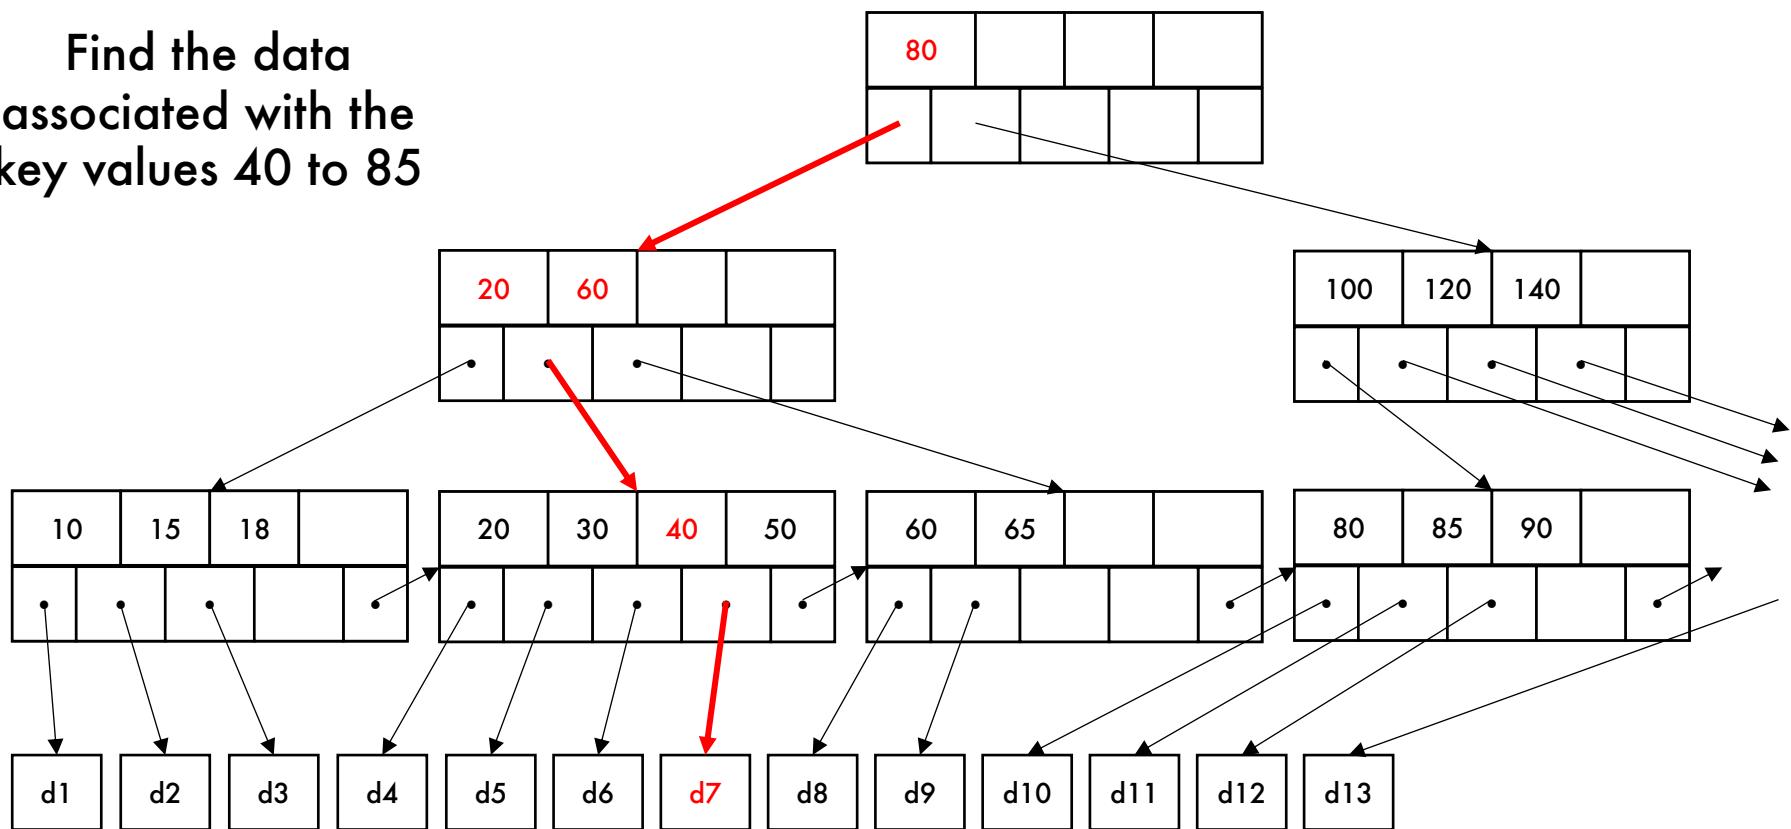
Find the data associated with the key values 40 to 85



Benefits of B+ Trees

- Range queries can be fast!
 - Filtering a value on a valid range is essentially looking up some portion of a B+ tree

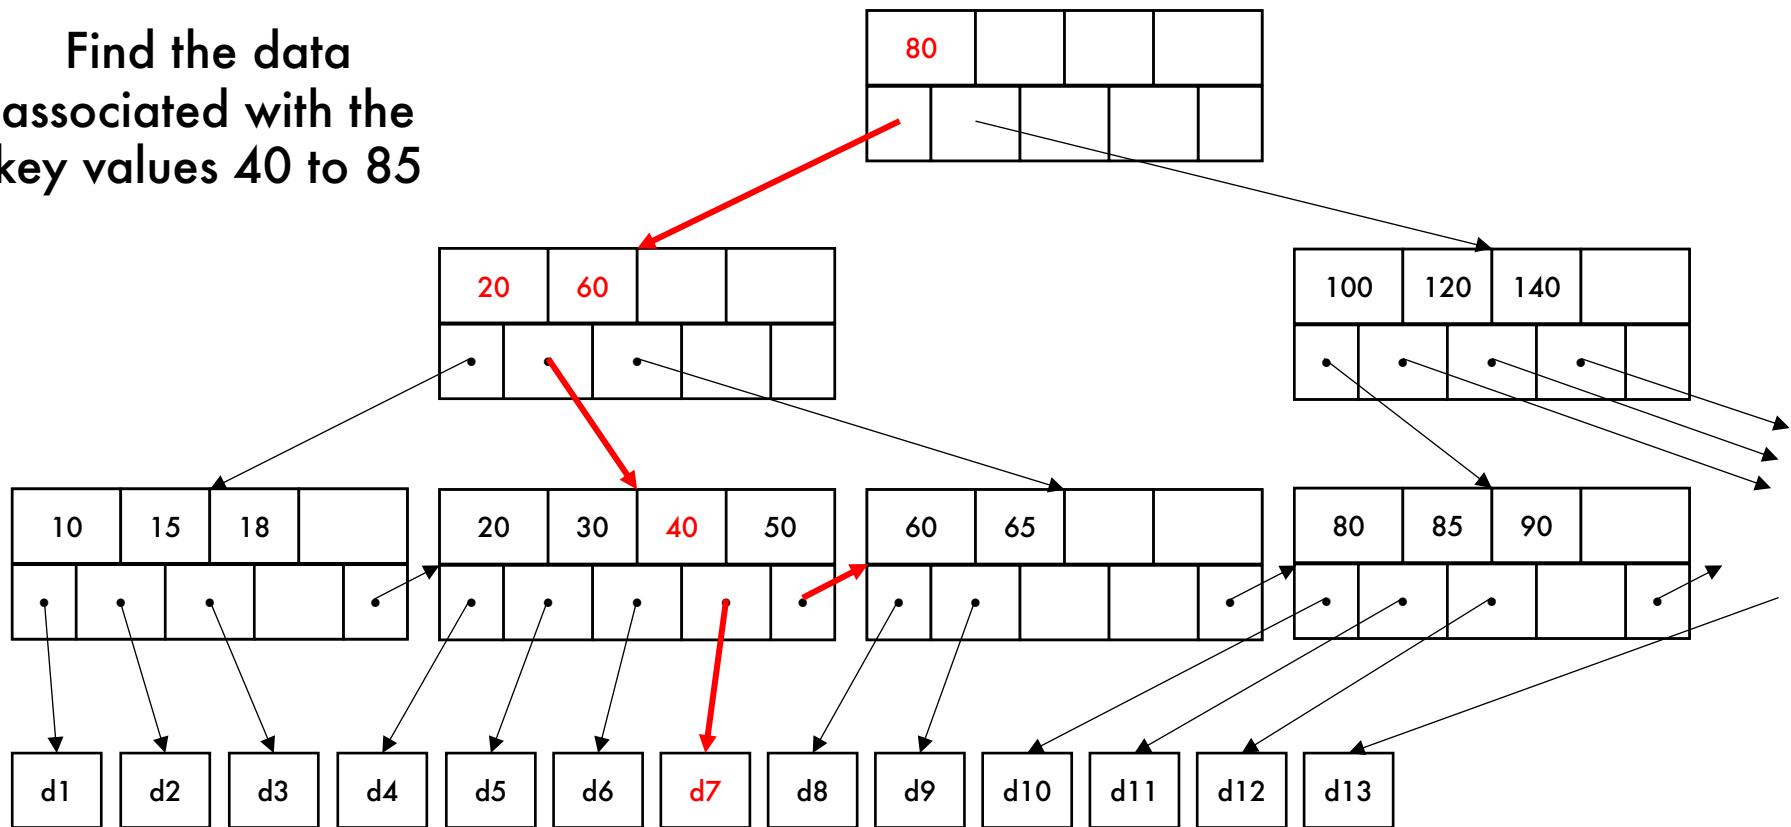
Find the data associated with the key values 40 to 85



Benefits of B+ Trees

- Range queries can be fast!
 - Filtering a value on a valid range is essentially looking up some portion of a B+ tree

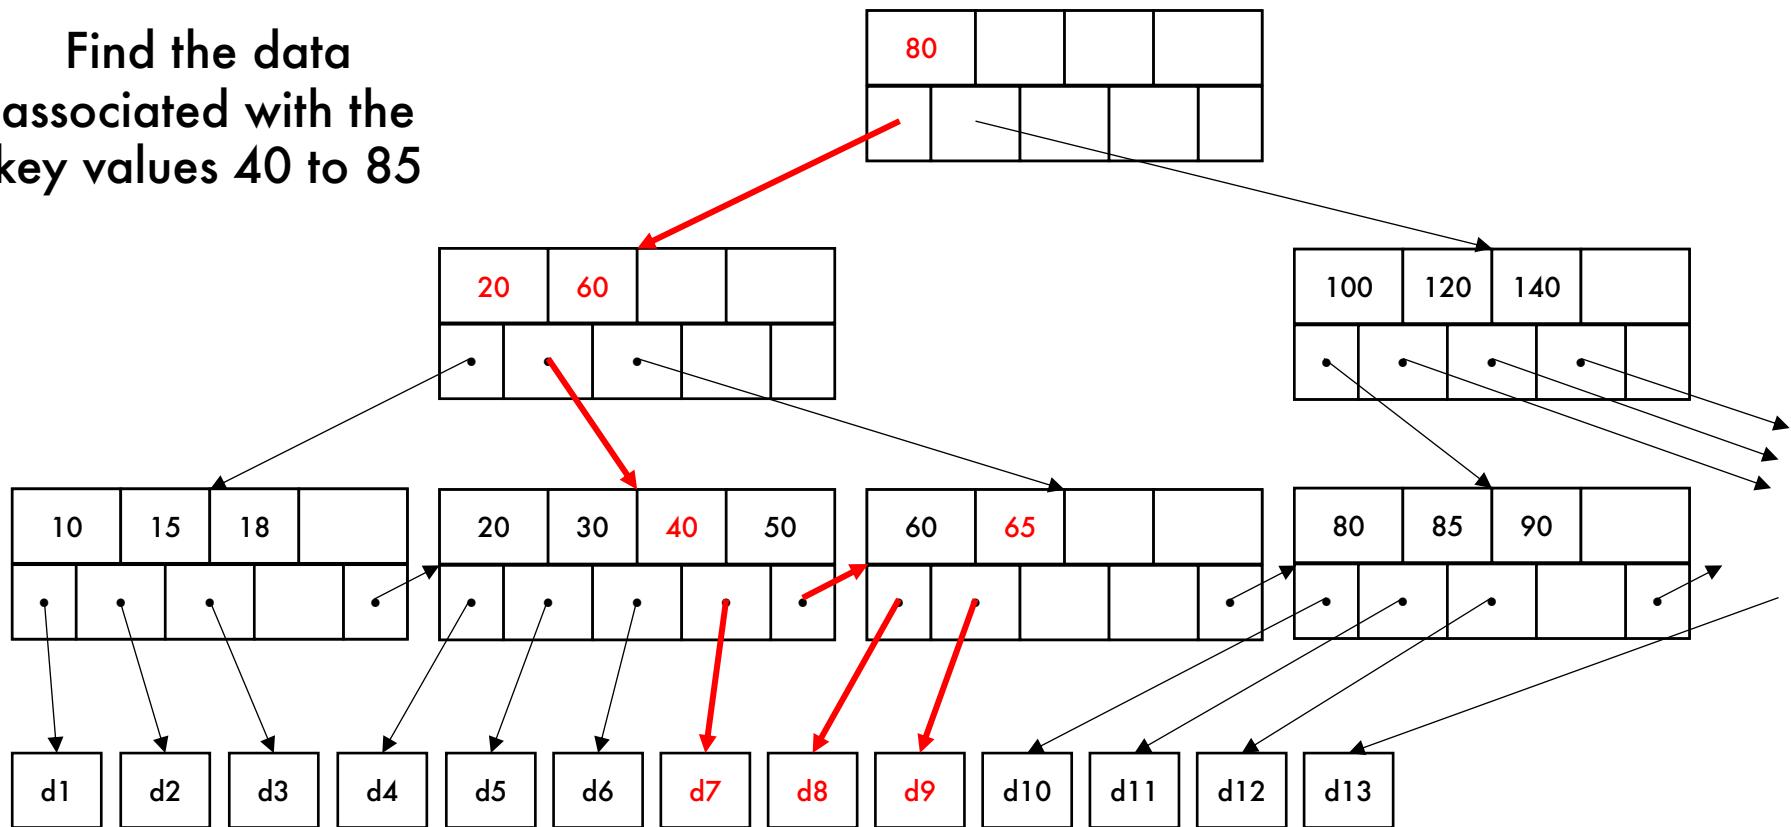
Find the data associated with the key values 40 to 85



Benefits of B+ Trees

- Range queries can be fast!
 - Filtering a value on a valid range is essentially looking up some portion of a B+ tree

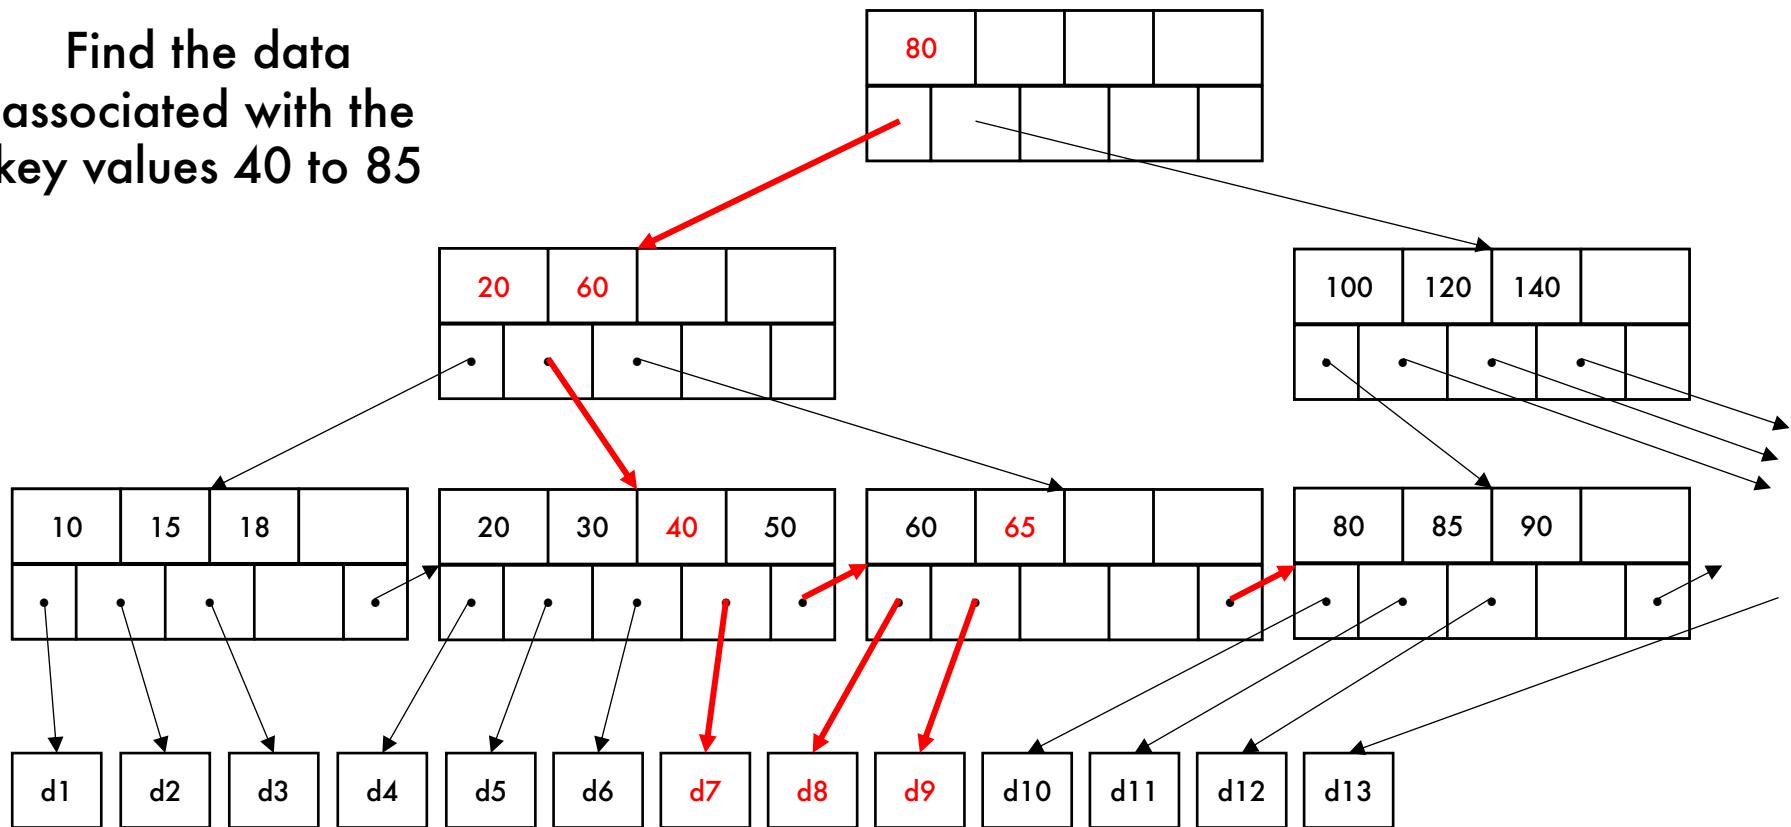
Find the data associated with the key values 40 to 85



Benefits of B+ Trees

- Range queries can be fast!
 - Filtering a value on a valid range is essentially looking up some portion of a B+ tree

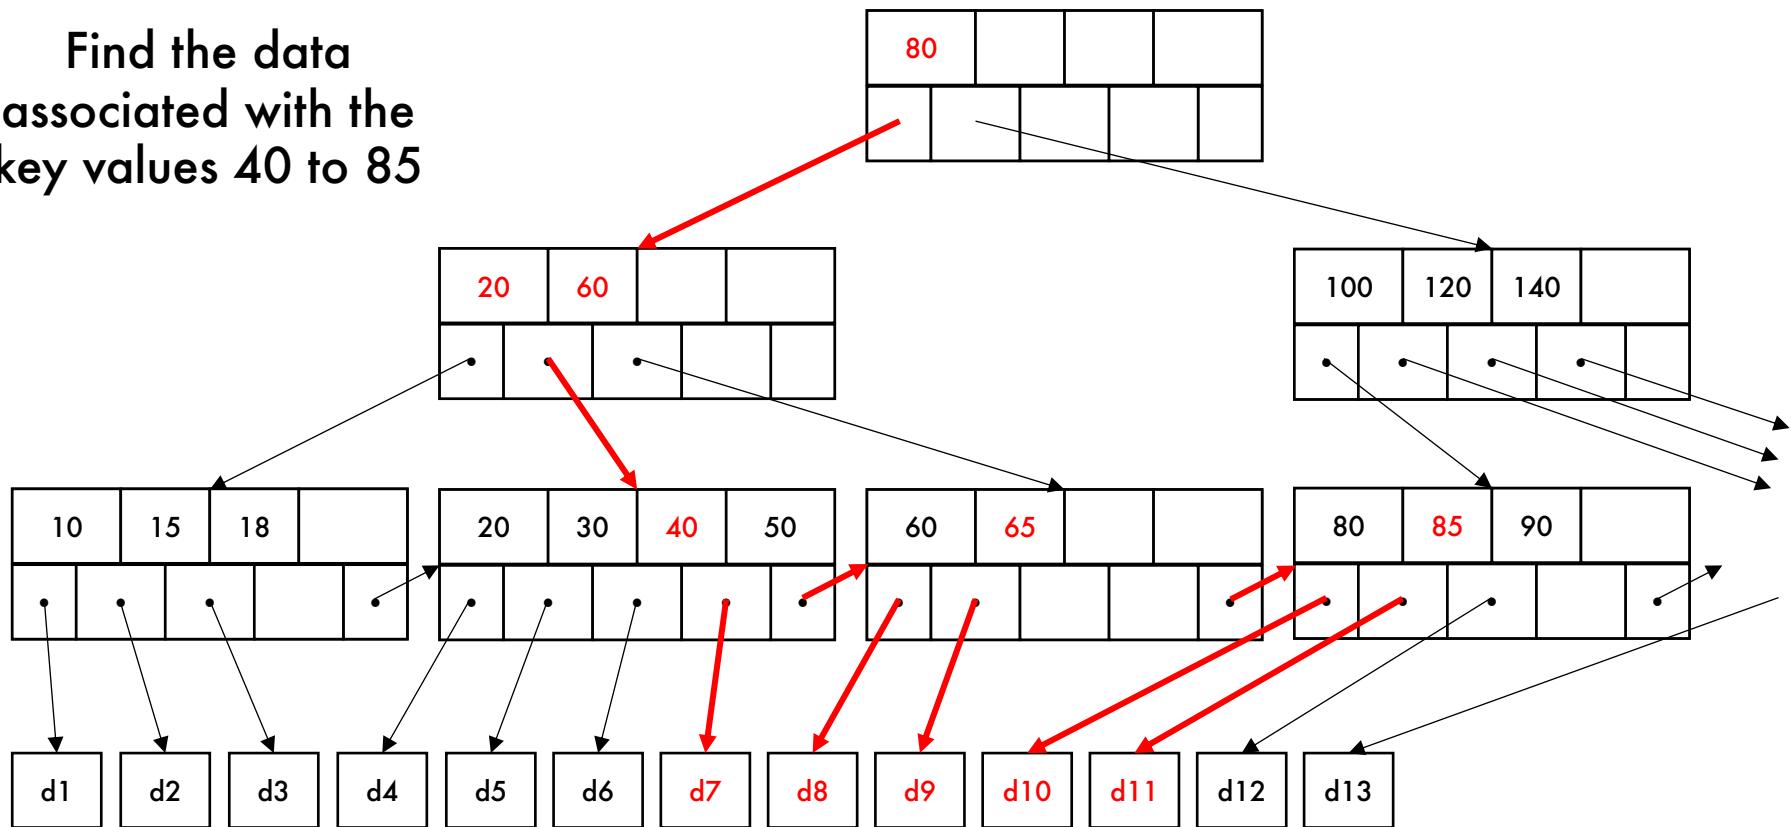
Find the data associated with the key values 40 to 85



Benefits of B+ Trees

- Range queries can be fast!
 - Filtering a value on a valid range is essentially looking up some portion of a B+ tree

Find the data associated with the key values 40 to 85



Estimating Amount of Data Read

- **Selectivity Factor (X)** → Proportion of total data needed
- Assuming uniform distribution of data values on numeric attribute a in table R , if the condition is:
 - $a = c \rightarrow X \cong \frac{1}{V(a,R)}$
 - $a < c \rightarrow X \cong \frac{c - \min(a,R)}{\max(a,R) - \min(a,R)}$
 - $c_1 < a < c_2 \rightarrow X \cong \frac{c_2 - c_1}{\max(a,R) - \min(a,R)}$
 - cond1 AND cond2 $\rightarrow X \cong X_1 * X_2$
- Disclaimer: More thorough selectivity estimation will use a histogram

Index-Based Selection

- For reference, a full sequential scan of data costs $B(R)$ IOs
- Provided some condition to read data:
 - Full **sequential scan** $\rightarrow B(R)$
 - Scan on **clustered index** $\rightarrow X^* B(R)$
 - Able to read a contiguous chunk of the file
 - Scan on **unclustered index** $\rightarrow X^* T(R)$
 - Worst case would read a different block everytime

Sequential Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

Without an index, finding a value must be done the “old fashioned way”

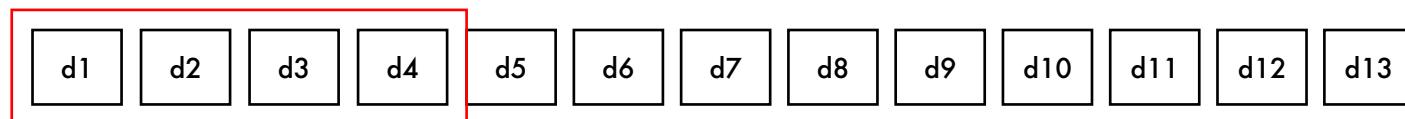


Disk

Sequential Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

Without an index, finding a value must be done the “old fashioned way”

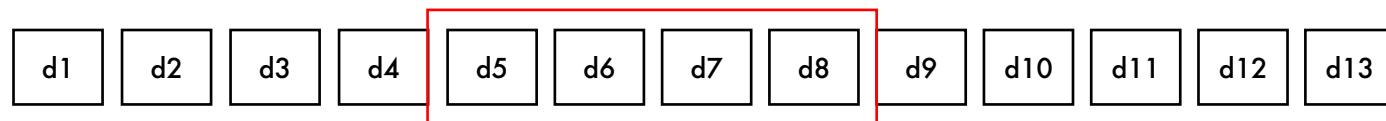


Disk

Sequential Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

Without an index, finding a value must be done the “old fashioned way”

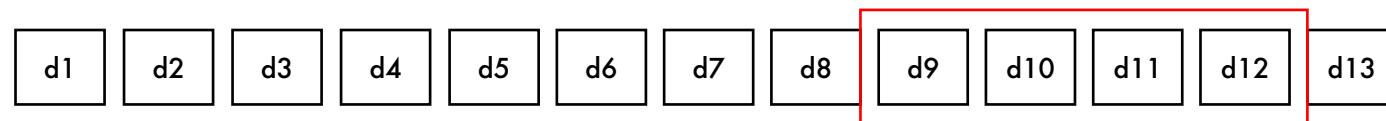


Disk

Sequential Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

Without an index, finding a value must be done the “old fashioned way”

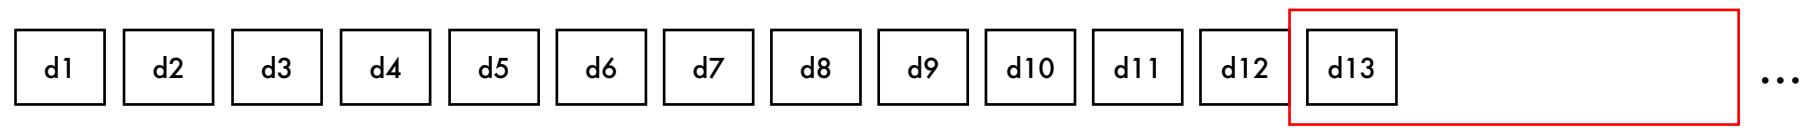


Disk

Sequential Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

Without an index, finding a value must be done the “old fashioned way”



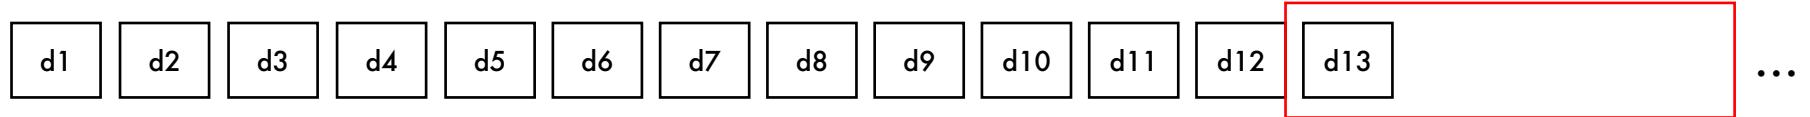
Disk

Sequential Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

Without an index, finding a value must be done the “old fashioned way”

Total cost: $B(R)$

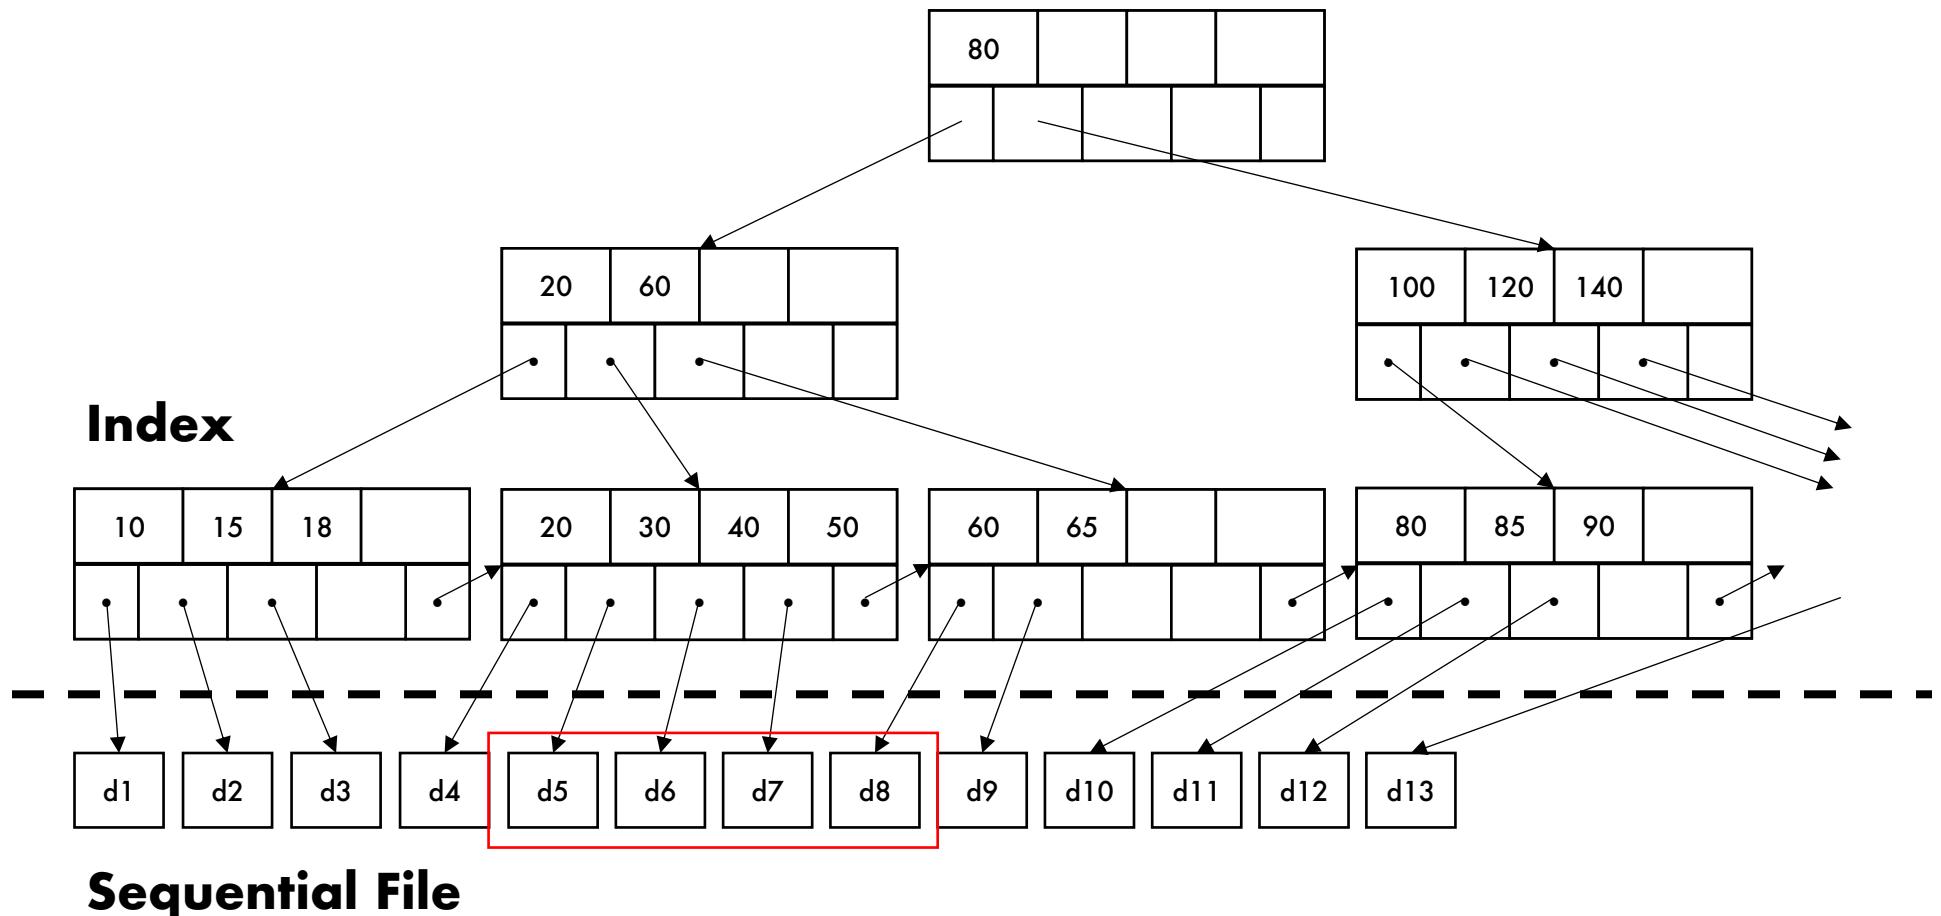


Disk

Clustered Index Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

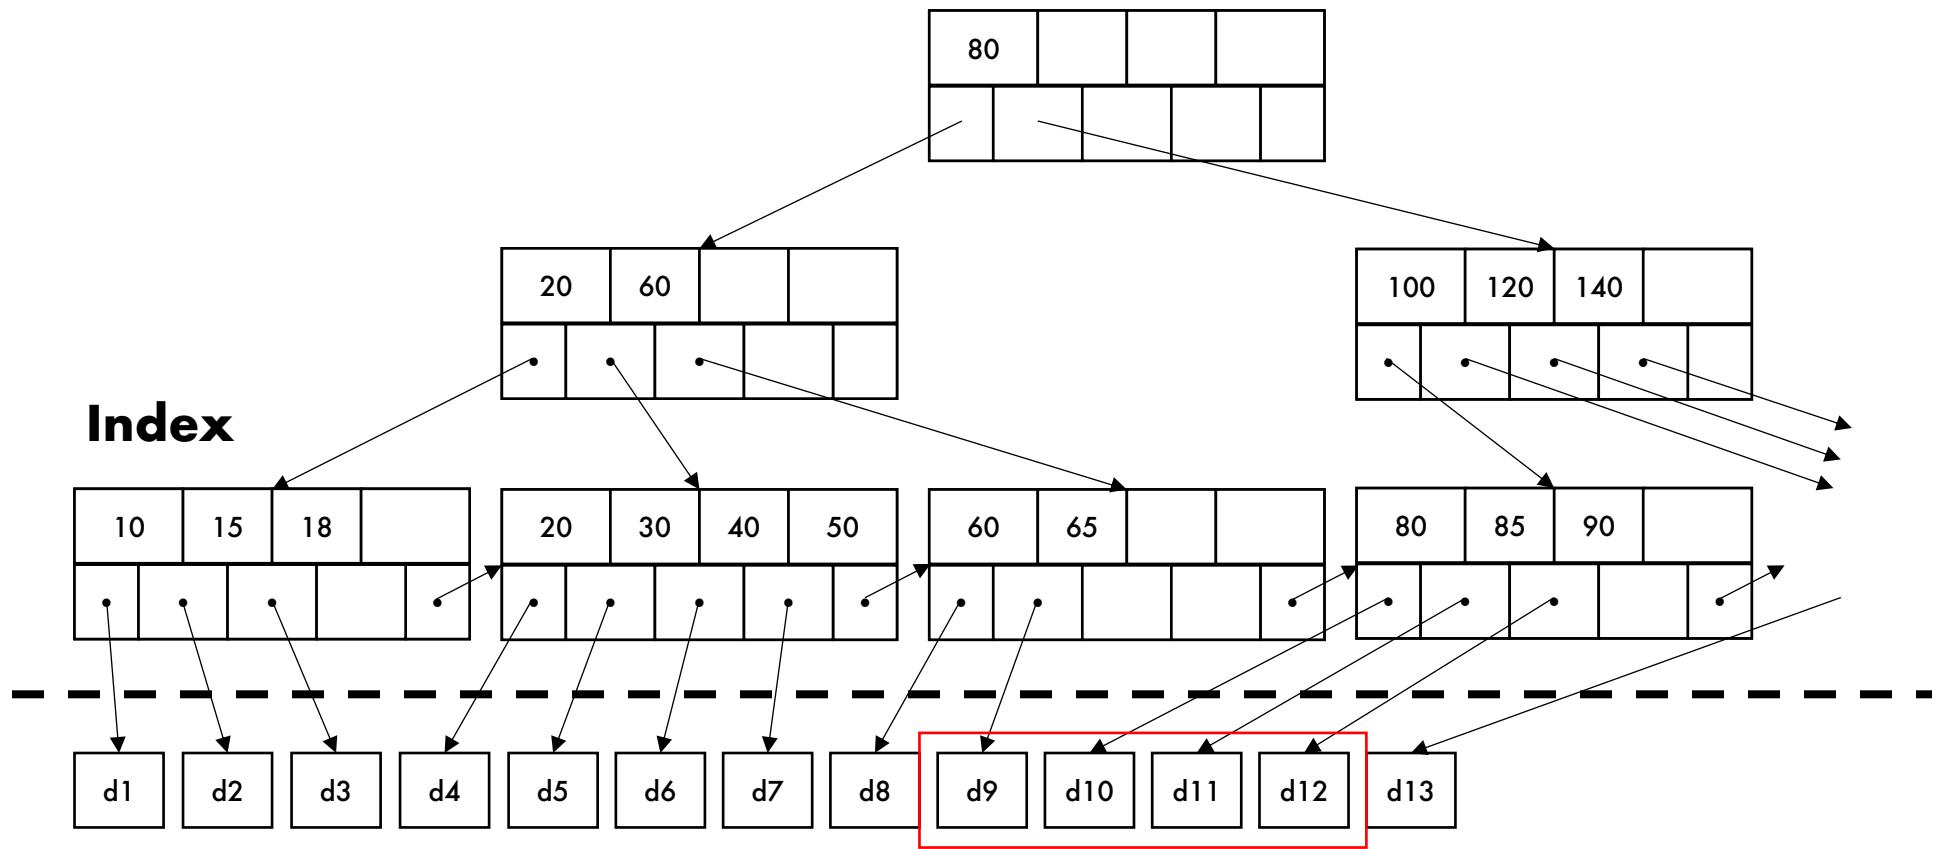
With a clustered index, I start scanning blocks in the range they are at



Clustered Index Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

With a clustered index, I start scanning blocks in the range they are at

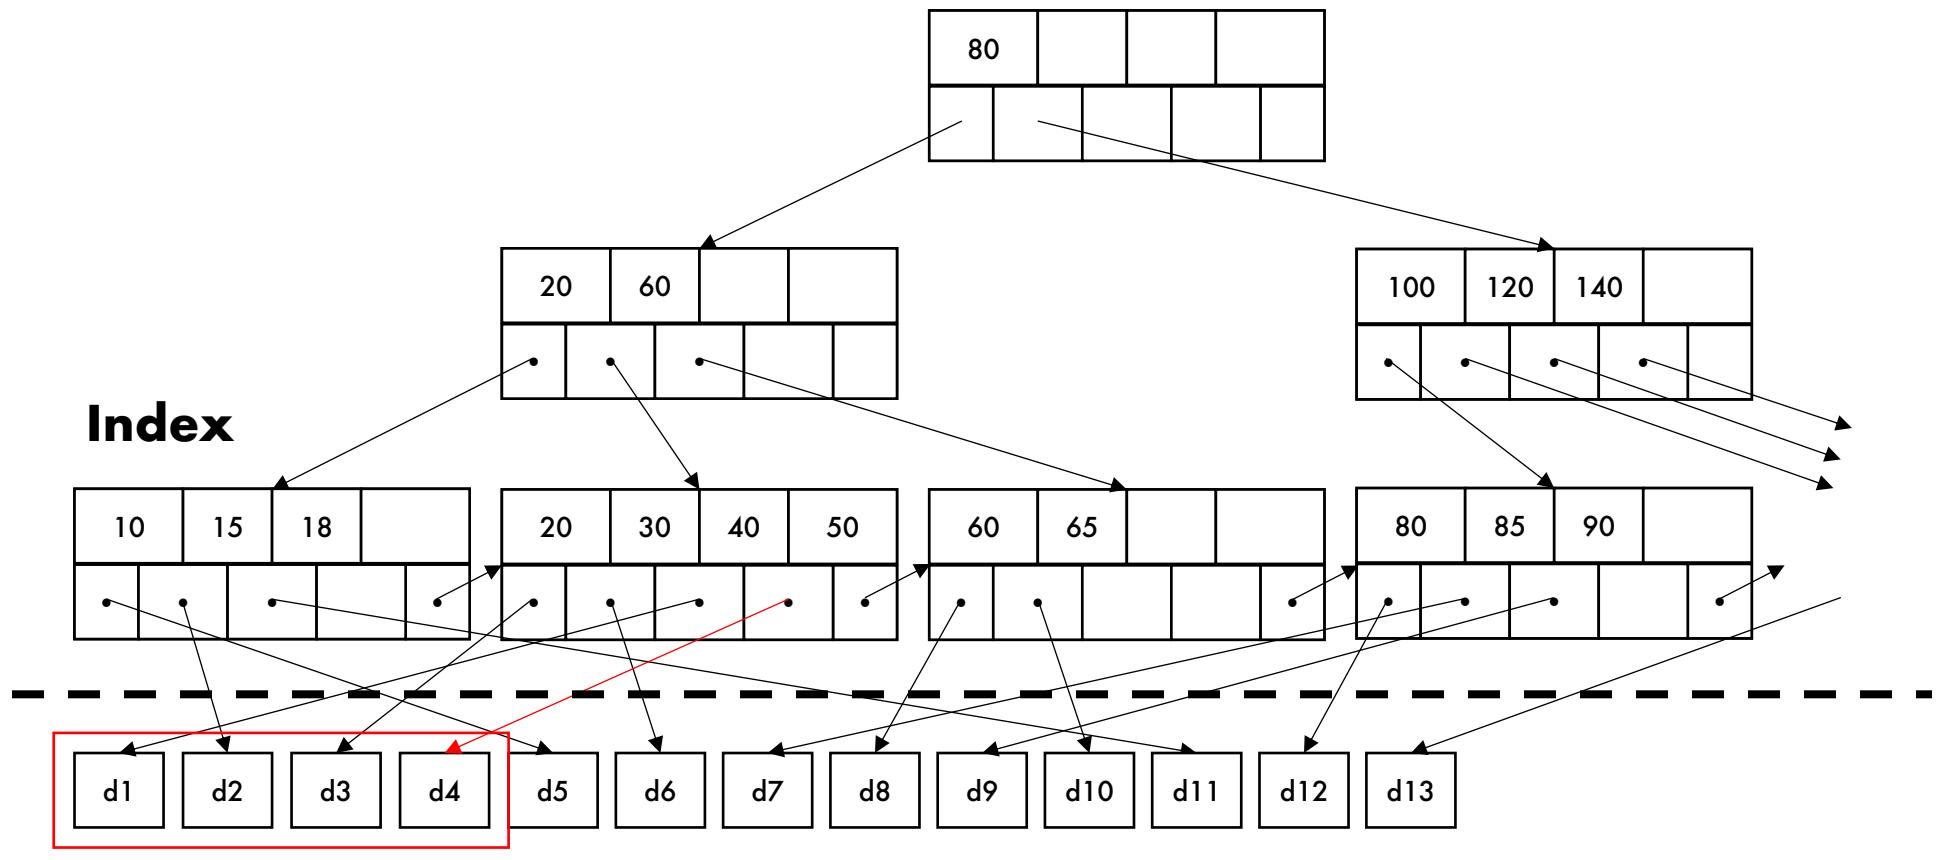


Sequential File

Unclustered Index Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

With an unclustered index, I start scanning tuples wherever they occur

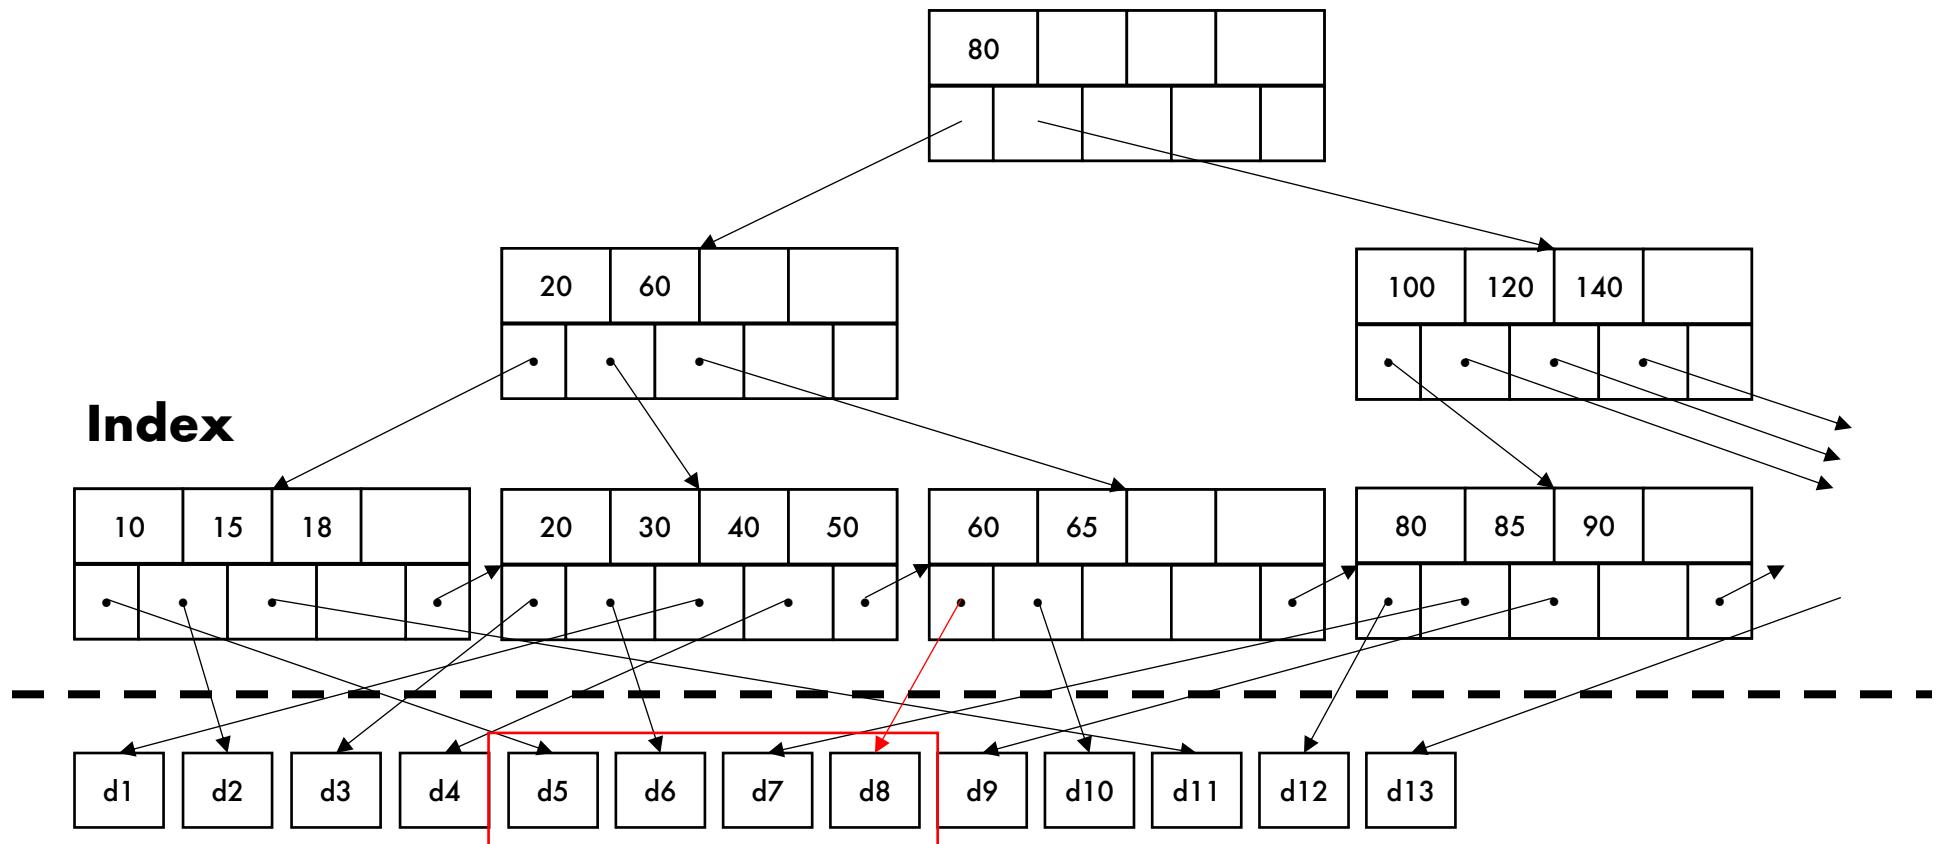


Sequential File with a different key or Heap File

Unclustered Index Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

With an unclustered index, I start scanning tuples wherever they occur

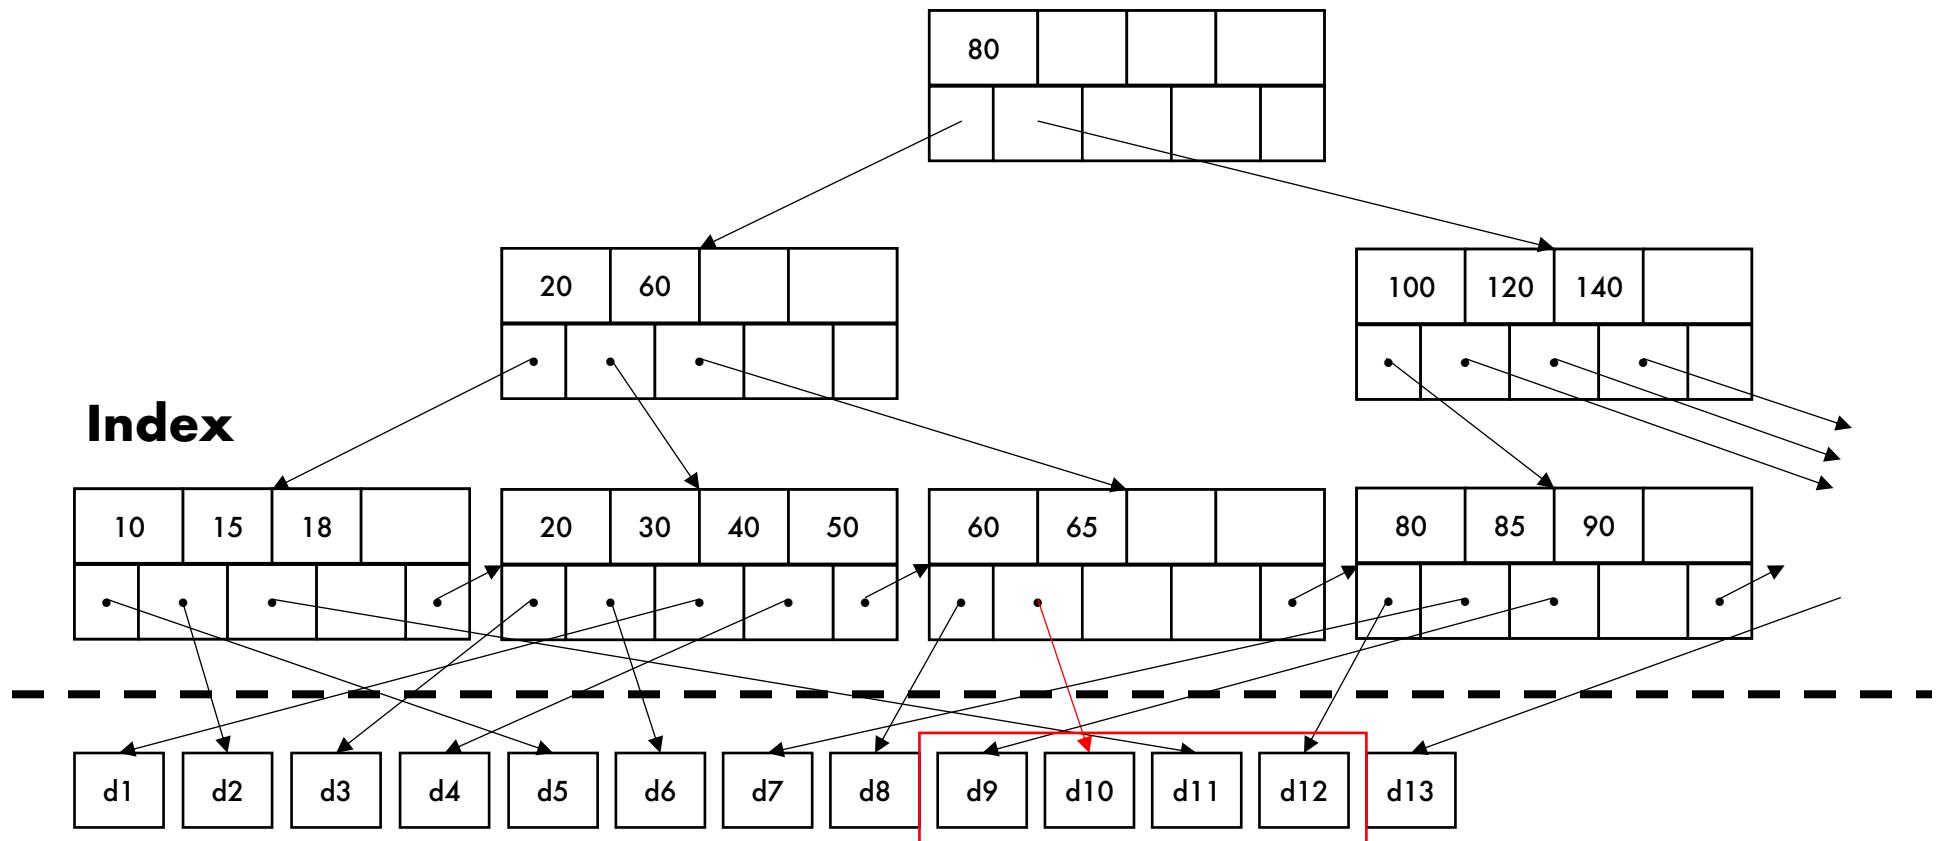


Sequential File with a different key or Heap File

Unclustered Index Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

With an unclustered index, I start scanning tuples wherever they occur

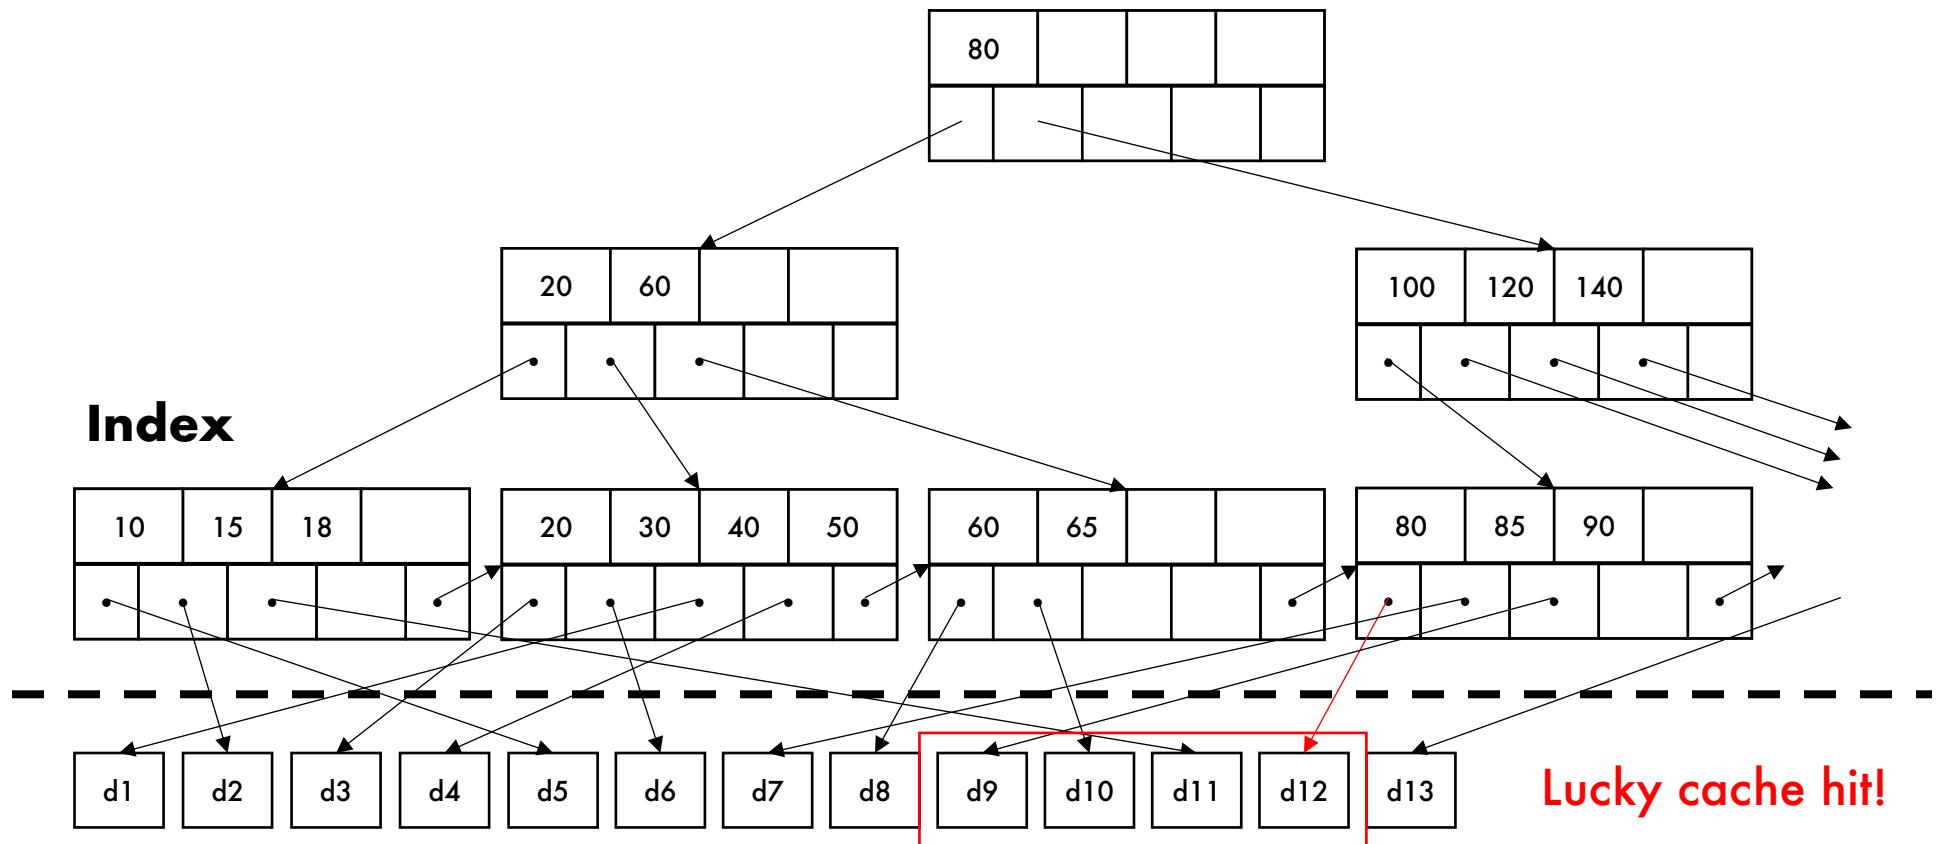


Sequential File with a different key or Heap File

Unclustered Index Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

With an unclustered index, I start scanning tuples wherever they occur

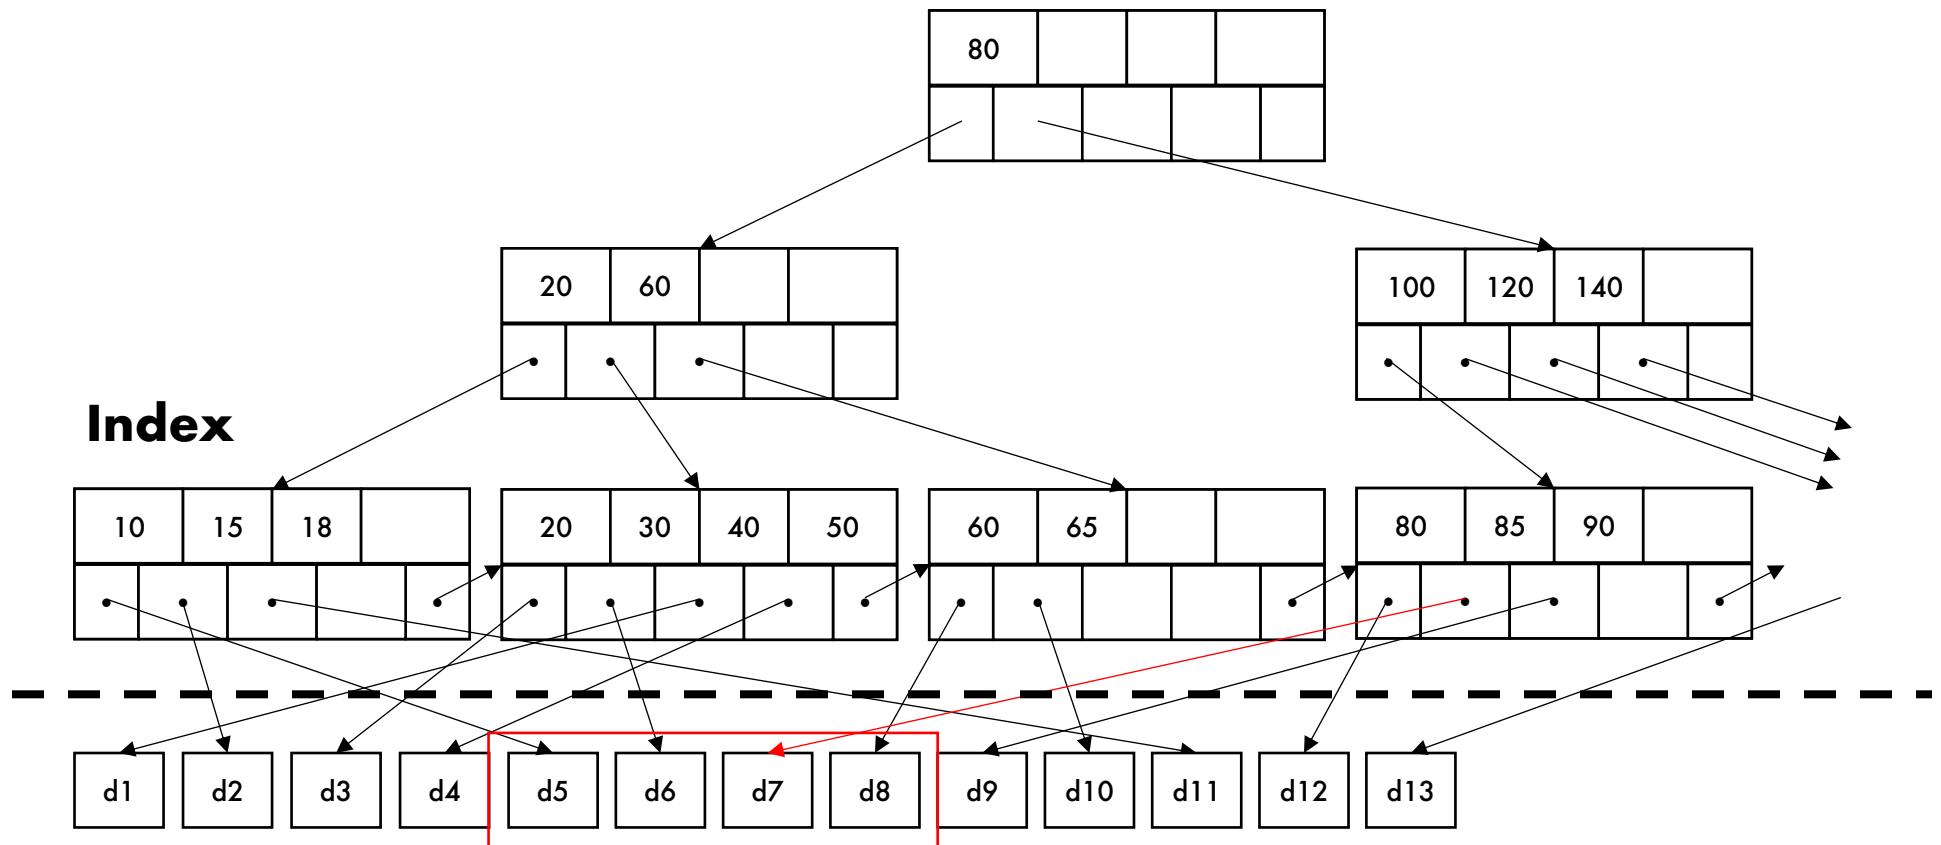


Sequential File with a different key or Heap File

Unclustered Index Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

With an unclustered index, I start scanning tuples wherever they occur



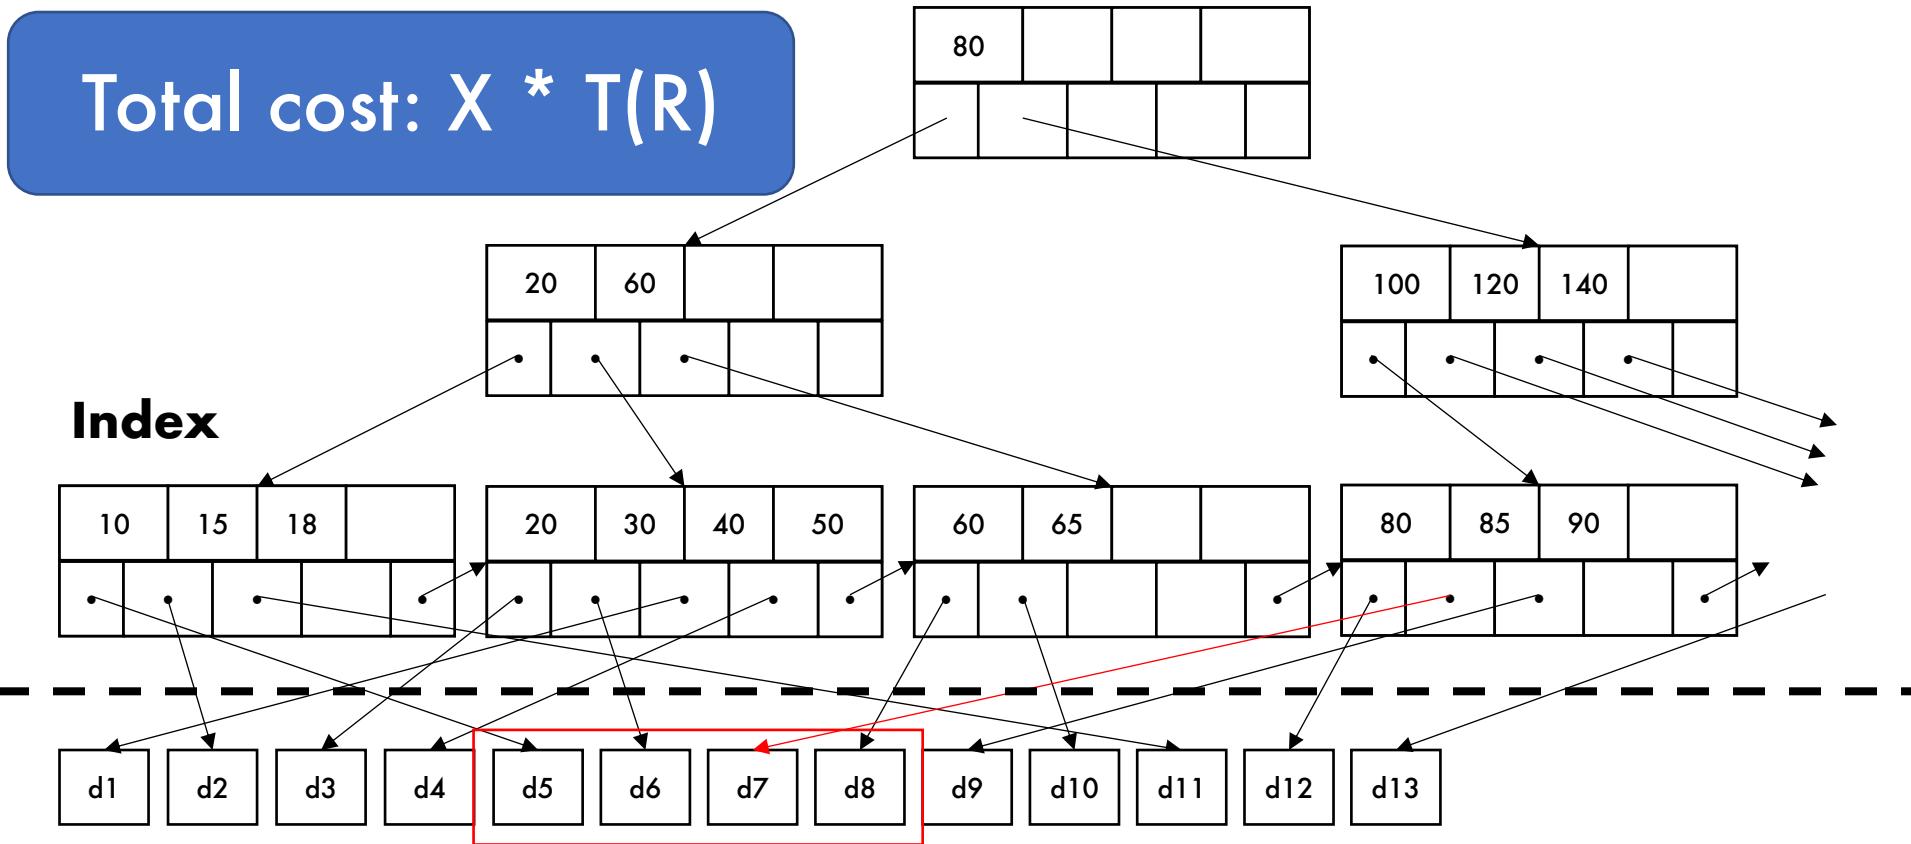
Sequential File with a different key or Heap File

Unclustered Index Scan

Assume a block holds 4 tuples. I want tuples associated with values 40-85.

With an unclustered index, I start scanning tuples wherever they occur

Total cost: $X * T(R)$



Index Expectations

- Using an index in the wrong scenario can lead to a slowdown!
- Common example: Full sequential scan vs unclustered index scan with **high X value** and/or **small tuple size** (large $T(R):B(R)$ ratio)

Known:

$$B(R) = 100$$

$$T(R) = 10000$$

Consider a query with $X=1/10$

Sequential scan

$$= B(R)$$

$$= 100$$

Index scan

$$= X * T(R)$$

$$= 1/10 * 10000$$

$$= 1000$$

Index Expectations

- Using an index in the wrong scenario can lead to a slowdown!
- Common example: Full sequential scan vs unclustered index scan with **high X value** and/or **small tuple size** (large $T(R):B(R)$ ratio)

Having indexes
doesn't mean you
will see a speedup!

Known:
 $B(R) = 100$
 $T(R) = 10000$

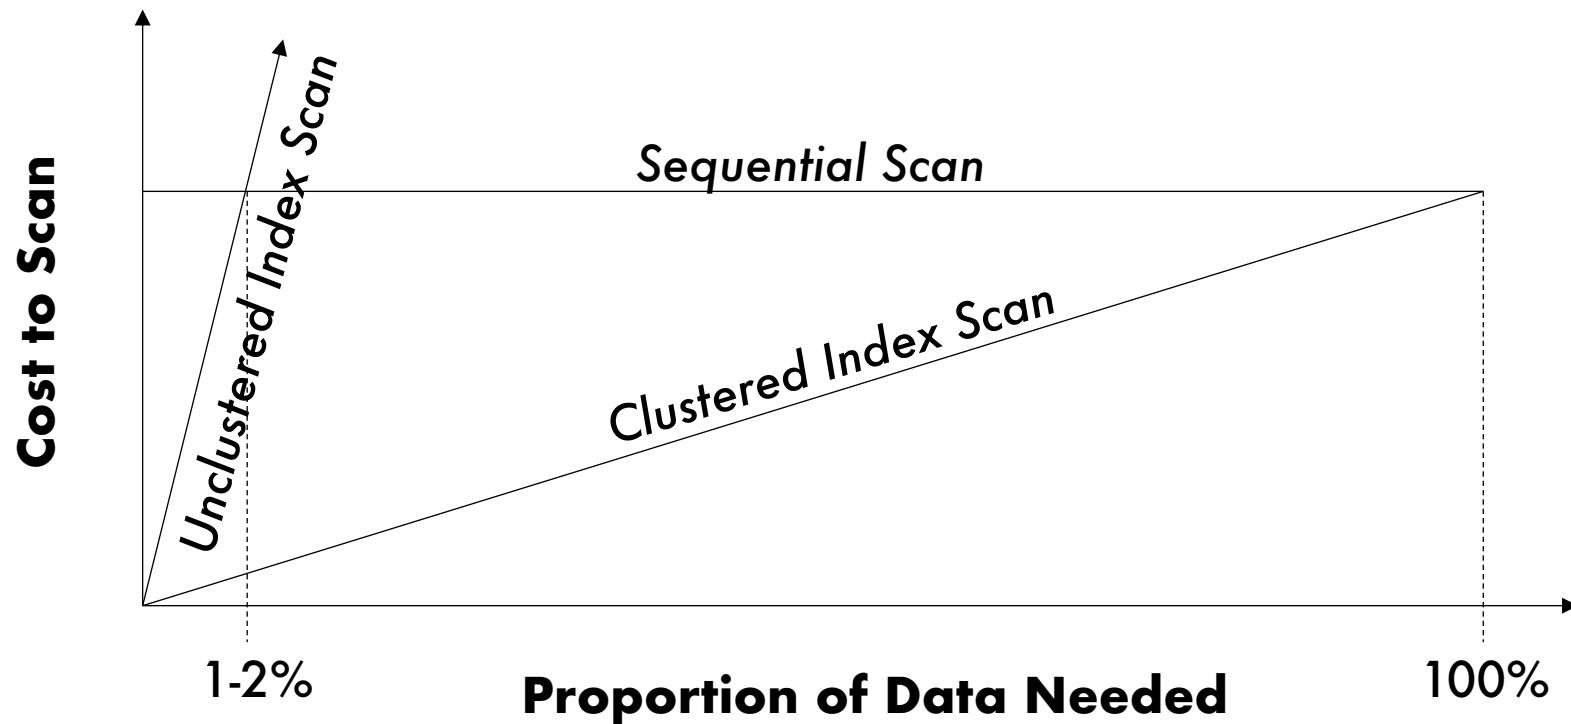
Consider a query with $X=1/10$

Sequential scan
 $= B(R)$
 $= 100$

Index scan
 $= X*T(R)$
 $= 1/10 * 10000$
 $= 1000$

Index Expectations

- **Sequential disk reads are faster than random ones**
 - Cost $\sim 1\text{-}2\%$ random scan = full sequential scan



Create Indexes in SQL

```
CREATE TABLE Users (  
    id INT,  
    age INT,  
    score INT);
```

```
CREATE INDEX U_age ON Users(age)
```

```
CREATE INDEX U_age_score ON Users(age, score)
```

```
CREATE CLUSTERED INDEX U_score_age ON Users(score, age)
```

Create Indexes in SQL

```
CREATE TABLE Users (  
    id INT,  
    age INT,  
    score INT);
```

```
CREATE INDEX U_age ON Users(age)
```

Unclustered
by default

```
CREATE INDEX U_age_score ON Users(age, score)
```

```
CREATE CLUSTERED INDEX U_score_age ON Users(score, age)
```

Create Indexes in SQL

```
CREATE TABLE Users (  
    id INT,  
    age INT,  
    score INT);
```

```
CREATE INDEX U_age ON Users(age)
```

Unclustered
by default

```
CREATE INDEX U_age_score ON Users(age, score)
```

Order specifies
precedence in sorting

```
CREATE CLUSTERED INDEX U_score_age ON Users(score, age)
```

Create Indexes in SQL

```
CREATE TABLE Users (  
    id INT,  
    age INT,  
    score INT);
```

```
CREATE INDEX U_age ON Users(age)
```

Unclustered
by default

```
CREATE INDEX U_age_score ON Users(age, score)
```

Order specifies
precedence in sorting

```
CREATE CLUSTERED INDEX U_score_age ON Users(score, age)
```

Reorders data on disk! (Fails if
another clustered index exists)

Create Indexes in SQL

```
CREATE TABLE Users (  
    id INT,  
    age INT,  
    score INT);
```

Fine to create indexes on non-numeric data
(just not discussed in this class)

```
CREATE INDEX U_age ON Users(age)
```

Unclustered
by default

```
CREATE INDEX U_age_score ON Users(age, score)
```

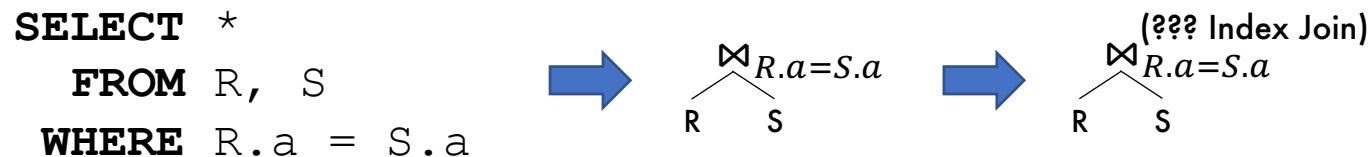
Order specifies
precedence in sorting

```
CREATE CLUSTERED INDEX U_score_age ON Users(score, age)
```

Reorders data on disk! (Fails if
another clustered index exists)

Index-Based Equijoin

- Assume index exists on the join attribute a of S



- **Clustered Index Join**

- Perform a clustered index scan for each tuple of R
- $B(R) + T(R)(X^* B(S)) = B(R) + T(R)(\underline{B(S)} / \underline{V(a, S)})$

- **Unclustered Index Join**

- Perform an unclustered index scan for each tuple of R
- $B(R) + T(R)(X^* T(S)) = B(R) + T(R)(\underline{T(S)} / \underline{V(a, S)})$

Index-Based Equijoin

- Assume index exists on the join attribute a of S

```
SELECT *
  FROM R, S
 WHERE R.a = S.a
```

↗
R S
↗
R S
(??? Index Join)

- **Clustered Index Join**

- Perform a clustered index scan for each tuple of R
- $B(R) + T(R)(X^*B(S)) = \underline{B(R) + T(R)(B(S)/V(a,S))}$

Can't scan per block of R
since tuples in blocks don't
have the same attribute values

- **Unclustered Index Join**

- Perform an unclustered index scan for each tuple of R
- $B(R) + T(R)(X^*T(S)) = \underline{B(R) + T(R)(T(S)/V(a,S))}$

Leveraging Indexes

- Often for applications, workloads can be well described
 - Lab 1 Flights application
 - Search method → query on city name values
 - Data visualization software (e.g. Tableau)
 - 2D plot → query on graph axis bounds
- **Create indexes to match expected query workload**

Leveraging Indexes

```
CREATE TABLE Users (
    id INT PRIMARY KEY,
    age INT,
    score INT, ...);
```

What indexes could we make on Users?

Expecting 1000 exec/day

```
SELECT *
    FROM Users, Assets
   WHERE Users.id = Assets.uid
```

Expecting 1000 exec/day

```
SELECT *
    FROM Users
   WHERE Users.score > 95
```

Expecting 10 exec/day

```
SELECT *
    FROM Users
   WHERE Users.age > 21
```

Leveraging Indexes

```
CREATE TABLE Users (
    id INT PRIMARY KEY,
    age INT,
    score INT, ...);
```

Expecting 1000 exec/day

```
SELECT *
    FROM Users, Assets
   WHERE Users.id = Assets.uid
```

Expecting 1000 exec/day

```
SELECT *
    FROM Users
   WHERE Users.score > 95
```

Expecting 10 exec/day

```
SELECT *
    FROM Users
   WHERE Users.age > 21
```

What indexes could we make on Users?

IDs are unique so an unclustered index would do fine.

Leveraging Indexes

```
CREATE TABLE Users (
    id INT PRIMARY KEY,
    age INT,
    score INT, ...);
```

Expecting 1000 exec/day

```
SELECT *
  FROM Users, Assets
 WHERE Users.id = Assets.uid
```

Expecting 1000 exec/day

```
SELECT *
  FROM Users
 WHERE Users.score > 95
```

Expecting 10 exec/day

```
SELECT *
  FROM Users
 WHERE Users.age > 21
```

What indexes could we make on Users?

IDs are unique so an unclustered index would do fine.

This range query would benefit from a clustered index on score



Only one can exist!

This range query would benefit from a clustered index on age

Leveraging Indexes

```
CREATE TABLE Users (
    id INT PRIMARY KEY,
    age INT,
    score INT, ...);
```

Expecting 1000 exec/day

```
SELECT *
  FROM Users, Assets
 WHERE Users.id = Assets.uid
```

Expecting 1000 exec/day

```
SELECT *
  FROM Users
 WHERE Users.score > 95
```

Expecting 10 exec/day

```
SELECT *
  FROM Users
 WHERE Users.age > 21
```

What indexes could we make on Users?

IDs are unique so an unclustered index would do fine.

Things to consider:

- What is the expected result size for each query?
- Do either of these queries need to be returned ASAP?

Leveraging Indexes

```
CREATE TABLE Users (
    id INT PRIMARY KEY,
    age INT,
    score INT, ...);
```

Expecting 1000 exec/day

```
SELECT *
    FROM Users, Assets
   WHERE Users.id = Assets.uid
```

Expecting 1000 exec/day

```
SELECT *
    FROM Users
   WHERE Users.score > 95
```

Expecting 10 exec/day

```
SELECT *
    FROM Users
   WHERE Users.age > 21
```

What indexes could we make on Users?

IDs are unique so an unclustered index would do fine.

Without more information, default to clustering on the index that will be used more (clustered index on score)

Leveraging Indexes

```
CREATE TABLE Users (
    id INT PRIMARY KEY,
    age INT,
    score INT, ...);
```

Expecting 1000 exec/day

```
SELECT *
    FROM Users, Assets
   WHERE Users.id = Assets.uid
```

Expecting 1000 exec/day

```
SELECT *
    FROM Users
   WHERE Users.score > 95
```

Expecting 10 exec/day

```
SELECT *
    FROM Users
   WHERE Users.age > 21
```

What indexes could we make on Users?

IDs are unique so an unclustered index would do fine.

Hack:

- Create a covering index primarily keyed on score
- Create a covering index primarily keyed on age

Leveraging Indexes

```
CREATE TABLE Users (
    id INT PRIMARY KEY,
    age INT,
    score INT, ...);
```

Expecting 1000 exec/day

```
SELECT *
  FROM Users, Assets
 WHERE Users.id = Assets.uid
```

Expecting 1000 exec/day

```
SELECT *
  FROM Users
 WHERE Users.score > 95
```

Expecting 10 exec/day

```
SELECT *
  FROM Users
 WHERE Users.age > 21
```

What indexes could we make on Users?

IDs
index
ed

Essentially a sorted copy of the table.
Fast but space inefficient and table updates are slow.

Hack:

- Create a covering index primarily keyed on score
- Create a covering index primarily keyed on age

“I’m not about that code monkey life”

- Choosing how to configure a database system is an interesting (i.e. hard) problem
- A database that is used by many people will often need one or more dedicated personnel to manage it (Database Administrator)
 - Logical design (multi-team coordination)
 - Physical design (hardware and system considerations)
 - Permission management (visibility and security)
 - Integration (company acquisitions and mergers)
 - ...

Multiple Joins

- **Pipelined Execution**

- Tuples are processed through the entire query plan
- Fast

- **Blocking Execution**

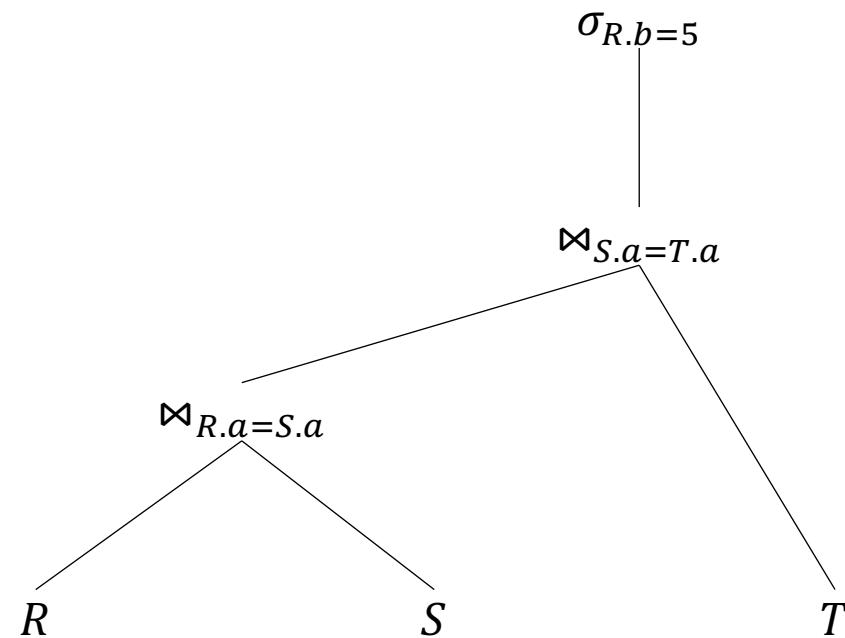
- Subplans are computed and stored before parent operation can start
- Simple

Pipelined Execution

- Iterator interface of RA operators (**Volcano Iterator Model**)
 - `open()` on every operator at start
 - `close()` on every operator at end
 - `next()` to get the next tuple from a child operator or input table

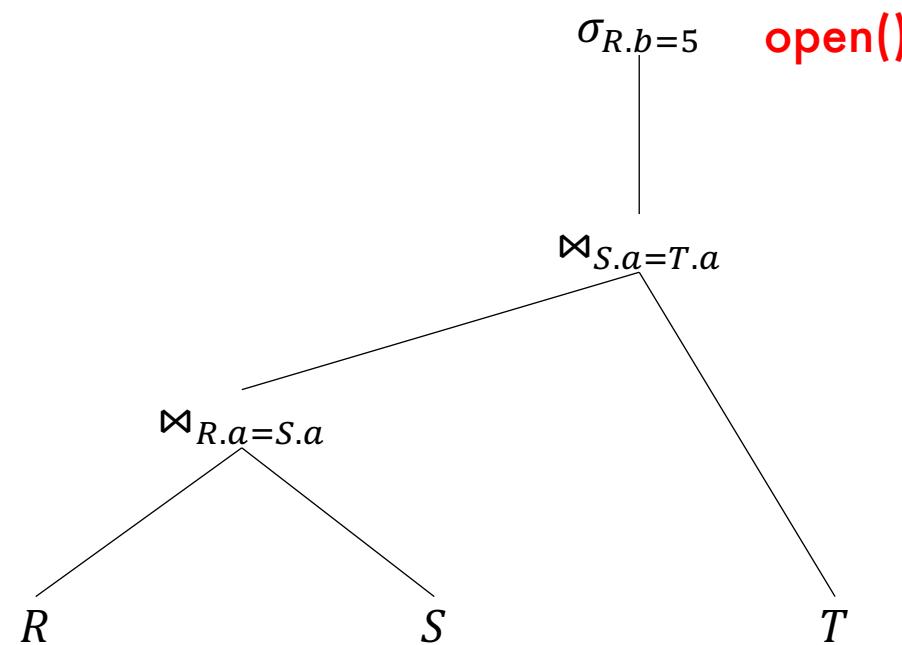
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



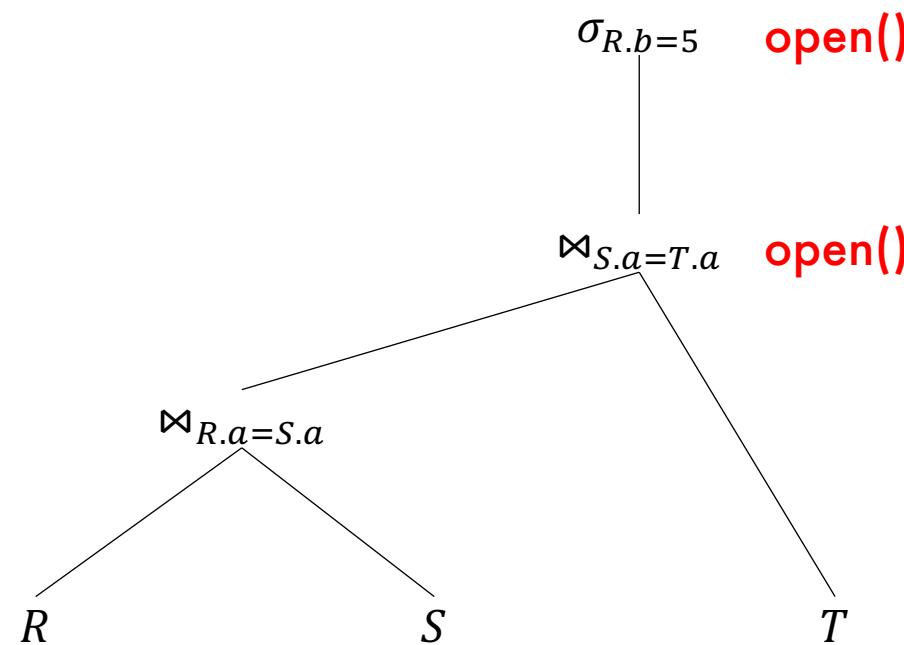
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



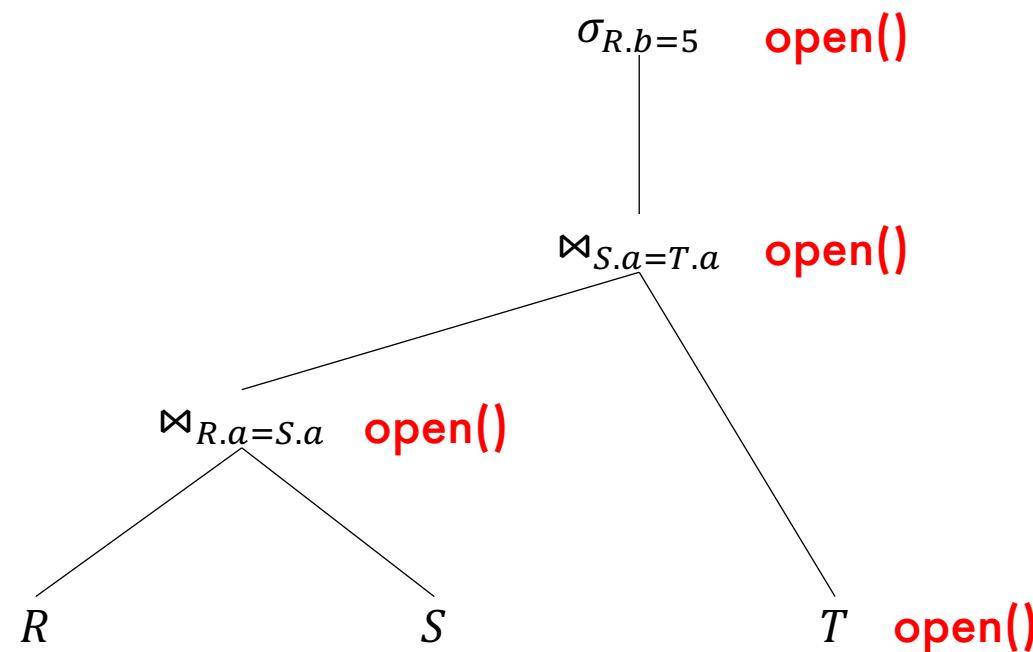
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



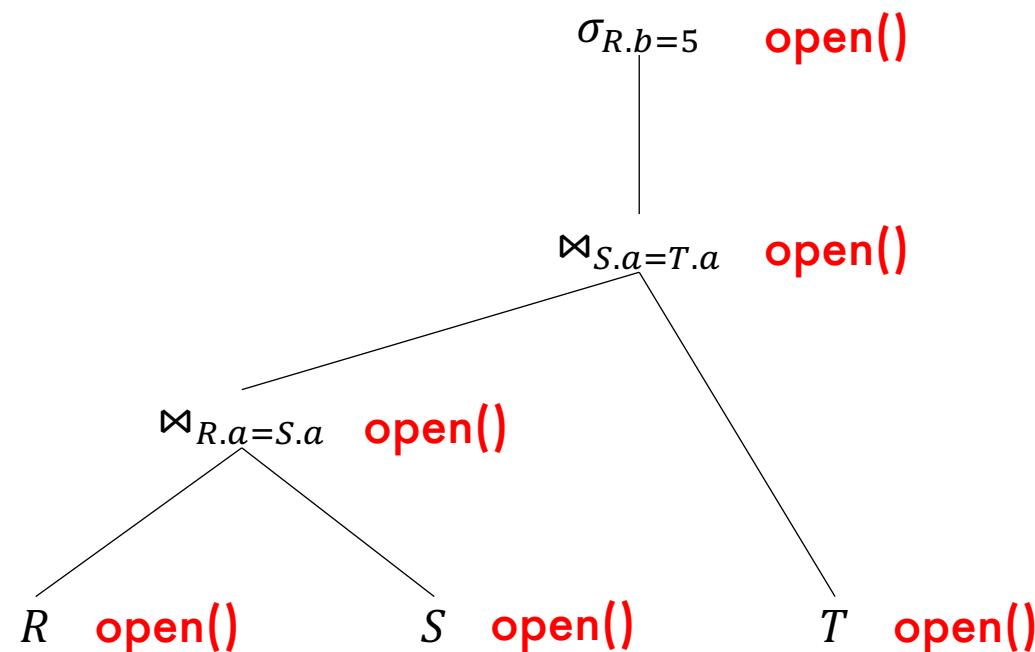
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



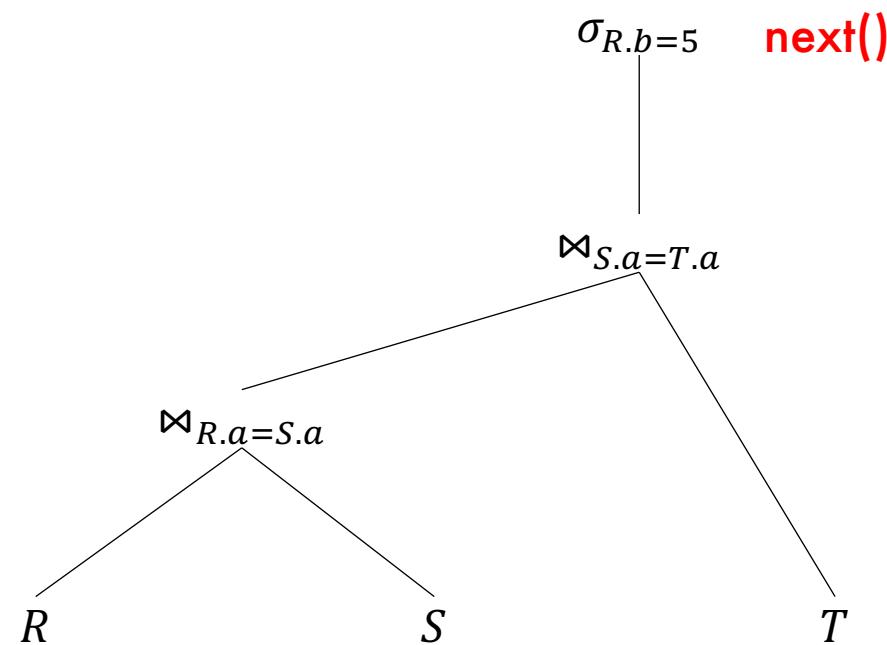
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



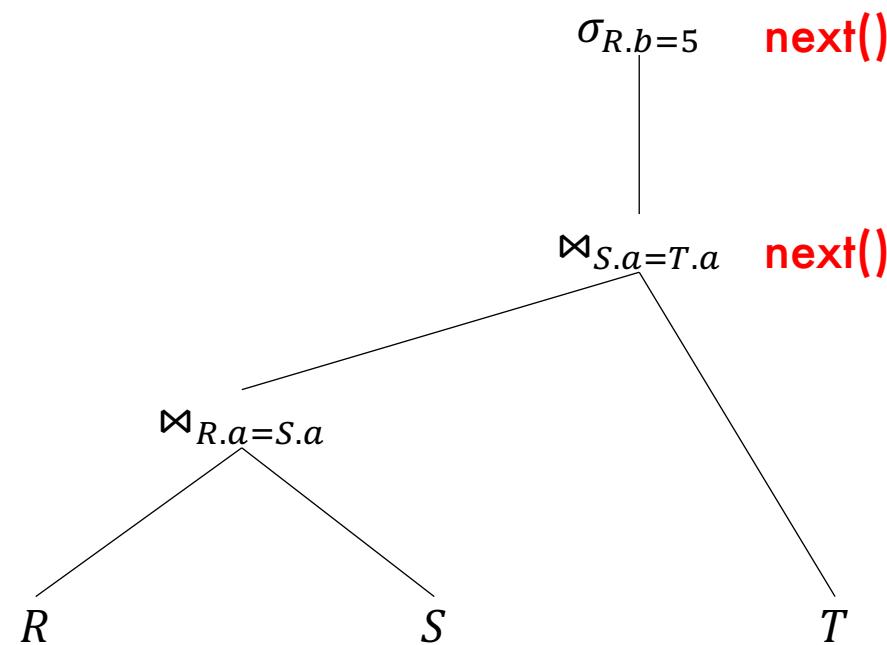
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



Pipelined Execution Example

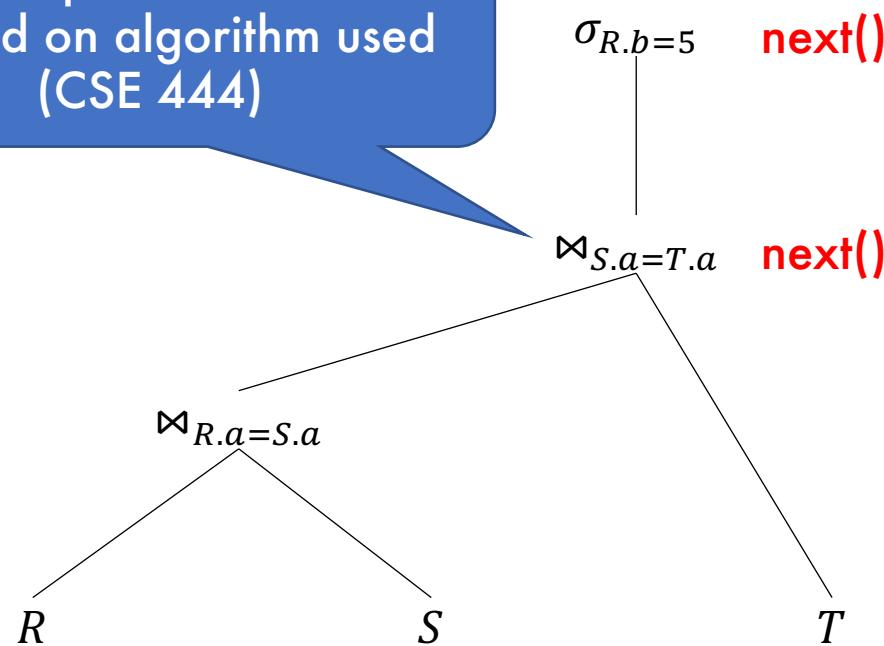
- Iterator interface of RA operators (**Volcano Iterator Model**)



Pipelined Execution Example

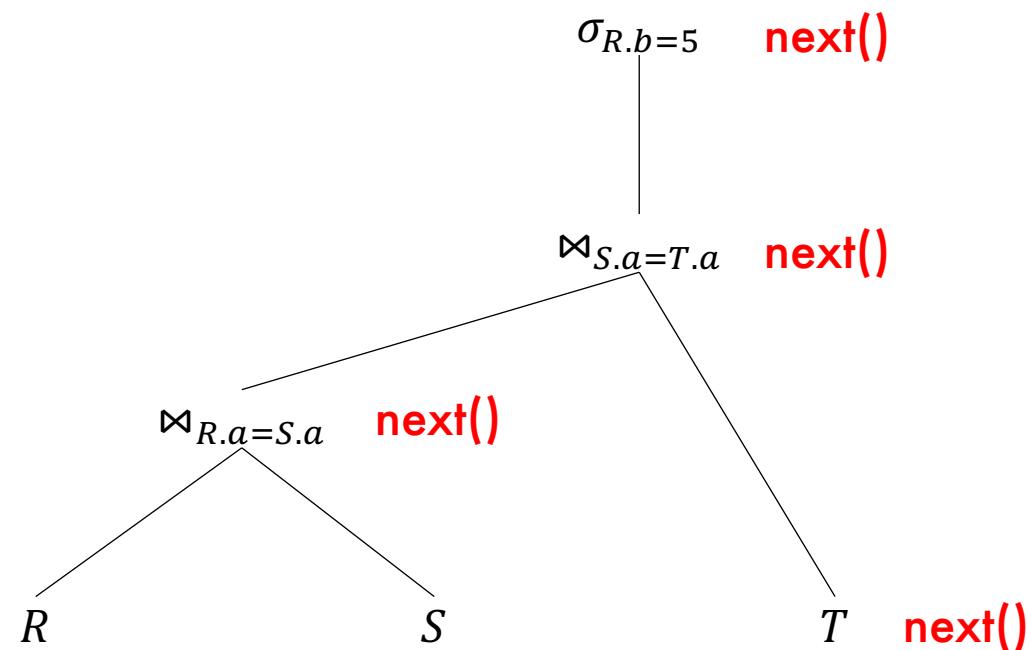
- Iterator interface of RA operators (**Volcano Iterator Model**)

next() implementation will depend on algorithm used
(CSE 444)



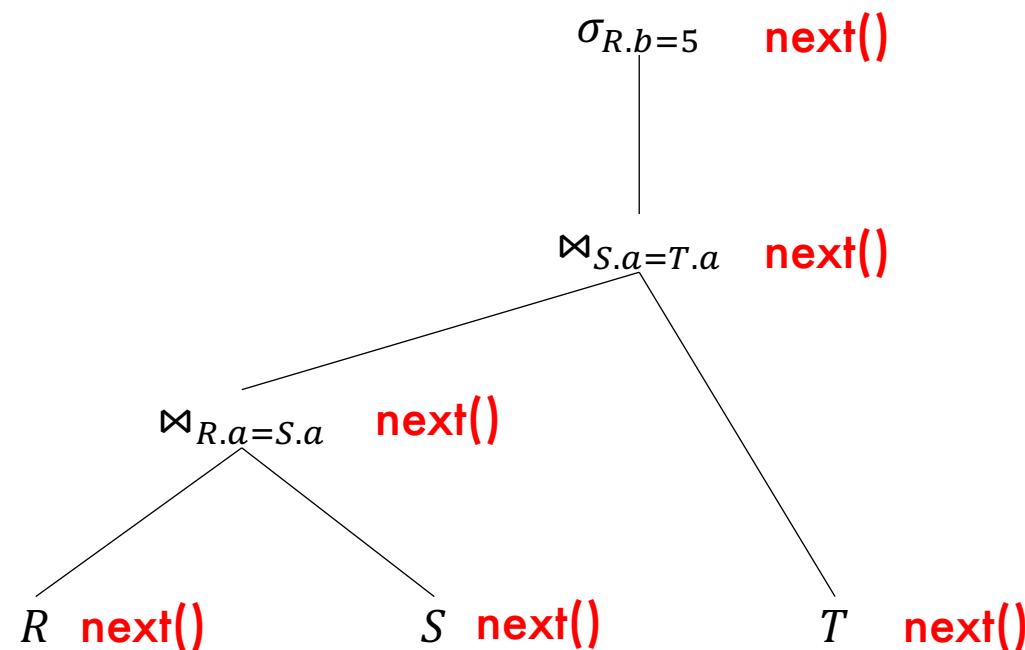
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



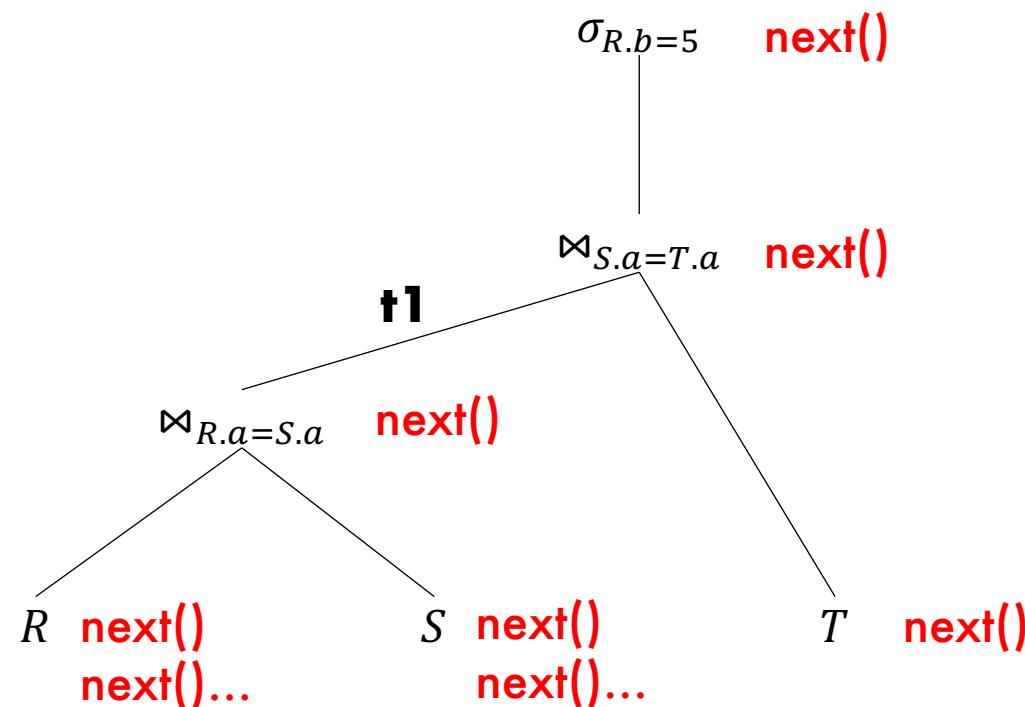
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



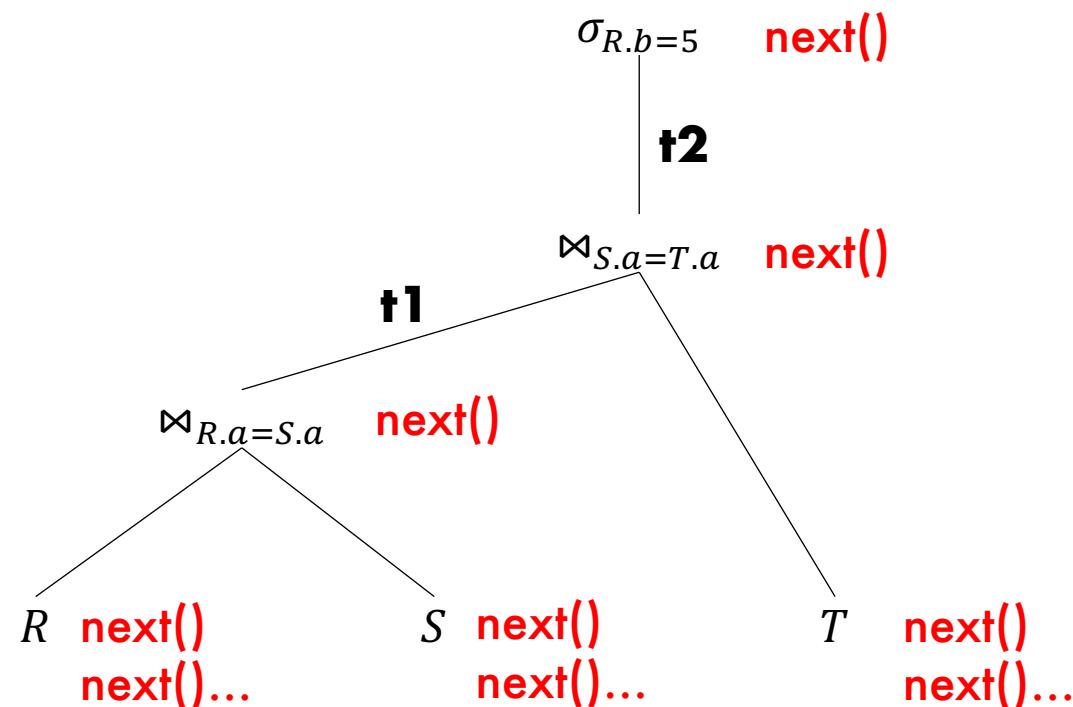
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



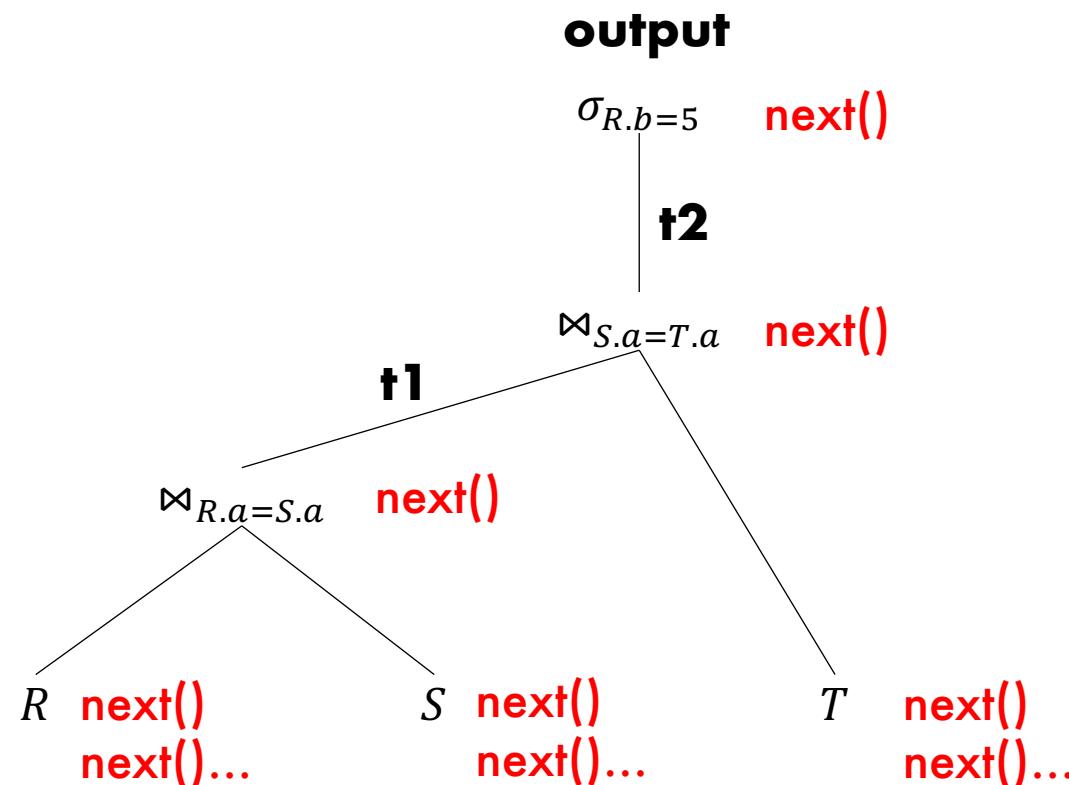
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



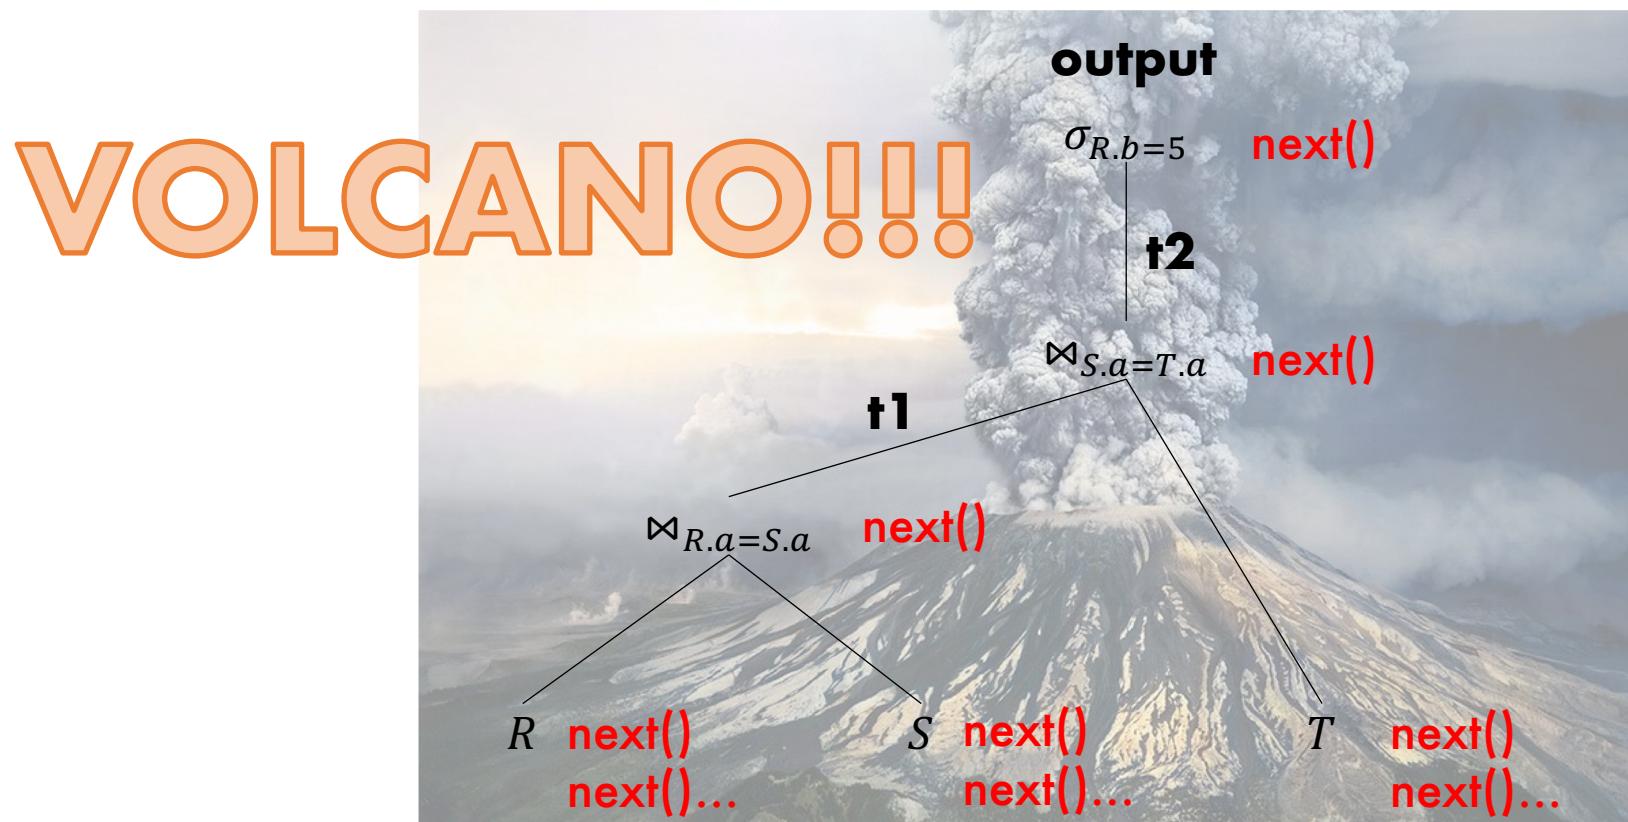
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



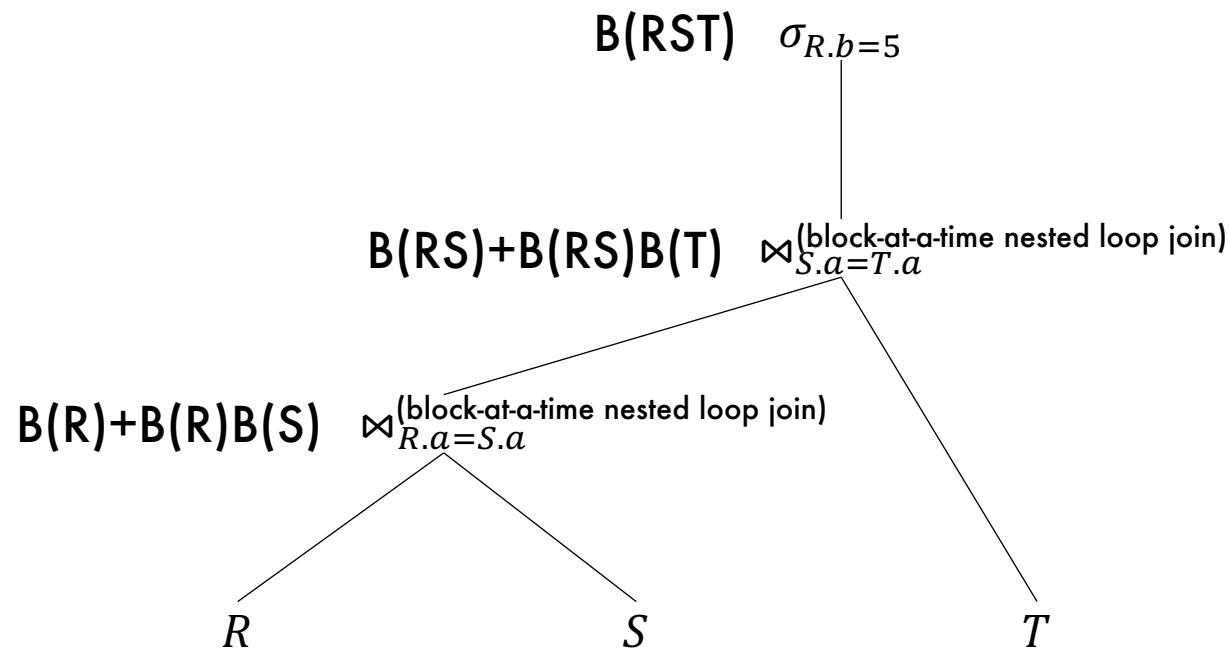
Pipelined Execution Example

- Iterator interface of RA operators (**Volcano Iterator Model**)



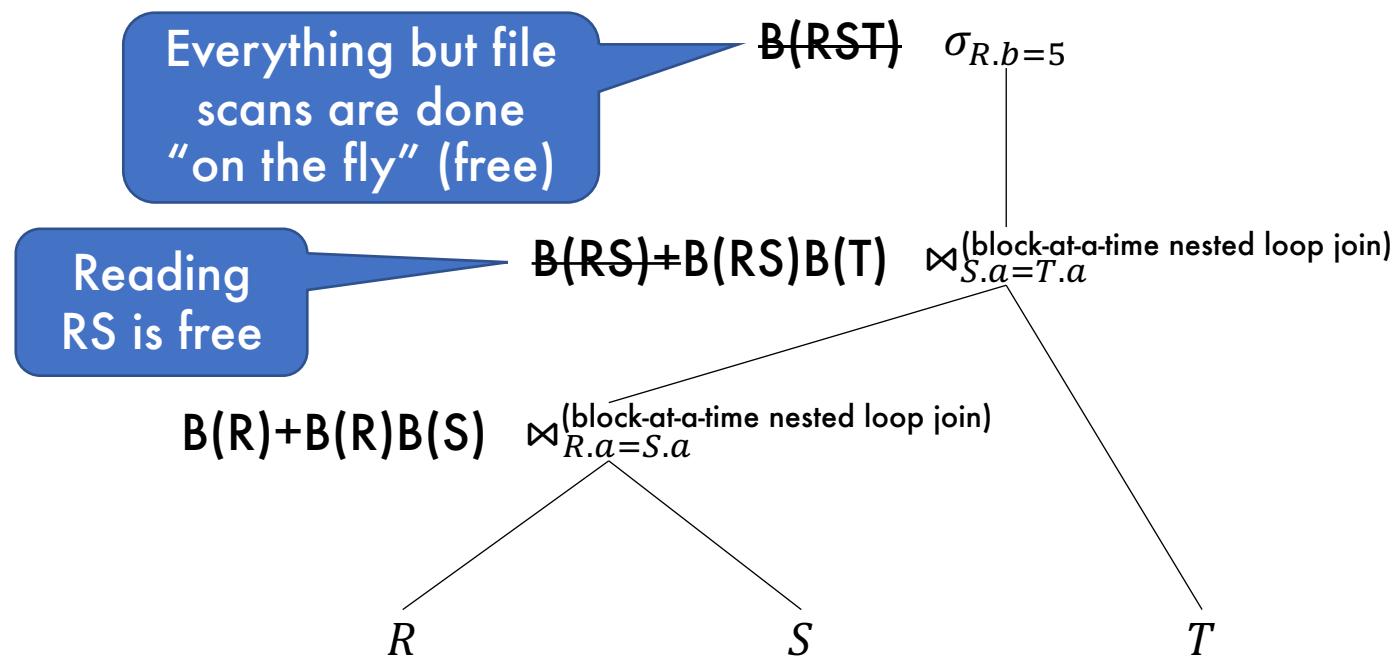
Full Plan Cost Estimation

- Estimated IO cost for a plan can be lowered under pipelined execution
- Generated tuples are in main memory so future access is “free”



Full Plan Cost Estimation

- Estimated IO cost for a plan can be lowered under pipelined execution
- Generated tuples are in main memory so future access is “free”



Full Plan Cost Estimation

- Estimated IO cost for a plan can be lowered under pipelined execution
- Generated tuples are in main memory so future access is “free”

