

Introduction to Data Management

Serializability

Paul G. Allen School of Computer Science and Engineering
University of Washington, Seattle

Recap: Leveraging Indexes

- Often for applications, workloads can be well described
 - HW 4 Flights application
 - Search method → query on city name values
 - Data visualization software (e.g. Tableau)
 - 2D plot → query on graph axis bounds
- **Create indexes to match expected query workload**

Transactions

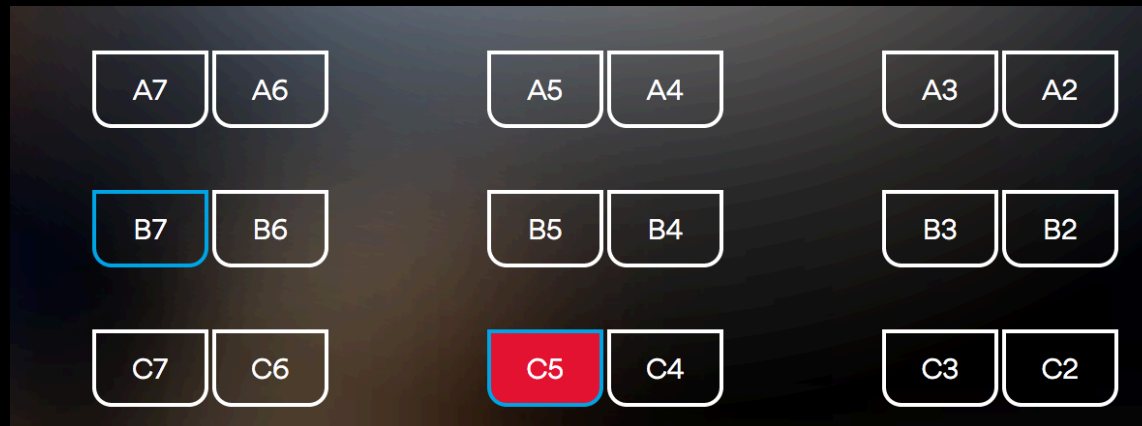

How do we support multiple people using a database at the same time?

- Multiple application users
- Multiple application programmers
- Multiple analysts
- Imagine a world where each person had to wait in line to use your database 😞

Common Concurrency Control Problems

- Non-Atomic Operations
- Lost Update
- Dirty/Inconsistent Read
- Unrepeatable Read
- Phantom Read

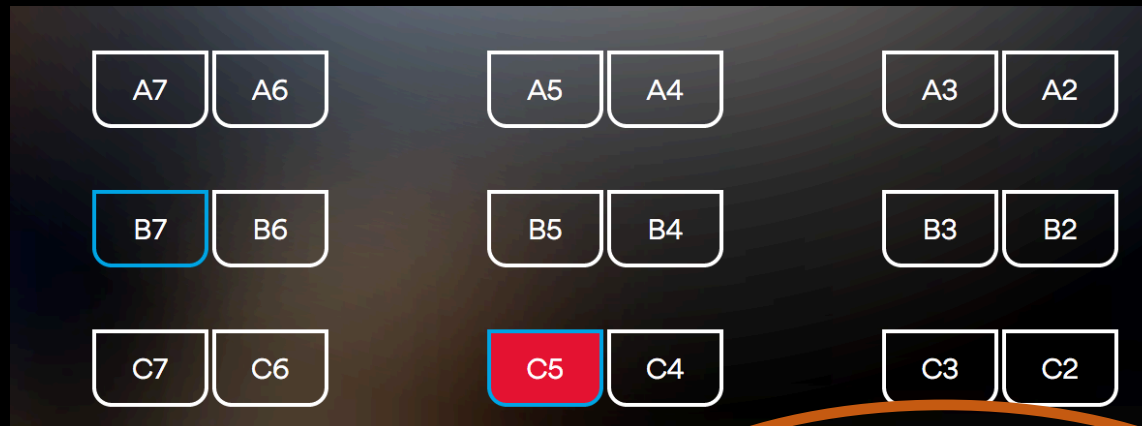
Non-Atomic Operations




Confirm Purchase

6:23 left 

Non-Atomic Operations



Confirm Purchase



Lost Update

- Write-Write (WW) conflict
- Consolidation scenario:

Account 1 = 100, Account 2 = 100

User 1 wants to pool
money into account 1

Set account 1 = 200

Set account 2 = 0

User 2 wants to pool
money into account 2

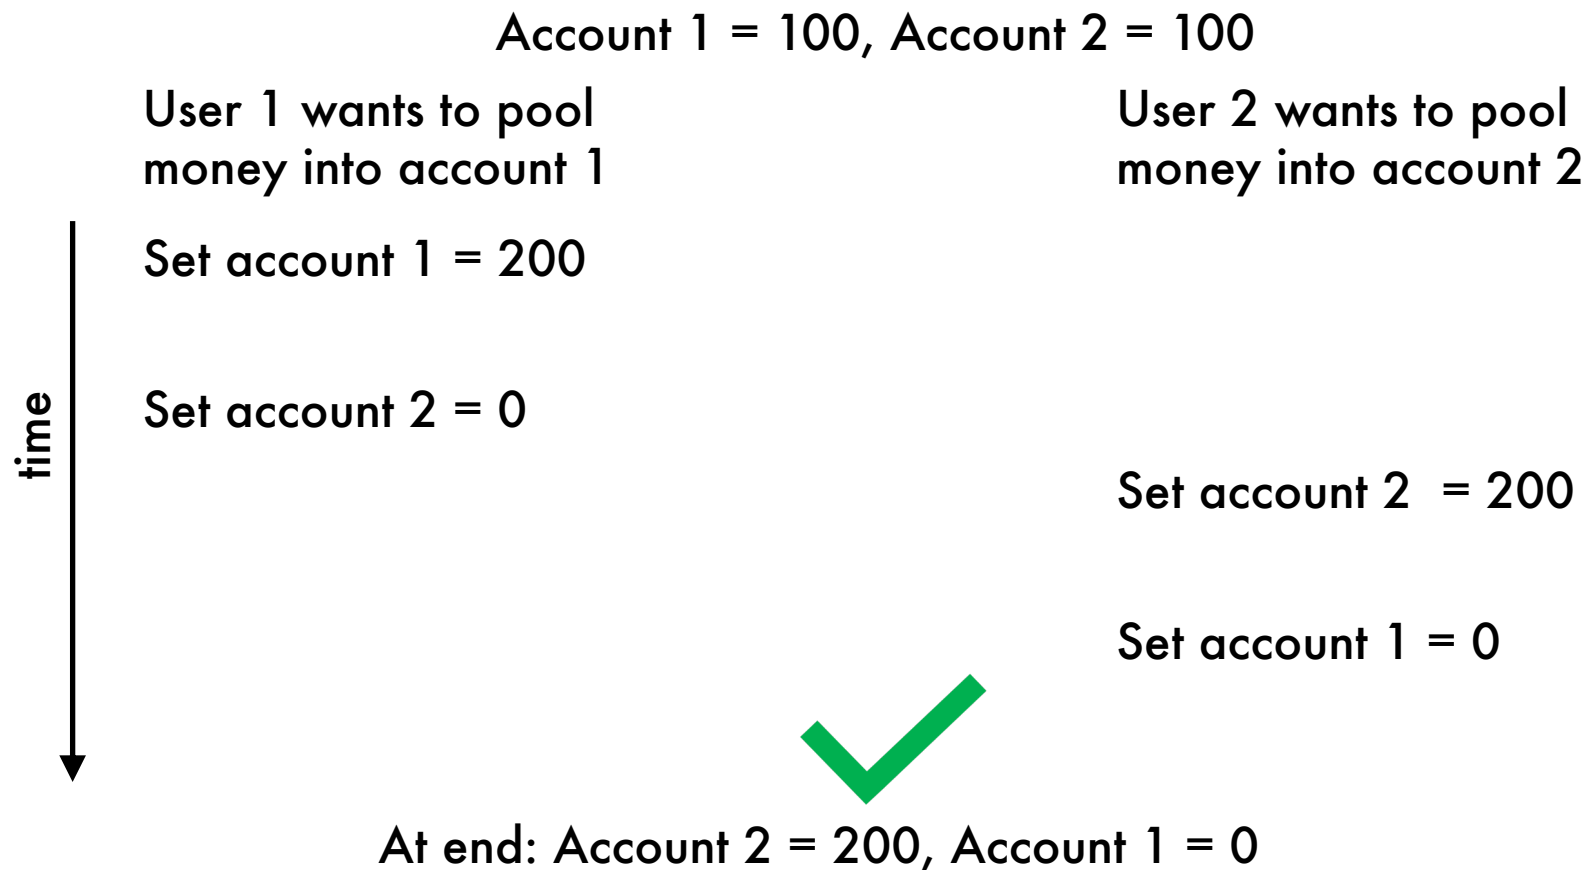
Set account 2 = 200

Set account 1 = 0

time
↓

Lost Update

- Write-Write (WW) conflict
- Consolidation scenario:



Lost Update

- Write-Write (WW) conflict
- Consolidation scenario:

Account 1 = 100, Account 2 = 100

User 1 wants to pool
money into account 1

Set account 1 = 200

Set account 2 = 0

User 2 wants to pool
money into account 2

Set account 2 = 200

Set account 1 = 0

Lost Update

- Write-Write (WW) conflict
- Consolidation scenario:

Account 1 = 100, Account 2 = 100

User 1 wants to pool
money into account 1

Set account 1 = 200

Set account 2 = 0

User 2 wants to pool
money into account 2

Set account 2 = 200

Set account 1 = 0



At end: Account 2 = 0, Account 1 = 0

Dirty/Inconsistent Read

- Write-Read (WR) conflict
- Budget management scenario:

Manager wants to
balance project budgets


-\$10mil from project A

+\$7mil to project B

+\$3mil to project C

CEO wants to check
company balance

`SELECT SUM(money) ...`



Dirty/Inconsistent Read

- Write-Read (WR) conflict
- Budget management scenario:


Manager wants to
balance project budgets

CEO wants to check
company balance

-\$10mil from project A

+\$7mil to project B

+\$3mil to project C



`SELECT SUM(money) ...`

Dirty/Inconsistent Read

- Write-Read (WR) conflict
- Budget management scenario:

Manager wants to
balance project budgets

CEO wants to check
company balance

-\$10mil from project A

+\$7mil to project B

+\$3mil to project C

SELECT SUM(money) ...



Unrepeatable Read

- Read-Write (RW) conflict
- Asset checking scenario:

Accountant wants to
check company assets

```
SELECT inventory  
FROM Products  
WHERE pid = 1
```

```
SELECT inventory*price  
FROM Products  
WHERE pid = 1
```

Application is automatically
updating inventories

```
UPDATE Products  
SET inventory = 0  
WHERE pid = 1
```

Unrepeatable Read

- Read-Write (RW) conflict
- Asset checking scenario:

Accountant wants to
check company assets

```
SELECT inventory  
FROM Products  
WHERE pid = 1
```

```
SELECT inventory*price  
FROM Products  
WHERE pid = 1
```

Application is automatically
updating inventories

```
UPDATE Products  
SET inventory = 0  
WHERE pid = 1
```

Might get a value that doesn't
correspond to previous read!

Phantom Read

- Same read has more rows
- Asset checking scenario:

Accountant wants to
check company assets

```
SELECT *  
FROM products  
WHERE price < 10.00
```

```
SELECT *  
FROM products  
WHERE price < 20.00
```

Warehouse catalogs new
products

```
INSERT INTO Products  
VALUES ('nuts', 10, 8.99)
```


Phantom Read

- Same read has more rows
- Asset checking scenario:

Accountant wants to
check company assets

```
SELECT *  
FROM products  
WHERE price < 10.00
```

```
SELECT *  
FROM products  
WHERE price < 20.00
```

Warehouse catalogs new
products

```
INSERT INTO Products  
VALUES ('nuts', 10, 8.99)
```



Gets a row that should
have been in the last read!

ACID

- Atomic
- Consistent
- Isolated
- Durable
- Ideally a DBMS follows these principles, but sacrificing good behavior for performance gains is common
- Definitely needs to follow these principles if you are dealing with \$\$\$

ACID

- Atomic
 - Consistent
 - Isolated
 - Durable
- Ideally a DBMS follows these principles, but sacrificing good behavior for performance gains is common
 - Definitely needs to follow these principles if you are dealing with \$\$\$

Atomic

- Operation encapsulation
- An operation is atomic if everything works or nothing happens

Consistent

- Integrity constraints and application specification
- Operations assume a valid database state and end in a valid database state

Isolated

- Concurrency management
- Isolated behavior is as if an operation ran as if it was the only one running

Durable

- Crash recovery
- CSE 444 topic, not discussed in this class

Transactions

- An application function may involve multiple different operations
- We want to make sure the parts of an operation execute properly **together as if it were a single action**
- We say that a transaction is one of these groups of executions
 - DBMS usually automatically treats each SQL statement as its own transaction unless otherwise specified

BEGIN TRANSACTION
[SQL Statements]
COMMIT – finalizes execution

BEGIN TRANSACTION
[SQL Statements]
ROLLBACK – undo everything

Concurrency Control Problems

- We've (sorta) solved the atomicity problem!
- **DBMS concurrency control is all based on specification**
- Merely specifying what your transactions are is good enough for the DBMS to take care of it as a single unit

Transaction Modeling

- Logical perspective → a database is a set of sets/bags of tuples
- Design perspective → a database is a schema that models information
- Physical perspective → a database is a catalog of organized files
- Transaction perspective → a database is a **collection of elements** that can be **written to** or **read from**
 - Element granularity can vary depending on DBMS and/or user specification
 - Transactions are sequence of element reads and/or writes

Schedules

- Transactions are sequence of element reads and/or writes
 - $R_i(A) \rightarrow$ **read** element A
 - $W_i(A) \rightarrow$ **update** element A
 - $I_i(A) \rightarrow$ **insert** an element A
 - $D_i(A) \rightarrow$ **delete** an element A
- Schedules are a sequence of interleaved actions from all transactions

Serial Schedules

- A **serial schedule** is a schedule where each transaction would be executed in some order
- A **serializable schedule** is a schedule where transaction reads and writes would be executed as if it were executed in serial order
 - If the schedule were executed and you were given a before and after, you would not be able to tell if there was interleaving

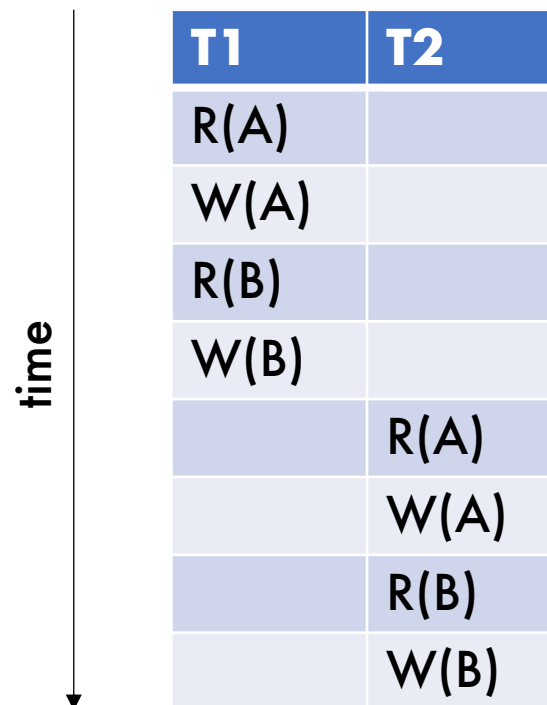
Transaction Schedule

T1	T2
R(A)	R(A)
W(A)	W(A)
R(B)	R(B)
W(B)	W(B)

Serial Schedule Example

- T1 then T2

$R_1(A), W_1(A), R_1(B), W_1(B), R_2(A), W_2(A), R_2(B), W_2(B)$



T1	T2
R(A)	
W(A)	
R(B)	
W(B)	
	R(A)
	W(A)
	R(B)
	W(B)

Serial Schedule Example

- T2 then T1

$R_2(A), W_2(A), R_2(B), W_2(B), R_1(A), W_1(A), R_1(B), W_1(B)$

T1	T2
	R(A)
	W(A)
	R(B)
	W(B)
R(A)	
W(A)	
R(B)	
W(B)	

Serializable Schedule

- Serializable to T1 then T2

$R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), R_2(B), W_2(B)$

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

Serializable Schedule

- Serializable to T1 then T2

$R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), R_2(B), W_2(B)$

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

Looks like T2
finished after T1
for each element

Serializable Schedule

- Not serializable to either order

$R_1(A), W_1(A), R_2(A), W_2(A), R_2(B), W_2(B), R_1(B), W_1(B)$

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
	R(B)
	W(B)
R(B)	
W(B)	

Serializable Schedule

- Not serializable to either order

$R_1(A), W_1(A), R_2(A), W_2(A), R_2(B), W_2(B), R_1(B), W_1(B)$

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
	R(B)
	W(B)
R(B)	
W(B)	

Looks like T1 finished last looking at B

Looks like T2 finished last looking at A

Checking Serializability

- How does the DBMS tell if some schedule is serializable?

Conflicts

- Most application concurrency problems are describable by conflicts
- Lost Update → Write-Write (WW) conflict
- Dirty Read → Write-Read (WR) conflict
- Unrepeatable Read → Read-Write (RW) conflict
- Phantom Read
 - We'll talk about this later...

Individual conflicts aren't "bad"!
Interleaving of conflicts can lead to trouble.

Types of Conflicts

- Changing the order of things in conflict will cause program behavior to behave badly
- **Intra-transaction conflicts**
 - Operations within a transaction cannot be swapped (you would be literally changing the program)
- **Inter-transaction conflicts**
 - WW conflicts $\rightarrow W1(X), W2(X)$
 - WR conflicts $\rightarrow W1(X), R2(X)$
 - RW conflicts $\rightarrow R1(X), W2(X)$

Conflict Serializability

- Showing program serializability is hard
 - Depending on the values read/written, many “bad” schedules can be serializable
 - Needs lots of extra information besides R, W, I, D
- Observation: Enforce something something simpler but stronger than serializability
- **Conflict serializability implies serializability**
- Serializability does not imply conflict serializability

Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
	R(A)
R(B)	
	W(A)
W(B)	
	R(B)
	W(B)

Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
R(B)	
	R(A)
	W(A)
W(B)	
	R(B)
	W(B)

Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
R(B)	
	R(A)
W(B)	
	W(A)
	R(B)
	W(B)

Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
R(B)	
W(B)	
	R(A)
	W(A)
	R(B)
	W(B)

Example

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
	R(B)
	W(B)
R(B)	
W(B)	

Example

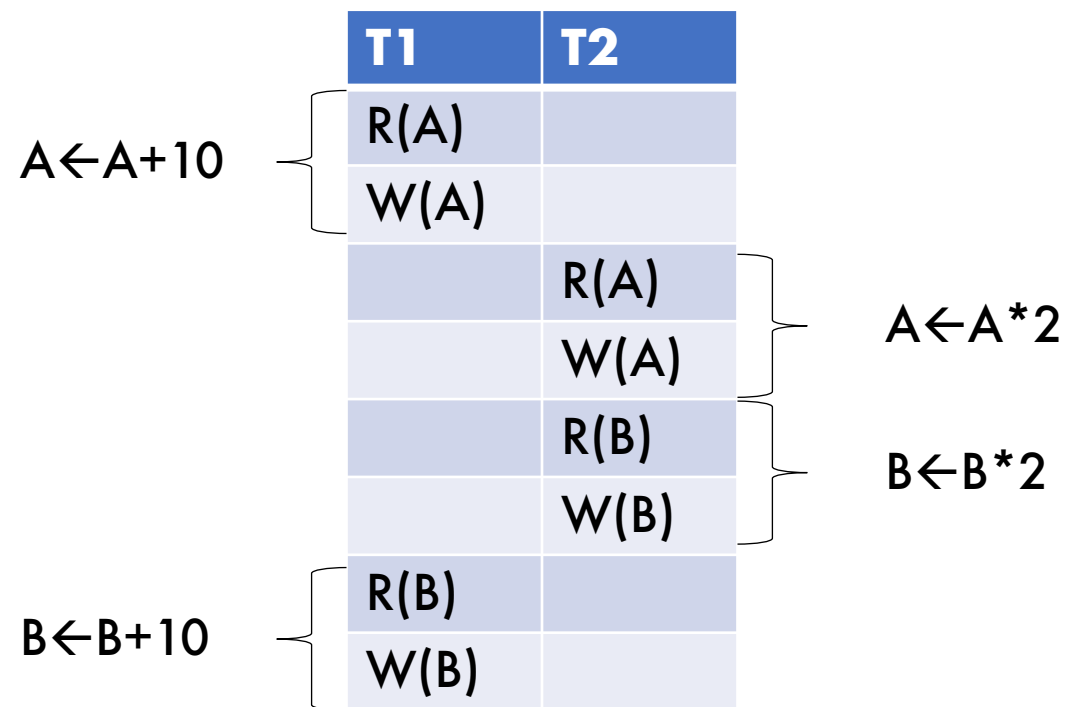
T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
	R(B)
R(B)	
	W(B)
W(B)	



Conflict rule broken!

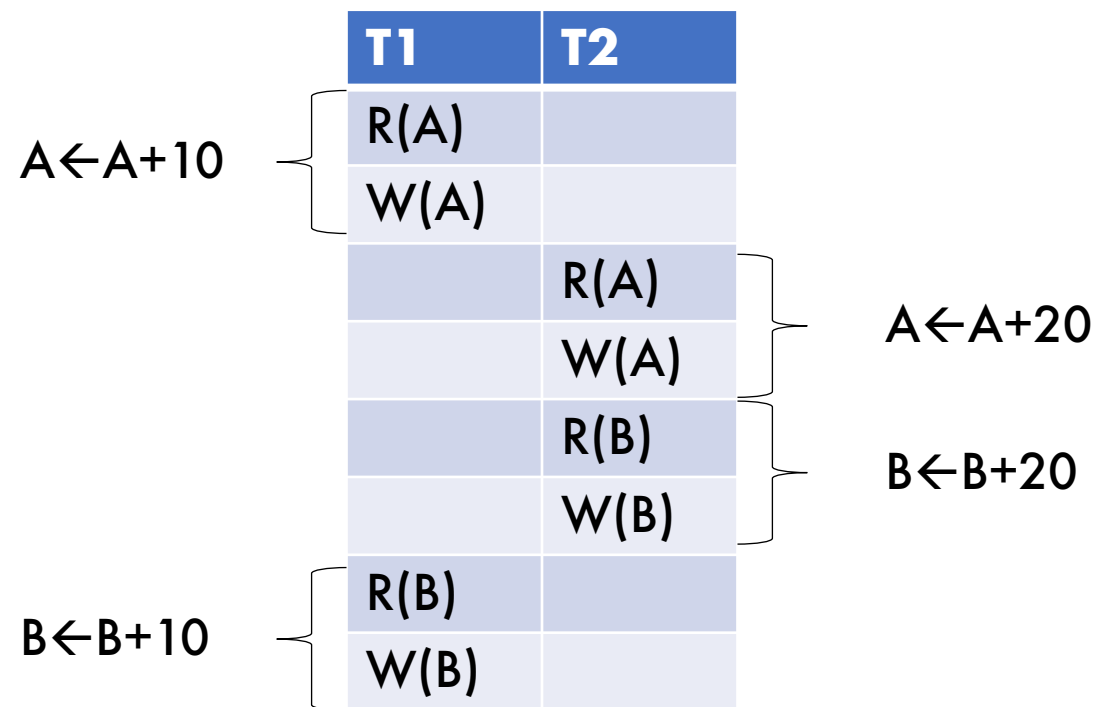
Serializable vs Conflict Serializable

Not serializable nor conflict serializable



Serializable vs Conflict Serializable

Serializable but not conflict serializable



Enforcing Conflict Serializability

- We only care if some conflict rule would be broken (no need to micromanage)
- Need an effective algorithm

Method:

- Model each transaction as a node
- Model a inter-transaction conflict as a directed edge
- If the resulting graph is a DAG then there is a serial order
- Conflict serializability enforcement turns into the graph cycle detection problem

Testing for Conflict-Serializability

Precedence graph:

- A node for each transaction T_i ,
- An edge from T_i to T_j whenever an action in T_i conflicts with, and comes before an action in T_j
- No edge for actions in the same transaction

Theorem:

The schedule is serializable iff the precedence graph is acyclic

Testing for Conflict-Serializability

Important:

Always draw the full graph, unless ONLY asked if (yes or no) the schedule is conflict serializable

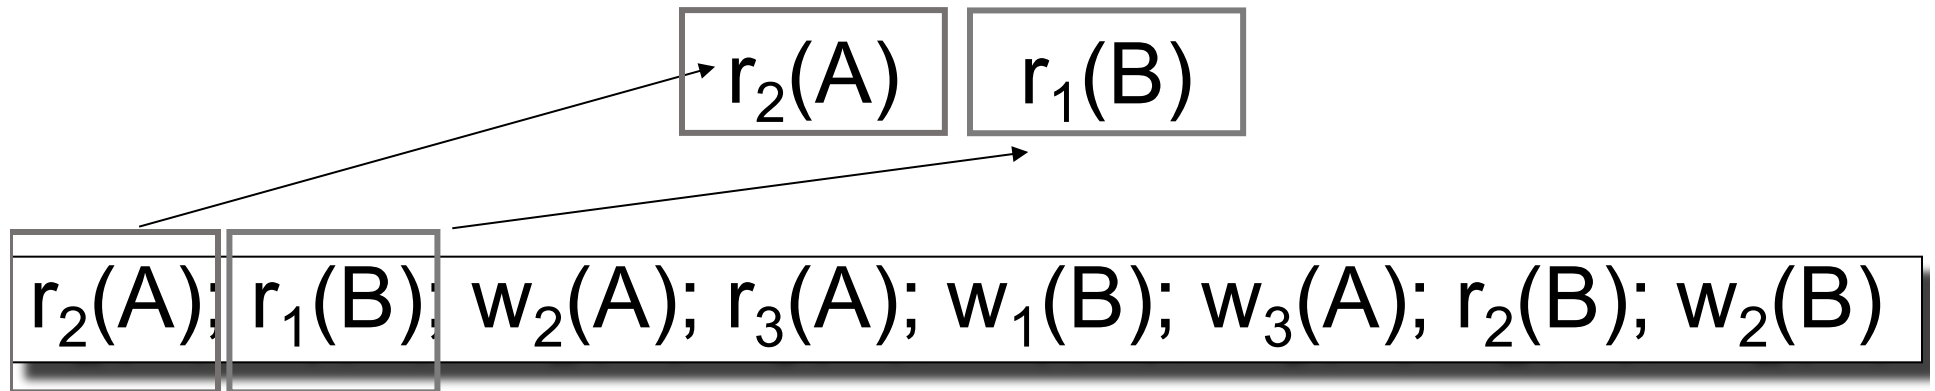
Example 1

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

①

②

③



①

②

③

$r_2(A)$

$r_1(B)$

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

①

②

③

$r_2(A)$

$r_1(B)$

No edge because
no conflict ($A \neq B$)

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

1

2

3

$r_2(A)$

$w_2(A)$

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

①

②

③

$r_2(A)$

$w_2(A)$

No edge because
same txn (2)

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

1

2

3

$r_2(A)$

$r_3(A)$

?

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

1

2

3

$r_2(A)$

$w_1(B)$

?

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

①

②

③

$r_2(A)$

$w_3(A)$

?

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

1

2

3

$r_2(A)$

$w_3(A)$

Edge! Conflict from
T2 to T3

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

1

2

3

$r_2(A)$

$w_3(A)$

Edge! Conflict from
T2 to T3

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

Edge label is optional

①

②

A

③

$r_2(A)$

$r_2(B)$

?

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

And so on until compared every pair of actions...


①

②

→ ③

Example 1


$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$



More edges, but repeats of the same directed edge not necessary

Example 1

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$



This schedule is **conflict-serializable**

Example 2


$r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B)$

①


②

③

Example 2



Example 2



This schedule is NOT conflict-serializable