

# Introduction to Data Management

## Isolation Levels

Jonathan Leang

Paul G. Allen School of Computer Science and Engineering  
University of Washington, Seattle

# Recap

- Schedules under 2PL are conflict serializable
  - Locking phase → unlocking phase
- Conflict serializable schedules follow the isolation principle of ACID
  - No dirty read (WR)
  - No unrepeatable read (RW)
  - No lost update (WW)
- Schedules under strict 2PL additionally provide recoverability
  - Locking phase → unlock with commit or rollback


# Outline

- Shared/Exclusive locks
- Isolation levels
- Implementing transactions in practice

# Practicality of Binary Locks

- Binary Locks → full control or no control
- Leads to excessive deadlocking

# Thrashing



# Shared/Exclusive Locks

- Observation: Reads don't conflict with each other
- Simple 3-tier lock hierarchy:
  - Exclusive/Write Lock  $\rightarrow \mathbf{X_i(A)}$ 
    - Full control
    - No other locks may exist
  - Shared/Read Lock  $\rightarrow \mathbf{S_i(A)}$ 
    - Shared control
    - May exist with other shared locks
  - Unlocked

# Shared/Exclusive Locks

Requested Lock	unlocked	S	X
S	Yes	Yes	No
X	Yes	No	No

# Practicality of Serializability

- **Easy to reason about**
  - Application programming is easier under serializability assumptions
- **Expensive to use**
  - Slow
  - Resource intensive
- **Applications often don't need serializability**
  - Application functionality may not depend on serializability
  - Financial/User experience cost is low enough for tradeoff considerations



# Isolation Levels

- **SET TRANSACTION ISOLATION LEVEL ...**
  - **READ UNCOMMITTED**
  - **READ COMMITTED**
  - **REPEATABLE READ**
  - **SERIALIZABLE**
  - **SNAPSHOT ISOLATION**
  - ...
- Default isolation level and configurability depends on the DBMS (read the docs)

# READ UNCOMMITTED

- Writes → Strict 2PL write locks
- Reads → No locks needed
- Dirty reads are possible

T1	T2
X(A) W(A)	
	R(A)
	COMMIT
ABORT U(A)	

# READ COMMITTED

- Writes → Strict 2PL write locks
- Reads → Short-duration read locks
  - Acquire lock before reading and release lock after (not 2PL)
- **Dirty reads are prevented**

T1	T2
X(A) W(A)	
	R(A)
	COMMIT
ABORT U(A)	



T1	T2
X(A) W(A)	
	S(A) blocked...
ABORT U(A)	...granted S(A)
	R(A)
	COMMIT U(A)

# READ COMMITTED

- **Unrepeatable reads are possible**

T1	T2
X(A) blocked...	S(A)
	<b>R(A)</b>
...granted X(A)	U(A)
<b>W(A)</b>	S(A) blocked...
COMMIT U(A)	...granted S(A)
	<b>R(A)</b>
	COMMIT U(A)

# REPEATABLE READ

- Writes → Strict 2PL write locks
- Reads → Strict 2PL read locks
- **Unrepeatable reads are prevented**

T1	T2
X(A) blocked...	S(A)
	<b>R(A)</b>
...granted X(A)	<del>U(A)</del>
<b>W(A)</b>	S(A) blocked...
COMMIT U(A)	...granted S(A)
	<b>R(A)</b>
	COMMIT U(A)

T1	T2
X(A) blocked...	S(A)
	<b>R(A)</b>
	<b>R(A)</b>
...granted X(A)	COMMIT U(A)
<b>W(A)</b>	
COMMIT U(A)	

# REPEATABLE READ

- Writes → Strict 2PL write locks
- Reads → Strict 2PL read locks
- **Phantom reads are possible...**

# Phantom Reads

- Conflict serializability implying serializability assumes a **static database**
  - Conflicts only matter for the same element
  - Inserting a new element (tuple-level granularity) means that the conflict model no longer is able to encapsulate it

	T1	T2	
SELECT * FROM Table;	R(A)		INSERT INTO Table VALUES (C...);
	R(B)		
		I(C)	
SELECT * FROM Table;	R(A)		
	R(B)		
	R(C)		

# Phantom Reads

- **Dynamic database** serializability needs either:
  - Table locking (prevent insertions) or
  - Predicate locking (lock based on query filters)



# SERIALIZABLE

- Write Lock → Strict 2PL
- Read Lock → Strict 2PL
- Plus predicate locks and/or table locks

# Isolation Level Summary

READ UNCOMMITTED → Dirty Read

READ COMMITTED → Unrepeatable Read

REPEATABLE READ → Phantom Read

SERIALIZABLE → No Anomalies

# Applying Transaction Logic

- Applications generally need to
  - **Check/Set isolation levels**
  - **Specify operations as transactions**
- Common mistakes/misconceptions:
  - You do not need to implement locking. The DBMS takes care of it.
  - You must **close all explicit transactions** with COMMIT or ROLLBACK. Not doing so will cause the application to hang (wait due to unfinished locking).

# Transaction Setup

```
conn.setTransactionIsolation(Connection.TRANSACTION_SERIALIZABLE);  
conn.setTransactionIsolation(Connection.TRANSACTION_READ_UNCOMMITTED);  
conn.setTransactionIsolation(Connection.TRANSACTION_READ_COMMITTED);  
conn.setTransactionIsolation(Connection.TRANSACTION_REPEATABLE_READ);  
conn.setAutoCommit(true);  
conn.setAutoCommit(false);
```

# DB Transaction Programming in Java

```
try {
    // Each Instance hold a unique conn
    PreparedStatement q = conn.prepareStatement("SELECT ...");
    PreparedStatement i = conn.prepareStatement("INSERT ...");
    // Make sure the statements don't execute separately
    conn.setAutoCommit(false);
    conn.execute("BEGIN TRANSACTION;");
    ResultSet rs = q.executeQuery();
    while(rs.next()) { ... } // Read out tuples from the ResultSet
    i.executeUpdate();
    conn.execute("COMMIT;");
    conn.setAutoCommit(true);
    return "success"
} catch (SQLException ex) {
    try {
        conn.execute("ROLLBACK;");
        conn.setAutoCommit(true);
        return "failed"
    } catch (SQLException e) {
        return "failed";
    }
}
```