

Introduction to Data Management

Semi-structured Data

Paul G. Allen School of Computer Science and Engineering
University of Washington, Seattle

Announcements

- HW 6:
 - due Sunday night
 - Up to 2 late days allowed
- HW 7: (out today, much less complex than normal HW)
 - due the Saturday after final exam
 - NO LATE DAYS ALLOWED
- Schedule for next week:
 - Today: Semi-structured data
 - Monday: SQL++ and HW 7
 - Wednesday: Misc. topics
 - Thursday section: exam review session
 - Friday: **Final exam**

Outline

- AsterixDB as a case study of Document Store
 - Semi-structured data model in JSON
 - Introducing AsterixDB and SQL++



Outline

- AsterixDB as a case study of Document Store
 - **Semi-structured data model in JSON**
 - Introducing AsterixDB and SQL++



What is a "document" anyways?

- Loose terminology
- Any "parsable" file qualifies
 - Ex: MongoDB can handle CSV files

Semi-Structured Documents

- Some notion of **tagging** to mark down semantics
- Examples:
 - XML
 - Protobuf
 - JSON


```
<?xml version="1.0" encoding="UTF-8"?>
<customers>
  <customer>
    <customer_id>1</customer_id>
    <first_name>John</first_name>
    <last_name>Doe</last_name>
    <email>john.doe@example.com</email>
  </customer>
  <customer>
    <customer_id>2</customer_id>
    <first_name>Sam</first_name>
    <last_name>Smith</last_name>
    <email>sam.smith@example.com</email>
  </customer>
  <customer>
    <customer_id>3</customer_id>
    <first_name>Jane</first_name>
    <last_name>Doe</last_name>
    <email>jane.doe@example.com</email>
  </customer>
</customers>
```



Tags surround the respective data

Semi-Structured Documents

- Some notion of **tagging** to mark down semantics
- Examples:
 - XML
 - **Protobuf**
 - JSON



Able to record field number
and type but not name

Semi-Structured Documents

- Some notion of **tagging** to mark down semantics
- Examples:
 - XML
 - Protobuf
 - **JSON**

```
{
  "orders": [
    {
      "orderno": "748745375",
      "date": "June 30, 2088 1:54:23 AM",
      "trackingno": "TN0039291",
      "custid": "11045",
      "customer": [
        {
          "custid": "11045",
          "fname": "Sue",
          "lname": "Hatfield",
          "address": "1409 Silver Street",
          "city": "Ashland",
          "state": "NE",
          "zip": "68003"
        }
      ]
    }
  ]
}
```

Tags introduce the respective data

Semi-Structured Documents

- Some notion of **tagging** to mark down semantics
- Examples:
 - XML
 - Protobuf
 - **JSON**

Many applications have phased out XML in favor of JSON

```
{
  "orders": [
    {
      "orderno": "748745375",
      "date": "June 30, 2088 1:54:23 AM",
      "trackingno": "TN0039291",
      "custid": "11045",
      "customer": [
        {
          "custid": "11045",
          "fname": "Sue",
          "lname": "Hatfield",
          "address": "1409 Silver Street",
          "city": "Ashland",
          "state": "NE",
          "zip": "68003"
        }
      ]
    }
  ]
}
```

Tags introduce the respective data

Relational vs Semi-Structured Tradeoffs

▪ Relational Model

- Fixed schema
- Flat data

▪ Semi-Structured

- Self-described schema
- Tree-structured data

Relational vs Semi-Structured Tradeoffs

▪ Relational Model

- Fixed schema
- Flat data

▪ Semi-Structured

- Self-described schema
- Tree-structured data



Less well-defined / More flexible

Relational vs Semi-Structured Tradeoffs

▪ Relational Model

- Fixed schema
- Flat data

▪ Semi-Structured

- Self-described schema
- Tree-structured data



Less well-defined / More flexible

• Basic retrieval process:

1. Retrieve table
2. Run through rows
3. Return data

• Basic retrieval process:

1. Retrieve document
2. Parse document tree
3. Return data

Relational vs Semi-Structured Tradeoffs

▪ Relational Model

- Fixed schema
- Flat data

▪ Semi-Structured

- Self-described schema
- Tree-structured data



Less well-defined / More flexible

• Basic retrieval process:

1. Retrieve table
2. Run through rows
3. Return data

• Basic retrieval process:

1. Retrieve document
2. Parse document tree
3. Return data



Inefficient encoding / Easy exchange of data



- No database paradigm is "better" than another
- One-size does **not** fit all (M. Stonebraker)
 - Excellent article on data management in 21st century
 - http://cs.brown.edu/research/db/publications/fits_all.pdf
- Everything is getting mixed up anyways

JSON Standard – Rules of the Game

- JavaScript Object Notation (JSON)
 - "Lightweight text-based open standard designed for **human-readable** data interchange"

```
{
  "book": [
    {
      "id": "01",
      "language": "Java",
      "author": "H. Javeson",
      "year": 2015
    },
    {
      "author": "E. Sepp",
      "id": "07",
      "language": "C++",
      "edition": null,
      "sale": true
    }
  ]
}
```

JSON Standard – Rules of the Game

- JavaScript Object Notation (JSON)
 - "Lightweight text-based open standard designed for **human-readable** data interchange"

```
{
  "book": [
    {
      "id": "01",
      "language": "Java",
      "author": "H. Javeson",
      "year": 2015
    },
    {
      "author": "E. Sepp",
      "id": "07",
      "language": "C++",
      "edition": null,
      "sale": true
    }
  ]
}
```

Types

Primitives include:

- String (in quotes)
- Numeric (unquoted number)
- Boolean (unquoted true/false)
- Null (literally just null)

JSON Standard – Rules of the Game

- JavaScript Object Notation (JSON)
 - "Lightweight text-based open standard designed for **human-readable** data interchange"

```
{
  "book": [
    {
      "id": "01",
      "language": "Java",
      "author": "H. Javeson",
      "year": 2015
    },
    {
      "author": "E. Sepp",
      "id": "07",
      "language": "C++",
      "edition": null,
      "sale": true
    }
  ]
}
```

Types

Objects are an *unordered* collection of name-value pairs:

- "name": <value>
- Values can be primitives, objects, or arrays
- **Enclosed by { }**

JSON Standard – Rules of the Game

- JavaScript Object Notation (JSON)
 - "Lightweight text-based open standard designed for **human-readable** data interchange"

```
{
  "book": [
    {
      "id": "01",
      "language": "Java",
      "author": "H. Javeson",
      "year": 2015
    },
    {
      "author": "E. Sepp",
      "id": "07",
      "language": "C++",
      "edition": null,
      "sale": true
    }
  ]
}
```

Types

Objects are an *unordered* collection of name-value pairs:

- "name": <value>
- Values can be primitives, objects, or arrays
- **Enclosed by { }**

JSON Standard – Rules of the Game

- JavaScript Object Notation (JSON)
 - "Lightweight text-based open standard designed for **human-readable** data interchange"

```
{
  "book": [
    {
      "id": "01",
      "language": "Java",
      "author": "H. Javeson",
      "year": 2015
    },
    {
      "author": "E. Sepp",
      "id": "07",
      "language": "C++",
      "edition": null,
      "sale": true
    }
  ]
}
```

Types

Objects are an *unordered* collection of name-value pairs:

- "name": <value>
- Values can be primitives, objects, or arrays
- Enclosed by { }

JSON Standard – Rules of the Game

- JavaScript Object Notation (JSON)
 - "Lightweight text-based open standard designed for **human-readable** data interchange"

```
{
  "book": [
    {
      "id": "01",
      "language": "Java",
      "author": "H. Javeson",
      "year": 2015
    },
    {
      "author": "E. Sepp",
      "id": "07",
      "language": "C++",
      "edition": null,
      "sale": true
    }
  ]
}
```

Types

Objects are an *unordered* collection of name-value pairs:

- "name": <value>
- Values can be primitives, objects, or arrays
- Enclosed by { }

JSON Standard – Rules of the Game

- JavaScript Object Notation (JSON)
 - "Lightweight text-based open standard designed for **human-readable** data interchange"

```
{
  "book": [
    {
      "id": "01",
      "language": "Java",
      "author": "H. Javeson",
      "year": 2015
    },
    {
      "author": "E. Sepp",
      "id": "07",
      "language": "C++",
      "edition": null,
      "sale": true
    }
  ]
}
```

Types

Arrays are an *ordered* list of values:

- Order is preserved in interpretation
- May contain any mix of types
- Enclosed by []

JSON Standard – Rules of the Game

- JavaScript Object Notation (JSON)
 - "Lightweight text-based open standard designed for **human-readable** data interchange"

```
{
  "book": [
    {
      "id": "01",
      "language": "Java",
      "author": "H. Javeson",
      "year": 2015
    },
    {
      "author": "E. Sepp",
      "id": "07",
      "language": "C++",
      "edition": null,
      "sale": true
    }
  ]
}
```

Types

Arrays are an *ordered* list of values:

- Order is preserved in interpretation
- May contain any mix of types
- Enclosed by []

Read as "book":[{object1}, {object2}]

Can have mix of types like
[{object1}, "string", 124, {object2}]

JSON Standard – Rules of the Game

- JSON Standard too expressive
 - Implementations **restrict syntax**
 - Ex: Duplicate fields

```
{  
  "id": "01",  
  "language": "Java",  
  "author": "H. Javeson",  
  "author": "D. Suciú",  
  "author": "A. Cheung",  
  "year": 2015  
}
```

JSON Standard – Rules of the Game

- JSON Standard too expressive
 - Implementations **restrict syntax**
 - Ex: Duplicate fields



```
{  
  "id": "01",  
  "language": "Java",  
  "author": "H. Javeson",  
  "author": "D. Suciu",  
  "author": "A. Cheung",  
  "year": 2015  
}
```



```
{  
  "id": "01",  
  "language": "Java",  
  "author": ["H. Javeson",  
            "D. Suciu",  
            "A. Cheung"],  
  "year": 2015  
}
```


Thinking About Semi-Structured Data

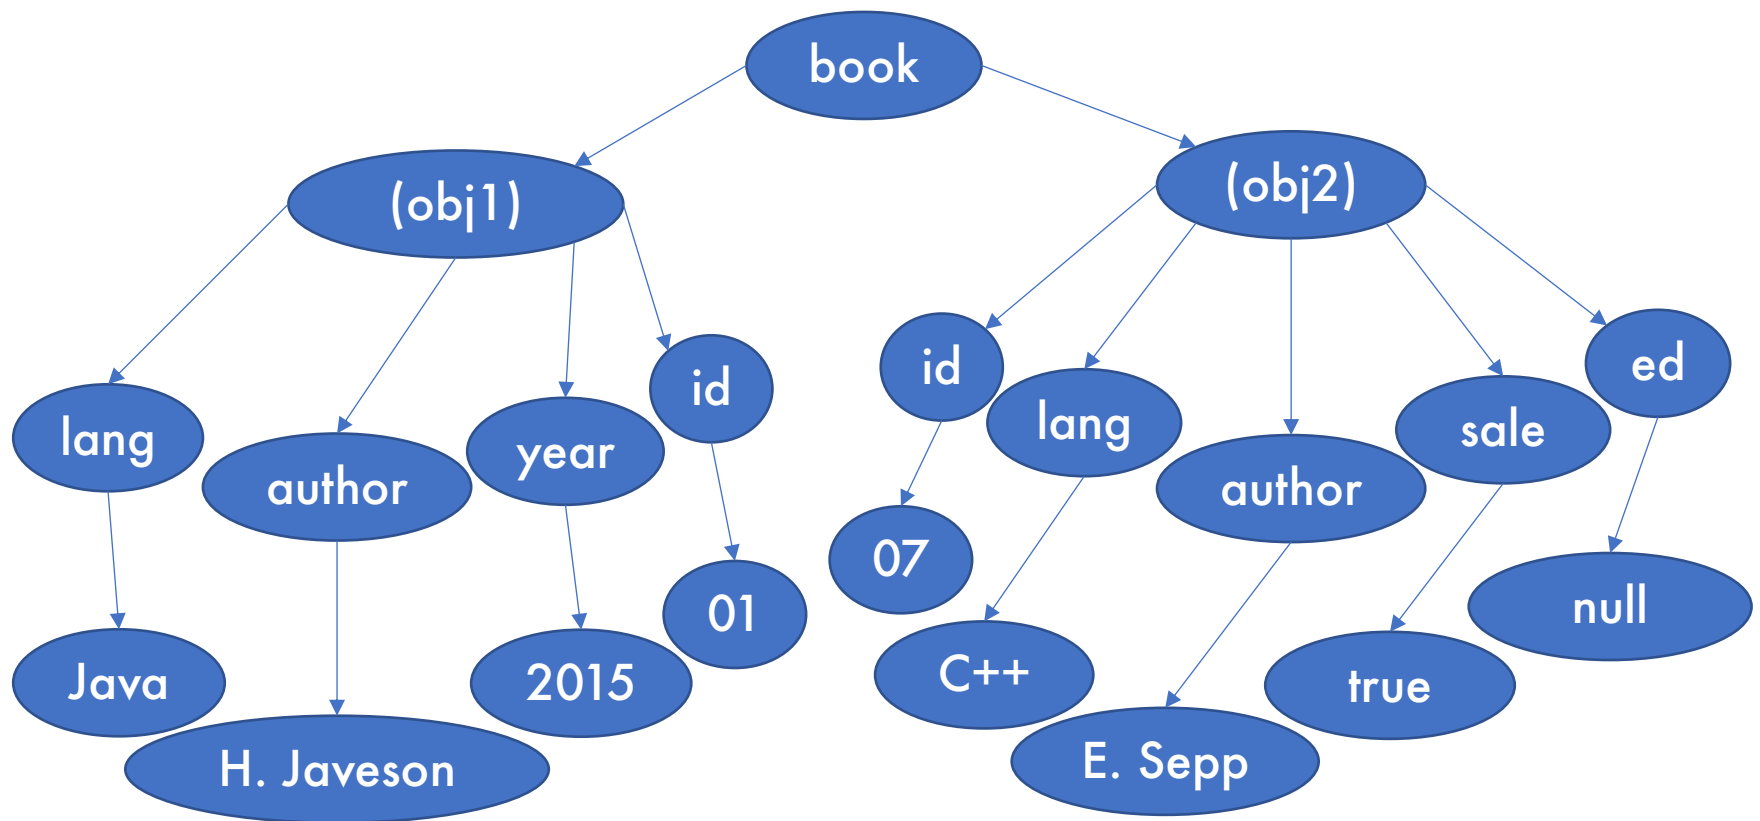
What does semi-structured data structure encode?

```
{
  "book": [
    {
      "id": "01",
      "language": "Java",
      "author": "H. Javeson",
      "year": 2015
    },
    {
      "author": "E. Sepp",
      "id": "07",
      "language": "C++",
      "edition": null,
      "sale": true
    }
  ]
}
```

Thinking About Semi-Structured Data

What does semi-structured data structure encode?

Tree semantics!

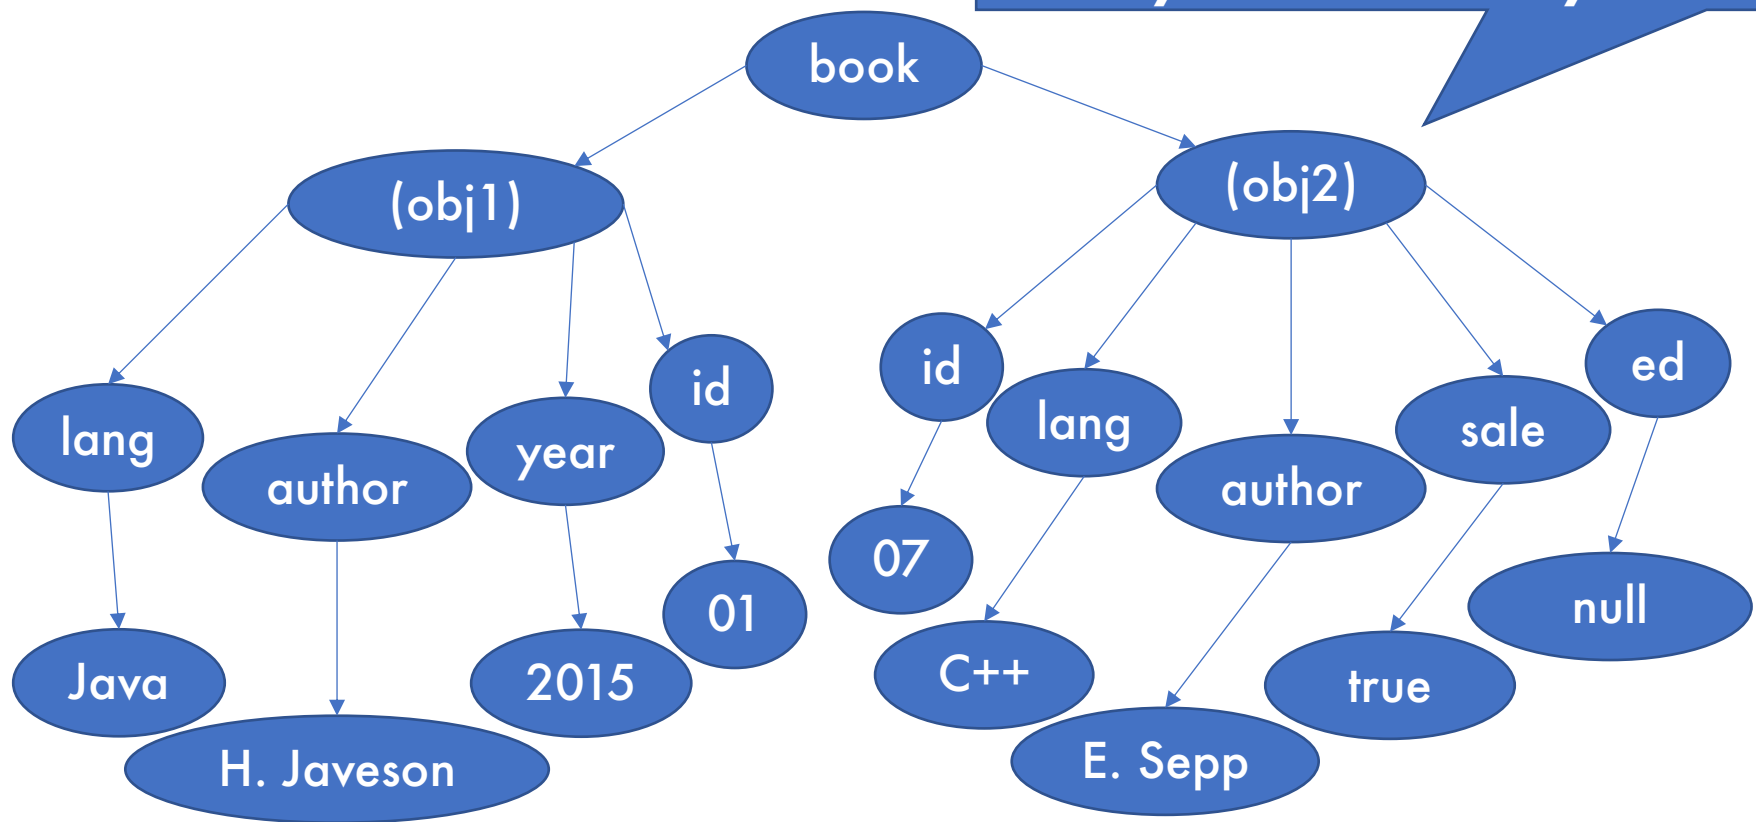


Thinking About Semi-Structured Data

What does semi-structured data structure encode?

Tree semantics!

These object don't have labels, as they are in an array



From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

What is a table in semi-structured land?

person

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

What is a table in semi-structured land?

person


From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

```
{  
  "person": [  
    {  
      ...  
    },  
    {  
      ...  
    },  
    {  
      ...  
    }  
  ]  
}
```

What is a table in semi-structured land?



Tables are just an array of elements (rows)


From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

```
{  
  "person": [  
    {  
      ...  
    },  
    {  
      ...  
    },  
    {  
      ...  
    }  
  ]  
}
```

What is a table in semi-structured land?




Tables are just an array of elements (rows)

Rows are just simple (unnested) objects

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789



```
{
  "person": [
    {
      "name": "Dan",
      "phone": "555-123-4567"
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678"
    },
    {
      "name": "Magda",
      "phone": "555-345-6789"
    }
  ]
}
```


From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

How can NULL
be represented?

```
{
  "person": [
    {
      "name": "Dan",
      "phone": "555-123-4567"
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678"
    },
    {
      "name": "Magda",
      "phone": "555-345-6789"
    }
  ]
}
```

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	NULL

How can NULL
be represented?

```
{
  "person": [
    {
      "name": "Dan",
      "phone": "555-123-4567"
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678"
    },
    {
      "name": "Magda",
      "phone": "555-345-6789"
    }
  ]
}
```

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	NULL

How can NULL
be represented?

```
{
  "person": [
    {
      "name": "Dan",
      "phone": "555-123-4567"
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678"
    },
    {
      "name": "Magda",
      "phone": null
    }
  ]
}
```

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	NULL

How can NULL
be represented?

```
{
  "person": [
    {
      "name": "Dan",
      "phone": "555-123-4567"
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678"
    },
    {
      "name": "Magda"
    }
  ]
}
```

OK for field to
be missing!

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Are there things that
the Relational Model
can't represent?

```
{
  "person": [
    {
      "name": "Dan",
      "phone": "555-123-4567"
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678"
    },
    {
      "name": "Magda",
      "phone": "555-345-6789"
    }
  ]
}
```

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Are there things that
the Relational Model
can't represent?

Non-flat data!

- Array data
- Multi-part data

```
{
  "person": [
    {
      "name": "Dan",
      "phone": "555-123-4567"
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678"
    },
    {
      "name": "Magda",
      "phone": "555-345-6789"
    }
  ]
}
```

From Relational to Semi-Structured

Person

Name	Phone
Dan	???
Alvin	555-234-5678
Magda	555-345-6789

Are there things that
the Relational Model
can't represent?

Non-flat data!

- **Array data**
- Multi-part data

Array with 2
objects

```
{
  "person": [
    {
      "name": "Dan",
      "phone": [
        "555-123-4567",
        "555-987-6543"
      ]
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678"
    },
    {
      "name": "Magda",
      "phone": "555-345-6789"
    }
  ]
}
```

From Relational to Semi-Structured

Person

Name	Phone
???	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Are there things that
the Relational Model
can't represent?

Non-flat data!


- Array data
- **Multi-part data**

Object with 2
key-value pairs

```
{
  "person": [
    {
      "name": {
        "fname": "Dan",
        "lname": "Suciu"
      },
      "phone": "555-123-4567"
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678"
    },
    {
      "name": "Magda",
      "phone": "555-345-6789"
    }
  ]
}
```


From Relational to Semi-Structured

Person



Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Orders

PName	Date	Product
Dan	1997	Furby
Alvin	2000	Furby
Alvin	2012	Magic8

How do we represent
foreign keys?

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Orders

PName	Date	Product
Dan	1997	Furby
Alvin	2000	Furby
Alvin	2012	Magic8

```
{
  "person": [
    {
      "name": "Dan",
      "phone": "555-123-4567",
      "orders": [
        {
          "date": 1997,
          "product": "Furby"
        }
      ]
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678",
      "orders": [
        {
          "date": 2000,
          "product": "Furby"
        },
        {
          "date": 2012,
          "product": "Magic8"
        }
      ]
    },
    {
      "name": "Magda",
      "phone": "555-345-6789",
      "orders": []
    }
  ]
}
```

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Orders

PName	Date	Product
Dan	1997	Furby
Alvin	2000	Furby
Alvin	2012	Magic8

Precomputed
equijoin!

```
{
  "person": [
    {
      "name": "Dan",
      "phone": "555-123-4567",
      "orders": [
        {
          "date": 1997,
          "product": "Furby"
        }
      ]
    },
    {
      "name": "Alvin",
      "phone": "555-234-5678",
      "orders": [
        {
          "date": 2000,
          "product": "Furby"
        },
        {
          "date": 2012,
          "product": "Magic8"
        }
      ]
    },
    {
      "name": "Magda",
      "phone": "555-345-6789",
      "orders": []
    }
  ]
}
```

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Is this many-to-many
relationship easily
convertible to JSON?

Orders

PName	Date	Product
Dan	1997	Furby
Alvin	2000	Furby
Alvin	2012	Magic8

Product

ProdName	Price
Furby	9.99
Magic8	15.99
Tomagachi	18.99

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Orders

PName	Date	Product
Dan	1997	Furby
Alvin	2000	Furby
Alvin	2012	Magic8

Product

ProdName	Price
Furby	9.99
Magic8	15.99
Tomagachi	18.99

Is this many-to-many
relationship easily
convertible to JSON?

Nest the data?
Person → Orders → Product

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Orders

PName	Date	Product
Dan	1997	Furby
Alvin	2000	Furby
Alvin	2012	Magic8

Product

ProdName	Price
Furby	9.99
Magic8	15.99
Tomagachi	18.99

Is this many-to-many
relationship easily
convertible to JSON?

Nest the data?
Person → Orders → Product

We might miss some products!
&
Product data will be duplicated!

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Orders

PName	Date	Product
Dan	1997	Furby
Alvin	2000	Furby
Alvin	2012	Magic8

Product

ProdName	Price
Furby	9.99
Magic8	15.99
Tomagachi	18.99

Is this many-to-many
relationship easily
convertible to JSON?

Nest the data?
Product → Orders → Person

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Orders

PName	Date	Product
Dan	1997	Furby
Alvin	2000	Furby
Alvin	2012	Magic8

Product

ProdName	Price
Furby	9.99
Magic8	15.99
Tomagachi	18.99

Is this many-to-many
relationship easily
convertible to JSON?

Nest the data?
Product → Orders → Person

We might miss some people!
&
People data will be duplicated!

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Orders

PName	Date	Product
Dan	1997	Furby
Alvin	2000	Furby
Alvin	2012	Magic8

Product

ProdName	Price
Furby	9.99
Magic8	15.99
Tomagachi	18.99

Is this many-to-many
relationship easily
convertible to JSON?

Convert each table to a
separate array/document?

From Relational to Semi-Structured

Person

Name	Phone
Dan	555-123-4567
Alvin	555-234-5678
Magda	555-345-6789

Orders

PName	Date	Product
Dan	1997	Furby
Alvin	2000	Furby
Alvin	2012	Magic8

Product

ProdName	Price
Furby	9.99
Magic8	15.99
Tomagachi	18.99

Is this many-to-many
relationship easily
convertible to JSON?

Convert each table to a
separate array/document?

We wanted to avoid joining
in the first place!

From Relational to Semi-Structured

Big ideas:

- Semi-structured data is **parsed**
 - Data model flexibility
 - Potentially lots of redundancy
- Semi-structured data expresses **unique patterns**
 - Collection/multi-part data
 - Precompute joins
- Semi-structured data **has limits**
 - Relies on relational-like patterns in some situations

A Semi-structured DBMS

- AsterixDB as a case study of Document Store
 - Semi-structured data model in JSON
 - **Introducing AsterixDB and SQL++**



The 5 W's of AsterixDB

- Who
 - M. J. Carey & co.
- What
 - "A Scalable, Open Source BDMS"
 - It is now also an Apache project
- Where
 - UC Irvine, Cloudera Inc, Google, IBM, ...
- When
 - 2014
- Why
 - To develop a next-gen system for managing semi-structured data

The 5 W's of SQL++

- Who
 - K. W. Ong & Y. Papakonstantinou
- What
 - A query language that is applicable to JSON native stores and SQL databases
- Where
 - UC San Diego
- When
 - 2015
- Why
 - Stand in for other semi-structured query languages that lack formal semantics.

Why We are Choosing SQL++

- Strong formal semantics
 - Original paper:
<https://arxiv.org/pdf/1405.3631.pdf>
 - Nested relational algebra:
<https://dl.acm.org/citation.cfm?id=588133>
- Many systems adopting or converging to SQL++
 - Apache AsterixDB
 - CouchBase (N1QL)
 - Apache Drill
 - Snowflake

Asterix Data Model (ADM)

- Nearly identical to the JSON standard
- Some additions
 - New primitive: **universally unique identifier (uuid)**
 - Ex: 123e4567-e89b-12d3-a456-426655440000
 - New derived type: **multiset**
 - Like an array but unordered
 - Encapsulated by double curly braces {{ }}
- Queried data must be a multiset or array

Introducing the New and Improved SQL++



SQL++ Mini Demo

General Installation (Details in HW7 spec)

Download from:

<https://asterixdb.apache.org/download.html>

Start local cluster from:

`<asterix root>/opt/local/bin/start-sample-cluster`

Use web browser for interaction, default address:
127.0.0.1:19002

Don't forget to stop cluster when you're done:

`<asterix root>/opt/local/bin/stop-sample-cluster`

SQL++ Mini Demo

General Usage:

Everything is running locally so make sure your computer doesn't die (advise against `SELECT *`)

Don't use attu, previous quarters people accidentally used other people's instance

Learn something! I dare say that SQL++ is a model for many future query languages.

SQL++ Hello World

```
SELECT x.phone
FROM [
    {"name": "Dan", "phone": [300, 150]},
    {"name": "Alvin", "phone": 420}
] AS x;
```

```
-- output, same for-loop semantics like in SQL
-- array data
/*
{ "phone": [300, 150] }
{ "phone": 420 }
*/
```

SQL++ Hello World

```
SELECT x.phone
FROM {{
    {"name": "Dan", "phone": [300, 150]},
    {"name": "Alvin", "phone": 420}
}} AS x;
```

```
-- same output as array data
-- multiset data
```

SQL++ Hello World

```
-- error
SELECT x.phone
  FROM {"name": "Dan", "phone": [300, 150]} AS x;

-- output
-- trying to query an object
/*
Type mismatch: function scan-collection expects its
1st input parameter to be type multiset or array,
but the actual input type is object
[TypeMismatchException]
*/
```

SQL++ Hello World

```
SELECT x.phone
FROM [
    {"name": "Dan", "phone": [300, 150]},
    {"name": "Alvin", "phone": null}
] AS x;
```

```
-- output, null works like in SQL
-- null values
/*
{ "phone": [300, 150] }
{ "phone": null }
*/
```

SQL++ Hello World

```
SELECT x.phone
FROM [
    {"name": "Dan", "phone": [300, 150]},
    {"name": "Alvin"}
] AS x;
```

```
-- output, missing data is simply passed over (beware of typos!)
-- missing values
/*
{ "phone": [300, 150] }
{ }
*/
```


SQL++ Hello World

```
SELECT x.fone -- intentional typo
FROM [
    {"name": "Dan", "phone": [300, 150]},
    {"name": "Alvin", "phone": 420}
] AS x;

-- output, beware of typos! No errors are thrown
/*
{ }
{ }
*/
```

SQL++ Hello World

```
FROM [
    {"name": "Dan", "phone": [300, 150]},
    {"name": "Alvin", "phone": 420}
] AS x
WHERE is_array(x.phone) OR x.phone > 100
GROUP BY x.name, x.phone
HAVING x.name = "Dan" OR x.name = "Alvin"
SELECT x.phone
ORDER BY x.name DESC;

-- output, finally the keyword order matches FWGHOS!
/*
{ "phone": [300, 150] }
{ "phone": 420 }
*/
```