

Natural Language Processing (CSE 447/547M): Neural Language Models

Noah Smith

© 2019

University of Washington
`nasmith@cs.washington.edu`

January 16, 2019

Quick Review

A language model is a probability distribution over \mathcal{V}^{\dagger} .

Typically p decomposes into probabilities $p(x_i \mid \mathbf{h}_i)$.

- ▶ n-gram: \mathbf{h}_i is $(n - 1)$ previous symbols
- ▶ Probabilities are estimated from data.

Two kinds of language models so far:

	representation?	estimation?	think about?
n-gram	\mathbf{h}_i is $(n - 1)$ previous symbols	count and normalize	smoothing
log-linear	featurized representation of $\langle \mathbf{h}_i, x_i \rangle$	follow gradients	features

Please review slides 51–58 from last time (added today).

Today: neural language models!

Neural Network: Definitions

Warning: there is no widely accepted standard notation!

A feedforward neural network n_{ν} is defined by:

- ▶ A function family that maps parameter values to functions of the form $n : \mathbb{R}^{d_{in}} \rightarrow \mathbb{R}^{d_{out}}$; typically:
 - ▶ Non-linear
 - ▶ Differentiable with respect to its inputs
 - ▶ “Assembled” through a series of affine transformations and non-linearities, composed together
 - ▶ Symbolic/discrete inputs handled through lookups.
- ▶ Parameter values ν
 - ▶ Typically a collection of scalars, vectors, and matrices
 - ▶ We often assume they are linearized into \mathbb{R}^D


A Couple of Useful Functions

- softmax : $\mathbb{R}^k \rightarrow \mathbb{R}^k$

$$\langle x_1, x_2, \dots, x_k \rangle \mapsto \left\langle \frac{e^{x_1}}{\sum_{j=1}^k e^{x_j}}, \frac{e^{x_2}}{\sum_{j=1}^k e^{x_j}}, \dots, \frac{e^{x_k}}{\sum_{j=1}^k e^{x_j}} \right\rangle$$

- tanh : $\mathbb{R} \rightarrow [-1, 1]$

$$x \mapsto \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



Generalized to be *elementwise*, so that it maps $\mathbb{R}^k \rightarrow [-1, 1]^k$.

- Others include: ReLUs, logistic sigmoids, PReLUs, ...

“One Hot” Vectors

Arbitrarily order the words in \mathcal{V} , giving each an index in $\{1, \dots, V\}$.

Let $\mathbf{e}_i \in \mathbb{R}^V$ contain all zeros, with the exception of a 1 in position i .

This is the “one hot” vector for the i th word in \mathcal{V} .

Feedforward Neural Network Language Model

(Bengio et al., 2003)

Define the n-gram probability as follows:

$$p(\cdot \mid \langle h_1, \dots, h_{n-1} \rangle) = n_{\nu}(\langle \mathbf{e}_{h_1}, \dots, \mathbf{e}_{h_{n-1}} \rangle) = \\ \text{softmax} \left(\underset{\nu}{\mathbf{b}} + \sum_{j=1}^{n-1} \underset{\nu}{\mathbf{e}_{h_j}}^\top \underset{\nu \times d}{\mathbf{M}} \underset{d \times \nu}{\mathbf{A}_j} + \underset{\nu \times H}{\mathbf{W}} \tanh \left(\underset{H}{\mathbf{u}} + \sum_{j=1}^{n-1} \underset{\nu}{\mathbf{e}_{h_j}}^\top \underset{d \times H}{\mathbf{M}} \underset{d \times H}{\mathbf{T}_j} \right) \right)$$

where each $\mathbf{e}_{h_j} \in \mathbb{R}^V$ is a one-hot vector and H is the number of “hidden units” in the neural network (a “hyperparameter”).

Parameters ν include:

- ▶ $\mathbf{M} \in \mathbb{R}^{V \times d}$, which are called “embeddings” (row vectors), one for every word in \mathcal{V}
- ▶ Feedforward NN parameters $\mathbf{b} \in \mathbb{R}^V$, $\mathbf{A} \in \mathbb{R}^{(n-1) \times d \times V}$, $\mathbf{W} \in \mathbb{R}^{V \times H}$, $\mathbf{u} \in \mathbb{R}^H$, $\mathbf{T} \in \mathbb{R}^{(n-1) \times d \times H}$

Breaking It Down

Look up each of the history words $h_j, \forall j \in \{1, \dots, n - 1\}$ in \mathbf{M} ; keep two copies.

$$\underset{v}{\mathbf{e}_{h_j}}^\top \underset{v \times d}{\mathbf{M}}$$
$$\underset{v}{\mathbf{e}_{h_j}}^\top \underset{v \times d}{\mathbf{M}}$$

Breaking It Down

Look up each of the history words $h_j, \forall j \in \{1, \dots, n-1\}$ in \mathbf{M} ; keep two copies.
Rename the embedding for h_j as \mathbf{m}_{h_j} .

$$\mathbf{e}_{h_j}^\top \mathbf{M} = \mathbf{m}_{h_j}$$

$$\mathbf{e}_{h_j}^\top \mathbf{M} = \mathbf{m}_{h_j}$$

Breaking It Down

Apply an affine transformation to the second copy of the history-word embeddings (\mathbf{u} , \mathbf{T})

$$\mathbf{m}_{h_j} + \sum_{j=1}^{n-1} \mathbf{m}_{h_j} \mathbf{T}_j$$

\mathbf{u}_H $\mathbf{T}_j_{d \times H}$

Breaking It Down

Apply an affine transformation to the second copy of the history-word embeddings (\mathbf{u} , \mathbf{T}) and a \tanh nonlinearity.

$$\mathbf{m}_{h_j} \tanh \left(\mathbf{u} + \sum_{j=1}^{n-1} \mathbf{m}_{h_j} \mathbf{T}_j \right)$$

Breaking It Down

Apply an affine transformation to everything (\mathbf{b} , \mathbf{A} , \mathbf{W}).

$$\mathbf{b}_v + \sum_{j=1}^{n-1} \mathbf{m}_{h_j} \mathbf{A}_j^{d \times v} + \mathbf{W}_{v \times H} \tanh \left(\mathbf{u} + \sum_{j=1}^{n-1} \mathbf{m}_{h_j} \mathbf{T}_j \right)$$

Breaking It Down

Apply a softmax transformation to make the vector sum to one.

$$\text{softmax} \left(\mathbf{b} + \sum_{j=1}^{n-1} \mathbf{m}_{h_j} \mathbf{A}_j \right. \\ \left. + \mathbf{W} \tanh \left(\mathbf{u} + \sum_{j=1}^{n-1} \mathbf{m}_{h_j} \mathbf{T}_j \right) \right)$$

Breaking It Down


$$\text{softmax} \left(\mathbf{b} + \sum_{j=1}^{n-1} \mathbf{m}_{h_j} \mathbf{A}_j \right. \\ \left. + \mathbf{W} \tanh \left(\mathbf{u} + \sum_{j=1}^{n-1} \mathbf{m}_{h_j} \mathbf{T}_j \right) \right)$$

Like a log-linear language model with two kinds of features:

- Concatenation of context-word embeddings vectors \mathbf{m}_{h_j}
- \tanh -affine transformation of the above

New parameters arise from (i) embeddings and (ii) affine transformation “inside” the nonlinearity.

Visualization



Number of Parameters

$$D = \underbrace{Vd}_{\mathbf{M}} + \underbrace{V}_{\mathbf{b}} + \underbrace{(n-1)dV}_{\mathbf{A}} + \underbrace{VH}_{\mathbf{W}} + \underbrace{H}_{\mathbf{u}} + \underbrace{(n-1)dH}_{\mathbf{T}}$$

For Bengio et al. (2003):

- ▶ $V \approx 18000$ (after OOV processing)
- ▶ $d \in \{30, 60\}$
- ▶ $H \in \{50, 100\}$
- ▶ $n - 1 = 5$

So $D = 461V + 30100$ parameters, compared to $O(V^n)$ for classical n-gram models.

- ▶ Forcing $\mathbf{A} = \mathbf{0}$ eliminated $300V$ parameters and performed a bit better, but was slower to converge.
- ▶ If we averaged \mathbf{m}_{h_j} instead of concatenating, we'd get to $221V + 6100$ (this is a variant of “continuous bag of words,” Mikolov et al., 2013).

Why does it work?


Why does it work?

- ▶ Historical answer: multiple layers and nonlinearities allow feature *combinations* a linear model can't get.

Why does it work?

- ▶ Historical answer: multiple layers and nonlinearities allow feature *combinations* a linear model can't get.
 - ▶ Suppose we want $y = \text{xor}(x_1, x_2)$; this can't be expressed as a linear function of x_1 and x_2 .

xor Example



Tuples where $y = \text{xor}(x_1, x_2)$ are red; tuples where $y \neq \text{xor}(x_1, x_2)$ are blue.

Why does it work?


- ▶ Historical answer: multiple layers and nonlinearities allow feature *combinations* a linear model can't get.
 - ▶ Suppose we want $y = \text{xor}(x_1, x_2)$; this can't be expressed as a linear function of x_1 and x_2 . But:

$$z = x_1 \cdot x_2$$

$$y = x_1 + x_2 - 2z$$

xor Example ($D = 13$)

Credit: Chris Dyer (<https://github.com/clab/cnn/blob/master/examples/xor.cc>)



$$\min_{\mathbf{v}, a, \mathbf{W}, \mathbf{b}} \sum_{x_1 \in \{0,1\}} \sum_{x_2 \in \{0,1\}} \left(\text{xor}(x_1, x_2) - \underset{3}{\mathbf{v}}^\top \left(\underset{3 \times 2}{\mathbf{W}} \underset{2}{\mathbf{x}} + \underset{3}{\mathbf{b}} \right) + a \right)^2$$

$$\min_{\mathbf{v}, a, \mathbf{W}, \mathbf{b}} \sum_{x_1 \in \{0,1\}} \sum_{x_2 \in \{0,1\}} \left(\text{xor}(x_1, x_2) - \underset{3}{\mathbf{v}}^\top \tanh \left(\underset{3 \times 2}{\mathbf{W}} \underset{2}{\mathbf{x}} + \underset{3}{\mathbf{b}} \right) + a \right)^2$$

Why does it work?

- ▶ Historical answer: multiple layers and nonlinearities allow feature *combinations* a linear model can't get.
 - ▶ Suppose we want $y = \text{xor}(x_1, x_2)$; this can't be expressed as a linear function of x_1 and x_2 . But:

$$z = x_1 \cdot x_2$$

$$y = x_1 + x_2 - 2z$$

- ▶ With high-dimensional inputs, there are a lot of conjunctive features to search through. For log-linear models, Della Pietra et al. (1997) did this, greedily.

Why does it work?

- ▶ Historical answer: multiple layers and nonlinearities allow feature *combinations* a linear model can't get.
 - ▶ Suppose we want $y = \text{xor}(x_1, x_2)$; this can't be expressed as a linear function of x_1 and x_2 . But:

$$z = x_1 \cdot x_2$$

$$y = x_1 + x_2 - 2z$$

- ▶ With high-dimensional inputs, there are a lot of conjunctive features to search through. For log-linear models, Della Pietra et al. (1997) did this, greedily.
- ▶ Neural models seem to smoothly explore lots of approximately-conjunctive features.

Why does it work?

- ▶ Historical answer: multiple layers and nonlinearities allow feature *combinations* a linear model can't get.
 - ▶ Suppose we want $y = \text{xor}(x_1, x_2)$; this can't be expressed as a linear function of x_1 and x_2 . But:

$$z = x_1 \cdot x_2$$

$$y = x_1 + x_2 - 2z$$

- ▶ With high-dimensional inputs, there are a lot of conjunctive features to search through. For log-linear models, Della Pietra et al. (1997) did this, greedily.
 - ▶ Neural models seem to smoothly explore lots of approximately-conjunctive features.
- ▶ Modern answer: representations of words and histories are tuned to the prediction problem.

Why does it work?

- ▶ Historical answer: multiple layers and nonlinearities allow feature *combinations* a linear model can't get.
 - ▶ Suppose we want $y = \text{xor}(x_1, x_2)$; this can't be expressed as a linear function of x_1 and x_2 . But:

$$z = x_1 \cdot x_2$$

$$y = x_1 + x_2 - 2z$$

- ▶ With high-dimensional inputs, there are a lot of conjunctive features to search through. For log-linear models, Della Pietra et al. (1997) did this, greedily.
 - ▶ Neural models seem to smoothly explore lots of approximately-conjunctive features.
- ▶ Modern answer: representations of words and histories are tuned to the prediction problem.
- ▶ Word embeddings: a powerful idea ...

Parameter Estimation

Bad news for neural language models:

- ▶ Log-likelihood function is not convex.
 - ▶ So any perplexity experiment is evaluating the model *and* an algorithm for estimating it.
- ▶ Calculating log-likelihood and its gradient is very expensive (5 epochs took 3 weeks on 40 CPUs).

Parameter Estimation

Bad news for neural language models:

- ▶ Log-likelihood function is not convex.
 - ▶ So any perplexity experiment is evaluating the model *and* an algorithm for estimating it.
- ▶ Calculating log-likelihood and its gradient is very expensive (5 epochs took 3 weeks on 40 CPUs).

Good news:

- ▶ n_{ν} is differentiable with respect to \mathbf{M} (from which its inputs come) and ν (its parameters), so gradient-based methods are available.

Lots more details in Bengio et al. (2003) and (for NNs more generally) in Goldberg (2015).

Observations about Neural Language Models (So Far)

- ▶ There's no knowledge built in that the most recent word h_{n-1} should generally be more informative than earlier ones.
 - ▶ This has to be learned.
- ▶ In addition to choosing n , also have to choose dimensionalities like d and H .
- ▶ Parameters of these models are hard to interpret.
- ▶ Architectures are not intuitive.
- ▶ Still, impressive perplexity gains got people's interest.

Observations about Neural Language Models (So Far)

- ▶ There's no knowledge built in that the most recent word h_{n-1} should generally be more informative than earlier ones.
 - ▶ This has to be learned.
- ▶ In addition to choosing n , also have to choose dimensionalities like d and H .
- ▶ Parameters of these models are hard to interpret.
 - ▶ Example: ℓ_2 -norm of $\mathbf{A}_{j,*,*}$ and $\mathbf{T}_{j,*,*}$ in the feedforward model correspond to the importance of history position j .
 - ▶ Individual word embeddings can be clustered and dimensions can be analyzed (e.g., Tsvetkov et al., 2015).
- ▶ Architectures are not intuitive.
- ▶ Still, impressive perplexity gains got people's interest.

Recurrent Neural Network

- ▶ Each input element is understood to be an element of a sequence: $\langle \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_\ell \rangle$
- ▶ At each timestep t :
 - ▶ The t th input element \mathbf{x}_t is processed alongside the previous state \mathbf{s}_{t-1} to calculate the new **state** (\mathbf{s}_t).
 - ▶ The t th output is a function of the state \mathbf{s}_t .
 - ▶ The *same functions* are applied at each iteration:

$$\mathbf{s}_t = f_{\text{recurrent}}(\mathbf{x}_t, \mathbf{s}_{t-1})$$

$$\mathbf{y}_t = f_{\text{output}}(\mathbf{s}_t)$$

In RNN language models, words *and* histories are represented as vectors (respectively, $\mathbf{x}_t = \mathbf{e}_{x_t}$ and \mathbf{s}_t).

RNN Language Model

The original version, by Mikolov et al. (2010) used a “simple” RNN architecture along these lines:


$$\mathbf{s}_t = f_{\text{recurrent}}(\mathbf{e}_{x_t}, \mathbf{s}_{t-1}) = \text{sigmoid} \left(\left(\mathbf{e}_{x_t}^\top \mathbf{M} \right)^\top \mathbf{A} + \mathbf{s}_{t-1}^\top \mathbf{B} + \mathbf{c} \right)$$

$$\mathbf{y}_t = f_{\text{output}}(\mathbf{s}_t) = \text{softmax} \left(\mathbf{s}_t^\top \mathbf{U} \right)$$


$$p(v \mid x_1, \dots, x_{t-1}) = [\mathbf{y}_t]_v$$

Note: this is *not* an n-gram (Markov) model!

Visualization



Visualization



Improvements to RNN Language Models

The simple RNN is known to suffer from two related problems:

- ▶ “Vanishing gradients” during learning make it hard to propagate error into the distant past.
- ▶ State tends to change a lot on each iteration; the model “forgets” too much.

Some variants:

- ▶ “Stacking” these functions to make deeper networks.
- ▶ Sundermeyer et al. (2012) use “long short-term memories” (LSTMs; see Olah, 2015) and Cho et al. (2014) use “gated recurrent units” (GRUs) to define $f_{\text{recurrent}}$.
- ▶ Mikolov et al. (2014) engineer the linear transformation in the simple RNN for better preservation.
- ▶ Jozefowicz et al. (2015) used randomized search to find even better architectures.

References I

- Yoshua Bengio, Réjean Ducharme, Pascal Vincent, and Christian Jauvin. A neural probabilistic language model. *Journal of Machine Learning Research*, 3(Feb):1137–1155, 2003. URL <http://www.jmlr.org/papers/volume3/bengio03a/bengio03a.pdf>.
- Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using RNN encoder–decoder for statistical machine translation. In *Proc. of EMNLP*, 2014.
- Scott C. Deerwester, Susan T. Dumais, Thomas K. Landauer, George W. Furnas, and Richard A. Harshman. Indexing by latent semantic analysis. *Journal of the American Society for Information Science*, 41(6):391–407, 1990.
- Stephen Della Pietra, Vincent Della Pietra, and John Lafferty. Inducing features of random fields. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 19(4):380–393, 1997.
- Yoav Goldberg. A primer on neural network models for natural language processing, 2015. URL <http://u.cs.biu.ac.il/~yogo/nnlp.pdf>.
- Rafal Jozefowicz, Wojciech Zaremba, and Ilya Sutskever. An empirical exploration of recurrent network architectures. In *Proc. of ICML*, 2015. URL <http://www.jmlr.org/proceedings/papers/v37/jozefowicz15.pdf>.
- Tomas Mikolov, Martin Karafiát, Lukas Burget, Jan Cernocký, and Sanjeev Khudanpur. Recurrent neural network based language model. In *Proc. of Interspeech*, 2010. URL http://www.fit.vutbr.cz/research/groups/speech/publi/2010/mikolov_interspeech2010_IS100722.pdf.

References II

- Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. In *Proceedings of ICLR*, 2013. URL <http://arxiv.org/pdf/1301.3781.pdf>.
- Tomas Mikolov, Armand Joulin, Sumit Chopra, Michael Mathieu, and Marc'Aurelio Ranzato. Learning longer memory in recurrent neural networks, 2014. arXiv:1412.7753.
- Christopher Olah. Understanding LSTM networks, 2015. URL <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>.
- Martin Sundermeyer, Ralf Schlüter, and Hermann Ney. LSTM neural networks for language modeling. In *Proc. of Interspeech*, 2012.
- Yulia Tsvetkov, Manaal Faruqui, Wang Ling, Guillaume Lample, and Chris Dyer. Evaluation of word vector representations by subspace alignment. In *Proc. of EMNLP*, 2015.
- Peter D. Turney and Patrick Pantel. From frequency to meaning: Vector space models of semantics. *Journal of Artificial Intelligence Research*, 37(1):141–188, 2010. URL <https://www.jair.org/media/2934/live-2934-4846-jair.pdf>.