

# Natural Language Processing (CSE 447/547M): Word Representations

Noah Smith

© 2019

University of Washington  
`nasmith@cs.washington.edu`

January 30, 2019

# Preliminaries

**Token:** an instance of a word observed in text

**Type:** the word in the abstract

# How to Represent an English Word in Code

# How to Represent an English Word in Code

1. Strings: apple  $\rightarrow$  "apple" ; banana  $\rightarrow$  "banana"

# How to Represent an English Word in Code

1. Strings: apple  $\rightarrow$  "apple" ; banana  $\rightarrow$  "banana"
2. Integers: apple  $\rightarrow$  211; banana  $\rightarrow$  633

# How to Represent an English Word in Code

1. Strings: apple  $\rightarrow$  "apple" ; banana  $\rightarrow$  "banana"
2. Integers: apple  $\rightarrow$  211; banana  $\rightarrow$  633
3. One-hot vector: apple  $\rightarrow [0, \dots, 0, 1, 0, \dots]$ ; banana  $\rightarrow [0, \dots, 0, 1, 0, \dots]$   

$\underbrace{\hspace{1.5cm}}$   
210 0s

$\underbrace{\hspace{1.5cm}}$   
632 0s

# How to Represent an English Word in Code

1. Strings: apple  $\rightarrow$  "apple" ; banana  $\rightarrow$  "banana"
2. Integers: apple  $\rightarrow$  211; banana  $\rightarrow$  633
3. One-hot vector: apple  $\rightarrow [0, \dots, 0, 1, 0, \dots]$ ; banana  $\rightarrow [0, \dots, 0, 1, 0, \dots]$   

$\underbrace{\hspace{1.5cm}}$   
210 0s

$\underbrace{\hspace{1.5cm}}$   
632 0s
4. Continuous vector

# How to Represent an English Word in Code

1. Strings: apple  $\rightarrow$  "apple" ; banana  $\rightarrow$  "banana"
2. Integers: apple  $\rightarrow$  211; banana  $\rightarrow$  633
3. One-hot vector: apple  $\rightarrow [0, \dots, 0, 1, 0, \dots]$ ; banana  $\rightarrow [0, \dots, 0, 1, 0, \dots]$   

$\underbrace{\hspace{1.5cm}}$   
210 0s

$\underbrace{\hspace{1.5cm}}$   
632 0s
4. Continuous vector
  - Learned as parameters of a neural language model (January 16 lecture)



# How to Represent an English Word in Code

1. Strings: apple  $\rightarrow$  "apple" ; banana  $\rightarrow$  "banana"
2. Integers: apple  $\rightarrow$  211; banana  $\rightarrow$  633
3. One-hot vector: apple  $\rightarrow [0, \dots, 0, 1, 0, \dots]$ ; banana  $\rightarrow [0, \dots, 0, 1, 0, \dots]$   

$\underbrace{\hspace{1.5cm}}$   
210 0s

$\underbrace{\hspace{1.5cm}}$   
632 0s
4. Continuous vector
  - ▶ Learned as parameters of a neural language model (January 16 lecture)
  - ▶ Learned as parameters of a neural text classifier (January 28 lecture)

# How to Represent an English Word in Code

1. Strings: apple  $\rightarrow$  "apple" ; banana  $\rightarrow$  "banana"
2. Integers: apple  $\rightarrow$  211; banana  $\rightarrow$  633
3. One-hot vector: apple  $\rightarrow [0, \dots, 0, 1, 0, \dots]$ ; banana  $\rightarrow [0, \dots, 0, 1, 0, \dots]$   

$\underbrace{\hspace{1.5cm}}$   
210 0s

$\underbrace{\hspace{1.5cm}}$   
632 0s
4. Continuous vector
  - ▶ Learned as parameters of a neural language model (January 16 lecture)
  - ▶ Learned as parameters of a neural text classifier (January 28 lecture)
  - ▶ More methods today

# The Word-Document Matrix

Let  $\mathbf{A} \in \mathbb{R}^{V \times C}$  contain statistics of association between words in  $\mathcal{V}$  and  $C$  documents.  
 $N$  is the total number of word tokens.

Tiny example, three documents:

- ▶ yes , we have no bananas
- ▶ say yes for bananas
- ▶ no bananas , we say

	1	2	3
,	1	0	1
bananas	1	1	1
for	0	1	0
have	1	0	0
no	1	0	1
say	0	1	1
we	1	0	1
yes	1	1	0

Count matrix:  $[\mathbf{A}]_{v,c} = c\mathbf{x}_c(v)$

## A Word-Context Matrix

Let  $\mathbf{A} \in \mathbb{R}^{V \times C}$  contain statistics of association between words in  $\mathcal{V}$  and symbols that occur just before them

Tiny example, three sentences:

- ▶ yes , we have no bananas
- ▶ say yes for bananas
- ▶ no bananas , we say

	○ _	, _	bananas _	for _	have _	no _	say _	we _	yes _
,	0	0	1	0	0	0	0	0	1
bananas	0	0	0	1	0	2	0	0	0
for	0	0	0	0	0	0	0	0	1
have	0	0	0	0	0	0	0	1	0
no	1	0	0	0	1	0	0	0	0
say	1	0	0	0	0	0	0	1	0
we	0	2	0	0	0	0	0	0	0
yes	1	0	0	0	0	0	1	0	0

Count matrix:  $[\mathbf{A}]_{v,v'} = \mathbf{c}_x(v'v)$

# Association Score

What we really want here is some way to get at “surprise.”

# Association Score

What we really want here is some way to get at “surprise.”

One way to think about this is, is the occurrence of word  $v$  in document  $c$  surprisingly high (or low), given what we'd expect due to chance?

# Association Score

What we really want here is some way to get at “surprise.”

One way to think about this is, is the occurrence of word  $v$  in document  $c$  surprisingly high (or low), given what we'd expect due to chance?

Chance would be  $\frac{c_{x_{1:C}}(v)}{N}$  words out of the  $\ell_c$  words in document  $c$ .

# Association Score

What we really want here is some way to get at “surprise.”

One way to think about this is, is the occurrence of word  $v$  in document  $c$  surprisingly high (or low), given what we'd expect due to chance?

Chance would be  $\frac{c_{\mathbf{x}_{1:C}}(v)}{N}$  words out of the  $\ell_c$  words in document  $c$ .

Intuition: consider the ratio of *observed* frequency ( $c_{\mathbf{x}_c}(v)$ ) to “chance” under independence ( $\frac{c_{\mathbf{x}_{1:C}}(v)}{N} \cdot \ell_c$ ).



# Pointwise Mutual Information

A common starting point is positive **pointwise mutual information**:

$$[\mathbf{A}]_{v,c} = \left[ \log \frac{c_{\mathbf{x}_c}(v)}{\frac{c_{\mathbf{x}_{1:C}}(v)}{N} \cdot \ell_c} \right]_+ = \left[ \log \frac{N \cdot c_{\mathbf{x}_c}(v)}{c_{\mathbf{x}_{1:C}}(v) \cdot \ell_c} \right]_+$$

From our example:

$$[\mathbf{A}]_{\text{bananas},1} = \log \frac{15 \cdot 1}{3 \cdot 6} \approx -0.18 \rightarrow 0$$

$$[\mathbf{A}]_{\text{for},2} = \log \frac{15 \cdot 1}{1 \cdot 4} \approx 1.32$$

	1	2	3
,	1	0	1
bananas	1	1	1
for	0	1	0
have	1	0	0
no	1	0	1
say	0	1	1
we	1	0	1
yes	1	1	0

# Pointwise Mutual Information

A common starting point is positive **pointwise mutual information**:

$$[\mathbf{A}]_{v,c} = \left[ \log \frac{c_{\mathbf{x}_c}(v)}{\frac{c_{\mathbf{x}_{1:C}}(v)}{N} \cdot \ell_c} \right]_+ = \left[ \log \frac{N \cdot c_{\mathbf{x}_c}(v)}{c_{\mathbf{x}_{1:C}}(v) \cdot \ell_c} \right]_+$$

Notes:

- If a word  $v$  appears with nearly the same frequency in every document, its row  $[\mathbf{A}]_{v,*}$  will be all nearly zero.

# Pointwise Mutual Information

A common starting point is positive **pointwise mutual information**:

$$[\mathbf{A}]_{v,c} = \left[ \log \frac{c_{\mathbf{x}_c}(v)}{\frac{c_{\mathbf{x}_{1:C}}(v)}{N} \cdot \ell_c} \right]_+ = \left[ \log \frac{N \cdot c_{\mathbf{x}_c}(v)}{c_{\mathbf{x}_{1:C}}(v) \cdot \ell_c} \right]_+$$

Notes:

- ▶ If a word  $v$  appears with nearly the same frequency in every document, its row  $[\mathbf{A}]_{v,*}$  will be all nearly zero.
- ▶ If a word  $v$  occurs *only* in document  $c$ , PMI will be large and positive.

# Pointwise Mutual Information

A common starting point is positive **pointwise mutual information**:

$$[\mathbf{A}]_{v,c} = \left[ \log \frac{c_{\mathbf{x}_c}(v)}{\frac{c_{\mathbf{x}_{1:C}}(v)}{N} \cdot \ell_c} \right]_+ = \left[ \log \frac{N \cdot c_{\mathbf{x}_c}(v)}{c_{\mathbf{x}_{1:C}}(v) \cdot \ell_c} \right]_+$$

Notes:

- ▶ If a word  $v$  appears with nearly the same frequency in every document, its row  $[\mathbf{A}]_{v,*}$  will be all nearly zero.
- ▶ If a word  $v$  occurs *only* in document  $c$ , PMI will be large and positive.
- ▶ PMI is very sensitive to rare occurrences; usually we smooth the frequencies and filter rare words.

# Pointwise Mutual Information

A common starting point is positive **pointwise mutual information**:

$$[\mathbf{A}]_{v,c} = \left[ \log \frac{c_{\mathbf{x}_c}(v)}{\frac{c_{\mathbf{x}_{1:C}}(v)}{N} \cdot \ell_c} \right]_+ = \left[ \log \frac{N \cdot c_{\mathbf{x}_c}(v)}{c_{\mathbf{x}_{1:C}}(v) \cdot \ell_c} \right]_+$$

Notes:

- ▶ If a word  $v$  appears with nearly the same frequency in every document, its row  $[\mathbf{A}]_{v,*}$  will be all nearly zero.
- ▶ If a word  $v$  occurs *only* in document  $c$ , PMI will be large and positive.
- ▶ PMI is very sensitive to rare occurrences; usually we smooth the frequencies and filter rare words.
- ▶ One way to think about PMI: it's telling us where a unigram model is most wrong.

# Pointwise Mutual Information

A common starting point is positive **pointwise mutual information**:

$$[\mathbf{A}]_{v,c} = \left[ \log \frac{c_{\mathbf{x}_c}(v)}{\frac{c_{\mathbf{x}_{1:C}}(v)}{N} \cdot \ell_c} \right]_+ = \left[ \log \frac{N \cdot c_{\mathbf{x}_c}(v)}{c_{\mathbf{x}_{1:C}}(v) \cdot \ell_c} \right]_+$$

Notes:

- ▶ If a word  $v$  appears with nearly the same frequency in every document, its row  $[\mathbf{A}]_{v,*}$  will be all nearly zero.
- ▶ If a word  $v$  occurs *only* in document  $c$ , PMI will be large and positive.
- ▶ PMI is very sensitive to rare occurrences; usually we smooth the frequencies and filter rare words.
- ▶ One way to think about PMI: it's telling us where a unigram model is most wrong.
- ▶ We could use  $\mathbf{A}$  as  $\mathbf{M}$  (though  $d$  is usually much smaller than  $C$ ) ...

# Topic Models: Latent Semantic Indexing/Analysis

(Deerwester et al., 1990)

LSI/A seeks to solve:

$$\underset{V \times C}{\mathbf{A}} \approx \underset{V \times d}{\hat{\mathbf{A}}} = \underset{V \times d}{\mathbf{M}} \times \underset{d \times d}{\text{diag}(\mathbf{s})} \times \underset{d \times C}{\mathbf{C}}^{\top}$$

where  $\mathbf{M}$  contains embeddings of words,  $\mathbf{C}$  contains embeddings of documents.

$$[\mathbf{A}]_{v,c} \approx \sum_{i=1}^d [\mathbf{m}_v]_i \cdot [\mathbf{s}]_i \cdot [\mathbf{c}_c]_i$$

# Topic Models: Latent Semantic Indexing/Analysis

(Deerwester et al., 1990)

LSI/A seeks to solve:

$$\underset{V \times C}{\mathbf{A}} \approx \underset{V \times d}{\hat{\mathbf{A}}} = \underset{V \times d}{\mathbf{M}} \times \underset{d \times d}{\text{diag}(\mathbf{s})} \times \underset{d \times C}{\mathbf{C}}^{\top}$$

where  $\mathbf{M}$  contains embeddings of words,  $\mathbf{C}$  contains embeddings of documents.

$$[\mathbf{A}]_{v,c} \approx \sum_{i=1}^d [\mathbf{m}_v]_i \cdot [\mathbf{s}]_i \cdot [\mathbf{c}_c]_i$$

This can be solved by applying singular value decomposition to  $\mathbf{A}$ , then truncating to  $d$  dimensions.

- ▶  $\mathbf{M}$  contains left singular vectors of  $\mathbf{A}$
- ▶  $\mathbf{C}$  contains right singular vectors of  $\mathbf{A}$
- ▶  $\mathbf{s}$  are singular values of  $\mathbf{A}$ ; they are nonnegative and conventionally organized in decreasing order.



# Truncated Singular Value Decomposition

SVD:

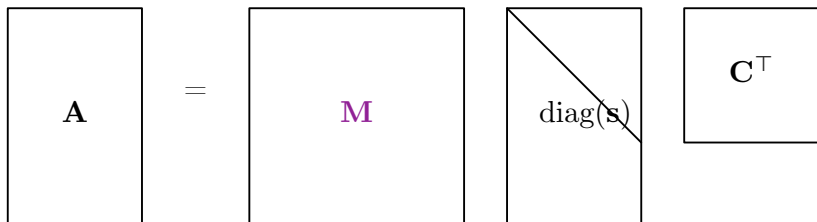


Diagram illustrating the full Singular Value Decomposition (SVD):

$$\mathbf{A} = \mathbf{M} \begin{matrix} \diagdown \\ \text{diag}(\mathbf{s}) \\ \diagup \end{matrix} \mathbf{C}^{\top}$$

The diagram shows three matrices:  $\mathbf{A}$  (a tall rectangle),  $\mathbf{M}$  (a square), and  $\mathbf{C}^{\top}$  (a wide rectangle). The matrix  $\mathbf{M}$  is colored purple. The diagonal matrix is represented by a rectangle with a diagonal line from the top-left to the bottom-right and the label  $\text{diag}(\mathbf{s})$  in the center.

truncated at  $d$ :

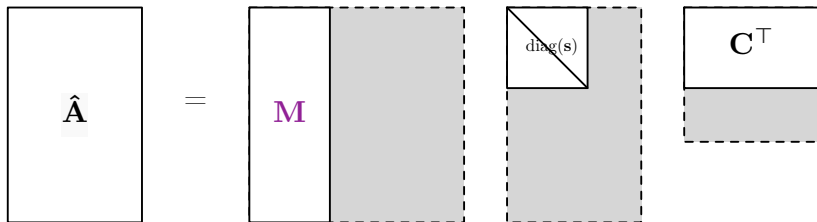


Diagram illustrating the truncated Singular Value Decomposition (SVD) at rank  $d$ :

$$\hat{\mathbf{A}} = \mathbf{M} \begin{matrix} \diagdown \\ \text{diag}(\mathbf{s}) \\ \diagup \end{matrix} \mathbf{C}^{\top}$$

The diagram shows three matrices:  $\hat{\mathbf{A}}$  (a tall rectangle),  $\mathbf{M}$  (a square), and  $\mathbf{C}^{\top}$  (a wide rectangle). The matrix  $\mathbf{M}$  is colored purple. The diagonal matrix is represented by a rectangle with a diagonal line from the top-left to the bottom-right and the label  $\text{diag}(\mathbf{s})$  in the center. The bottom-right portion of the  $\mathbf{M}$  matrix and the bottom portion of the  $\mathbf{C}^{\top}$  matrix are shaded gray and enclosed in dashed lines, indicating they are truncated.

# A Nod to Linear Algebra

For (not truncated) singular value decomposition  $\mathbf{A} = \mathbf{M} \times \text{diag}(\mathbf{s}) \times \mathbf{C}^\top$ :


- ▶ The columns of  $\mathbf{M}$  form an orthonormal basis,  $\mathbf{M}$  are eigenvectors of  $\mathbf{A}\mathbf{A}^\top$ , with eigenvalues  $s^2$ .
- ▶ The columns of  $\mathbf{C}$  form an orthonormal basis,  $\mathbf{C}$  are eigenvectors of  $\mathbf{A}^\top\mathbf{A}$ , with eigenvalues  $s^2$ .

If some elements of  $\mathbf{s}$  are zero, then  $\mathbf{A}$  is “low rank.”

Approximating  $\mathbf{A}$  by truncating  $\mathbf{s}$  equates to a “low rank approximation.”

# LSI/A Example

$k = 2$




	1	2	3
,	1	0	1
bananas	1	1	1
for	0	1	0
have	1	0	0
no	1	0	1
say	0	1	1
we	1	0	1
yes	1	1	0

Words and documents in two dimensions.

# LSI/A Example

$k = 2$



	1	2	3
,	1	0	1
bananas	1	1	1
for	0	1	0
have	1	0	0
no	1	0	1
say	0	1	1
we	1	0	1
yes	1	1	0

Words and documents in two dimensions.

Note how no, we, and , are all in the exact same spot. Why?

# Understanding LSI/A

- ▶ Mapping words and documents into the same  $k$ -dimensional space.
- ▶ Bag of words assumption (Salton et al., 1975): a document is nothing more than the distribution of words it contains.
- ▶ Distributional hypothesis (Harris, 1954; Firth, 1957): words are nothing more than the distribution of contexts (here, documents) they occur in. Words that occur in similar contexts have similar meanings.
- ▶  $\mathbf{A}$  is sparse and noisy; LSI/A “fills in” the zeroes and tries to eliminate the noise.
  - ▶ It finds the best rank- $k$  approximation to  $\mathbf{A}$ .

## Local Contexts: Distributional Semantics

Within NLP, emphasis has shifted from word/document associations to the relationship between  $v \in \mathcal{V}$  and more local contexts.

For example: LSI/A, but replace documents with “nearby words.” This is a way to recover word vectors that capture distributional similarity.

These models are designed to “guess” a word at position  $i$  given a word at a position in  $\{i - w, \dots, i - 1\} \cup \{i + 1, \dots, i + w\}$ . (My example on slide 12 used only position  $i - 1$ .)

Sometimes such methods are used to “pre-train” word vectors used in other, richer models (like neural language models).

# word2vec

(Mikolov et al., 2013a,b)

Two models for word vectors designed to be computationally efficient.

- ▶ Continuous bag of words (CBOW):  $p(v \mid c)$ 
  - ▶ Similar in spirit to the feedforward neural language model we saw before (Bengio et al., 2003)
- ▶ Skip-gram:  $p(c \mid v)$

It turns out these are closely related to matrix factorization as in LSI/A (Levy and Goldberg, 2014).

# Skip-Gram Model

$$p(C = c \mid X = v) = \frac{1}{Z_v} \exp \mathbf{c}_c^\top \mathbf{v}_v$$

- ▶ Two different vectors for each element of  $\mathcal{V}$ : one when it is “ $v$ ” ( $\mathbf{v}$ ) and one when it is “ $c$ ” ( $\mathbf{c}$ ).
- ▶ Normalization term  $Z_v$  is expensive, so approximations are required for efficiency.
- ▶ Can vary the context window, or expand this to be over the whole sentence or document, or otherwise choose which words “count” as contexts.



# Word Vector Evaluations

Several popular methods for *intrinsic* evaluations:

# Word Vector Evaluations

Several popular methods for *intrinsic* evaluations:

- ▶ How well do (cosine) similarities of pairs of words' vectors correlate with judgments of word similarity by humans?

# Word Vector Evaluations

Several popular methods for *intrinsic* evaluations:

- ▶ How well do (cosine) similarities of pairs of words' vectors correlate with judgments of word similarity by humans?
- ▶ TOEFL-like synonym tests, e.g.,  $\text{rug} \xrightarrow{?} \{\text{sofa}, \text{ottoman}, \text{carpet}, \text{hallway}\}$

# Word Vector Evaluations

Several popular methods for *intrinsic* evaluations:

- ▶ How well do (cosine) similarities of pairs of words' vectors correlate with judgments of word similarity by humans?
- ▶ TOEFL-like synonym tests, e.g.,  $\text{rug} \xrightarrow{?} \{\text{sofa}, \text{ottoman}, \text{carpet}, \text{hallway}\}$
- ▶ Syntactic analogies, e.g., “*walking* is to *walked* as *eating* is to what?” Solved via:

$$\min_{v \in \mathcal{V}} \cos(\mathbf{v}_v, \mathbf{v}_{\text{walking}} - \mathbf{v}_{\text{walked}} + \mathbf{v}_{\text{eating}})$$

# Extrinsic Evaluations

1. Use large unannotated corpus to get your word vectors (called **pretraining**).
2. Use them in a text classifier (or some other NLP system, more examples to be introduced in future lectures). Two options:
  - ▶ Plug in word vectors as “frozen” features.
  - ▶ Treat them as parameters of the text classifier; pretraining gives initial values, but they get updated, or “finetuned” during supervised learning.
3. Does that system’s performance improve?

# Stepping Back

Big ideas:

- ▶ This is *unsupervised* learning: all you need is lots of raw text (no labels!)
- ▶ Large corpora → powerful word representations (the distributional hypothesis from linguistics, brought to life through engineering)
- ▶ It's all about the relationship between *words* and their *contexts*

# Embeddings from Language Models (ELMo)

(Peters et al., 2018)

Why not give every word *token* (in context) its own vector?

- ▶ To do that, we need a function that maps contexts to vectors.
- ▶ An RNN language model can do that, for the entire left context.
  - ▶  $\mathbf{s}_i$  is  $x_i$ 's history.  $\mathbf{s}_{i+1} = f_{\text{recurrent}}(\mathbf{e}_{x_i}, \mathbf{s}_i)$  is a (left-) *contextualized* embedding of the token  $x_i$ .
  - ▶ To get the right-context, run an RNN language model from right to left, and extract the analogous state vector just after reading  $x_i$ . Concatenate the two.
- ▶ On extrinsic evaluations, this method gave big improvements to state of the art systems.

# References I

- Yoshua Bengio, Réjean Ducharme, Pascal Vincent, and Christian Jauvin. A neural probabilistic language model. *Journal of Machine Learning Research*, 3(Feb):1137–1155, 2003. URL <http://www.jmlr.org/papers/volume3/bengio03a/bengio03a.pdf>.
- Scott C. Deerwester, Susan T. Dumais, Thomas K. Landauer, George W. Furnas, and Richard A. Harshman. Indexing by latent semantic analysis. *Journal of the American Society for Information Science*, 41(6):391–407, 1990.
- J. R. Firth. A synopsis of linguistic theory 1930–1955. In *Studies in Linguistic Analysis*, pages 1–32. Blackwell, 1957.
- Zellig Harris. Distributional structure. *Word*, 10(23):146–162, 1954.
- Omer Levy and Yoav Goldberg. Neural word embedding as implicit matrix factorization. In *NIPS*, 2014.
- Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. In *Proceedings of ICLR*, 2013a. URL <http://arxiv.org/pdf/1301.3781.pdf>.
- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S. Corrado, and Jeff Dean. Distributed representations of words and phrases and their compositionality. In *NIPS*, 2013b. URL <http://papers.nips.cc/paper/5021-distributed-representations-of-words-and-phrases-and-their-compositionality.pdf>.
- Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. Deep contextualized word representations. In *Proceedings of NAACL*, 2018.
- Gerard Salton, Anita Wong, and Chung-Shu Yang. A vector space model for automatic indexing. *Communications of the ACM*, 18(11):613–620, 1975.