

CSE 573: Artificial Intelligence

Problem Spaces & Search

Dan Weld


With slides from
Dan Klein, Stuart Russell, Andrew Moore, Luke Zettlemoyer, Dana Nau...

Outline

- Search Problems
- Uninformed Search Methods
 - Depth-First Search
 - Breadth-First Search
 - Iterative Deepening Search
 - Uniform-Cost Search
- Heuristic Search Methods
- Heuristic Generation



Agent vs. Environment

- An **agent** is an entity that *perceives* and *acts*.
- A **rational agent** selects actions that maximize its **utility function**.
- Characteristics of the **percepts**, **environment**, and **action space** dictate techniques for selecting rational actions.




Goal Based Agents

- Plan ahead
- Ask "what if"
- Decisions based on (hypothesized) consequences of actions
- Must have a model of how the world evolves in response to actions
- Act on how the world **WOULD BE**



Search: It's not just for Agents

Hardware verification



Planning optimal repair sequences



11

Search thru a Problem Space (aka State Space)

• Input:

- Set of states
- Operators ~~and costs~~
- Start state
- Goal state *[or test]*

Functions: States \rightarrow States

Aka "Successor Function"

• Output:

- Path: start \Rightarrow a state satisfying goal test
[May require shortest path]
[Sometimes just need a state that passes test]


Example: Simplified Pac-Man

▪ Input:

- A state space



- Successor function



- A start state
- A goal test

▪ Output:

Ex: Route Planning: Arad \rightarrow Bucharest

▪ Input:

- Set of states
- Operators [and costs]
- Start state
- Goal state (test)


Different operators may be applicable in different states

▪ Output:



Ex: Dock Worker Robots

- A harbor with several locations
 - e.g., docks, docked ships, storage areas, parking areas
- Containers
 - going to/from ships
- Robot vehicles
 - can move containers
- Cranes
 - can load and unload containers
- Multiple robots can operate at the same time
- Move, load & other actions have **different durations**




Slide adapted from Dana Nau "Automated Planning", licensed under creative commons NC share alike

Dock Worker 2


Input:

- Set of states
 - Partially specified plans
- Operators [and costs]
 - Plan modification operators
- Start state
 - The null plan (no actions)
- Goal test
 - A plan which provably achieves the desired world configuration



Slide adapted from Dana Nau "Automated Planning", licensed under creative commons NC share alike

Plan Space




Blue boxes are plans = states in search space
 Operators modify plans
 Successors(p) = all possible ways of modifying p

18

Multiple Problem Spaces

Real World

States of the world (e.g. loading dock configurations)
 Actions (take one world-state to another)



Robot's Head

• Problem Space 1

- PS states =
 - models of world states
- Operators =
 - models of actions

• Problem Space 2

- PS states =
 - partially spec. plan
- Operators =
 - plan modificat'n ops



Algebraic Simplification


$$\begin{aligned} \partial_t^2 u &= - \left[E'' - \frac{l(l+1)}{r^3} - r^2 \right] u(r) \\ e^{-2s} (\partial_s^2 - \partial_s) u(s) &= - \left[E'' - l(l+1)e^{-2s} - e^{2s} \right] u(s) \\ e^{-2s} \left[e^{\frac{1}{2}s} \left(e^{-\frac{1}{2}s} u(s) \right)'' - \frac{1}{4} u \right] &= - \left[E'' - l(l+1)e^{-2s} - e^{2s} \right] u(s) \\ e^{-2s} \left[e^{\frac{1}{2}s} \left(e^{-\frac{1}{2}s} u(s) \right)' \right]' &= - \left[E'' - \left(l + \frac{1}{2} \right)^2 e^{-2s} - e^{2s} \right] u(s) \\ v'' &= -e^{2s} \left[E'' - \left(l + \frac{1}{2} \right)^2 e^{-2s} - e^{2s} \right] v \end{aligned}$$

- **Input:**
 - Set of states
 - Operators [and costs]
 - Start state
 - Goal state (test)
- **Output:**

20

State Space Graphs


- **State space graph:**
 - Each node is a state
 - The operators are represented by arcs
 - Edges may be labeled with costs



- In contrast to other areas of CS... *Ridiculously tiny search graph for a tiny search problem*
- We can rarely build this graph in memory (so we don't try)
- ... just a conceptual tool

State Space Sizes?

- Search Problem:
Eat all of the food
- Pacman positions:
 $10 \times 12 = 120$
- Pacman facing:
up, down, left, right
- Food configurations: 2^{30}
- Ghost1 positions: 12
- Ghost 2 positions: 11




$$120 \times 4 \times 2^{30} \times 12 \times 11 = 6.8 \times 10^{13}$$

Search Methods

- **Blind Search**
 - Depth first search
 - Breadth first search
 - Iterative deepening search
 - Uniform cost search
- **Local Search**
- **Informed Search**
- **Constraint Satisfaction**
- **Adversary Search**

Search Trees




▪ A search tree:

- Start state at the root node
- Children correspond to successors
- Nodes **contain** states, correspond to **PLANS** to those states
- Edges are labeled with actions and costs
- For most problems, we can never actually build the whole tree


Example: Tree Search

State graph:



What is the search tree?

State Graphs vs. Search Trees




Each **NODE** in the search tree denotes an entire **PATH** in the problem graph.

We construct both on demand – and we construct as little as possible.

Why duplicate nodes?

States vs. Nodes

- Vertices in state space graphs are problem states
 - Represent an abstracted state of the world
 - Have successors, can be goal / non-goal, have multiple predecessors
- Vertices in search trees ("Nodes") are plans
 - Contain a **problem state** and one parent, a path length, a depth & a cost
 - Represent a plan (sequence of actions) which results in the node's state
 - The same problem state may be achieved by multiple search tree nodes



Building Search Trees



Search:

- Expand out possible nodes (plans) in the tree
- Maintain a **fringe** of **unexpanded** nodes
- Try to expand as few nodes as possible

General Tree Search

```

function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
  
```


Important ideas:

- Fringe** (leaves of tree)
- Expansion** (adding successors of a leaf)
- Exploration **strategy**
which fringe node to expand next?

Detailed pseudocode is in the book!

Review: Breadth First Search



Strategy: expand **shallowest** node first
Implementation:
 Fringe is a queue - FIFO



Review: Breadth First Search


Expansion order:

(S, d, e, p, b, c, e, h, r, q, a, a, h, r, p, q, f, p, q, f, q, c, G)



Review: Depth First Search



Strategy: expand **deepest** node first
Implementation:
 Fringe is a stack - LIFO



Wolog: assume we choose children in lexicographic order

Review: Depth First Search

Expansion ordering:
 (d, b, a, c, a, e, h, p, q, q, r, f, c, a, G)



Search Algorithm Properties

- **Complete?** Guaranteed to find a solution if one exists?
- **Optimal?** Guaranteed to find the least cost path?
- **Time complexity?**
- **Space complexity?**

Variables:

| | |
|-------|--|
| n | Number of states in the problem |
| b | The maximum branching factor B (the maximum number of successors for a state) |
| C^* | Cost of least cost solution |
| d | Depth of the shallowest solution |
| m | Max depth of the search tree |


Depth-First Search

Assuming finite tree

| Algorithm | Complete | Optimal | Time | Space |
|---------------------------|----------|---------|----------|----------------|
| DFS Depth First Search | No | No | $O(b^m)$ | $O(b \cdot m)$ |


- Infinite paths make DFS incomplete...

- How can we fix this?
- Check new nodes against **path** from S



d depth of solution
 m max depth of tree

DFS Search (w/ cycle checking)




| Algorithm | Complete | Optimal | Time | Space |
|-----------|-------------|---------|----------|----------------|
| DFS | Y if finite | N | $O(b^m)$ | $O(b \cdot m)$ |

Only if finite tree

BFS Tree Search

| Algorithm | Complete | Optimal | Time | Space |
|-----------|--------------------|---------|----------|----------|
| DFS | N w/ Path Checking | N | $O(b^m)$ | $O(bm)$ |
| BFS | Y* | Y* | $O(b^d)$ | $O(b^d)$ |




* Assuming finite branching factor

Memory a Limitation?

- **Suppose:**
 - 4 GHz CPU
 - 32 GB main memory
 - 100 instructions / expansion
 - 5 bytes / node
- 40 M expansions / sec
 - Memory filled in ... 3 min

Iterative Deepening Search

- DFS with limit; incrementally grow limit
- Evaluation



Puzzles



15 puzzle



3x3 Rubiks cube

44

Speed

Assuming 10M nodes/sec & sufficient memory

| | BFS | | Iter. Deep. | |
|---------------|-----------|---------|-------------|---------------|
| | Nodes | Time | Nodes | Time |
| 8 Puzzle | 10^5 | .01 sec | 10^5 | .01 sec |
| 2x2x2 Rubik's | 10^6 | .2 sec | 10^6 | .2 sec |
| 15 Puzzle | 10^{13} | 6 days | 10^{17} | 20k yrs |
| 3x3x3 Rubik's | 10^{19} | 68k yrs | 10^{20} | 574k yrs |
| 24 Puzzle | 10^{25} | 12B yrs | 10^{37} | 10^{23} yrs |

Why the difference?


Rubik has higher branch factor
15 puzzle has greater depth

of duplicates

Slide adapted from Richard Korf presentation

Tree vs Graph Search

In BFS, for example, we shouldn't bother expanding the circled nodes (why?)



Graph Search

- Very simple fix: never expand a state type twice

```

function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE(node)) then return node
    if STATE(node) is not in closed then
      add STATE(node) to closed
      fringe ← INSERTALL(EXPAND(node, problem), fringe)
  end
  
```



Some Hints

- On small problems
 - Graph search almost always better than tree search
 - Implement your closed list as a dict or set!
- On many real problems
 - Storage space is a huge concern
 - Graph search impractical

Search Methods

- Depth first search (DFS)
- Breadth first search (BFS)
- Iterative deepening depth-first search (IDS)

Blind search

49

Search Methods

- Depth first search (DFS)
- Breadth first search (BFS)
- Iterative deepening depth-first search (IDS)
- Best first search
- Uniform cost search (UCS)
- Greedy search
- A*
- Iterative Deepening A* (IDA*)
- Beam search
- Hill climbing

Heuristic search

50


Blind vs Heuristic Search

- Costs on Actions
- Heuristic Guidance

Separable issues, but usually linked.


51

Costs on Actions



Objective: Path with smallest overall cost

Costs on Actions



What will BFS return?

... finds the shortest path in terms of number of transitions.
It does **not** find the least-cost path.

Best-First Search

- Generalization of breadth-first search
- Fringe = **Priority** queue of nodes to be explored
- Cost function $f(n)$ applied to each node

54



Priority Queue Refresher

A priority queue is a data structure in which you can insert and retrieve (key, value) pairs with the following operations:

| | |
|----------------------------------|---|
| <code>pq.push(key, value)</code> | inserts (key, value) into the queue. |
| <code>pq.pop()</code> | returns the key with the lowest value, and removes it from the queue. |

- You can decrease a key's priority by pushing it again
- Unlike a regular queue, insertions aren't constant time, usually $O(\log n)$
- We'll need priority queues for cost-sensitive search methods

Best-First Search

- Generalization of breadth-first search
- Fringe = **Priority** queue of nodes to be explored
- Cost function $f(n)$ applied to each node

Add initial state to priority queue

While queue not empty

Node = head(queue)

If goal?(node) then return node

Add new children of node to queue

← "expanding the node"

56

Old Friend

Breadth First = Best First
with $f(n) = \text{depth}(n)$

57

Uniform Cost Search

Best first, where


$f(n) = \text{"cost from start to n"}$

aka "Dijkstra's Algorithm"

Uniform Cost Search


Expansion order:

S, p, d, b, e, a, r, f, e, G



Uniform Cost Search

| Algorithm | Complete | Optimal | Time | Space |
|-----------|------------------|-------------|-----------------------|-----------------------|
| DFS | w/ Path Checking | Y if finite | N | $O(b^m)$ |
| BFS | | Y | $O(b^d)$ | $O(b^d)$ |
| UCS | | Y* | $O(b^{C^*/\epsilon})$ | $O(b^{C^*/\epsilon})$ |




C^* = Optimal cost

ϵ = Minimum cost of an action

* We'll keep assuming finite branching factor

Uniform Cost Issues


- Remember: explores increasing cost contours



- The good: UCS is complete and optimal!


- The bad:

- Explores options in every "direction"
- No information about **goal location**




Uniform Cost: Pac-Man

- Cost of 1 for each action
- Explores all of the states, but one



What is a *Heuristic*?

- An *estimate* of how close a state is to a goal
- Designed for a particular search problem



- Examples: Manhattan distance: $10+5 = 15$
Euclidean distance: 11.2

A* Search

Hart, Nilsson & Rafael 1968

Best first search with $f(n) = g(n) + h(n)$

- $g(n)$ = sum of costs from start to n
- $h(n)$ = (admissible) estimate of lowest cost path from n \rightarrow goal note: $h(\text{goal}) = 0$

Can view as cross-breed:

$g(n)$ ~ uniform cost search

$h(n)$ ~ greedy search

Best of both worlds...

68

*Nomenclature
Confusion:
Is h the heuristic or is f?*

A* Search

Hart, Nilsson & Rafael 1968

Best first search with $f(n) = g(n) + h(n)$

- $g(n)$ = sum of costs from start to n
- $h(n)$ = admissible estimate of lowest cost path from n \rightarrow goal $h(\text{goal}) = 0$

If $h(n)$ is **admissible** and **monotonic**
then A* is optimal


Underestimates (\leq) cost
of reaching goal from node

f values never decrease
From node to descendants
(triangle inequality)

69

Is Manhattan distance **admissible**?


- Underestimate?



70


Is Manhattan distance **monotonic**?

- f values **increase** from node to children
- (triangle inequality)



71

Monotonicity




$$F(a) \geq F(b)$$

$$G(a)+H(a) \geq G(b)+H(b)$$

72

Euclidean Distance

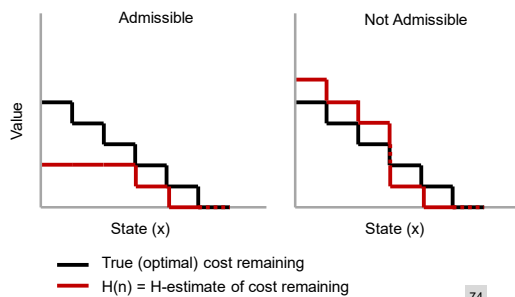
- Admissible?
- Monotonic?



73

Admissible Heuristics

$$F(n) = G(n) + H(n)$$




Slide credit: Travis Mandel

74

Monotonic/Consistent Heuristics

$$F(n) = G(n) + H(n)$$

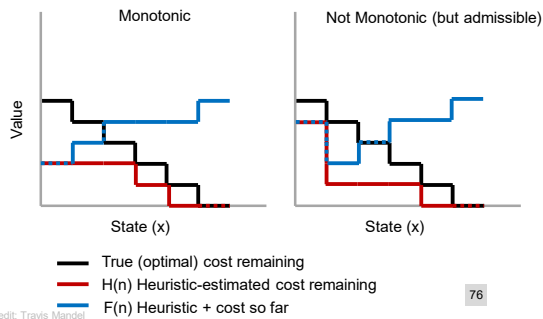


Slide credit: Travis Mandel

75

Monotonic/Consistent Heuristics

$$F(n) = G(n) + h(n)$$




Slide credit: Travis Mandel

76

76

Optimality of A* (tree search)

Suppose some suboptimal goal G_2 has been generated and is in the queue.
Let n be an unexpanded node on a shortest path to an optimal goal G_1 .



$$\begin{aligned} f(G_2) &= g(G_2) && \text{since } h(G_2) = 0 \\ &> g(G_1) && \text{since } G_2 \text{ is suboptimal} \\ &\geq f(n) && \text{since } h \text{ is admissible} \end{aligned}$$

Since $f(G_2) > f(n)$, A* will never select G_2 for expansion


Monotonicity required for proof in graph search version

77

Optimality Continued

Lemma: A* expands nodes in order of increasing f value*

Gradually adds " f -contours" of nodes (cf. breadth-first adds layers)
Contour i has all nodes with $f = f_i$, where $f_i < f_{i+1}$



A* Example




79

A* Example




80

A* Example




81

A* Example




82

A* Example




83

A* Example



84

European Example



85

A* Summary

Pros

Produces optimal cost solution!

Does so quite quickly (focused)

Cons

Maintains priority queue...

Which can get exponentially big ☹


86

Iterative-Deepening A*

Like iterative-deepening depth-first, but...

Depth bound modified to be an **f-limit**

- Start with $f\text{-limit} = h(\text{start})$
- Prune any node if $f(\text{node}) > f\text{-limit}$
- Next $f\text{-limit} = \text{min-cost of any node pruned}$



87

IDA* Analysis

- Complete & Optimal (ala A*)
- Space usage \propto depth of solution
- Each iteration is DFS - no priority queue!
- # nodes expanded relative to A*
 - Depends on # unique values of heuristic function
 - In 8 puzzle: few values \Rightarrow close to # A* expands
 - In traveling salesman: each f value is unique
 - $\Rightarrow 1+2+\dots+n = O(n^2)$ where n =nodes A* expands
 - if n is too big for main memory, n^2 is too long to wait!

89

Forgetfulness

- A* used exponential memory
- How much does IDA* use?
 - During a run?
 - In between runs?
 - SMA*

© Daniel S. Weld

90