

The Big Data Quadfecta

Brian O'Neill
Lead Architect, Health Market Science
@boneill42, bone@alumni.brown.edu


Taylor Goetz
Development Lead, Health Market Science
@ptgoetz, ptgoetz@gmail.com

Quadfecta?

1. Quadfecta

- A legendary beirut/beer pong shot that lands on the tops of four cups simultaneously. Considered the rarest shot in the game, topping even the trifecta, 2-cup knockover-and-sink, and simultaneous 6-cup game-ending double bounce-in

- Kafka
- Storm
- Elastic Search
- Cassandra



3 V's



Volume



Variety



Velocity

The Use Case

Our Mission

**THE SCIENCE OF BETTER
MARKET INTELLIGENCE**

Gain insights on the most
valuable physicians

[LEARN MORE »](#)

525624657824657893285256246578
98745332217533524678934689221
46829252385924634567468292219
874533221544487947898743332115
465789523383263456234657895238
332215532625678987459322155326
392528194834562246982925281946
292246524648254246982922465246
678987453357865235256789874533
563246862922185326345622468629
3878952315223281946487478952352
322155326396789874533221553263
292819464845622468292528194648
VERIFIED 465778452819464879424657

- 🎬 Prescriber eligibility and remediation
- 🎬 Eliminate fraud, waste and abuse
- 🎬 Insights into the healthcare space

The Business

Master Data Solutions

Health Care Provider & Facilities

Variety/Velocity

- >12000 of sources
- 6 Million unique HCPs
- 10+ years history

Data Challenges

- Constant change in real world data
- Conflicting & partial info
- Frequent changes to source structure
- Authoritative sources vs. crowdsource
- Predicting source quality

Business Solutions

CompleteView, Expense Manager, CompleteSpend

Prescriber Eligibility/Remediation

Analytics (Influencer Networks)

Medical Claims Data

Medical Procedures & Diagnosis

Volume/Velocity

- ~1B claims annually
- +5B records annually
- 5+ years history

Data Challenges

- Sources have incomplete capture
- Overlapping source data
- Statistical projections & biases
- Social media type relationships

Our Solutions

Business Needs



Sales & Marketing



Compliance




Business Systems




Finance & Legal

Solutions




Provider Data



Data Assessment, Integration & Enrichment Services





Compliance



Market Intelligence

Advanced Technology



Master Data Management


HMS Authoritative Sources



PDC



Medical Claims



Federal



State



Web



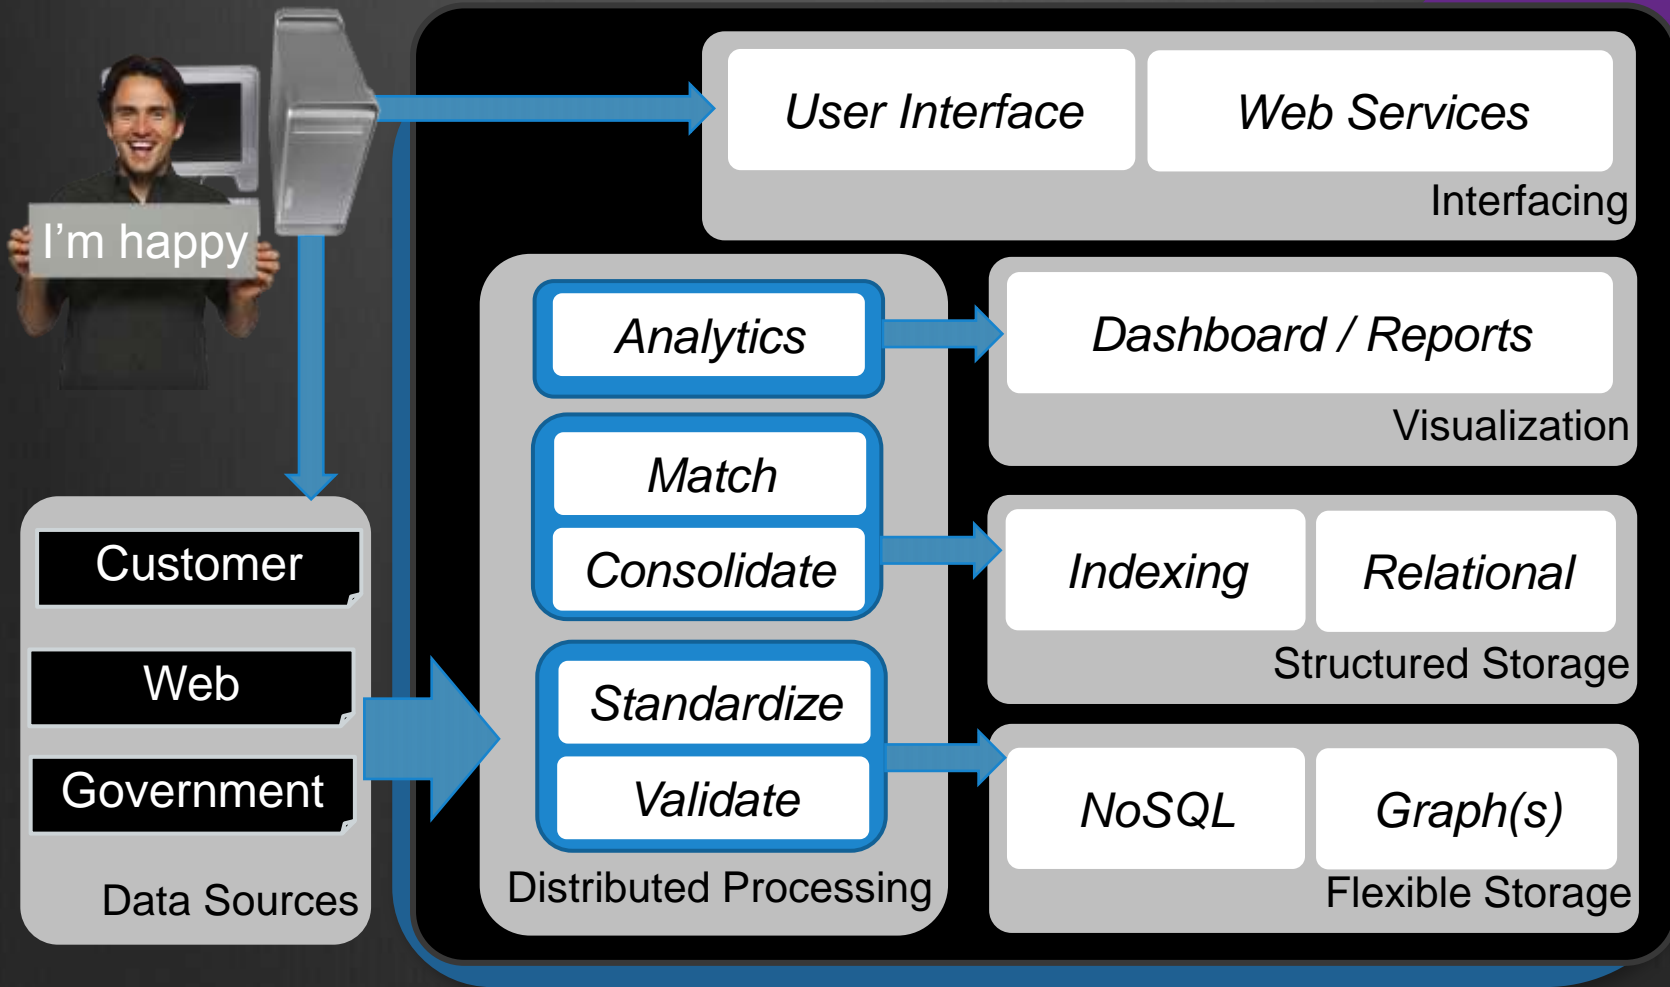
Derived

Datacenter


- $\frac{3}{4}$ Petabytes of raw storage
- Virtualized (VMware)
- On a SAN
- Should we go physical???



Under the Hood



Master Data Management



Harvested

$$f_{address} \hat{=} F@t_0$$

Government

$$f_{license} \hat{=} F@t_5$$

Private

$$f_{sanction} \hat{=} F@t_1$$

$$f_{sanction} \hat{=} F@t_4$$

Schema Change!

The Design

System of Record



Flexibility (Variety)
Scalability (Velocity + Volume)

Design Principles

⌚ Patterns

⌚ Idempotent Operations

- ⌚ Elegantly handle replay

⌚ Immutable data

- ⌚ Assertions of facts over time


⌚ Anti-Patterns

⌚ Transactions / Locking



State / Counting

- Exactly-once semantics for state
- Create small batches
- Order batches




Batch	Total
1	4
3	4 (wait)
2	10 (+6)
3	23 (+13)
3'	23 (+0)

What we did wrong...



- ❌ Could not react to transactional changes
- ❌ Needed extra logic to track what changed
- ❌ Took too long

What we did wrong... (II)



- ❶ AOP-based triggers
 - ❶ Worked well initially.
 - ❶ Business Processes captured as side-effects.

What we did right.

- REST APIs for Loose Coupling

- See Virgil:


- <https://github.com/hmsonline/virgil>

- But really... watch out for Intravert

- <https://github.com/zznate/intravert-ug>

Kafka

- Millions of Messages
- Replay Enabled
- No transactions / Lightning Fast




Elastic Search

- Edit Distance / Soundex
- Native Scalability
- Fuzzy Search
- Geospatial
- Facets

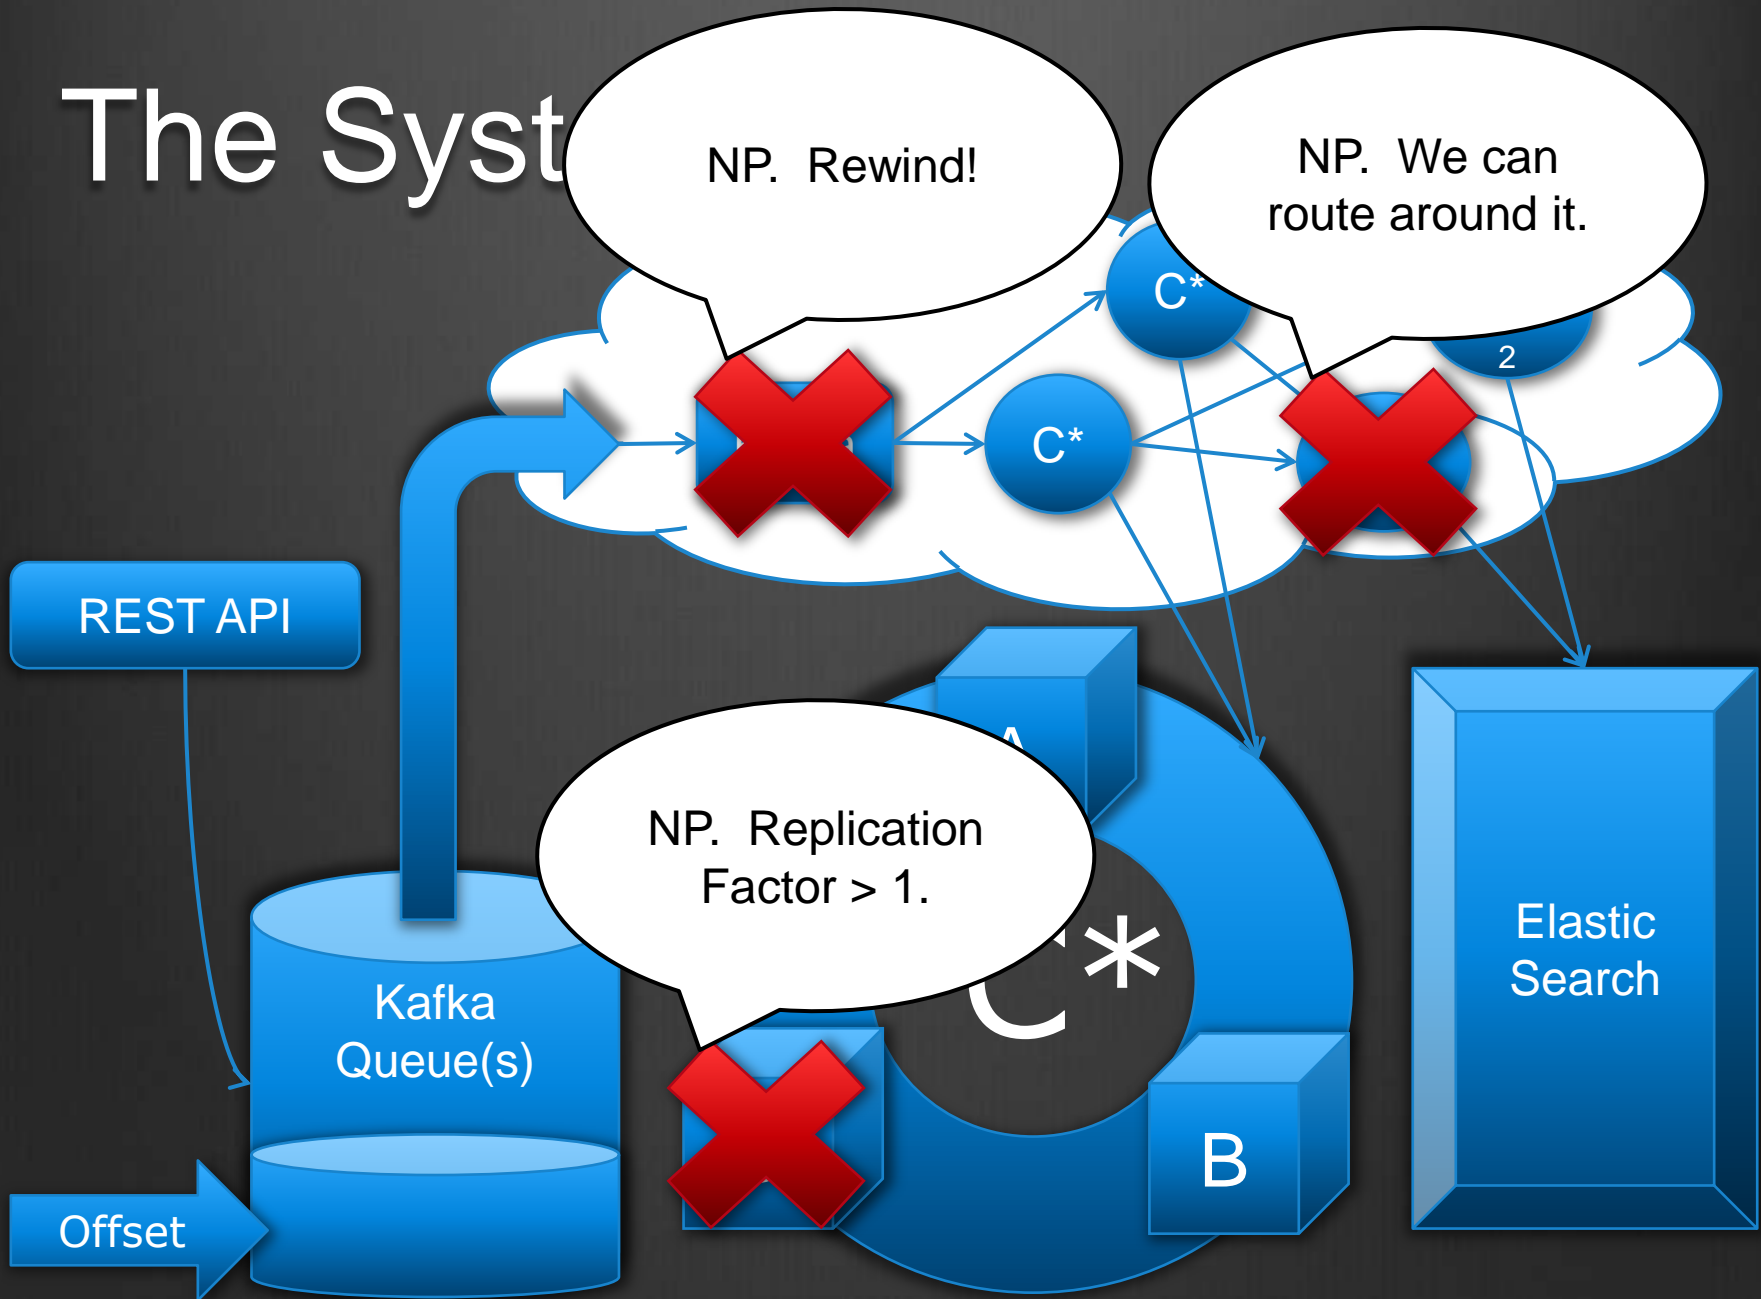


Storm


- Guaranteed once semantics
- Well-designed processing abstraction
- Beats BYODP
- Momentum



The System



What comes after Quadfecta?



?

Real-Time Integration

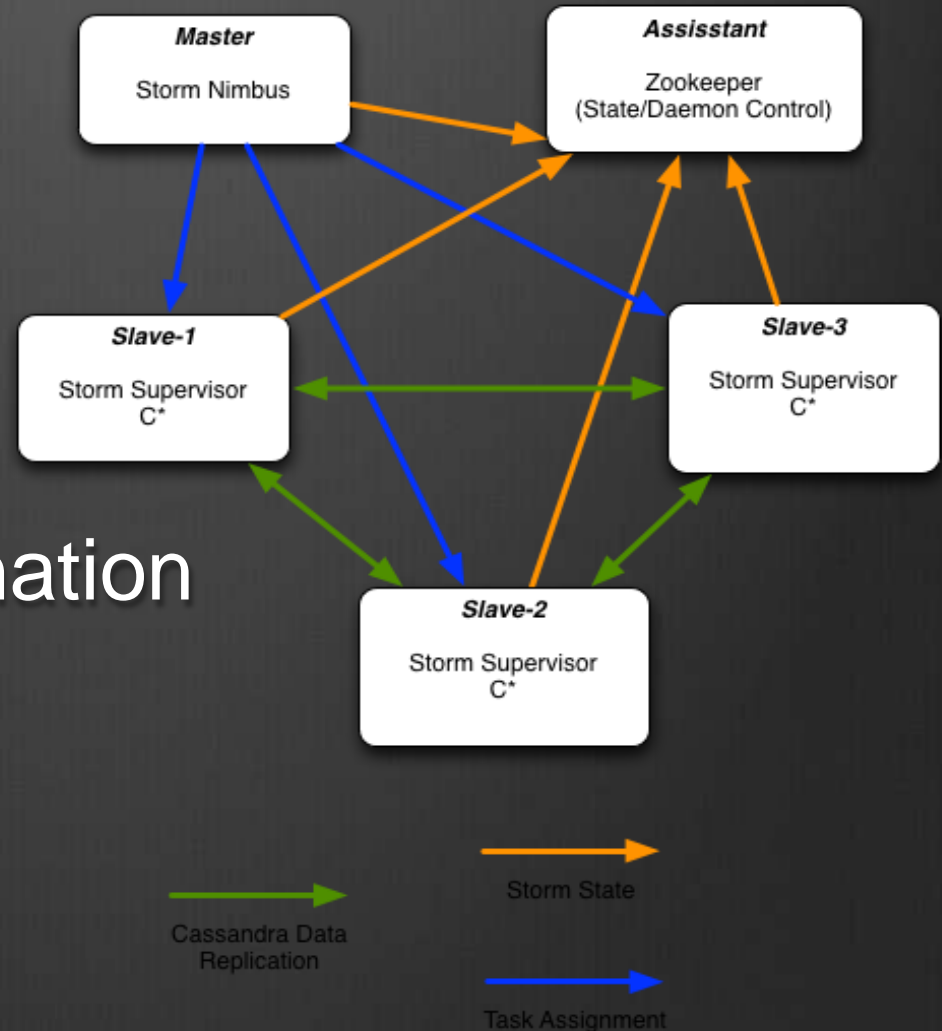
- ⦿ Real-time CRUD via Web Services
 - ⦿ DRPC
 - ⦿ “Real-time” Queue

Not quite sure?

The Storm/C* Bridge

Anatomy of a Storm Cluster

- ❶ Nimbus
 - ❶ Master Node
- ❷ Zookeeper
 - ❶ Cluster Coordination
- ❸ Supervisors
 - ❶ Worker Nodes



Storm Primitives

- Streams
 - Unbounded sequence of tuples
- Spouts
 - Stream Sources
- Bolts
 - Unit of Computation
- Topologies
 - Combination of n Spouts and n Bolts
 - Defines the overall “Computation”

Storm Spouts

- Represents a source (stream) of data
 - Queues (JMS, Kafka, Kestrel, etc.)
 - Twitter Firehose
 - Sensor Data
- Emits “Tuples” (Events) based on source
 - Primary Storm data structure
 - Set of Key-Value pairs



Storm Bolts

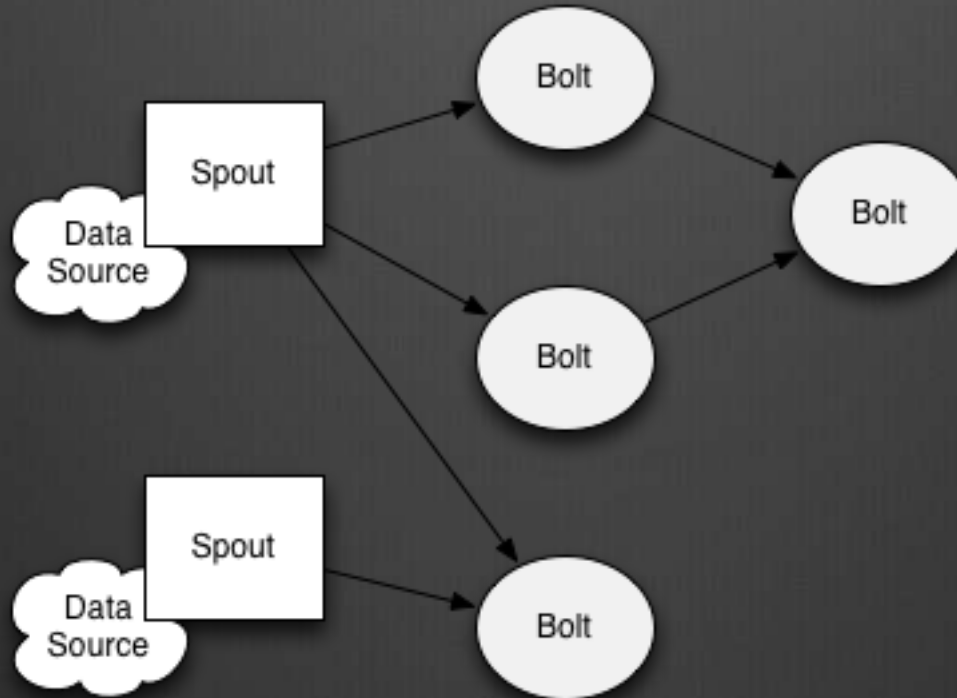
- ⦿ Receive Tuples from Spouts or other Bolts
- ⦿ Operate on, or React to Data
 - ⦿ Functions/Filters/Joins/Aggregations
 - ⦿ Database writes/lookups
- ⦿ Optionally emit additional Tuples



Storm Topologies

- Data flow between spouts and bolts
- Routing of Tuples between spouts/bolts
 - Stream “Groupings”
- Parallelism of Components
- Long-Lived

Storm Topologies



Storm and Cassandra

● Use Cases:

● Write Storm Tuple data to C*

- Computation Results
- Pre-compute indices

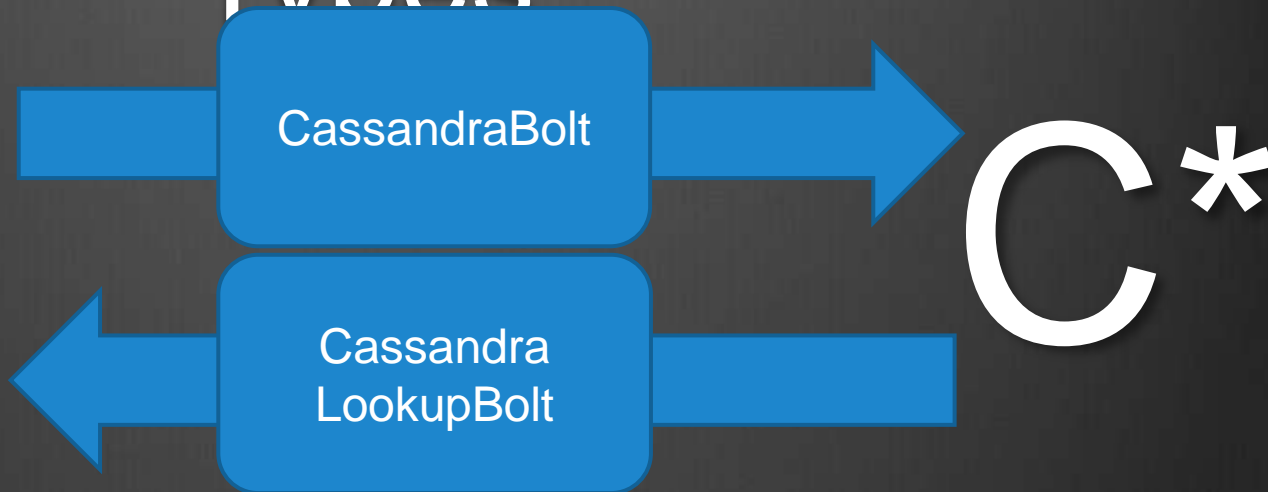
● Read data from C* and emit Storm Tuples

- Dynamic Lookups

<http://github.com/hmsonline/storm-cassandra>

Storm Cassandra Bolt

Types

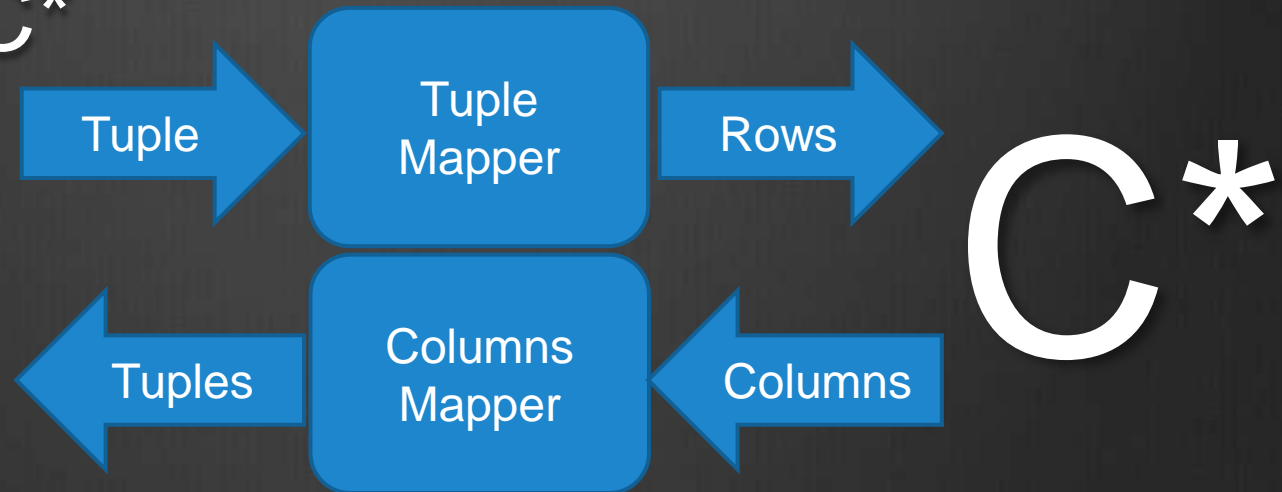


- ❶ CassandraBolt
 - ❶ Writes data to Cassandra
 - ❶ Available in Batching and Non-Batching
- ❶ CassandraLookupBolt
 - ❶ Reads data from Cassandra

<http://github.com/hmsonline/storm-cassandra>

Storm-Cassandra Project

- Provides generic Bolts for writing/reading Storm Tuples to/from C*



<http://github.com/hmsonline/storm-cassandra>

Storm-Cassandra Project

- TupleMapper Interface
 - Tells the CassandraBolt how to write a tuple to an arbitrary data model
- Given a Storm Tuple:
 - Map to Column Family
 - Map to Row Key
 - Map to Columns

<http://github.com/hmsonline/storm-cassandra>

Storm-Cassandra Project

• ColumnsMapper Interface

- Tells the CassandraLookupBolt how to transform a C* row into a Storm Tuple

• Given a C* Row Key and list of Columns:

- Return a list of Storm Tuples

<http://github.com/hmsonline/storm-cassandra>

Storm-Cassandra Project

- Current State:
 - Version 0.4.0
 - Uses Astyanax Client
 - Several out-of-the-box *Mapper Implementations:
 - Basic Key-Value Columns
 - Value-less Columns
 - Counter Columns
 - Lookup by row key
 - Lookup by range query
 - Composite Key/Column Support
 - Trident support

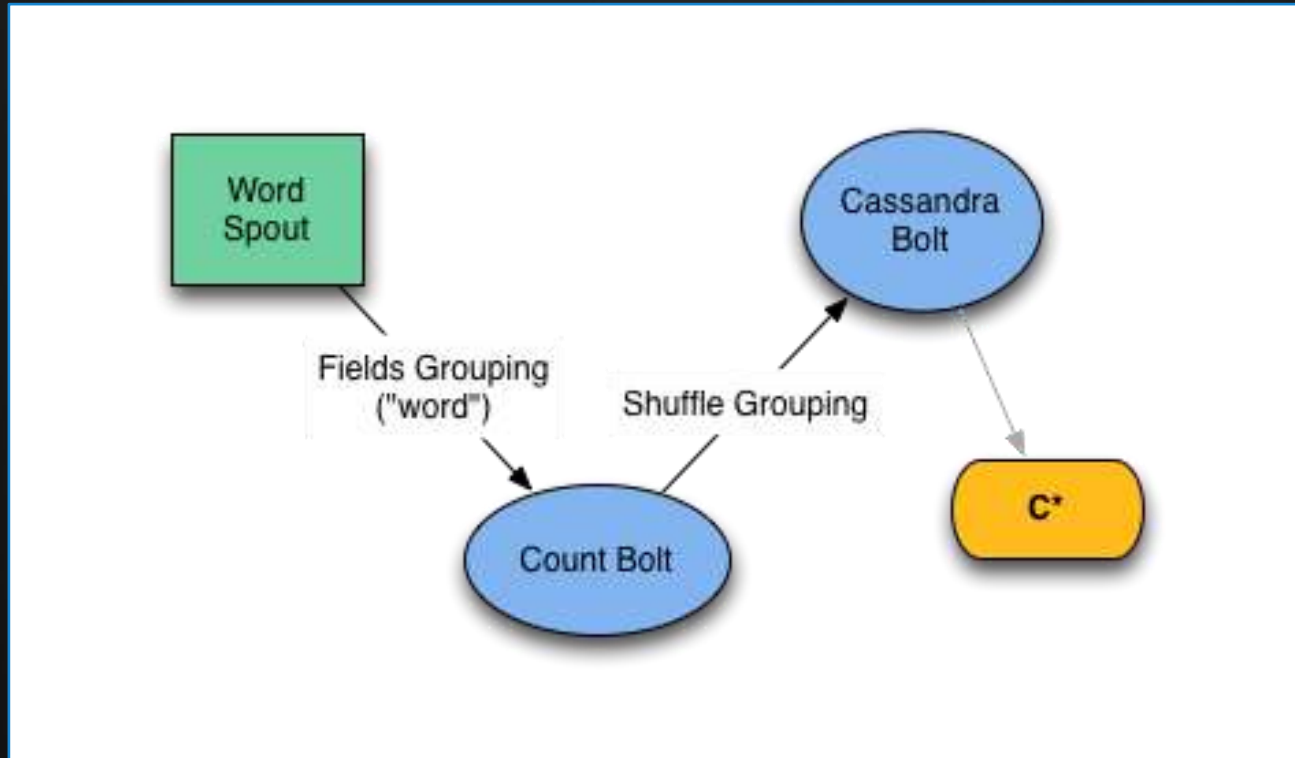
<http://github.com/hmsonline/storm-cassandra>

Storm-Cassandra Project

- Future Plans:
 - Switch to CQL
 - Enhanced Trident Support

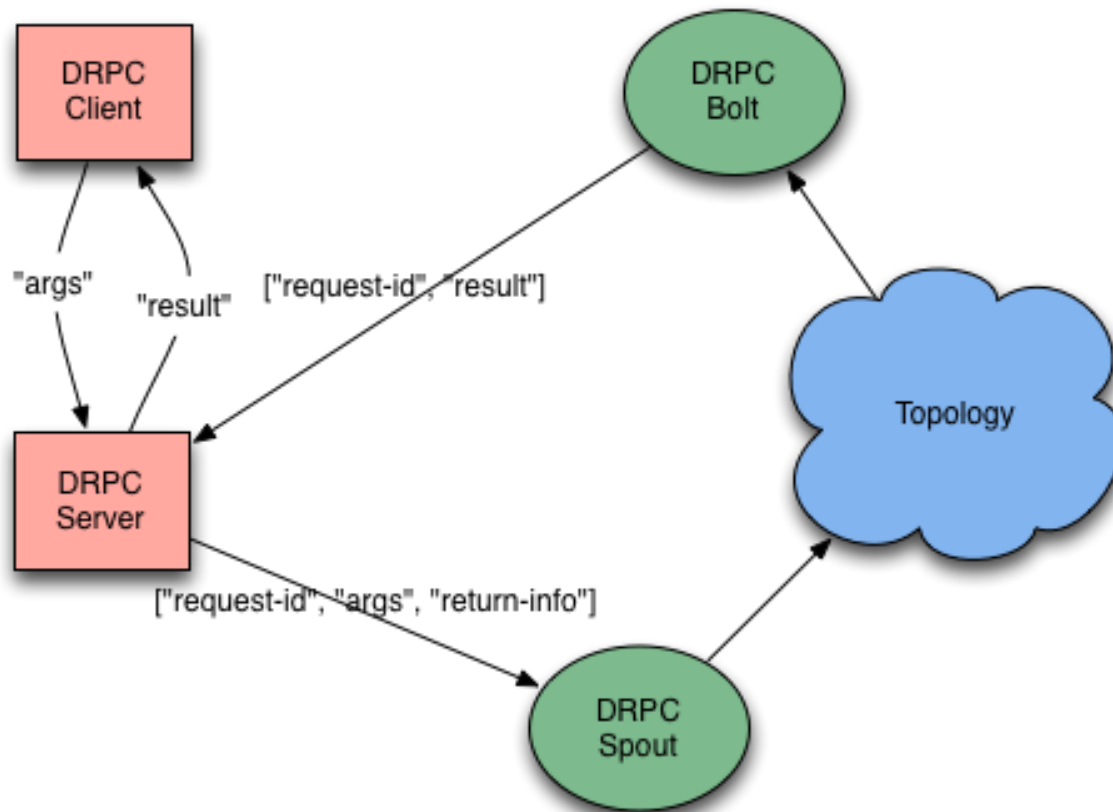
<http://github.com/hmsonline/storm-cassandra>

Persistent Word Count

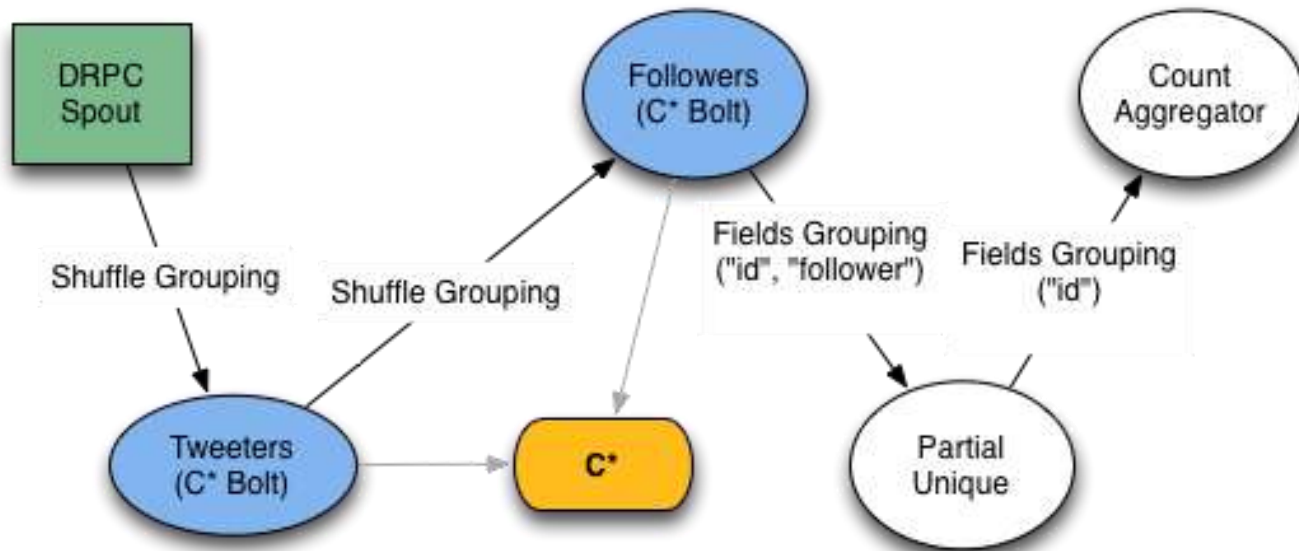


<http://github.com/hmsonline/storm-cassandra>

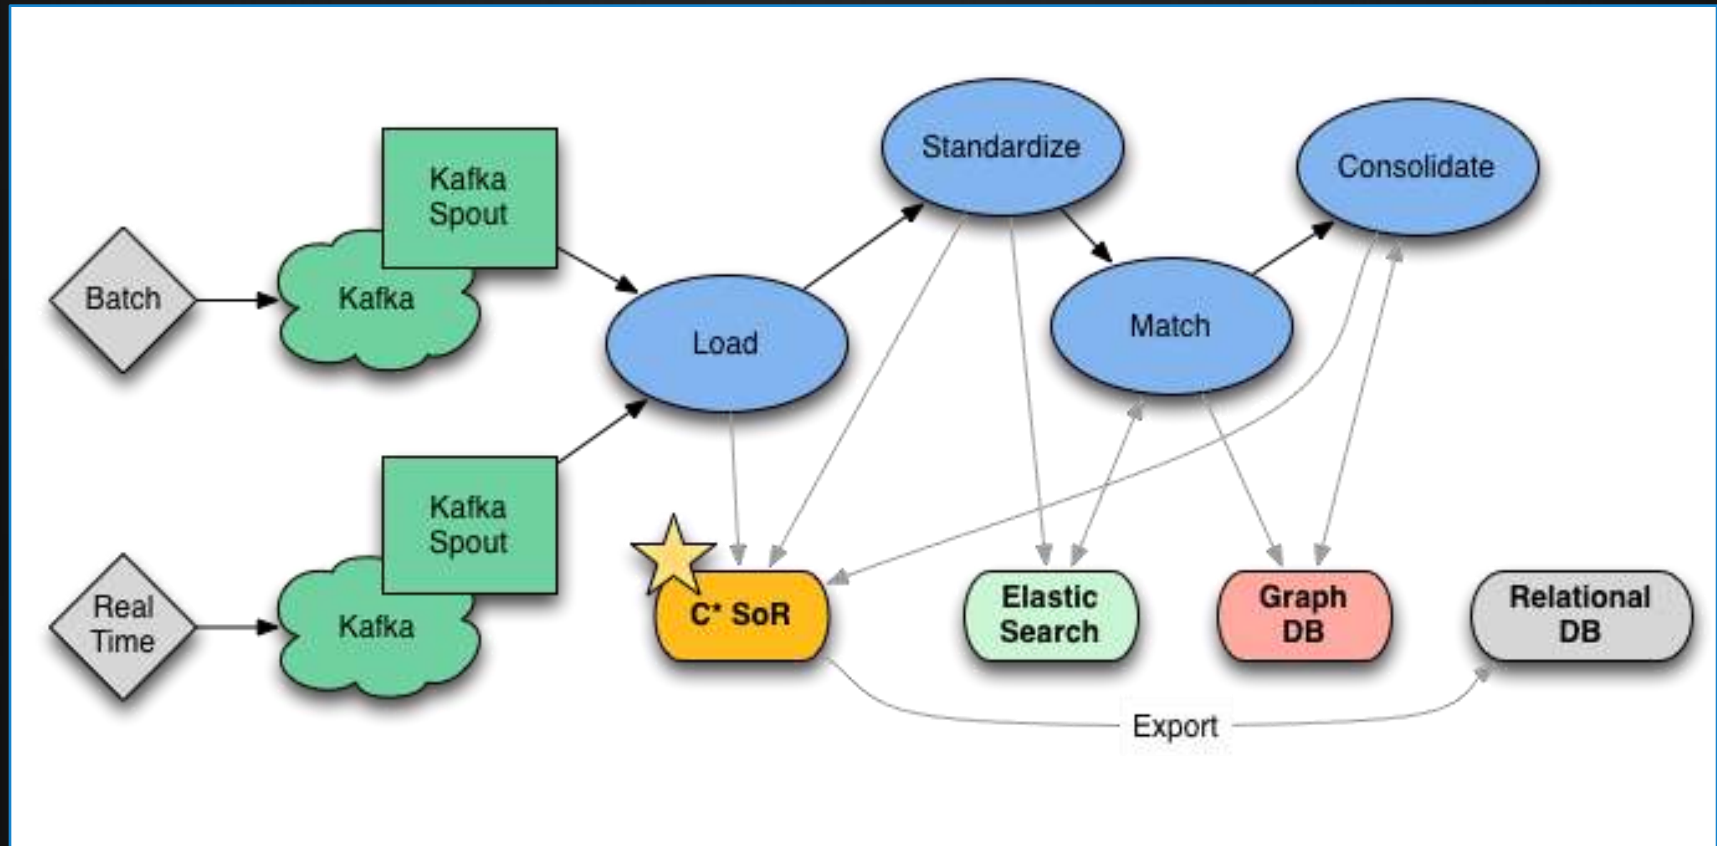
DRPC



“Reach” Computation

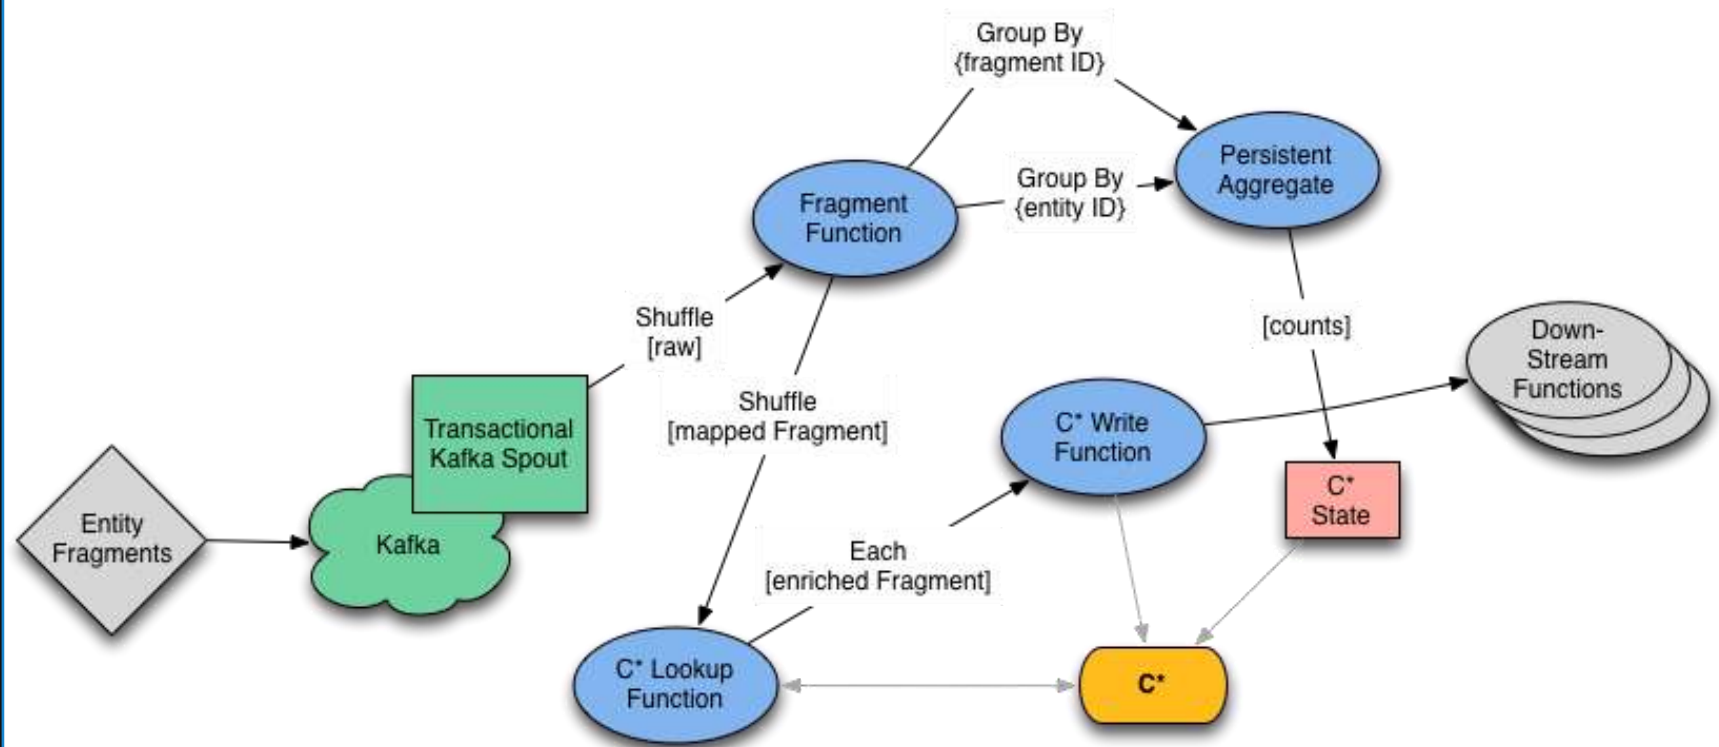


MDM Topology*



**Notional*

Load Topology



Shameless Shoutouts

- ⦿ HMS (<https://github.com/hmsonline/>)
 - ⦿ storm-cassandra
 - ⦿ storm-elastic-search
 - ⦿ storm-jdbi (coming soon)
- ⦿ ptgoetz (<https://github.com/ptgoetz>)
 - ⦿ storm-jms
 - ⦿ storm-signals

Next Level : Trident

Trident

- Provides a higher-level abstraction for stream processing
 - Constructs for state management and Batching
- Adds additional primitives that abstract away common topological patterns
- Deprecates transactional topologies
- Distributes with Storm

Sample Trident Operations

● Partition Local

● Functions (`execute(x) → x + y`)

● Filters (`isKeep(x) → 0,x`)

● PartitionAggregate

● Combiner (`pairwise combining`)

● Reducer (`iterative accumulation`)

● Aggregator (`byoa`)

A sample topology

```
TridentTopology topology = new TridentTopology();

TridentState wordCounts =

    topology.newStream("spout1", spout)

        .each(new Fields("sentence"),

            new Split(),

            new Fields("word"))

        .groupBy(new Fields("word"))

        .persistentAggregate(

            MemcachedState.opaque(serverLocations),

            new Count(),

            new Fields("count"))

        .parallelismHint(6);
```

Trident State

Sequenced writes by batch/transaction id.

❁ Spouts

❁ Transactional

- ❁ Batch contents never change

❁ Opaque

- ❁ Batch contents can change

❁ State

❁ Transactional

- ❁ Store tx_id with counts to maintain sequencing of writes.

❁ Opaque

- ❁ Store previous value in order to overwrite the current value when contents of a batch change.